



oscon.com/java

The Ghost in the Virtual Machine

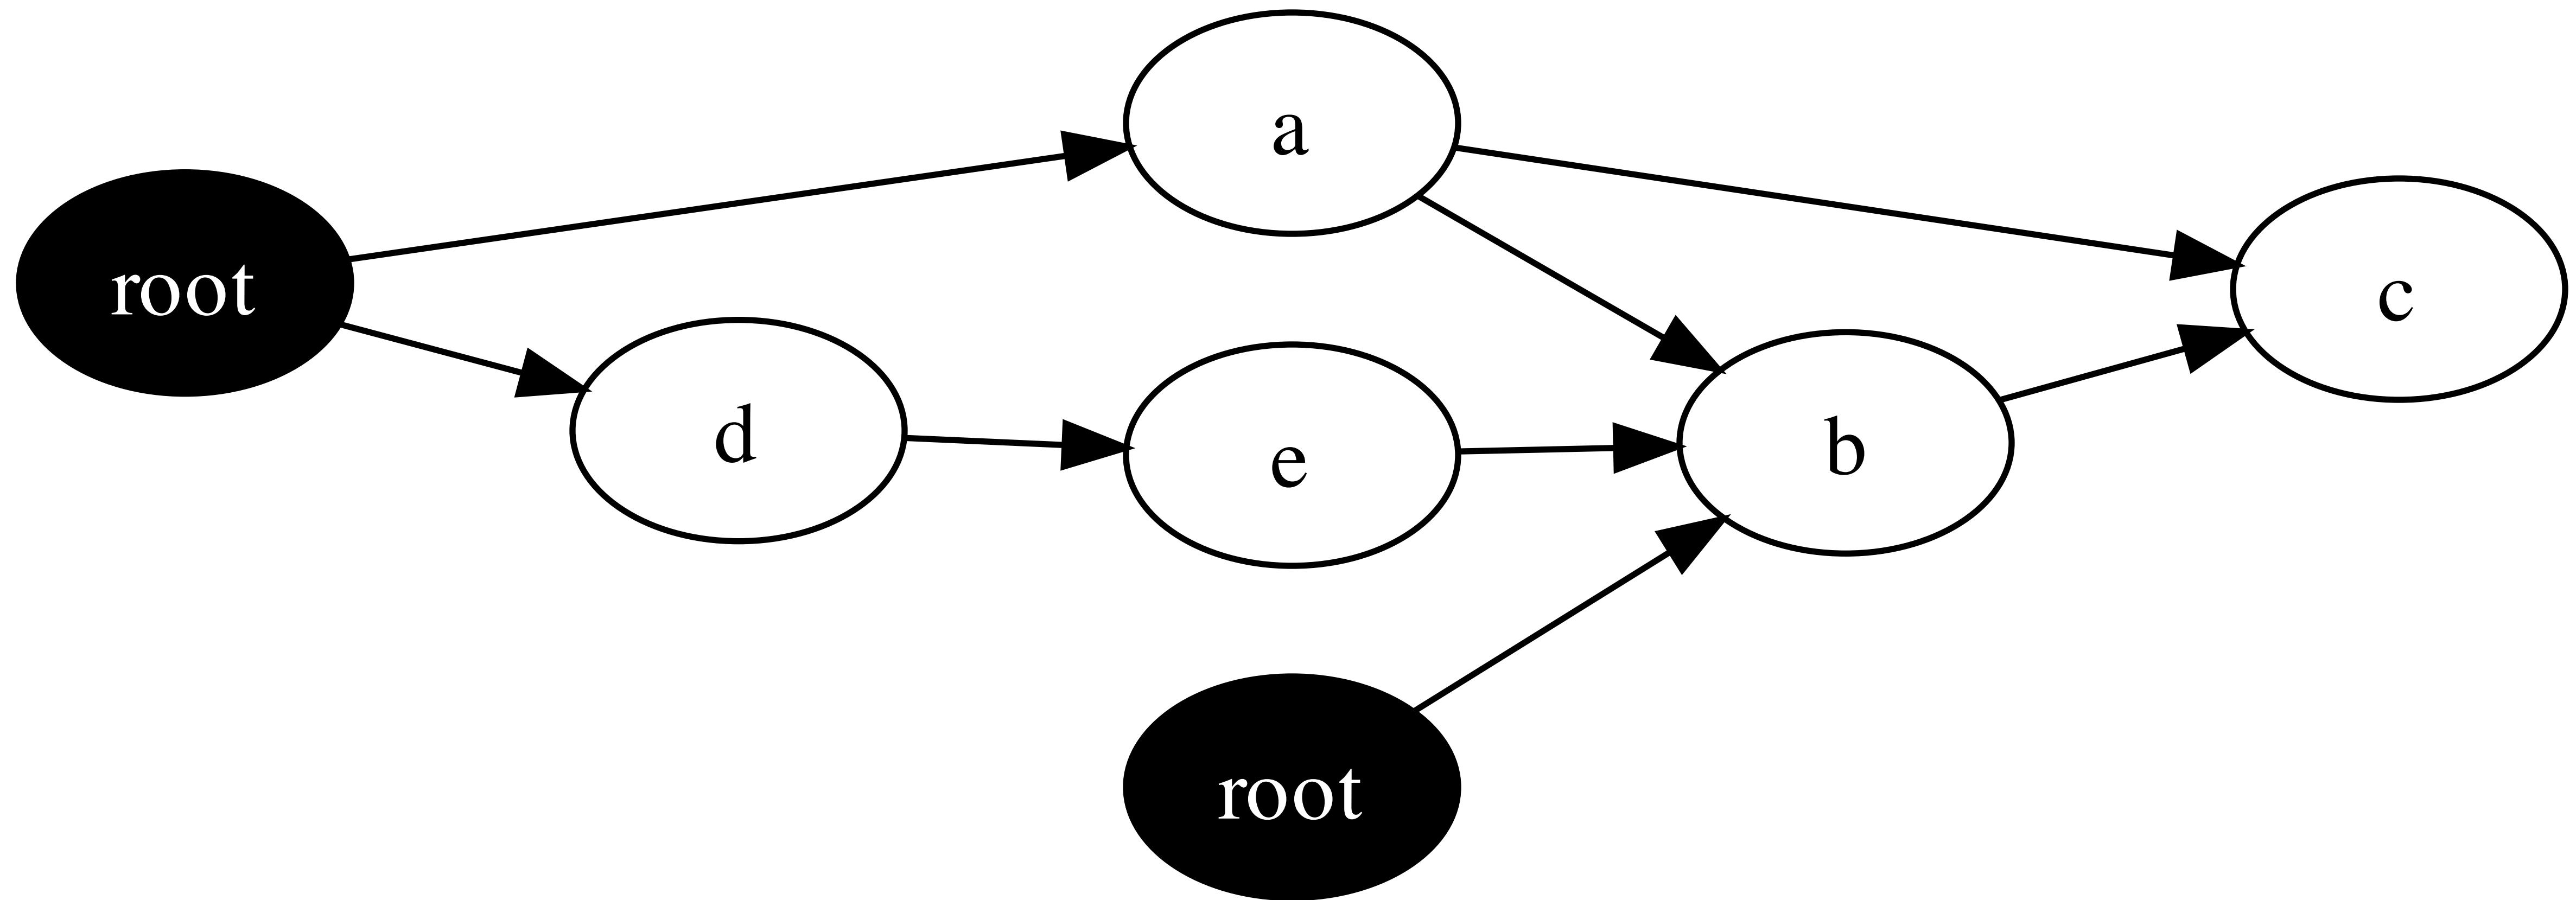
A Reference to References

Bob Lee
Square Inc.

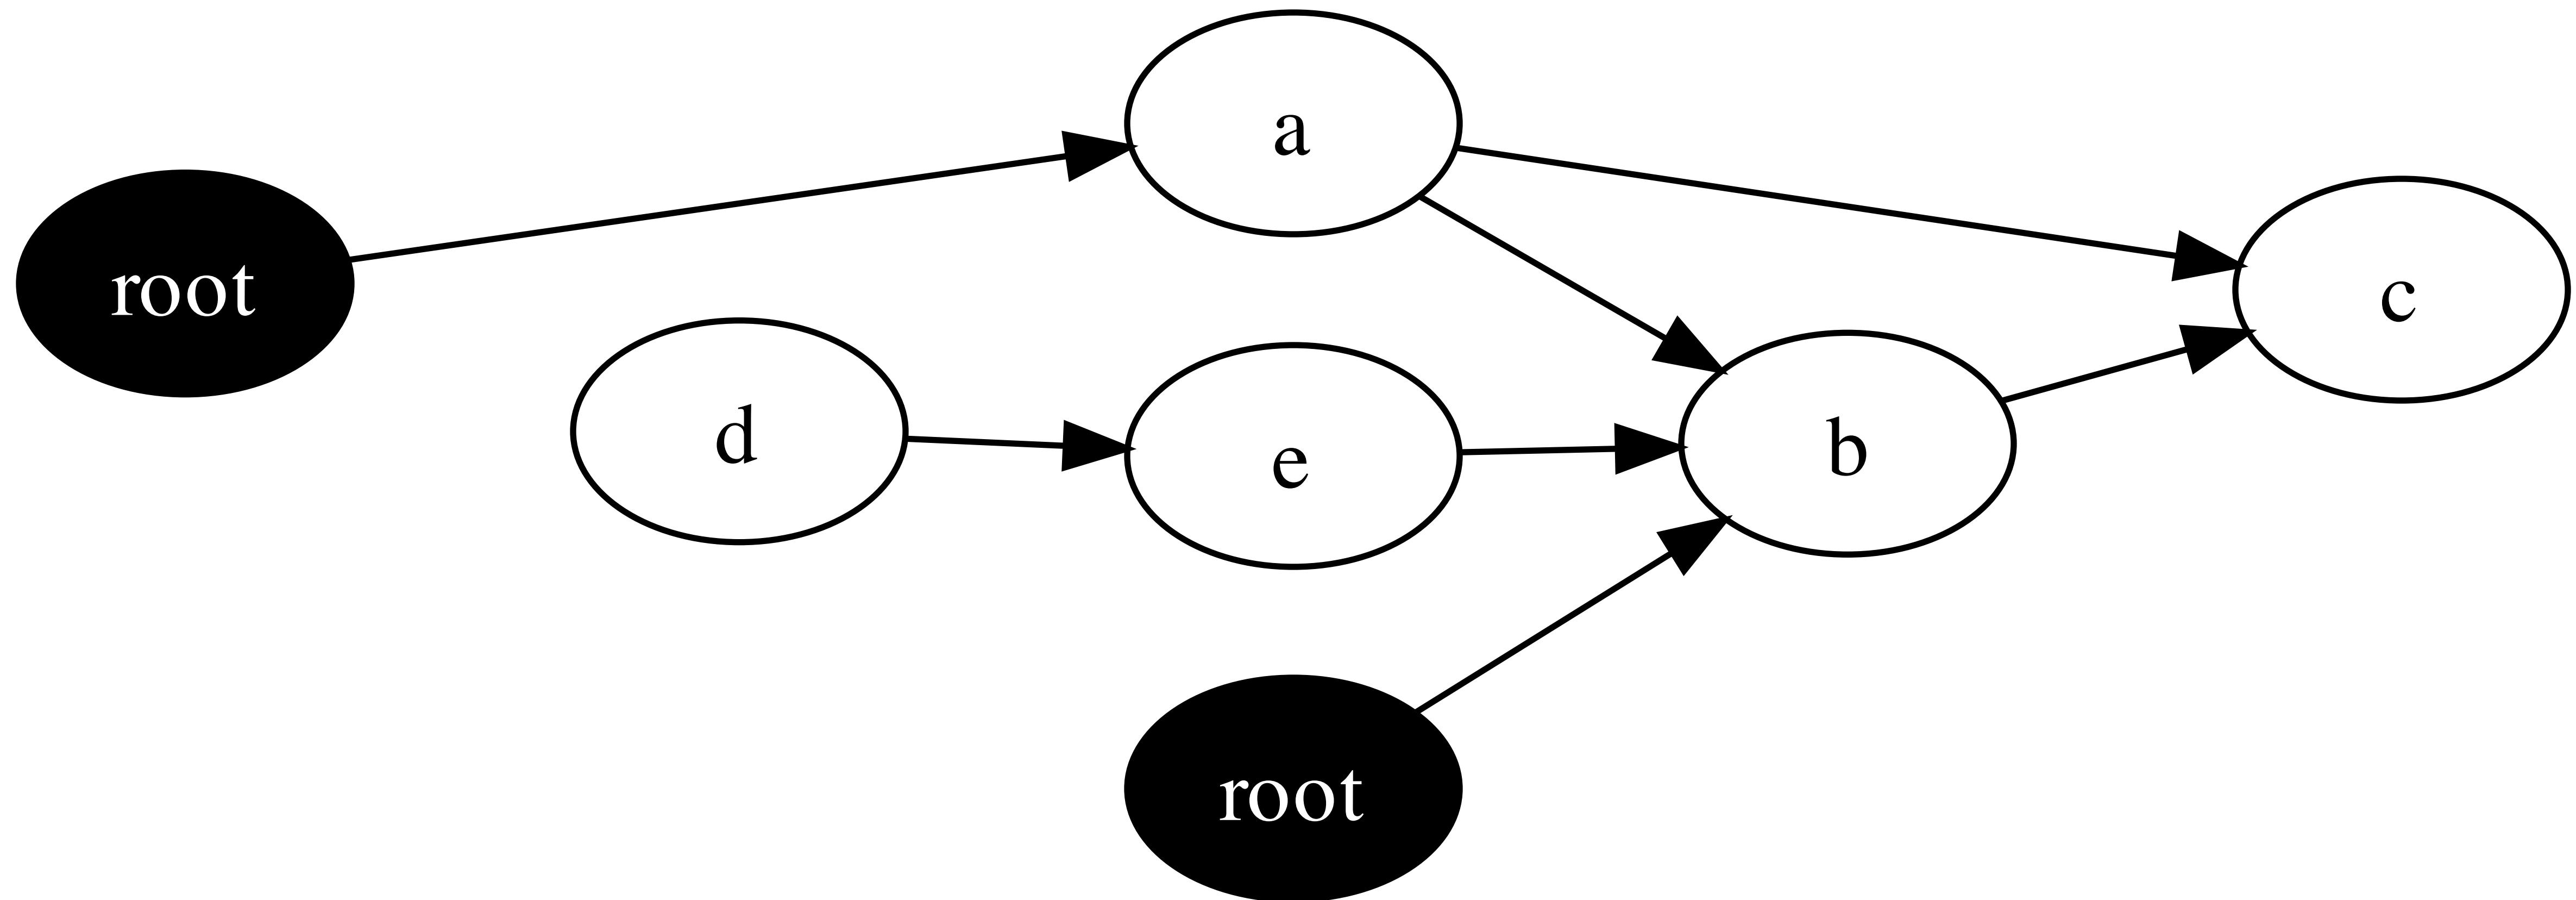
Goals

- Take the mystery out of garbage collection.
- Perform manual cleanup the Right way.
- Become honorary VM sanitation engineers.

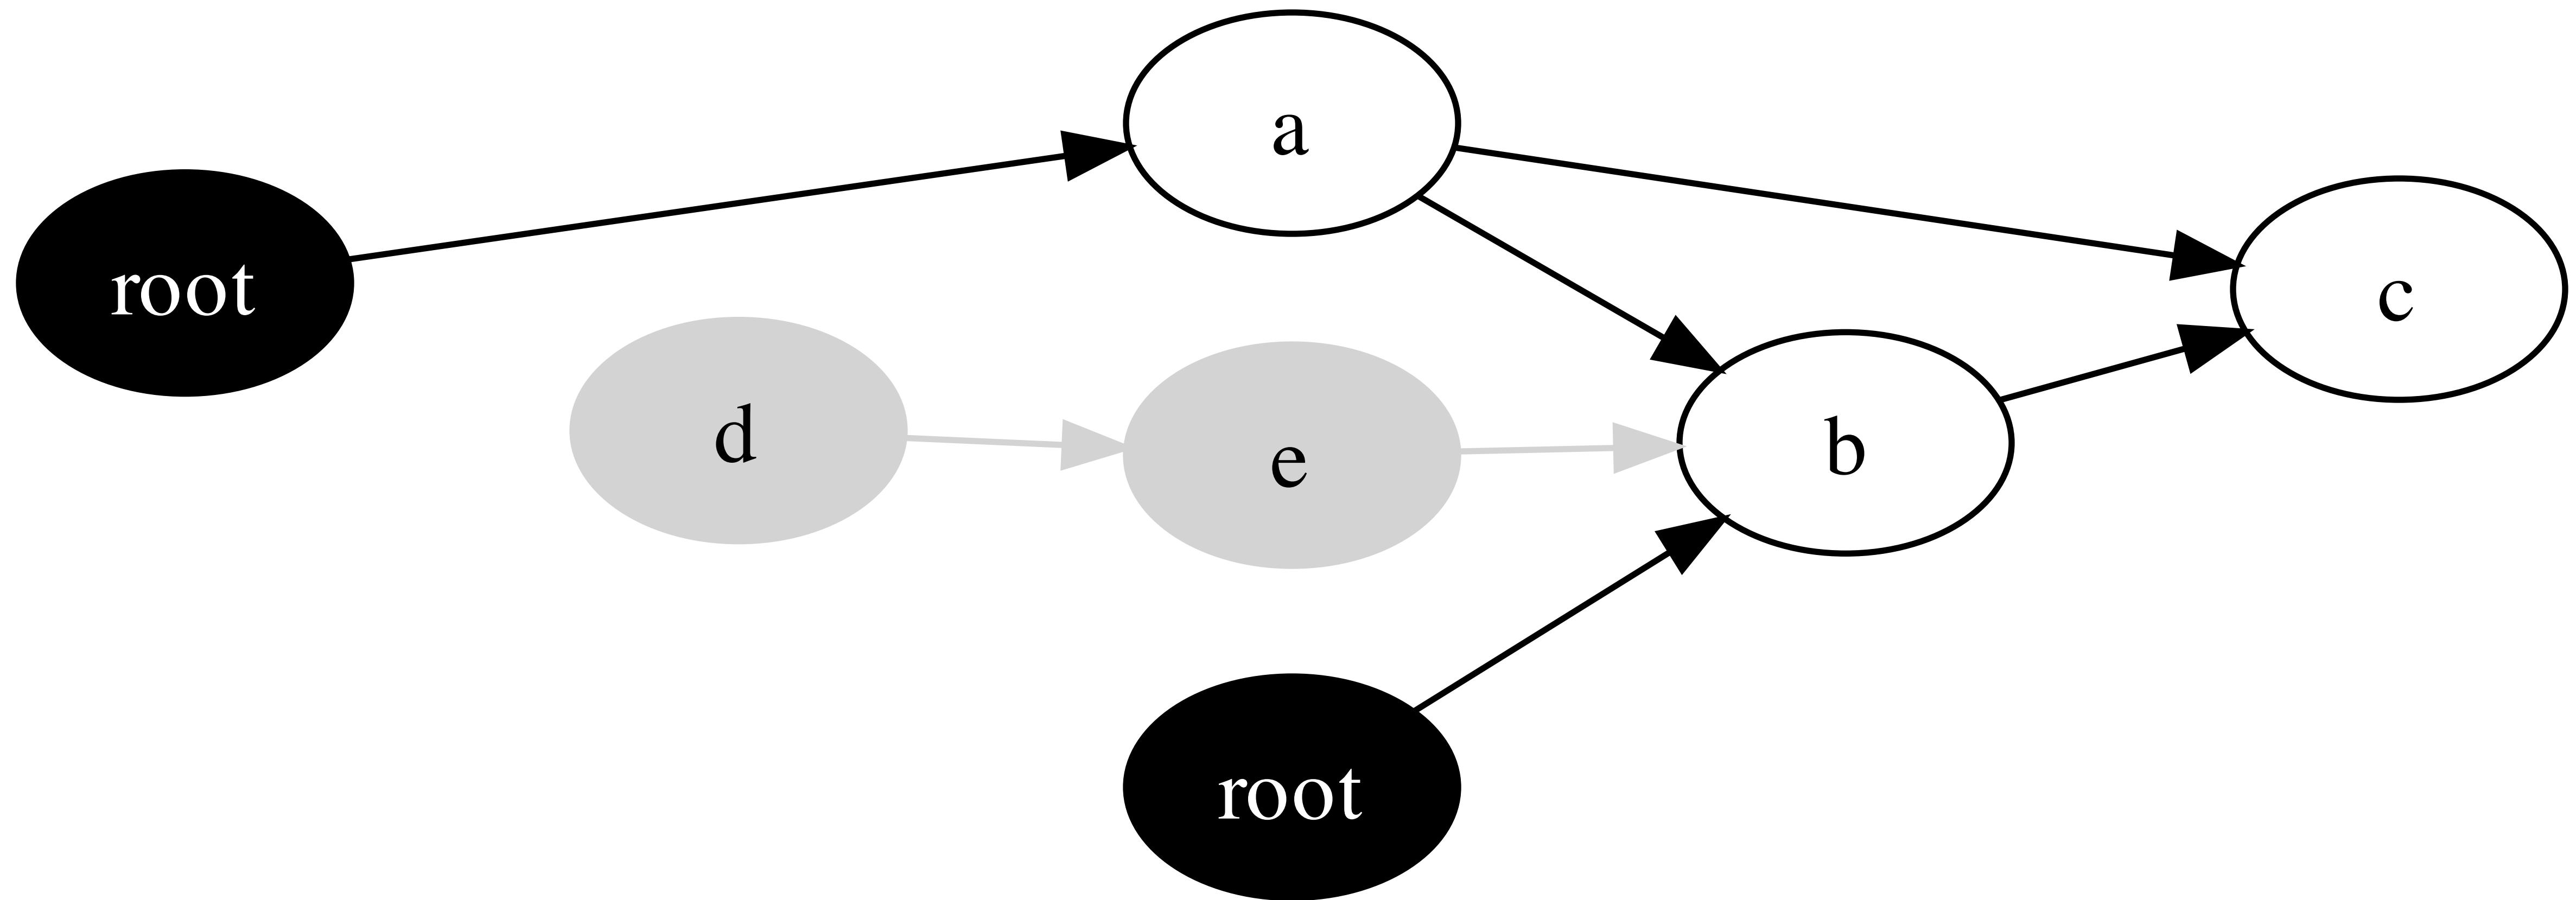
How does garbage collection work?



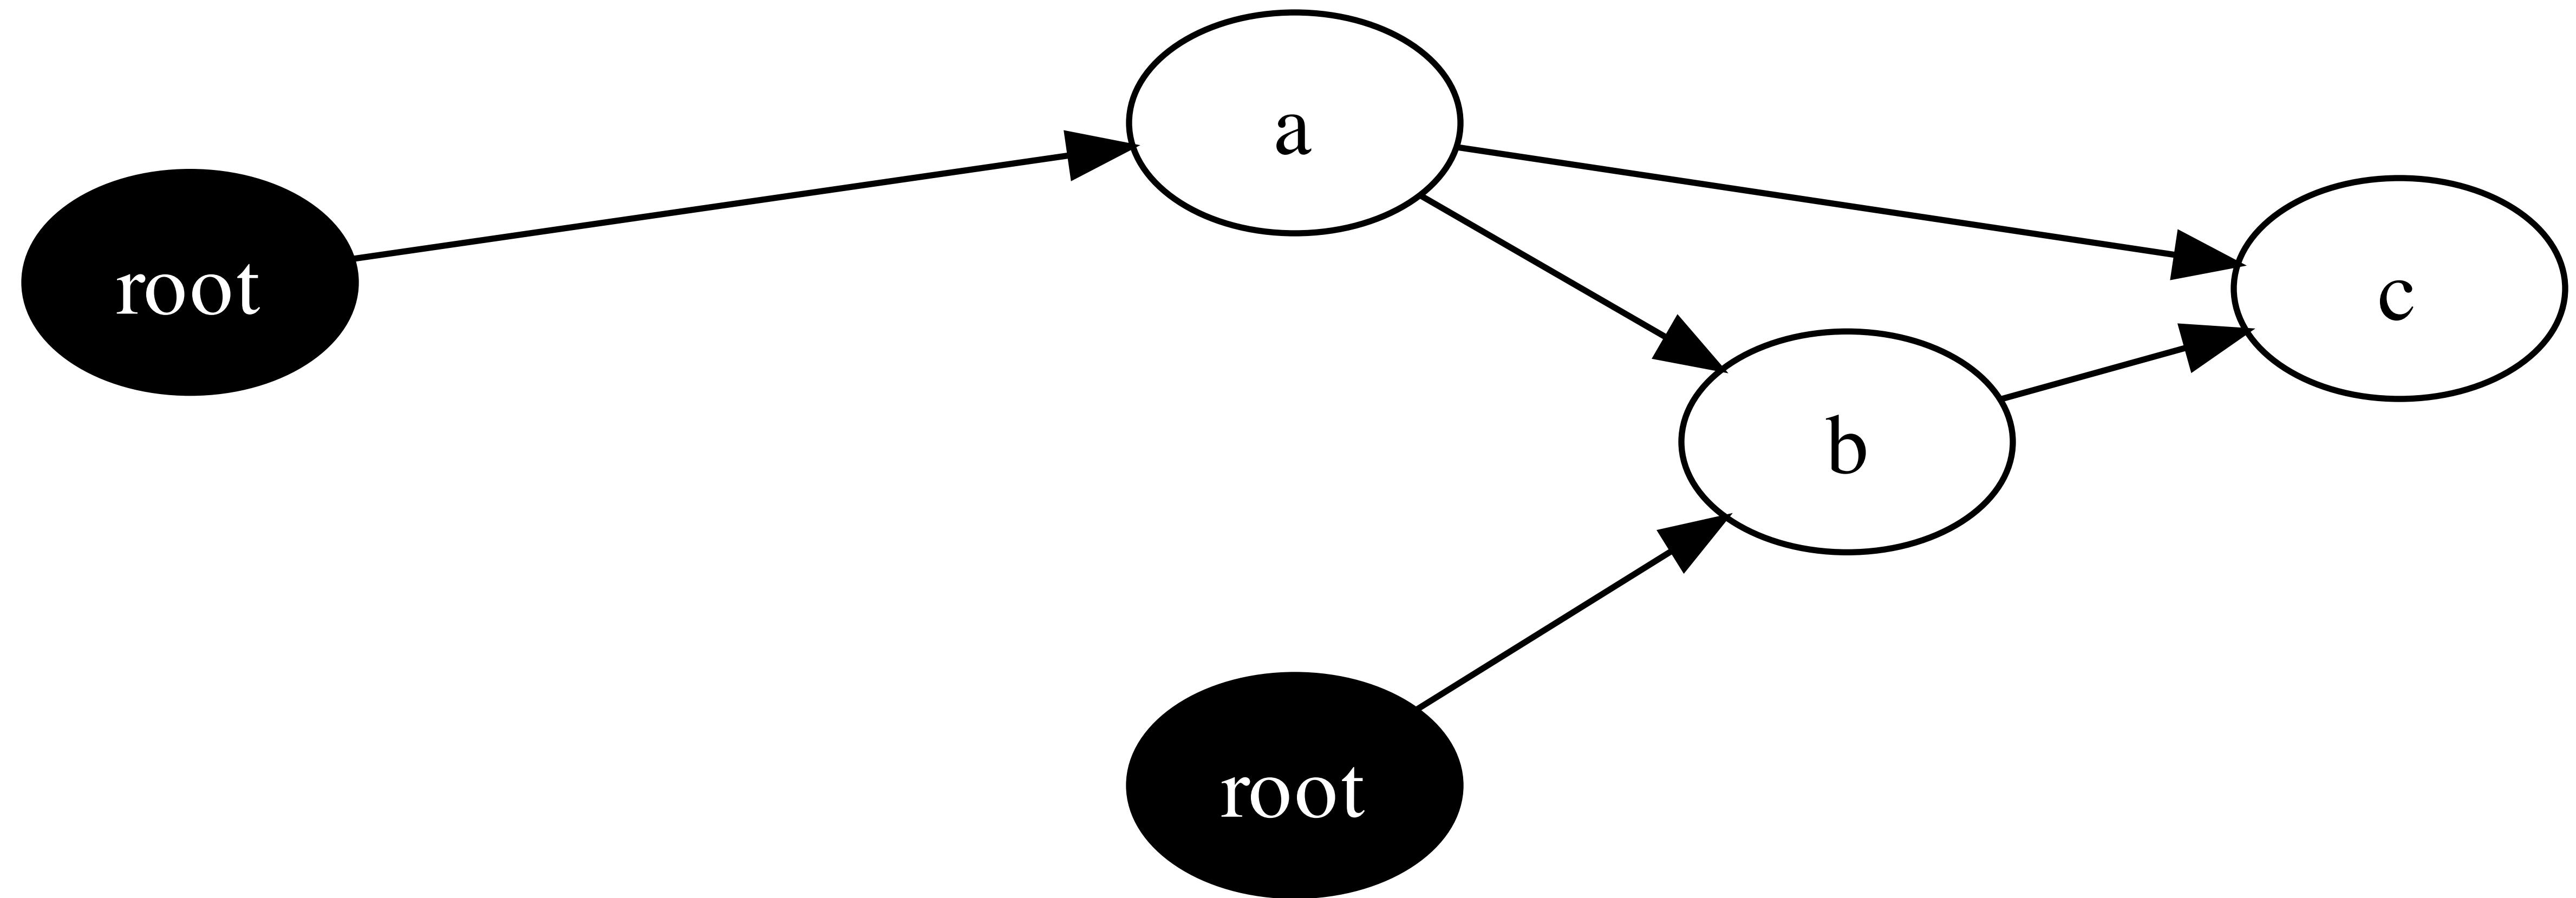
If the reference to D goes away...



We can no longer reach D or E.



So the collector reclaims them.



Reachability

- An object is *reachable* if a live thread can access it.
- Examples of heap roots:
 - System classes (which have static fields)
 - Thread stacks
 - In-flight exceptions
 - JNI global references
 - The finalizer queue
 - The interned String pool
 - etc. (VM-dependent)

Garbage collection isn't a magic bullet.

- Some things require manual cleanup.
 - Listeners
 - File descriptors
 - Native memory
 - External state (`IdentityHashMap`)
- Tools at your disposal:
 - `finally`
 - Overriding `Object.finalize()`
 - References (and reference queues)

Try finally first.

- Pros:
 - More straightforward
 - Handles exceptions in main thread
 - Ensures cleanup keeps pace
- Cons:
 - More work for programmers
 - More error prone
 - Cleanup happens in main thread
- ARM will help.

What is a finalizer?

A callback used by the garbage collector to notify an object when it is about to be reclaimed:

```
public class Foo extends Bar {  
    @Override protected void finalize() throws Throwable {  
        try {  
            ... // Clean up Foo.  
        } finally {  
            super.finalize(); // Clean up Bar.  
        }  
    }  
}
```

Finalizers are seductively simple, but...

- They're not guaranteed to run, especially not timely.
- Undefined threading model; they can run concurrently!
- You must remember to call `super.finalize()`.
- Exceptions are ignored (per spec).
- You can resurrect references.
- They keep objects alive longer than necessary.
- They can make allocation/reclamation 430X slower!
(Bloch, *Effective Java*)
- Worst of all, they messed up the reference API.

Example

```
public class NativeMemory {  
    final int address = allocate();  
    /** Allocates native memory. */  
    static native int allocate();  
  
    /** Writes to native memory. */  
    public void write(byte[] data) {  
        write(address, data);  
    }  
    static native void write(int address, byte[] data);  
  
    /** Frees native memory. */  
    @Override protected void finalize() {  
        free(address);  
    }  
    static native void free(int address);  
}
```

Let's play War!

SegfaultFactory can cause a segfault if its finalizer executes after **NativeMemory**'s:

```
public class SegfaultFactory {  
    private final NativeMemory nm;  
  
    public SegfaultFactory(NativeMemory nm) {  
        this.nm = nm;  
    }  
  
    @Override protected void finalize() {  
        // 50/50 chance of failure  
        nm.write("I'm taking the VM with me!".getBytes());  
    }  
}
```



Always use protection.

```
public class NativeMemory {  
    final int address = allocate();  
    /** Allocates native memory. */  
    static native int allocate();  
  
    /** Writes to native memory. */  
    boolean finalized;  
    public synchronized void write(byte[] data) {  
        if (!finalized) write(address, data);  
        else /* do nothing? */;  
    }  
    static native void write(int address, byte[] data);  
  
    /** Frees native memory. */  
    @Override protected synchronized void finalize() {  
        finalized = true;  
        free(address);  
    }  
    static native void free(int address);  
}
```

Basically, finalizers are good for one thing.

Logging warnings:

```
public class Connection {  
    ...  
    boolean closed;  
    public synchronized void close() {  
        reallyClose();  
        closed = true;  
    }  
    private native void reallyClose();  
  
    @Override protected synchronized void finalize() {  
        if (!closed) {  
            Logger.getLogger(Connection.class.getName())  
                .warning("You forgot to close me!!!!");  
            close();  
        }  
    }  
}
```

Basically, finalizers are good for one thing.

Logging warnings:

```
public class Connection {  
    ...  
    boolean closed;  
    public synchronized void close() {  
        reallyClose();  
        closed = true;  
    }  
    private native void reallyClose();  
  
    @Override protected synchronized void finalize() {  
        if (!closed) {  
            Logger.getLogger(Connection.class.getName())  
                .warning("You forgot to close me!!!!");  
            close();  
        }  
    }  
}
```

Unless you want to disable the warnings.

The alternative: The Reference API

- `@since 1.2`
- Reference types
 - **Soft**: for caching
 - **Weak**: for fast cleanup (pre-finalizer)
 - **Phantom**: for safe cleanup (post-finalizer)
- **Reference queues**: for notifications

package java.lang.ref

```
public abstract class Reference<T> {
    public T get() { ... }
}

public class SoftReference<T> extends Reference<T> {
    public SoftReference(T referent) { ... }
    public SoftReference(T referent, ReferenceQueue<? super T> q) { ... }
}

public class WeakReference<T> extends Reference<T> {
    public WeakReference(T referent) { ... }
    public WeakReference(T referent, ReferenceQueue<? super T> q) { ... }
}

public class PhantomReference<T> extends Reference<T> {
    public PhantomReference(T referent, ReferenceQueue<? super T> q) { ... }
}

public class ReferenceQueue<T> {
    public ReferenceQueue() { ... }
    public Reference<? extends T> poll() { ... }
    public Reference<? extends T> remove() { ... }
}
```

Soft references

- Cleared when the VM runs low on memory
 - *Hopefully* in LRU fashion
- Tuned with **-XX:SoftRefLRUPolicyMSPerMB**
 - How long to retain soft refs in *ms per free MB of heap*
 - **Default:** 1000ms

Use soft references judiciously.

- For quick-and-dirty caching only
- Soft refs have no notion of *weight*.
 - Memory usage
 - Computation time
 - CPU usage
- Soft refs can exacerbate low memory conditions.

Caching a file

```
public class CachedFile {  
    final File file;  
    public CachedFile(File file) {  
        this.file = file;  
    }  
    volatile SoftReference<byte[]> dataReference  
        = new SoftReference<byte[]>(null);  
    /** Gets file contents, reading them if necessary. */  
    public byte[] getData() {  
        byte[] data = dataReference.get();  
        if (data != null) return data;  
        data = readData();  
        dataReference = new SoftReference<byte[]>(data);  
        return data;  
    }  
    /** Reads file contents. */  
    byte[] readData() {  
        ...  
    }  
}
```

Weak references

- Cleared as soon as no strong or soft refs remain.
- Cleared ASAP, before the finalizer runs.
- **Not for caching!** Use soft references, as intended:

“Virtual machine implementations are encouraged to bias against clearing recently-created or recently-used soft references.”

- The `SoftReference` documentation

Can you hear me now?

```
public class Button {  
    public interface Listener {  
        void onClick();  
    }  
    private final List<WeakReference<Listener>> listeners  
        = new ArrayList<WeakReference<Listener>>();  
    public void add(Listener l) {  
        listeners.add(new WeakReference<Listener>(l));  
    }  
    public void click() {  
        Iterator<WeakReference<Listener>> i  
            = listeners.iterator();  
        while (i.hasNext()) {  
            Listener l = i.next().get();  
            if (l == null) i.remove();  
            else l.onClick();  
        }  
    }  
}
```

Phantom references

- Enqueued after no other references remain, *post-finalizer*.
 - Can suffer similar problems to finalizers.
- Must be cleared manually, for no good reason.
- `get()` always returns `null`.
 - So you must use a reference queue.

Let's replace a finalizer!

```
public class NativeMemory {  
    final int address = allocate();  
    /** Allocates native memory. */  
    static native int allocate();  
    NativeMemory() {}  
  
    /** Writes to native memory. */  
    public void write(byte[] data) {  
        write(address, data);  
    }  
    static native void write(int address, byte[] data);  
  
    /** Frees native memory. */  
    @Override protected void finalize() {  
        free(address);  
    }  
    static native void free(int address);  
}
```

The reference

```
class NativeMemoryReference
    extends PhantomReference<NativeMemory> {
    final int address;
    NativeMemoryReference(NativeMemory referent,
        ReferenceQueue<NativeMemory> rq) {
        super(referent, rq);
        address = referent.address;
    }
}
```

The manager

```
public class NativeMemoryManager {  
    private static final Set<Reference<?>> refs  
        = Collections.synchronizedSet(new HashSet<Reference<?>>());  
    private static final ReferenceQueue<NativeMemory> rq  
        = new ReferenceQueue<NativeMemory>();  
    public static NativeMemory allocate() {  
        NativeMemory nm = new NativeMemory();  
        refs.add(new NativeMemoryReference(nm, rq));  
        cleanUp();  
        return nm;  
    }  
    private static void cleanUp() {  
        NativeMemoryReference ref;  
        while ((ref = (NativeMemoryReference) rq.poll()) != null) {  
            NativeMemory.free(ref.address);  
            refs.remove(ref);  
        }  
    }  
}
```

The manager *with* the Guava Libraries

```
public class NativeMemoryManager {  
    private static final Set<Reference<?>> refs  
        = Collections.synchronizedSet(new HashSet<Reference<?>>());  
    private static final FinalizableReferenceQueue frq  
        = new FinalizableReferenceQueue();  
    public static NativeMemory allocate() {  
        NativeMemory nm = new NativeMemory();  
        final int address = nm.address;  
        refs.add(new FinalizablePhantomReference<NativeMemory>(nm, frq) {  
            public void finalizeReferent() {  
                NativeMemory.free(address);  
                refs.remove(this);  
            }  
        });  
        return nm;  
    }  
}
```

Tip: accessing a phantom referent

```
public class WeakPhantomReference<T> extends PhantomReference<T> {  
    final WeakReference<T> weakReference;  
  
    public WeakPhantomReference(T referent,  
        ReferenceQueue<? super T> q) {  
        super(referent, q);  
        weakReference = new WeakReference<T>(referent);  
    }  
  
    /** Returns referent so long as it's weakly-reachable. */  
    @Override public T get() {  
        return weakReference.get();  
    }  
}
```

Don't forget...

The GC runs concurrently with your code:

```
public class RaceTheCollector {  
    public <T> T dereference(WeakReference<T> referent) {  
        T t = referent.get();  
        if (t == null) {  
            throw new NullPointerException("Reference is cleared.");  
        }  
        ... // The garbage collector runs.  
        return referent.get(); // Can return null!!!  
    }  
}
```

`java.util.WeakHashMap`

- Useful for emulating additional fields
- Keeps weak refs to keys, strong refs to values
- Not concurrent
- Uses `equals()` when it should use `==`

Guava MapMaker

- Near drop-in replacement for `WeakHashMap`
- Strong, soft, or weak key and/or value references
- Concurrent, cleans up in background thread
- Uses `==` to compare weak and soft referents
- Supports on-demand computation of values
- **Supports size limiting**

Guava MapMaker

```
public class GetterMethods {  
    final static Map<Class<?>, List<Method>> cache = new MapMaker()  
        .weakKeys()  
        .softValues()  
        .makeComputingMap(new Function<Class<?>, List<Method>>() {  
            public List<Method> apply(Class<?> clazz) {  
                List<Method> getters = new ArrayList<Method>();  
                for (Method m : clazz.getMethods())  
                    if (m.getName().startsWith("get"))  
                        getters.add(m);  
                return getters;  
            }  
        });  
    public static List<Method> on(Class<?> clazz) {  
        return cache.get(clazz);  
    }  
}
```

Usage: `List<Method> l = GetterMethods.on(Foo.class);`

Recap: The Levels of Reachability

- Strong
- Soft
- Weak
- Finalizer
- Phantom, JNI weak
- Unreachable

Recap: The Levels of Reachability

- **Strong**
- Soft
- Weak
- Finalizer
- Phantom, JNI weak
- Unreachable

Recap: The Levels of Reachability

- Strong
- **Soft**
- Weak
- Finalizer
- Phantom, JNI weak
- Unreachable

Recap: The Levels of Reachability

- Strong
- Soft
- **Weak**
- Finalizer
- Phantom, JNI weak
- Unreachable

Recap: The Levels of Reachability

- Strong
- Soft
- Weak
- **Finalizer**
- Phantom, JNI weak
- Unreachable

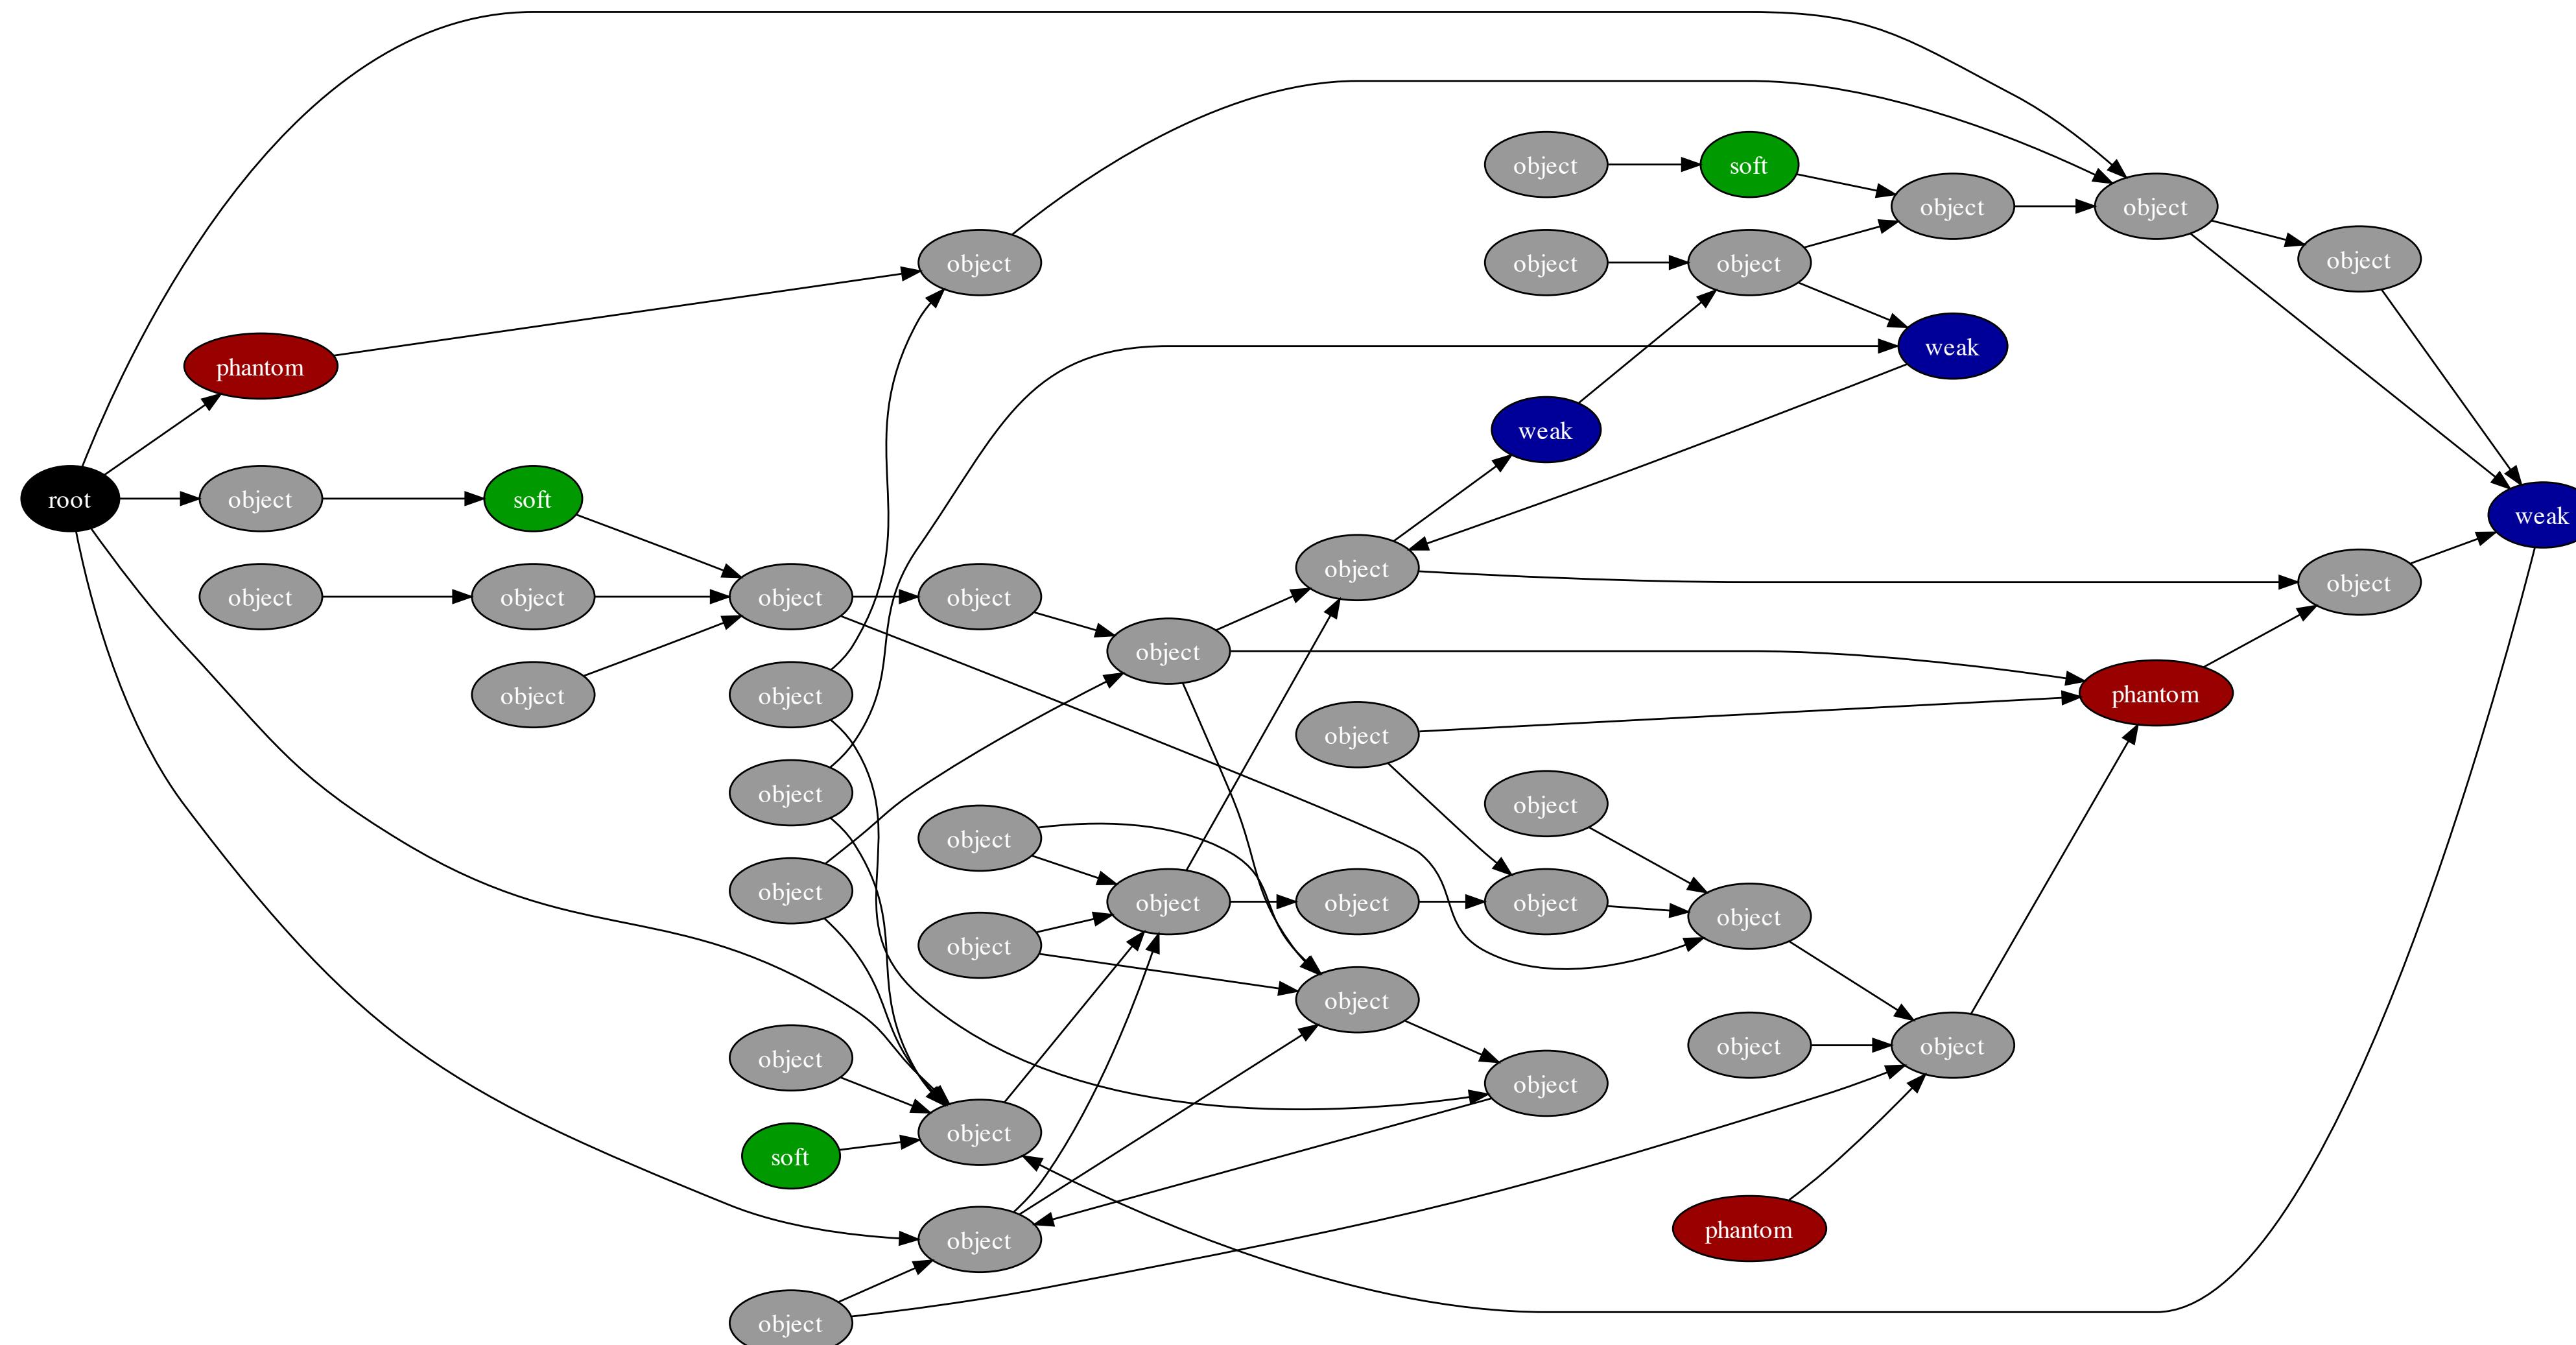
Recap: The Levels of Reachability

- Strong
- Soft
- Weak
- Finalizer
- **Phantom, JNI weak**
- Unreachable

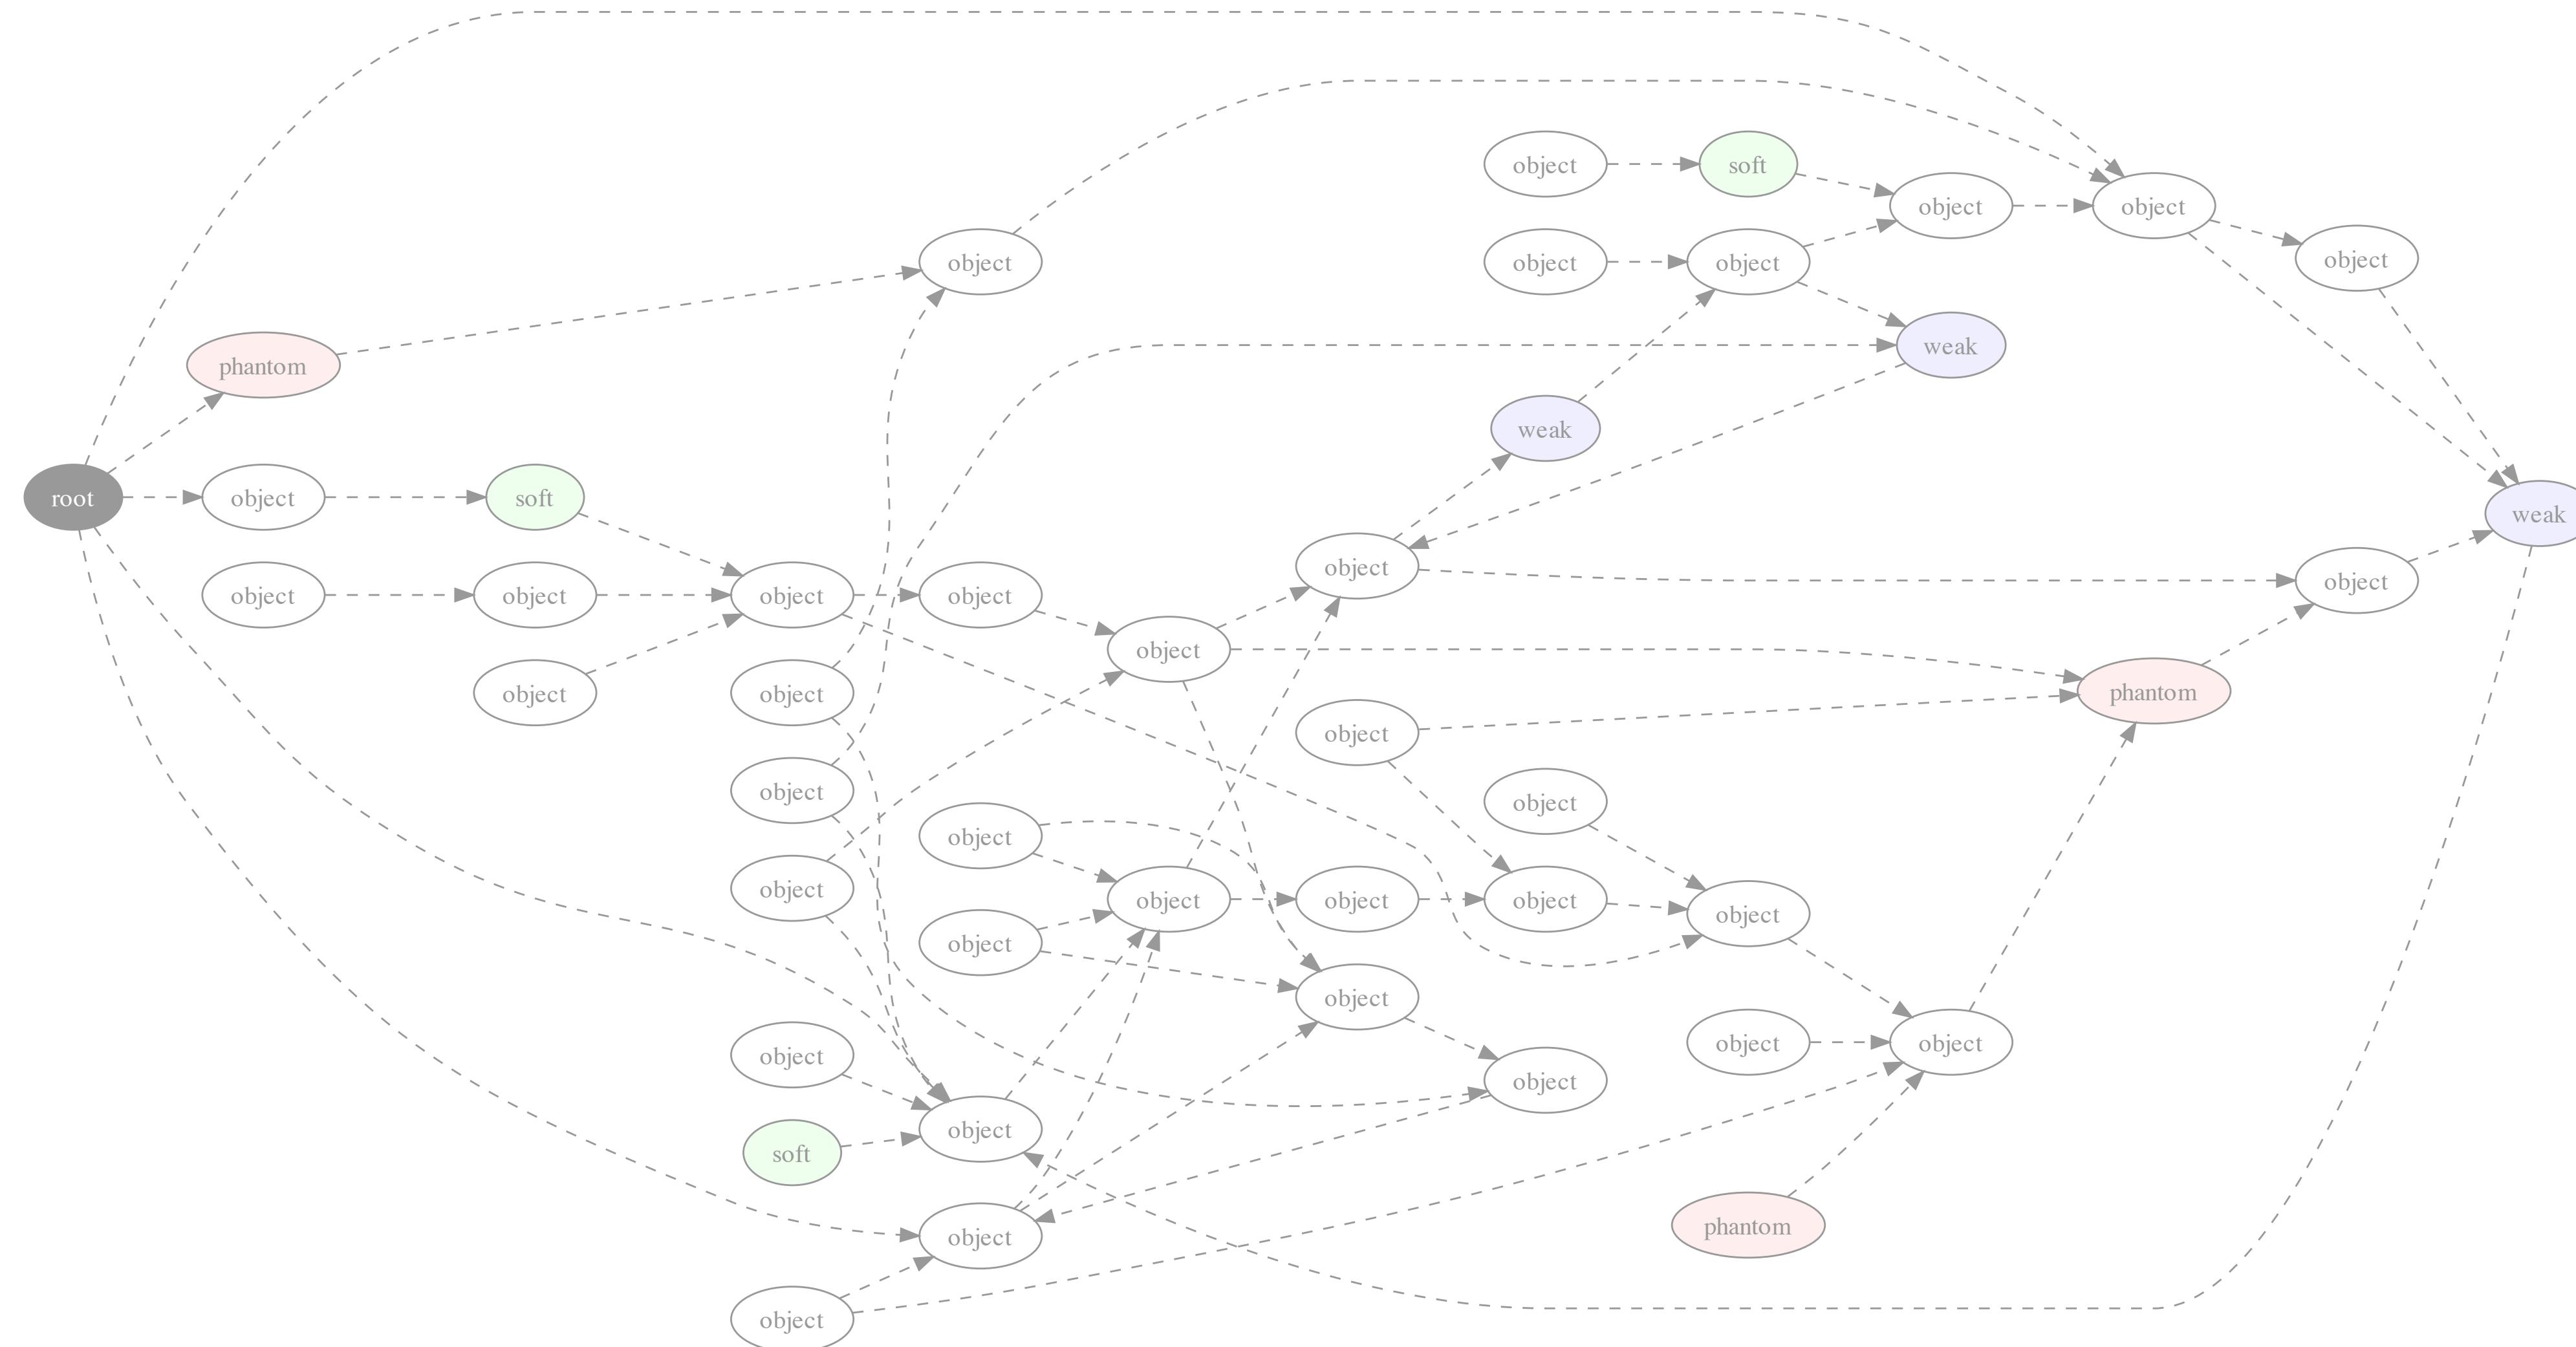
Recap: The Levels of Reachability

- Strong
- Soft
- Weak
- Finalizer
- Phantom, JNI weak
- **Unreachable**

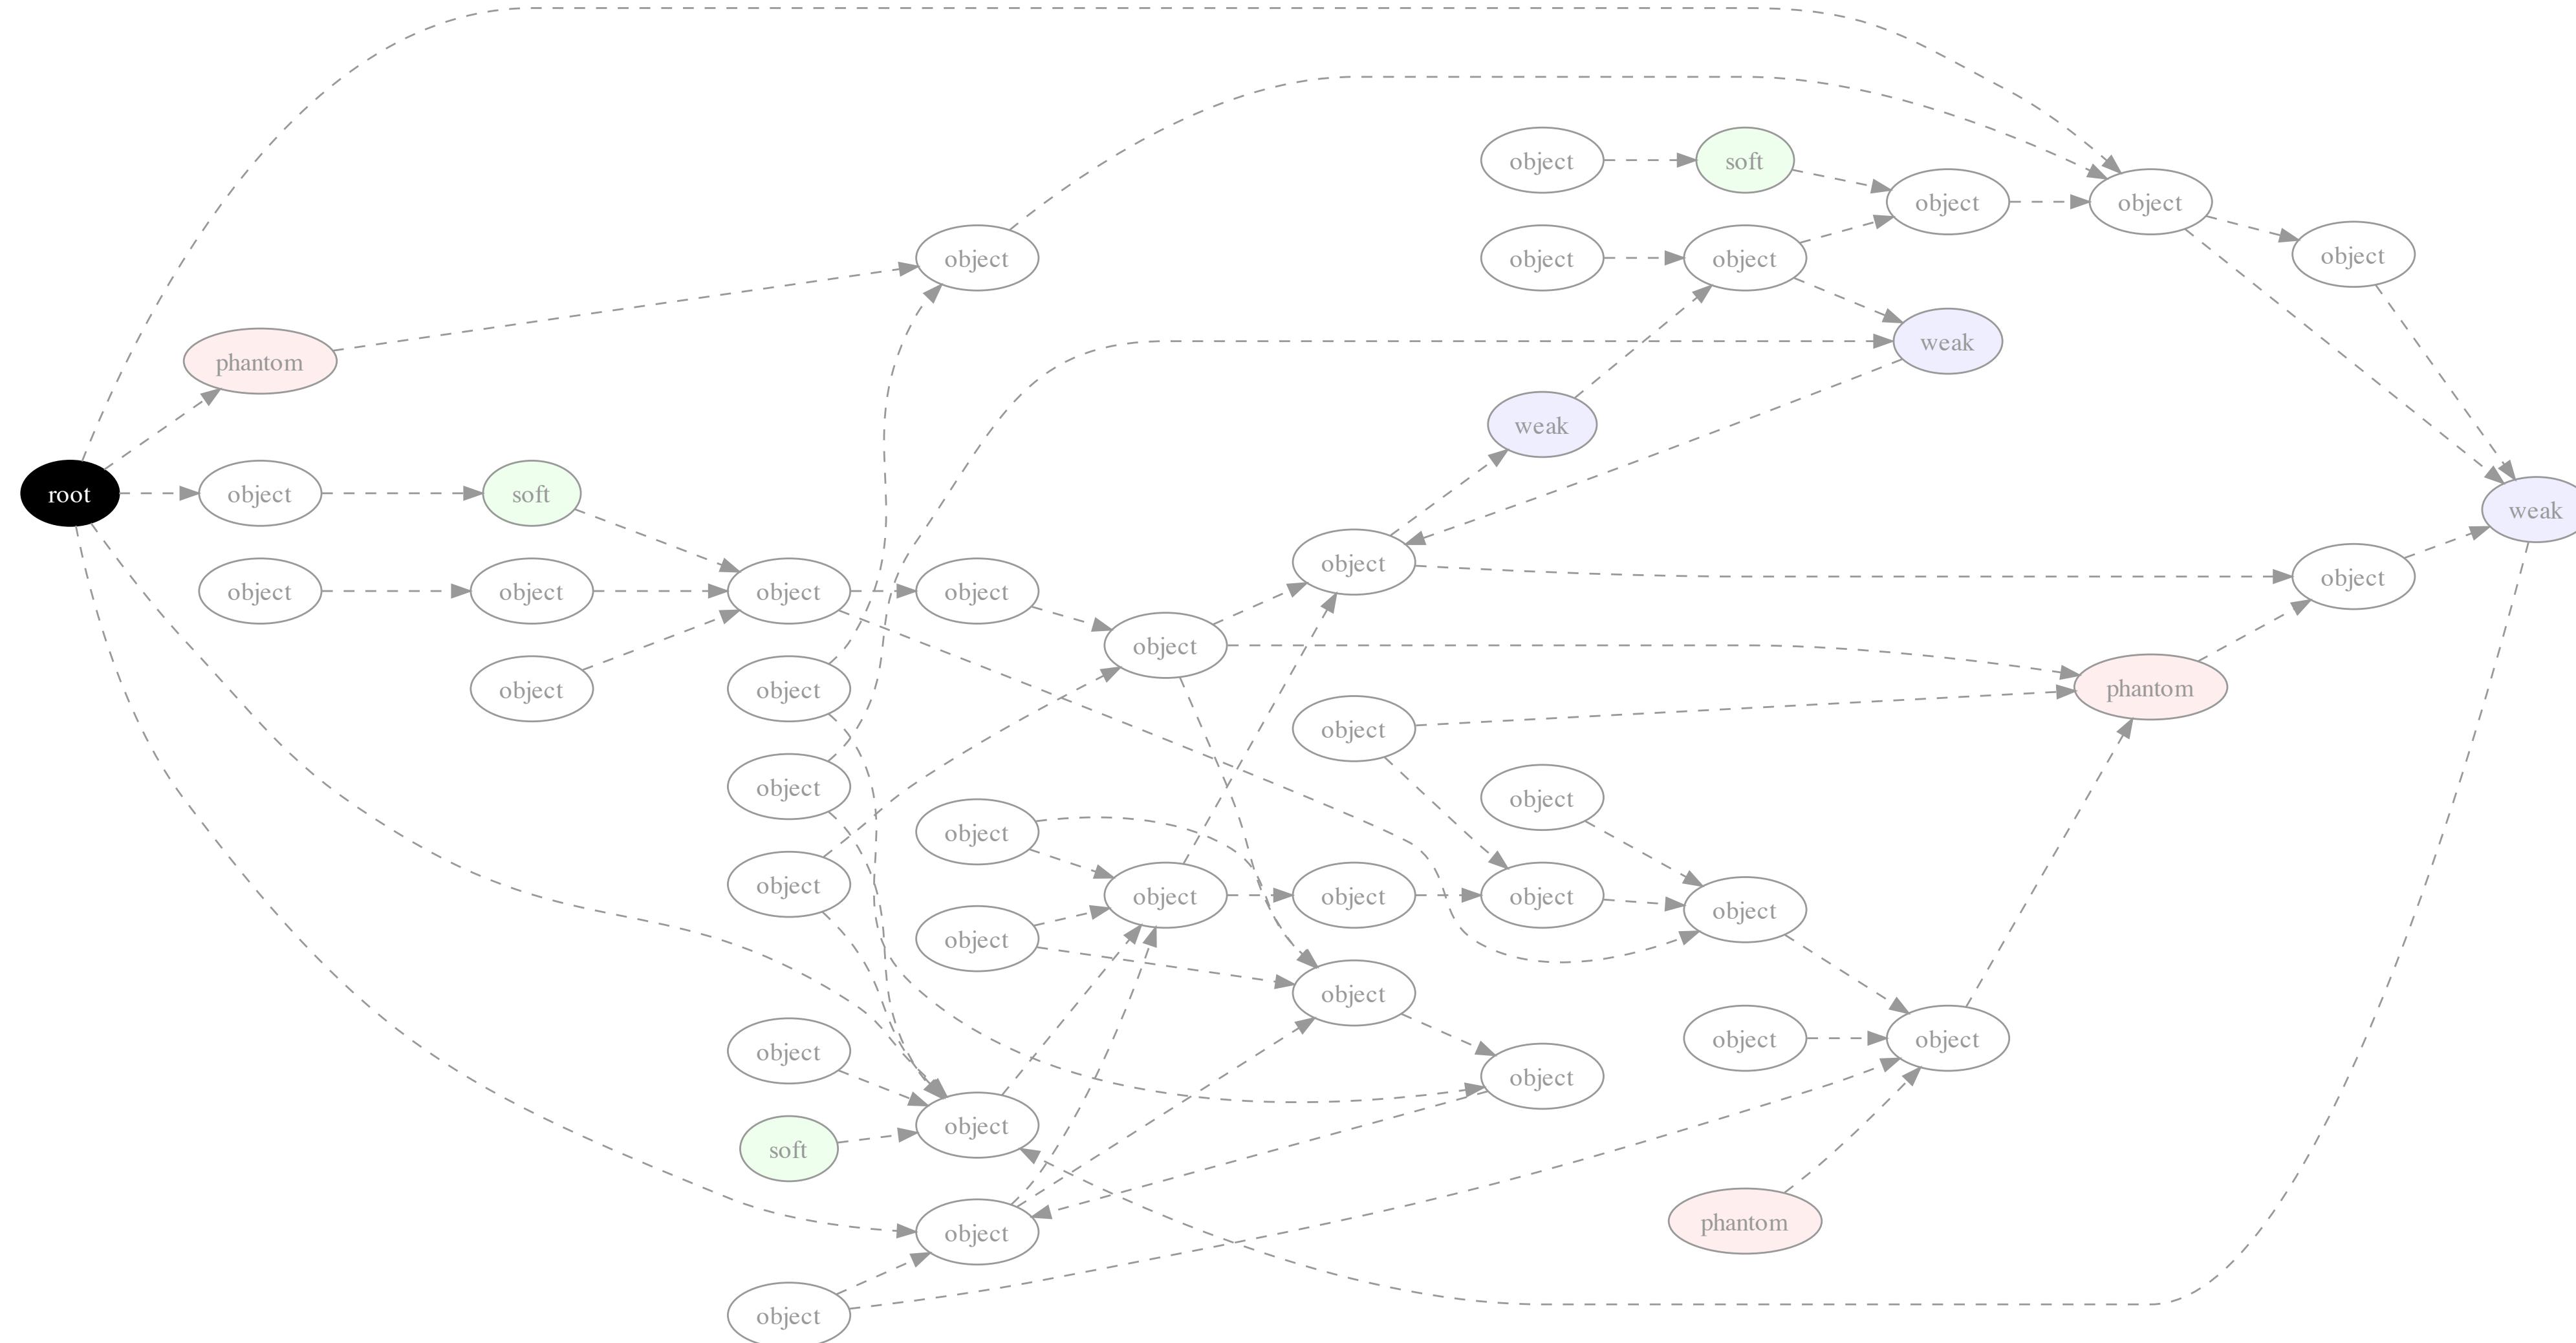
Let's mark and sweep a heap!



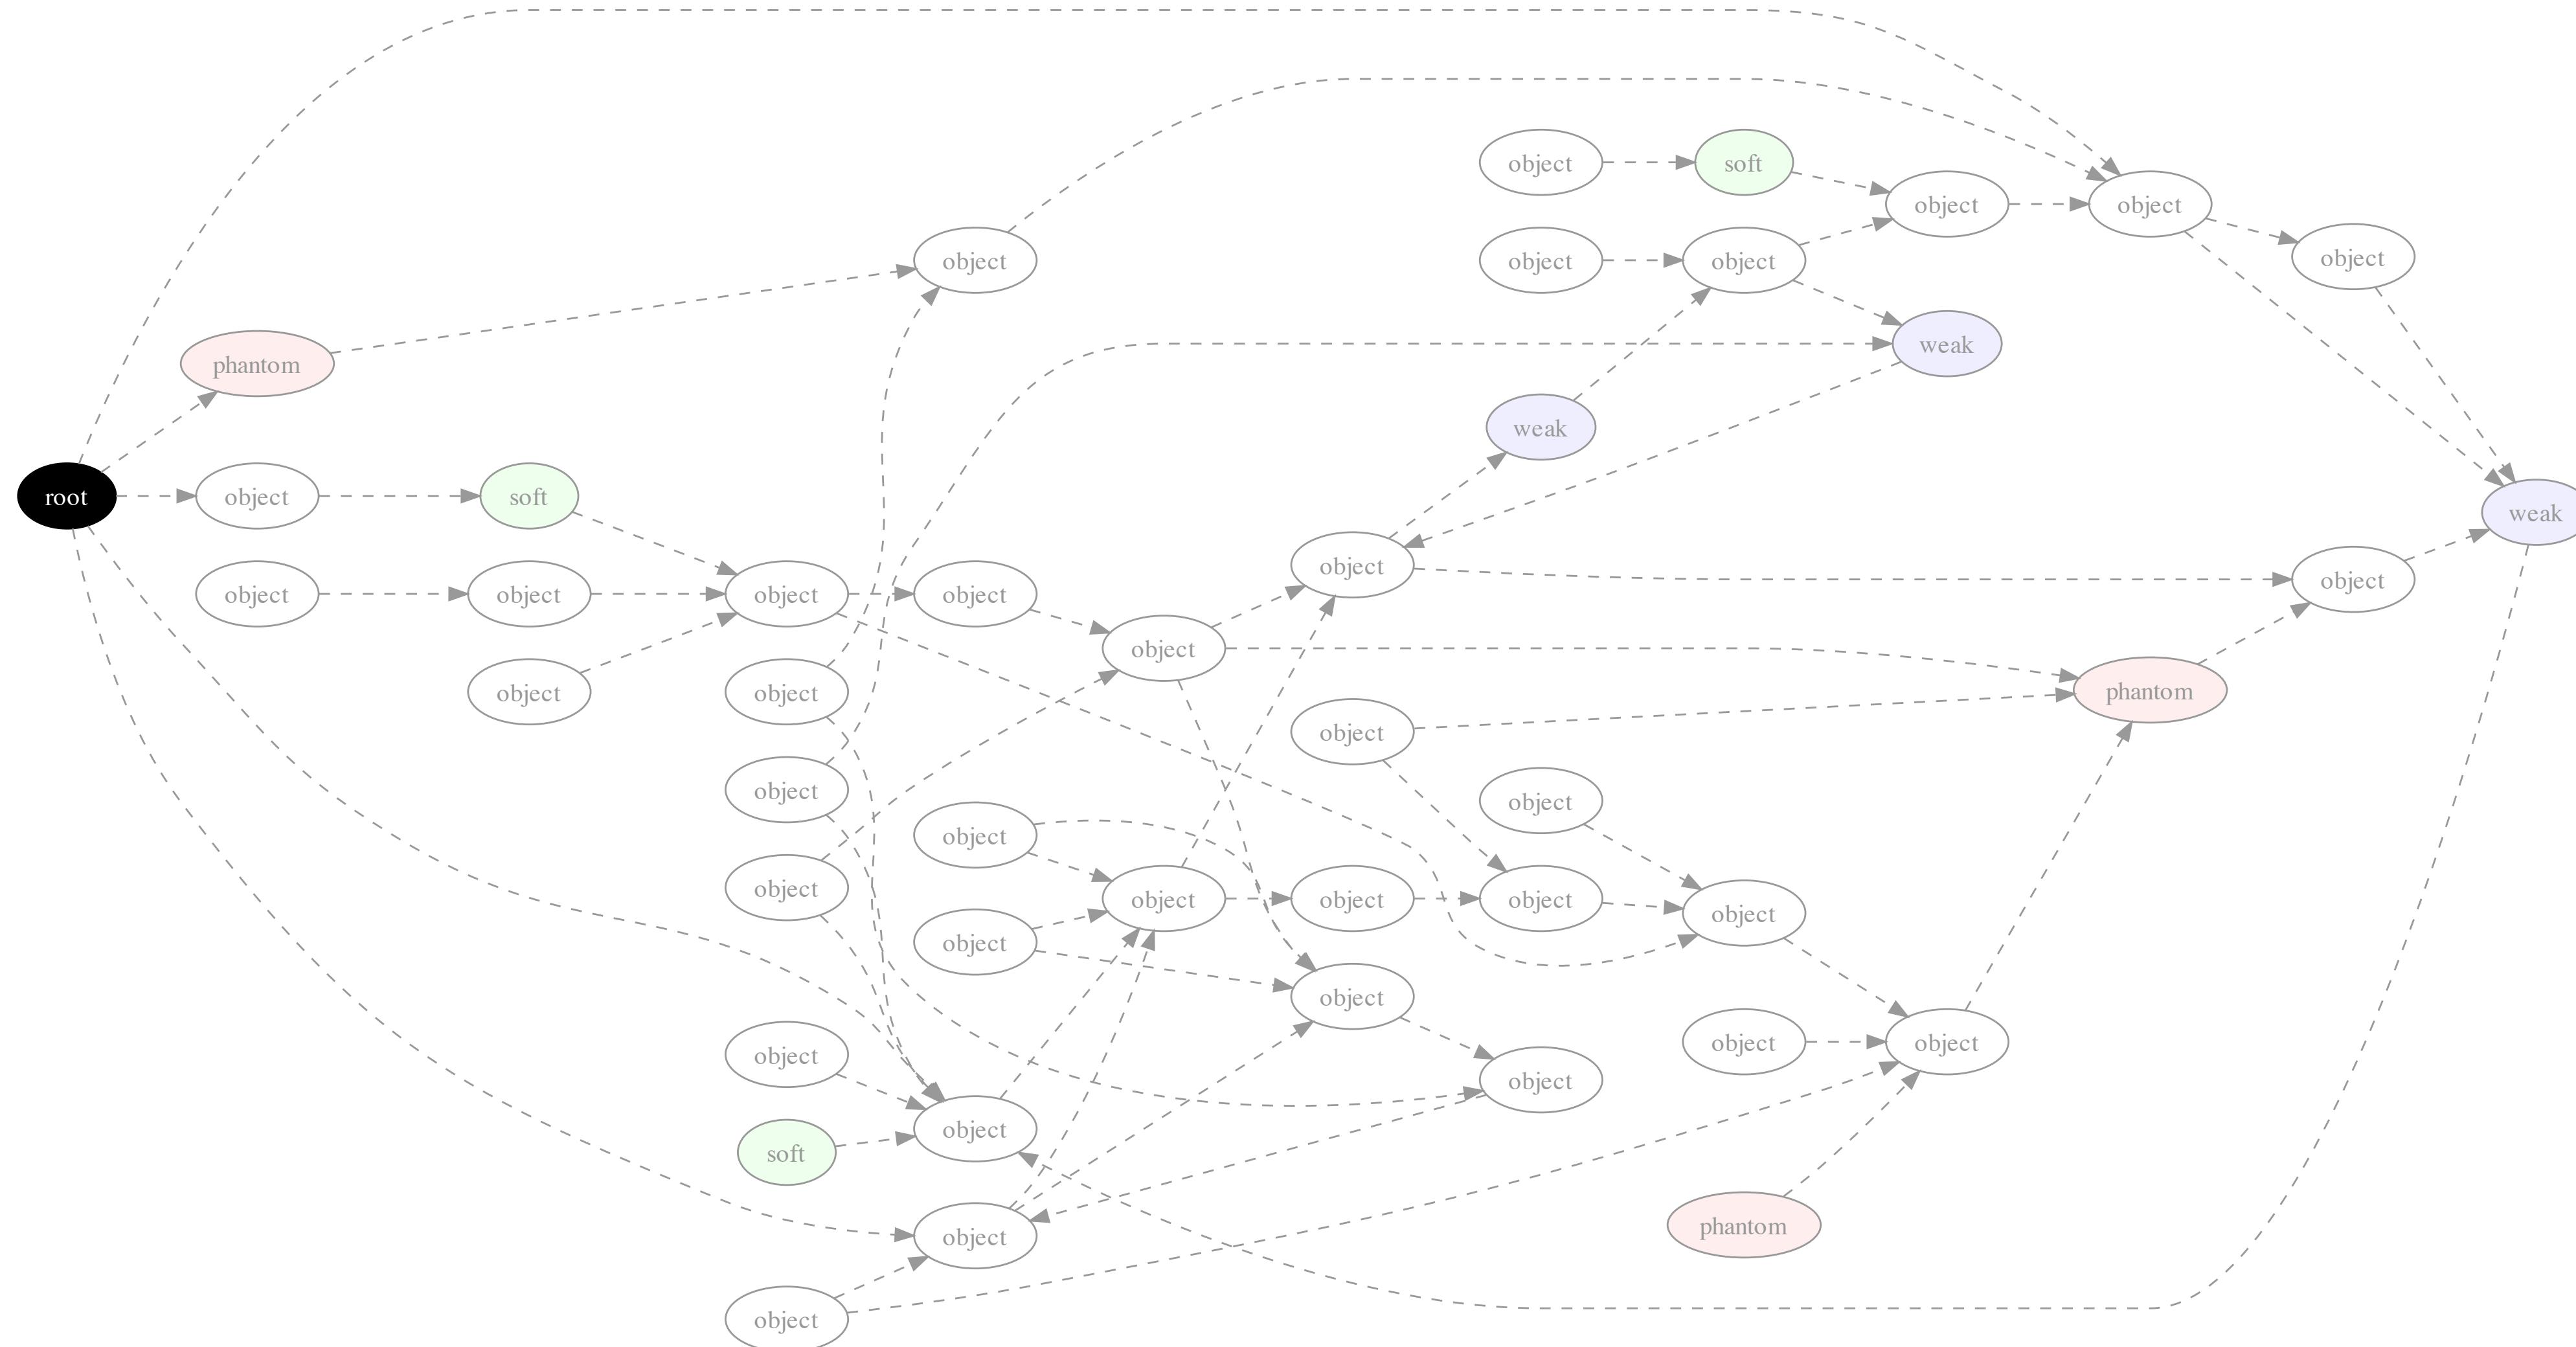
No objects are marked at first.



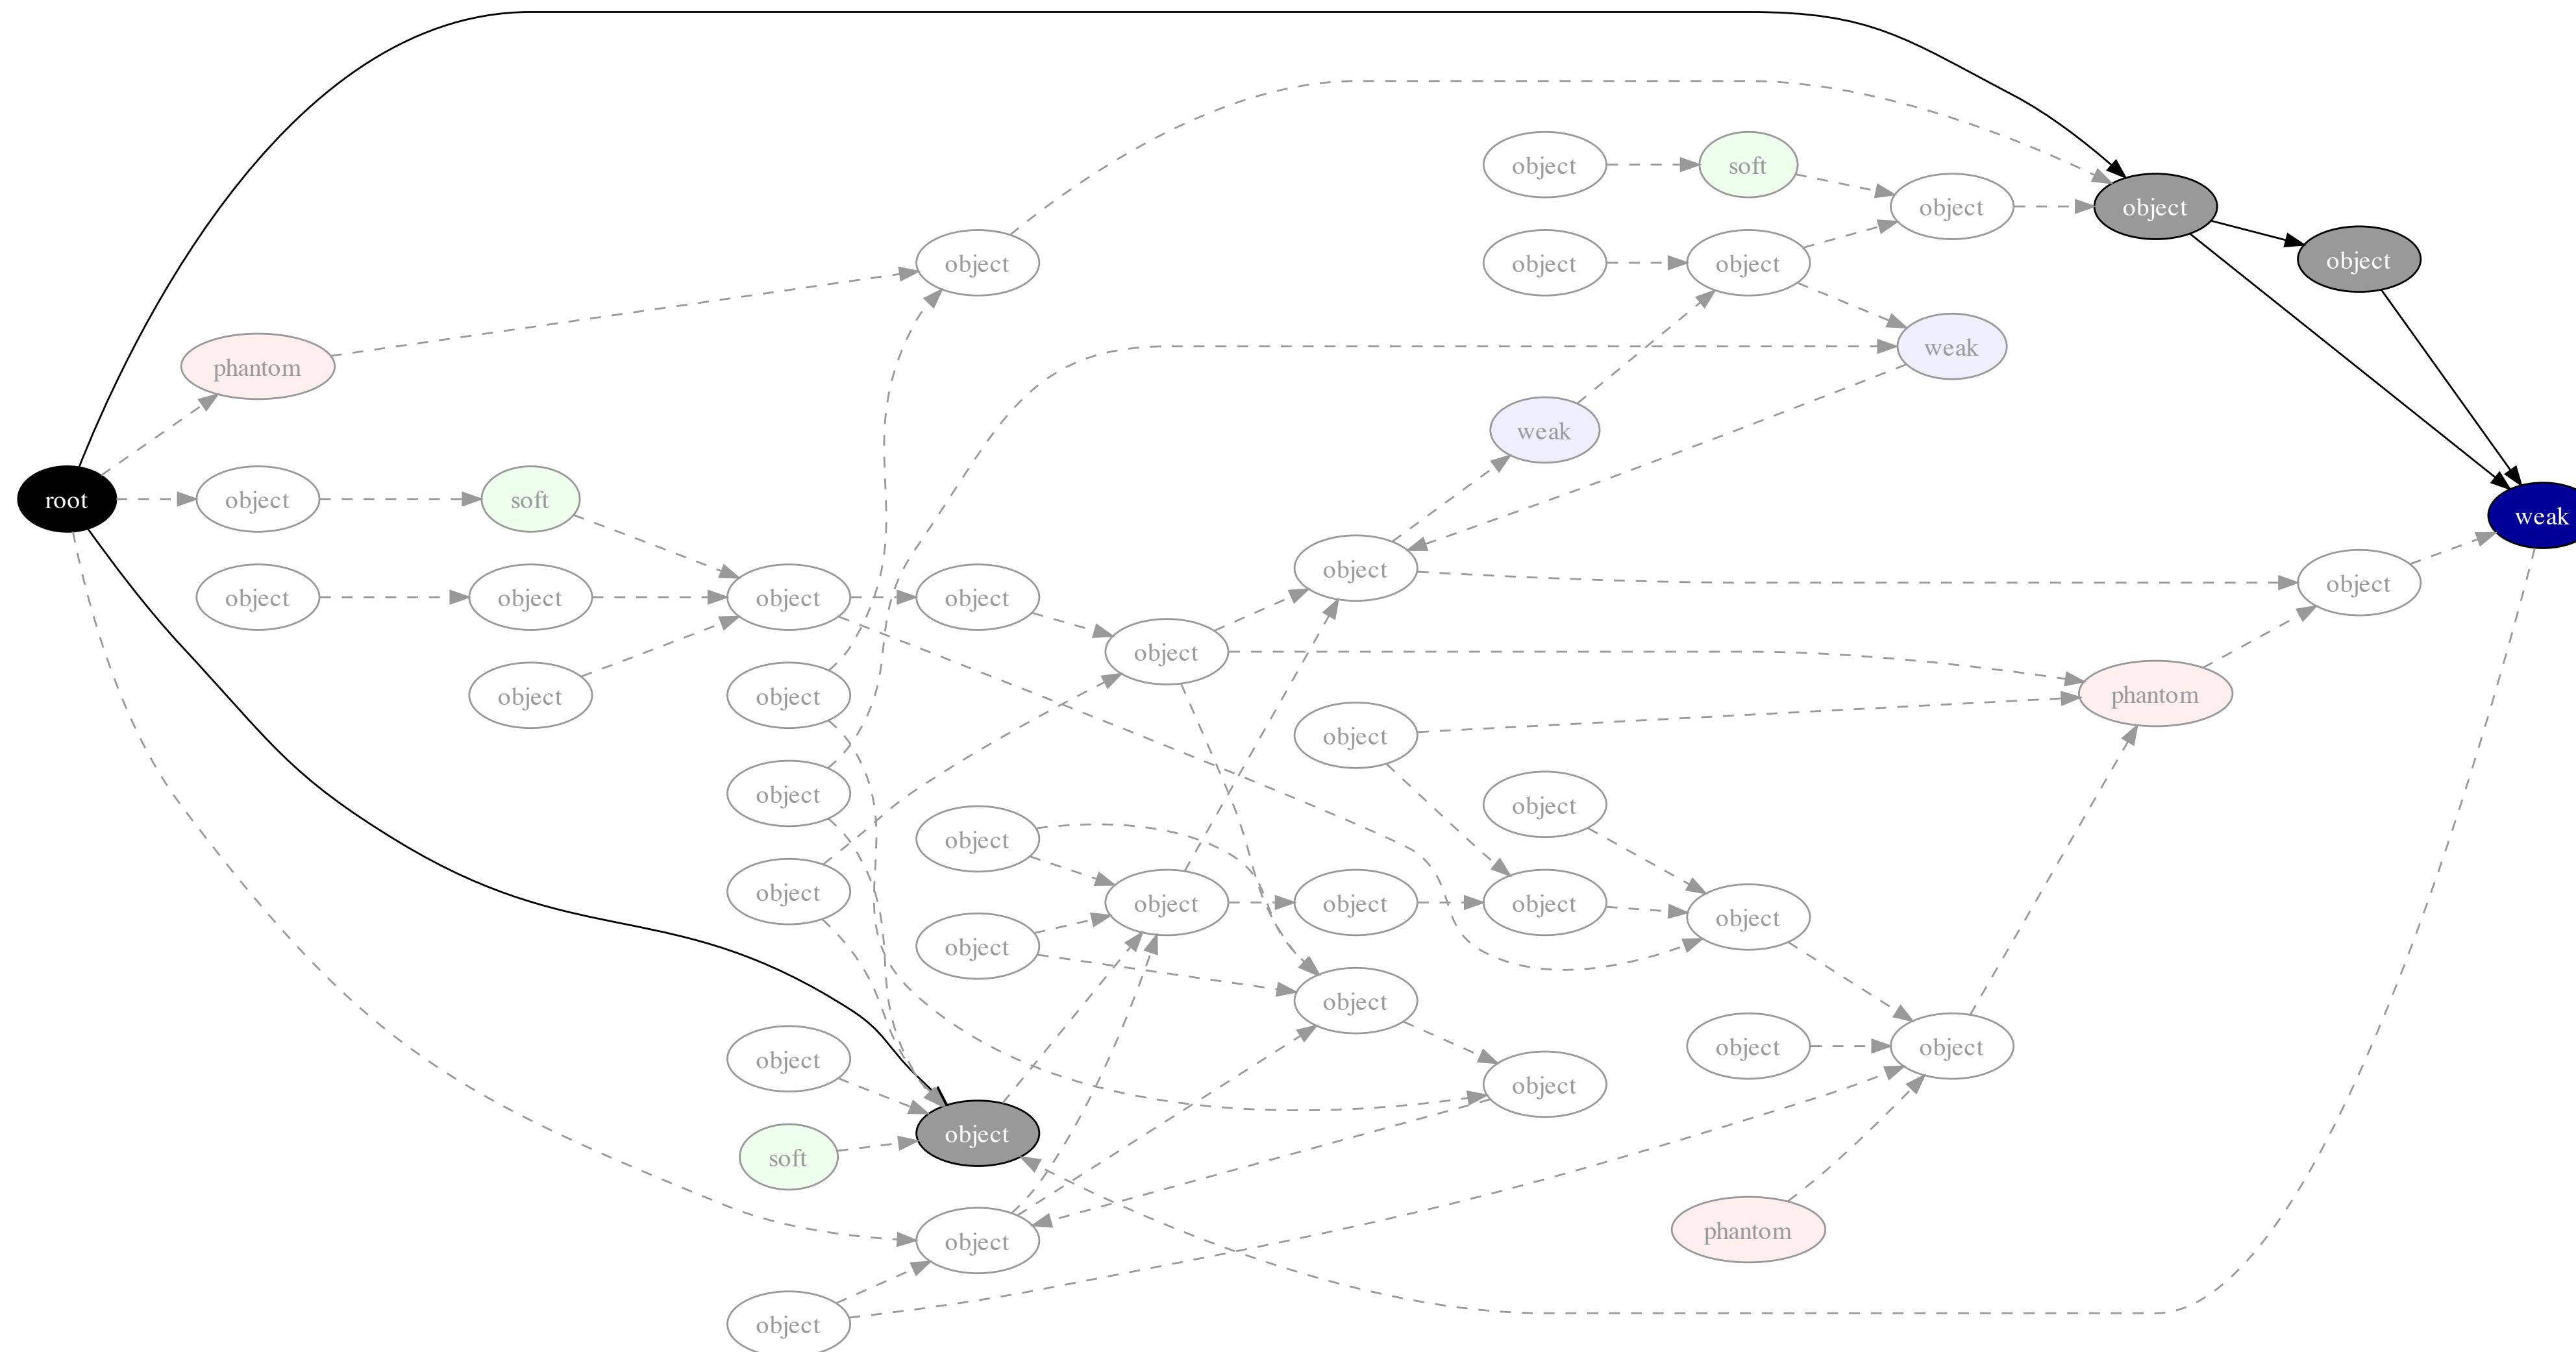
1. Start at a root.



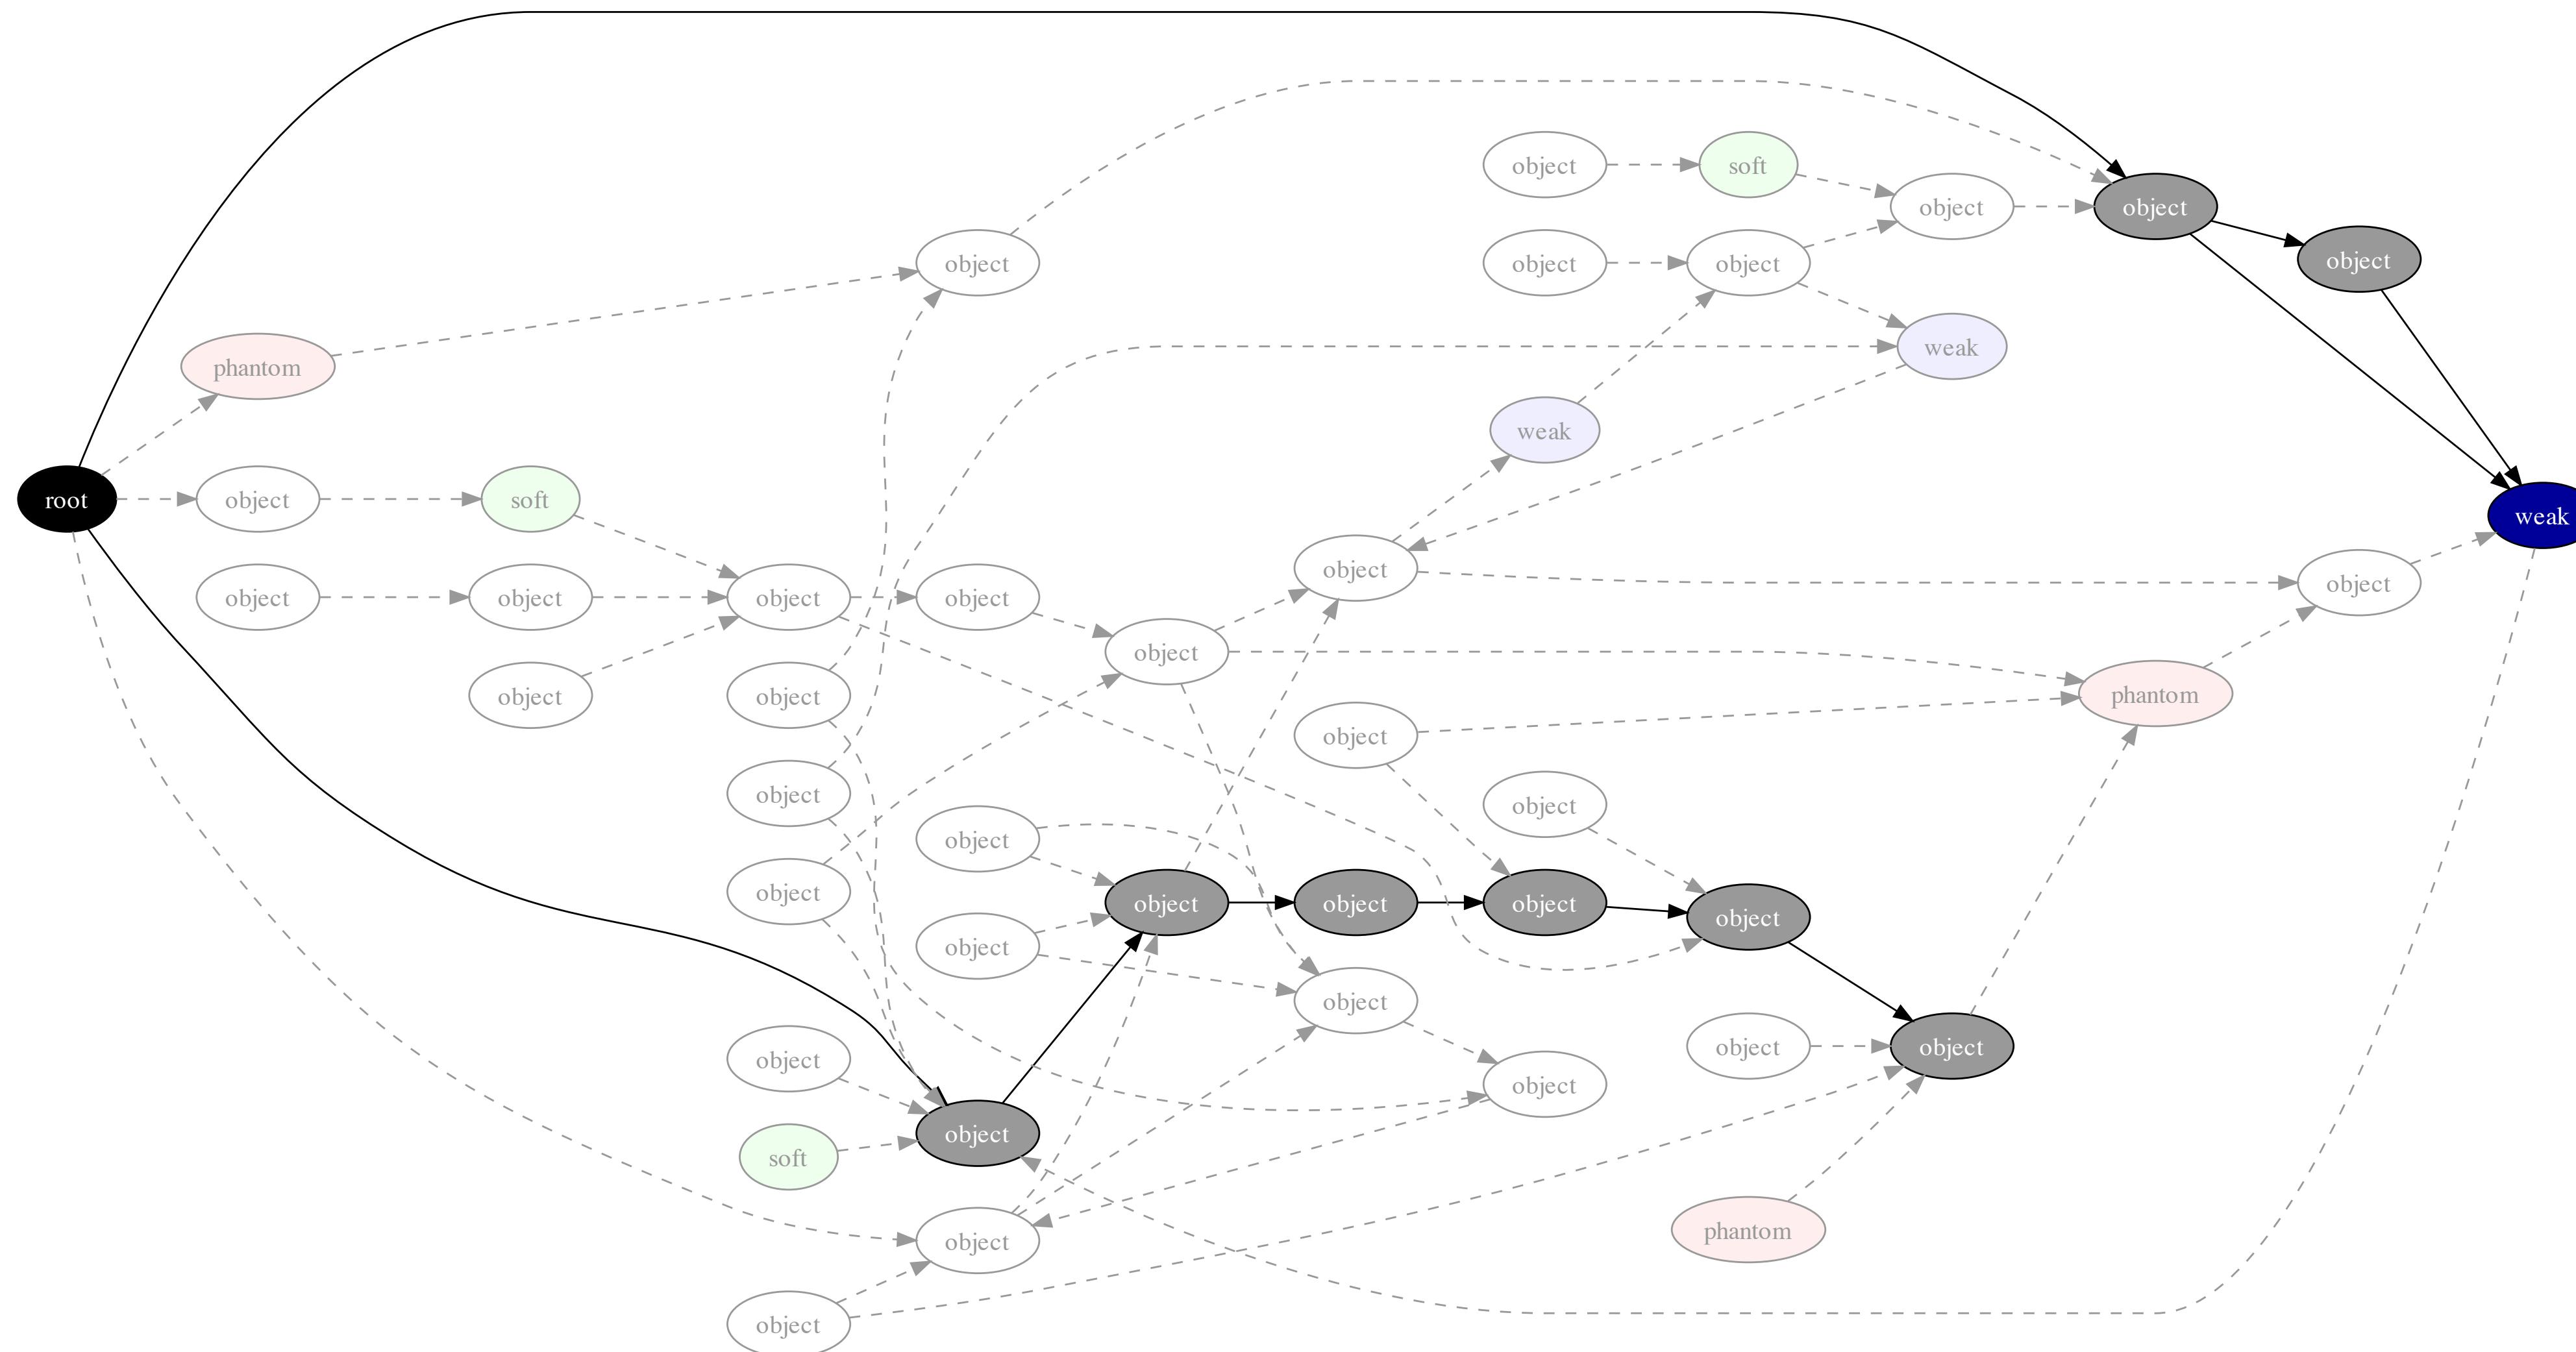
2. Trace and mark strongly-referenced objects.



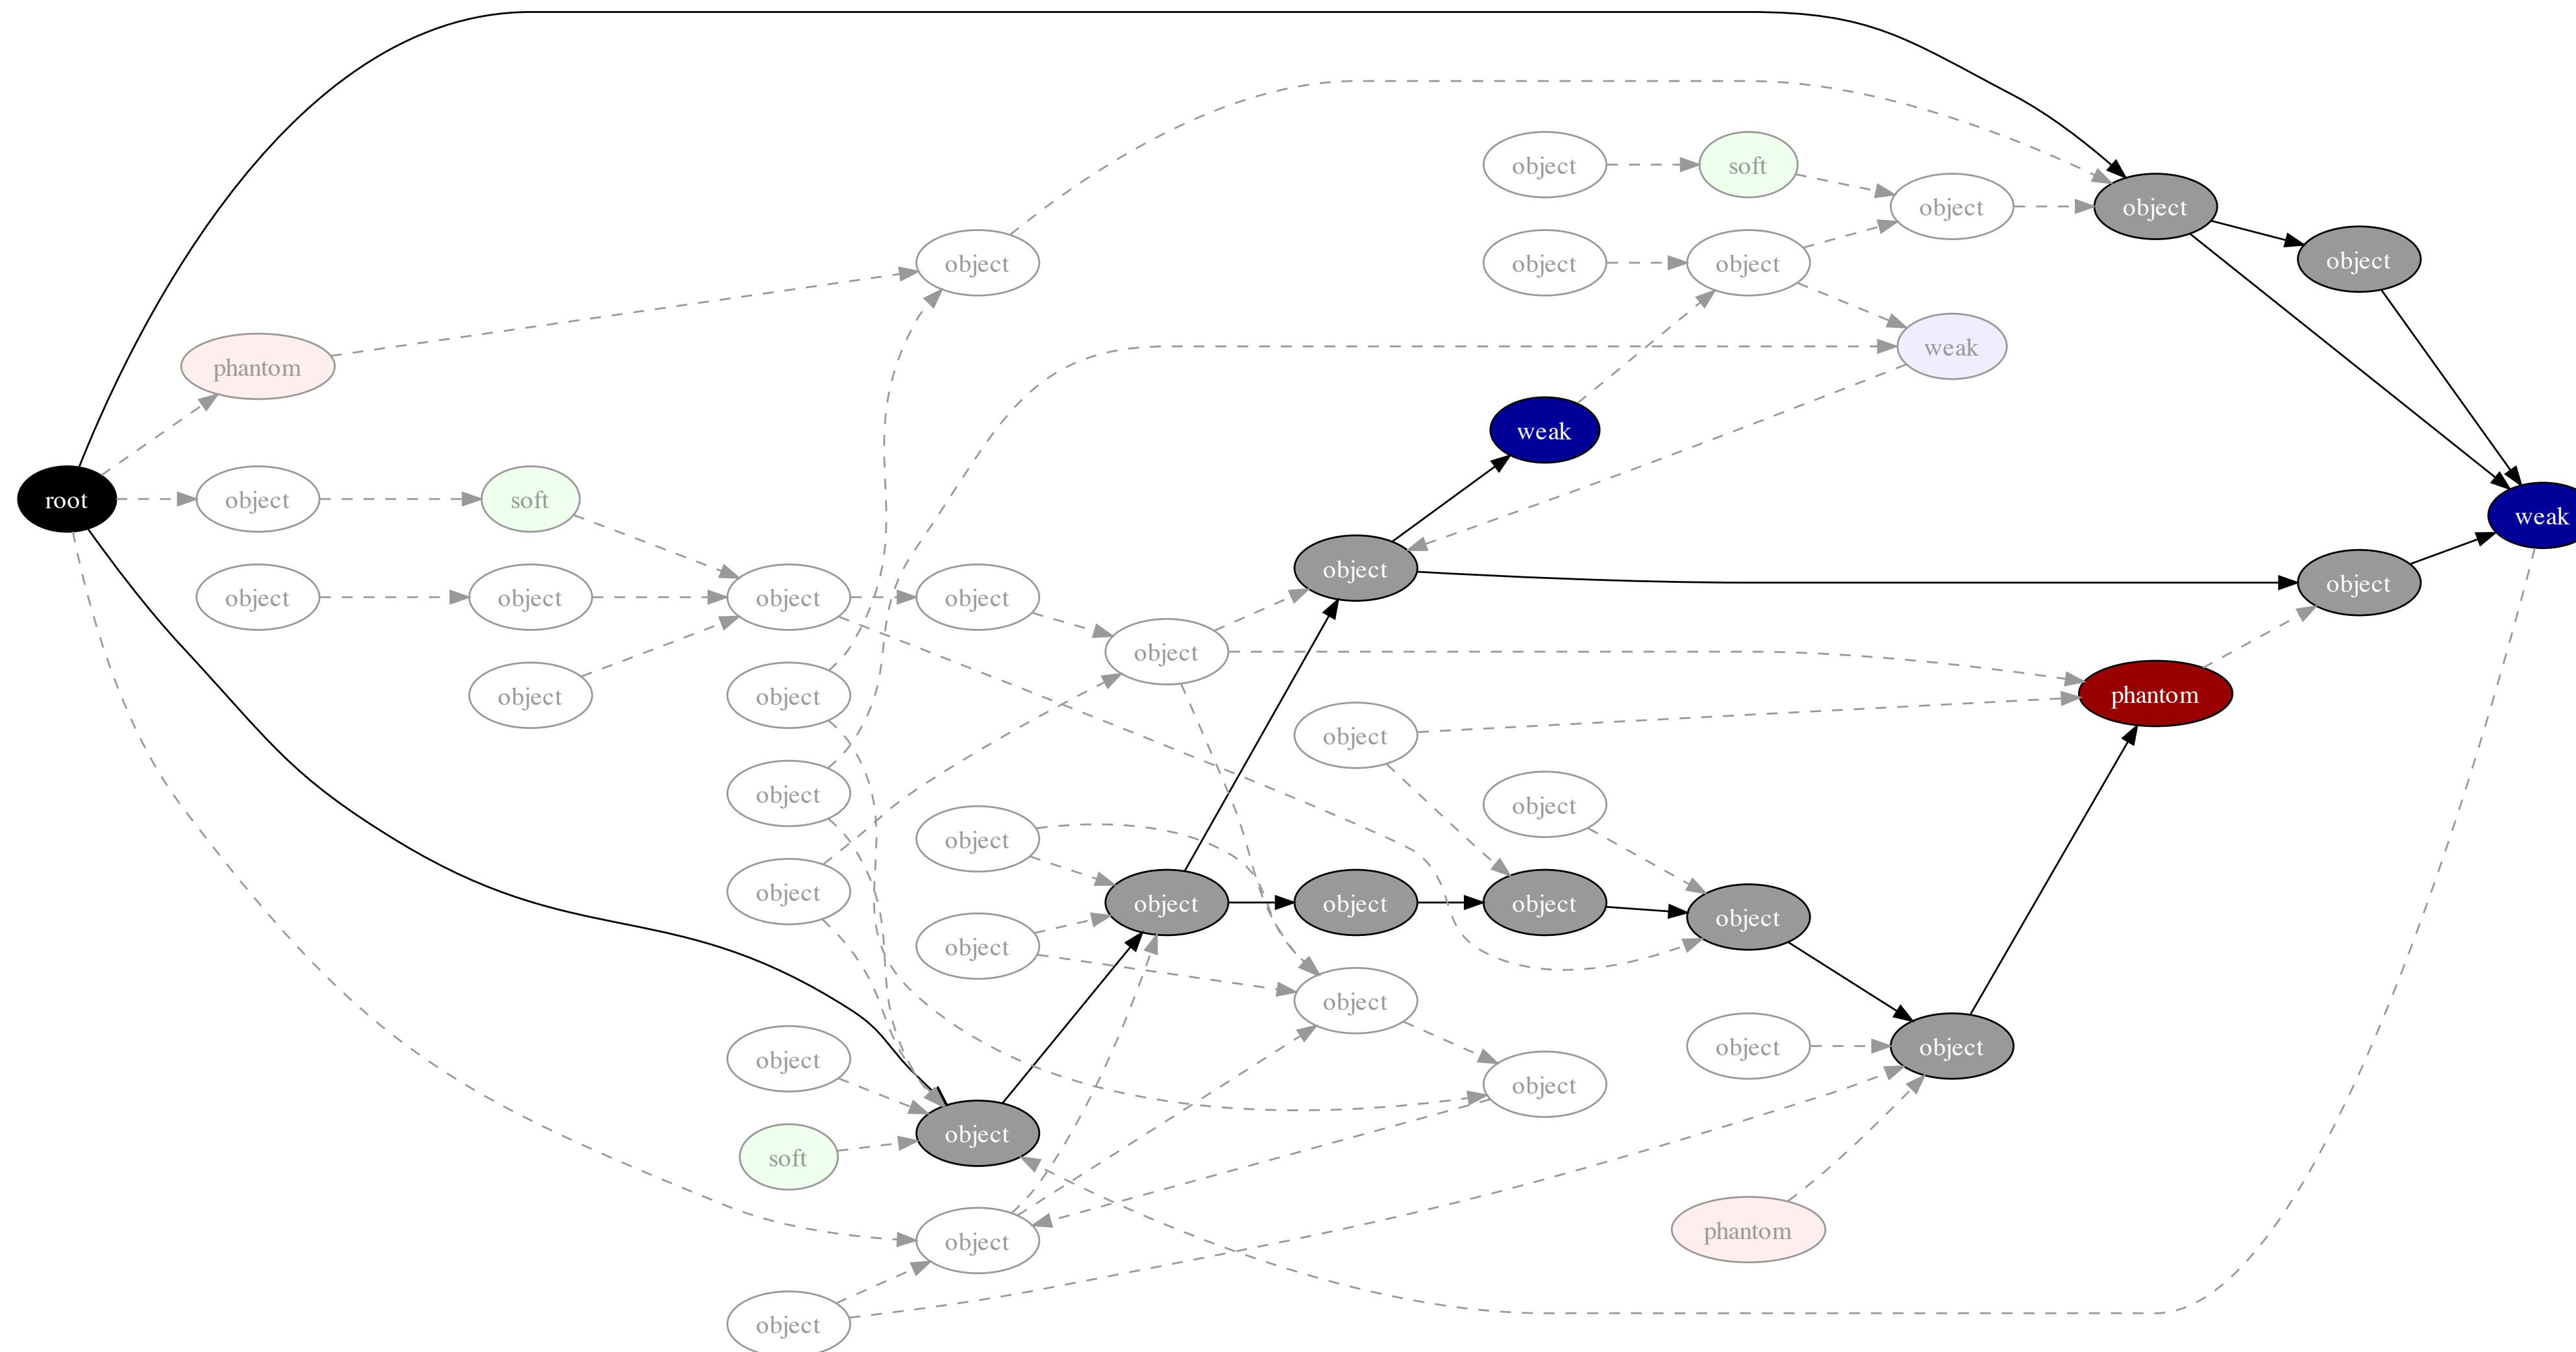
2. Trace and mark strongly-referenced objects.



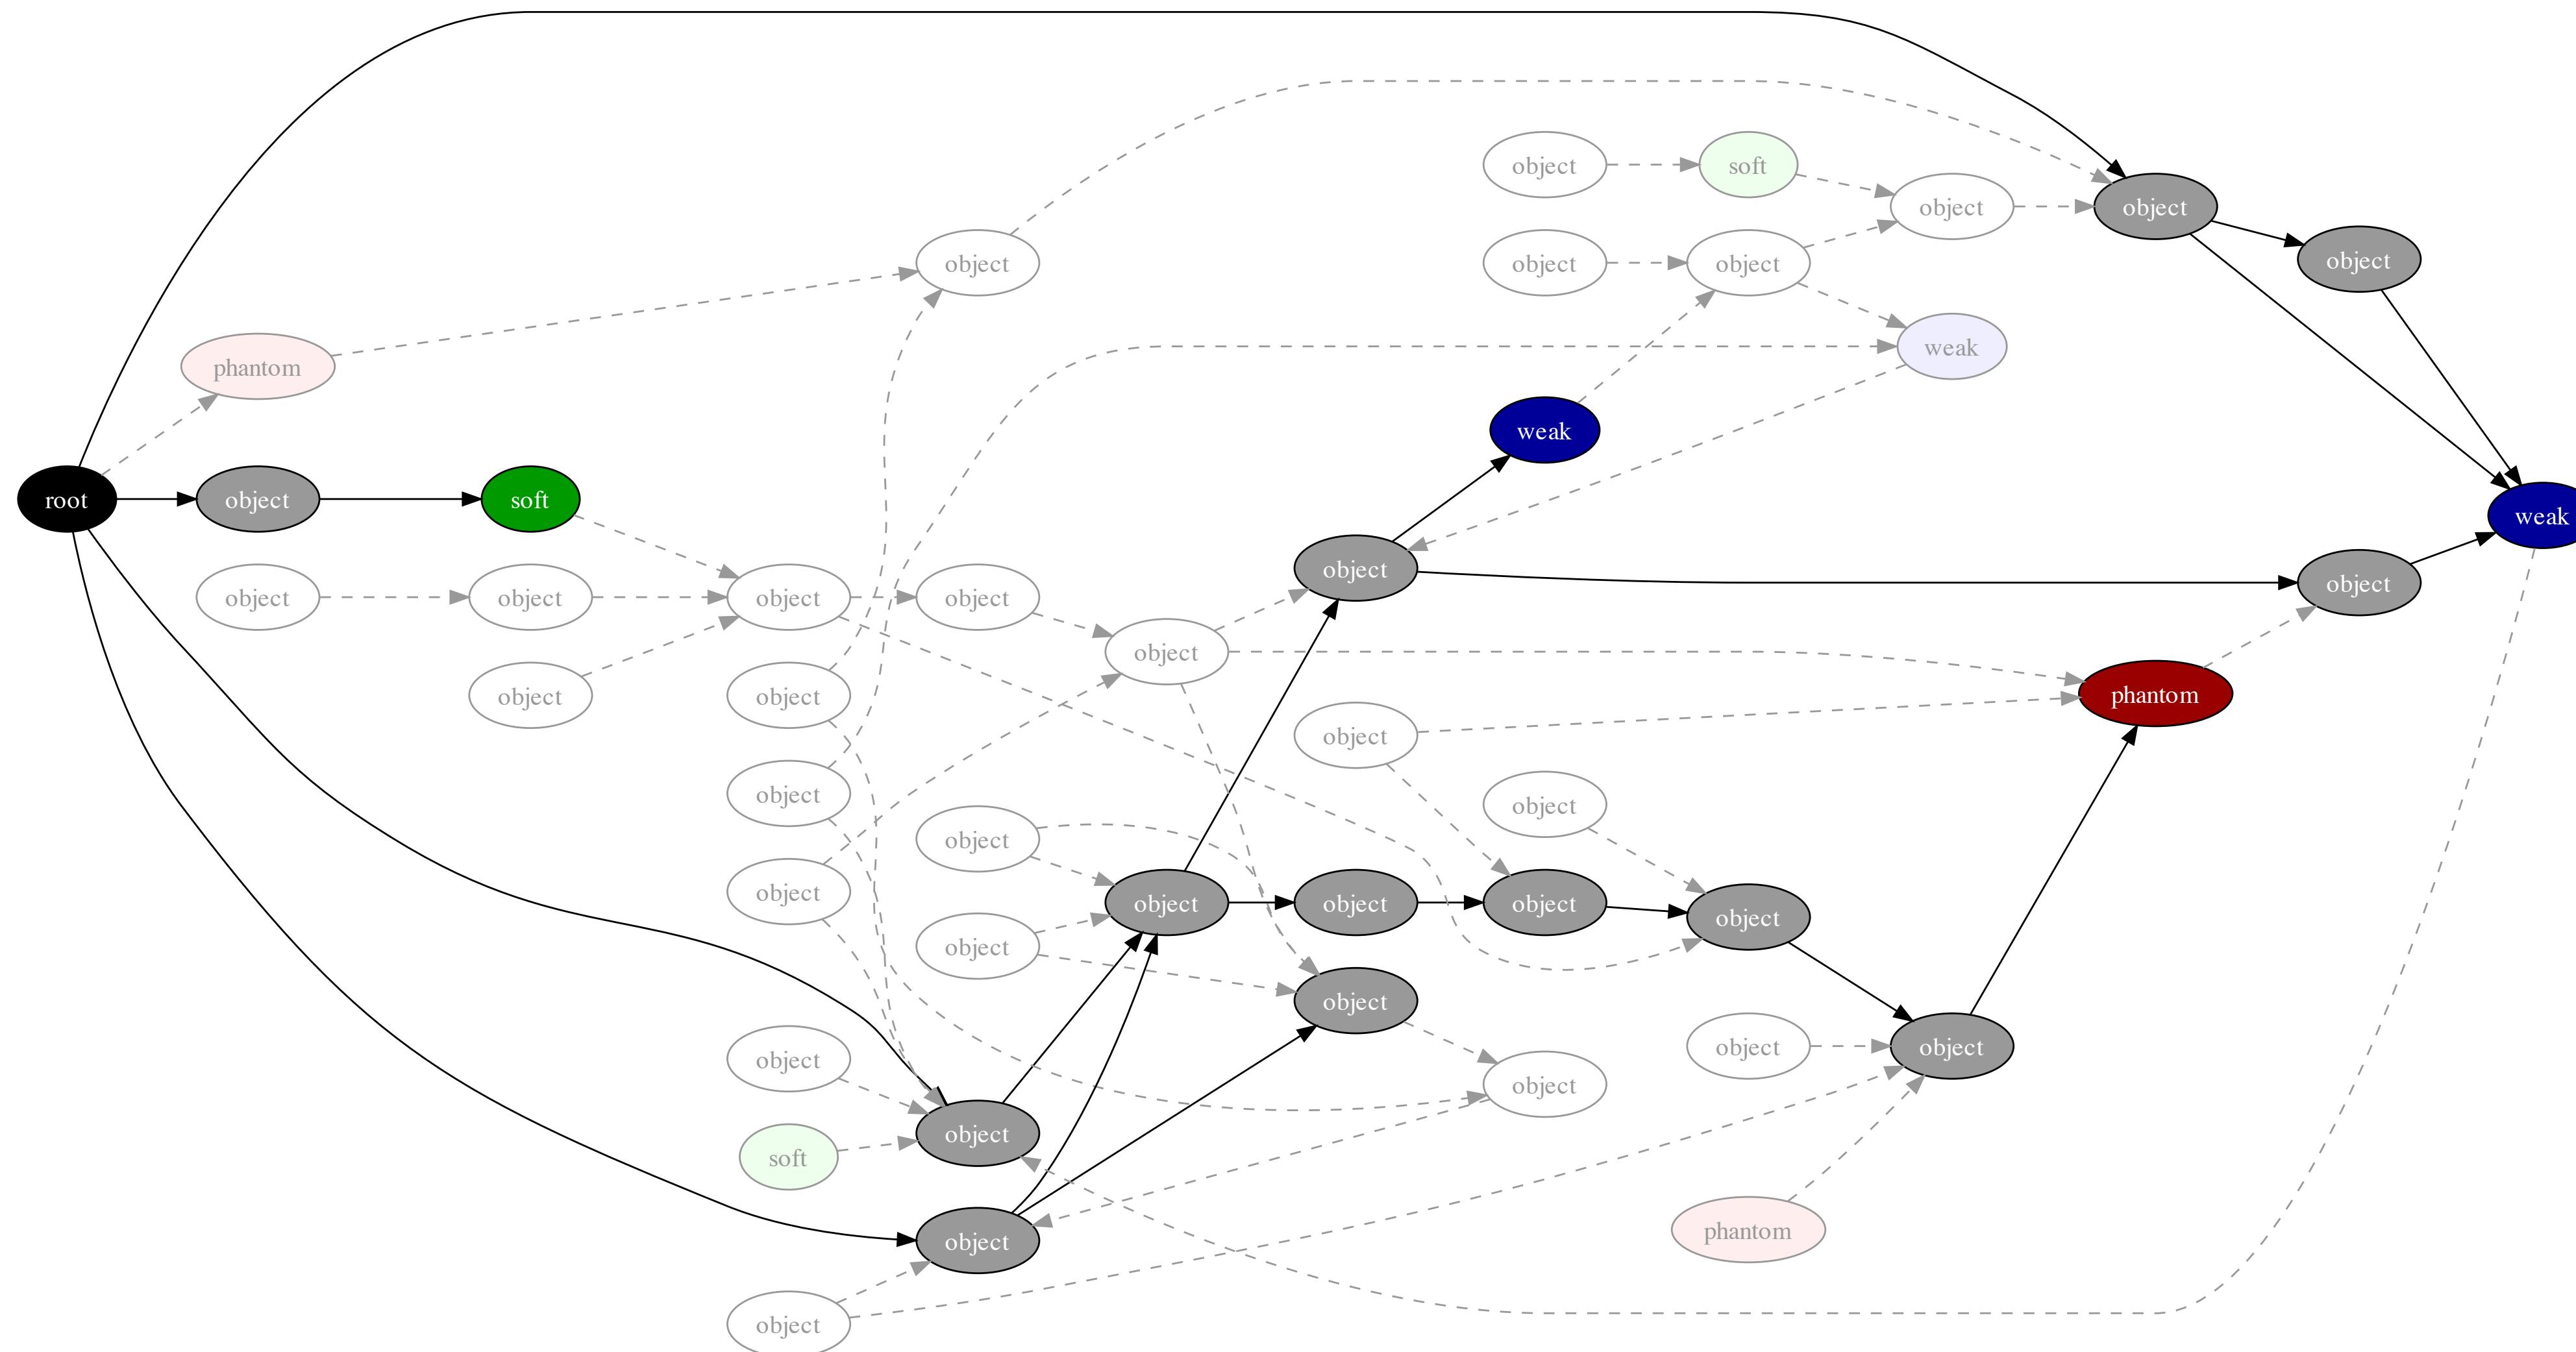
2. Trace and mark strongly-referenced objects.



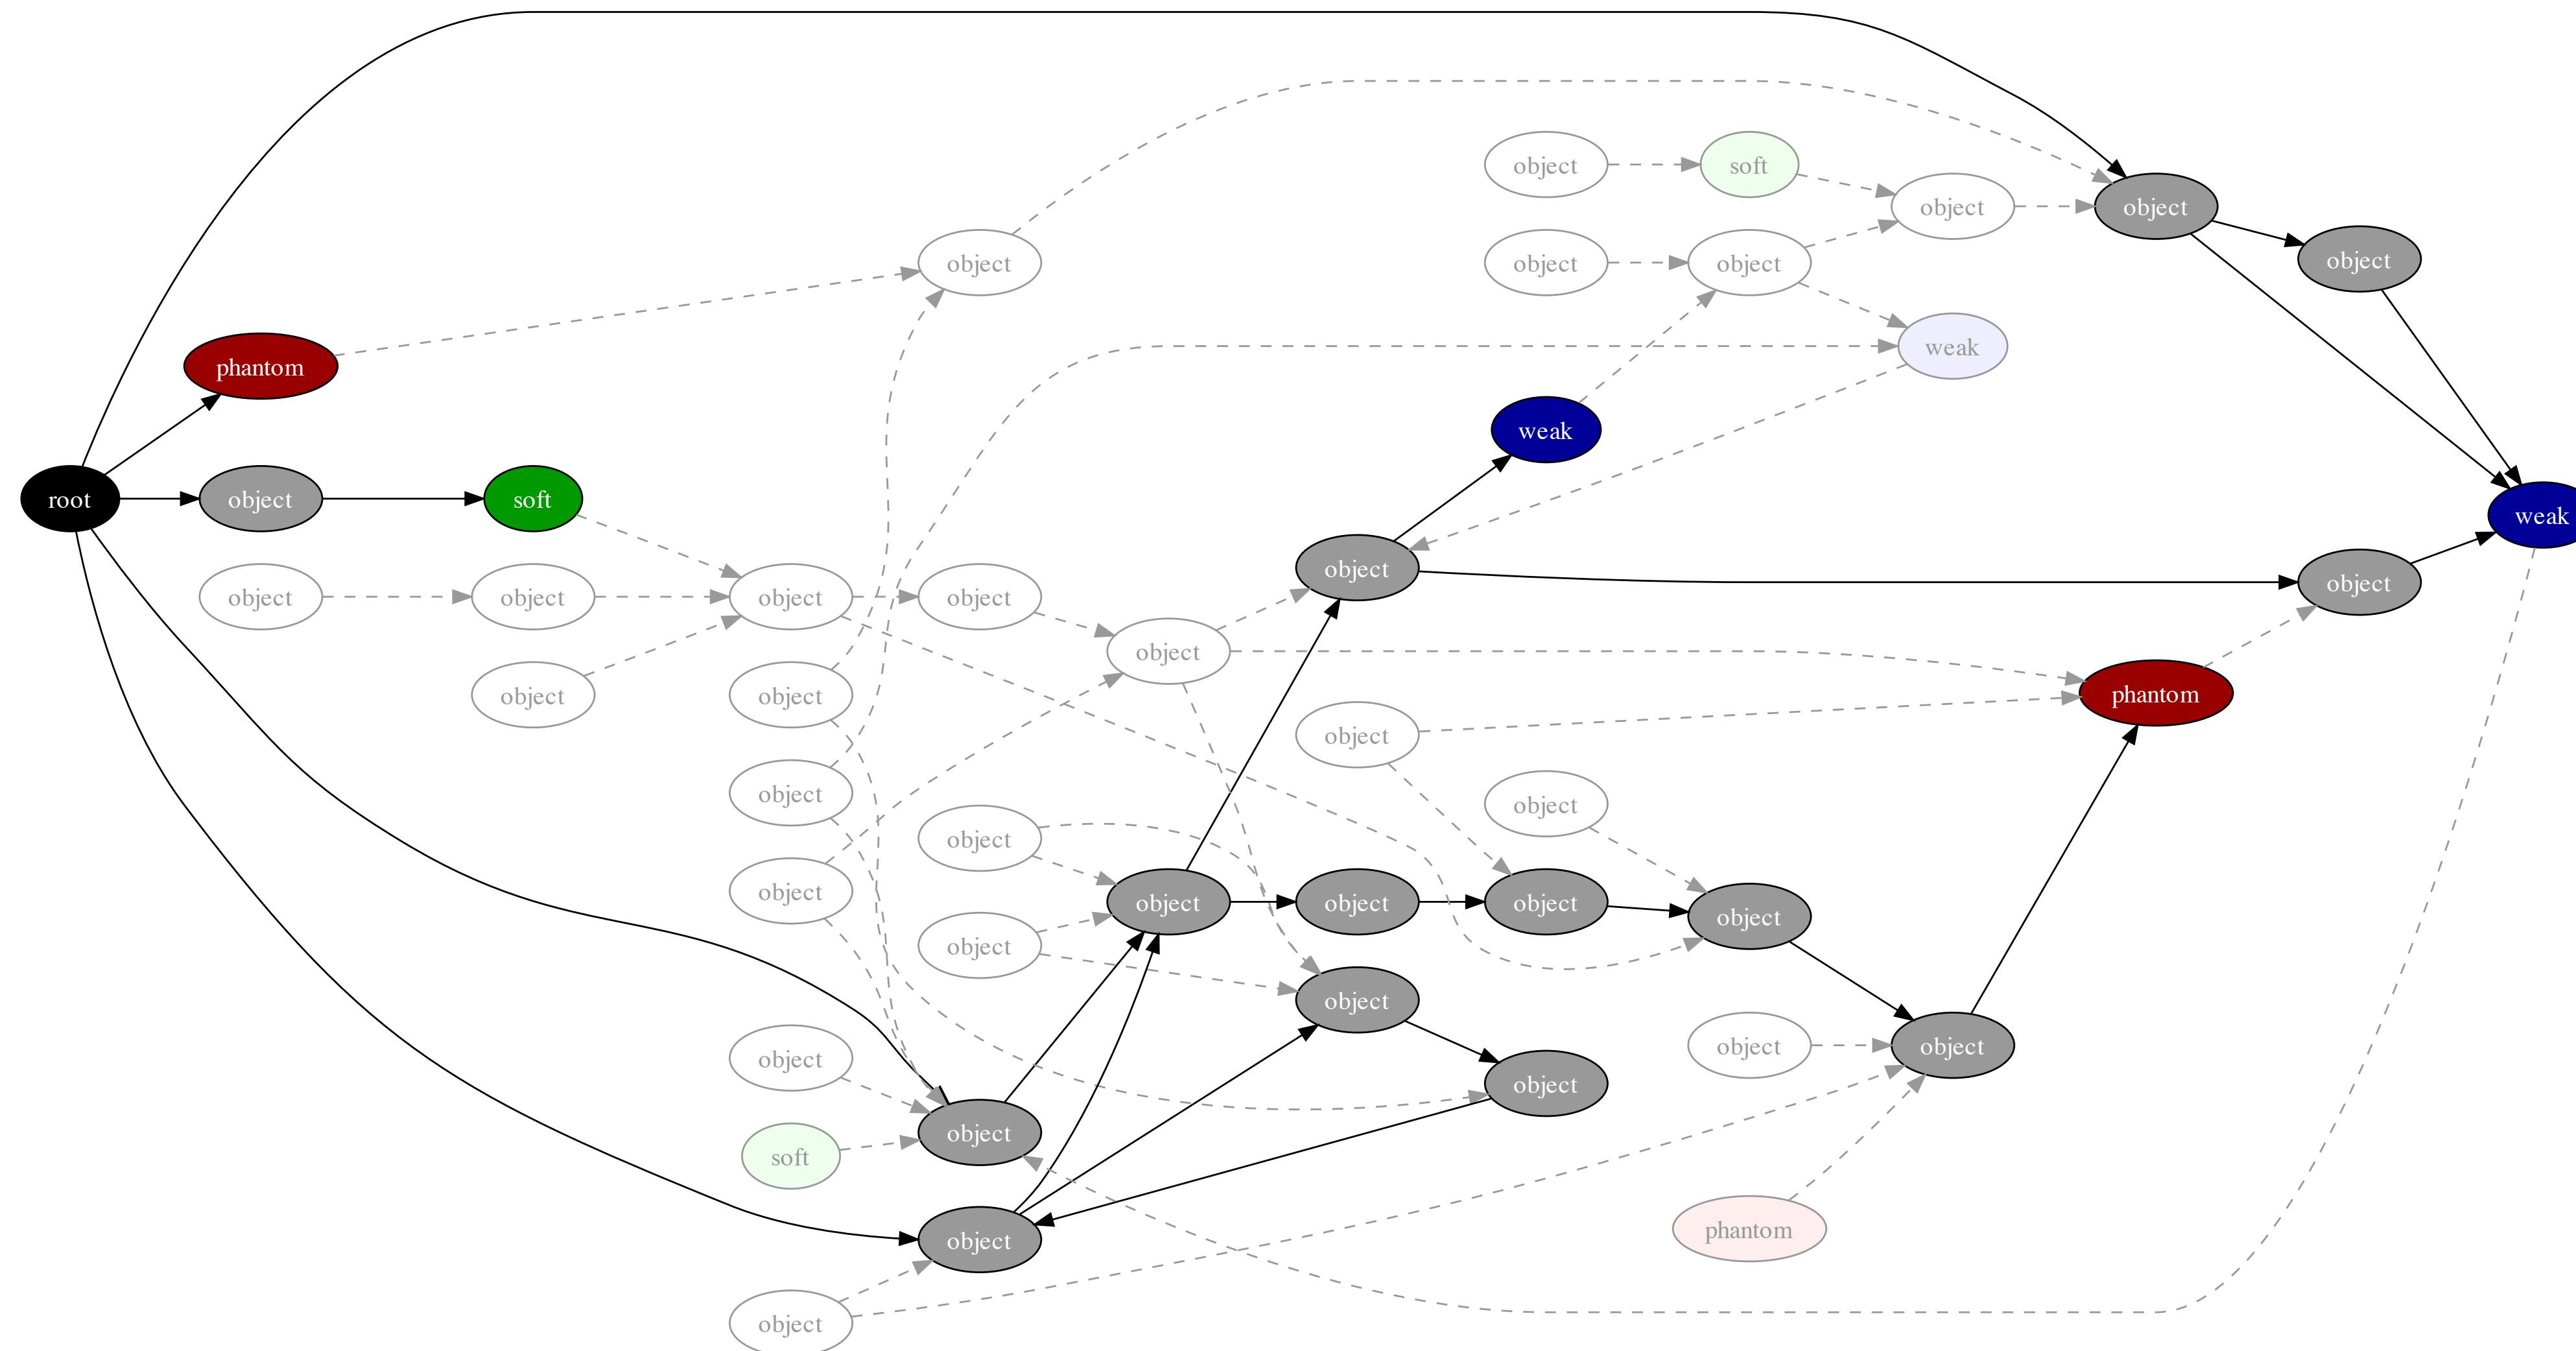
2. Trace and mark strongly-referenced objects.



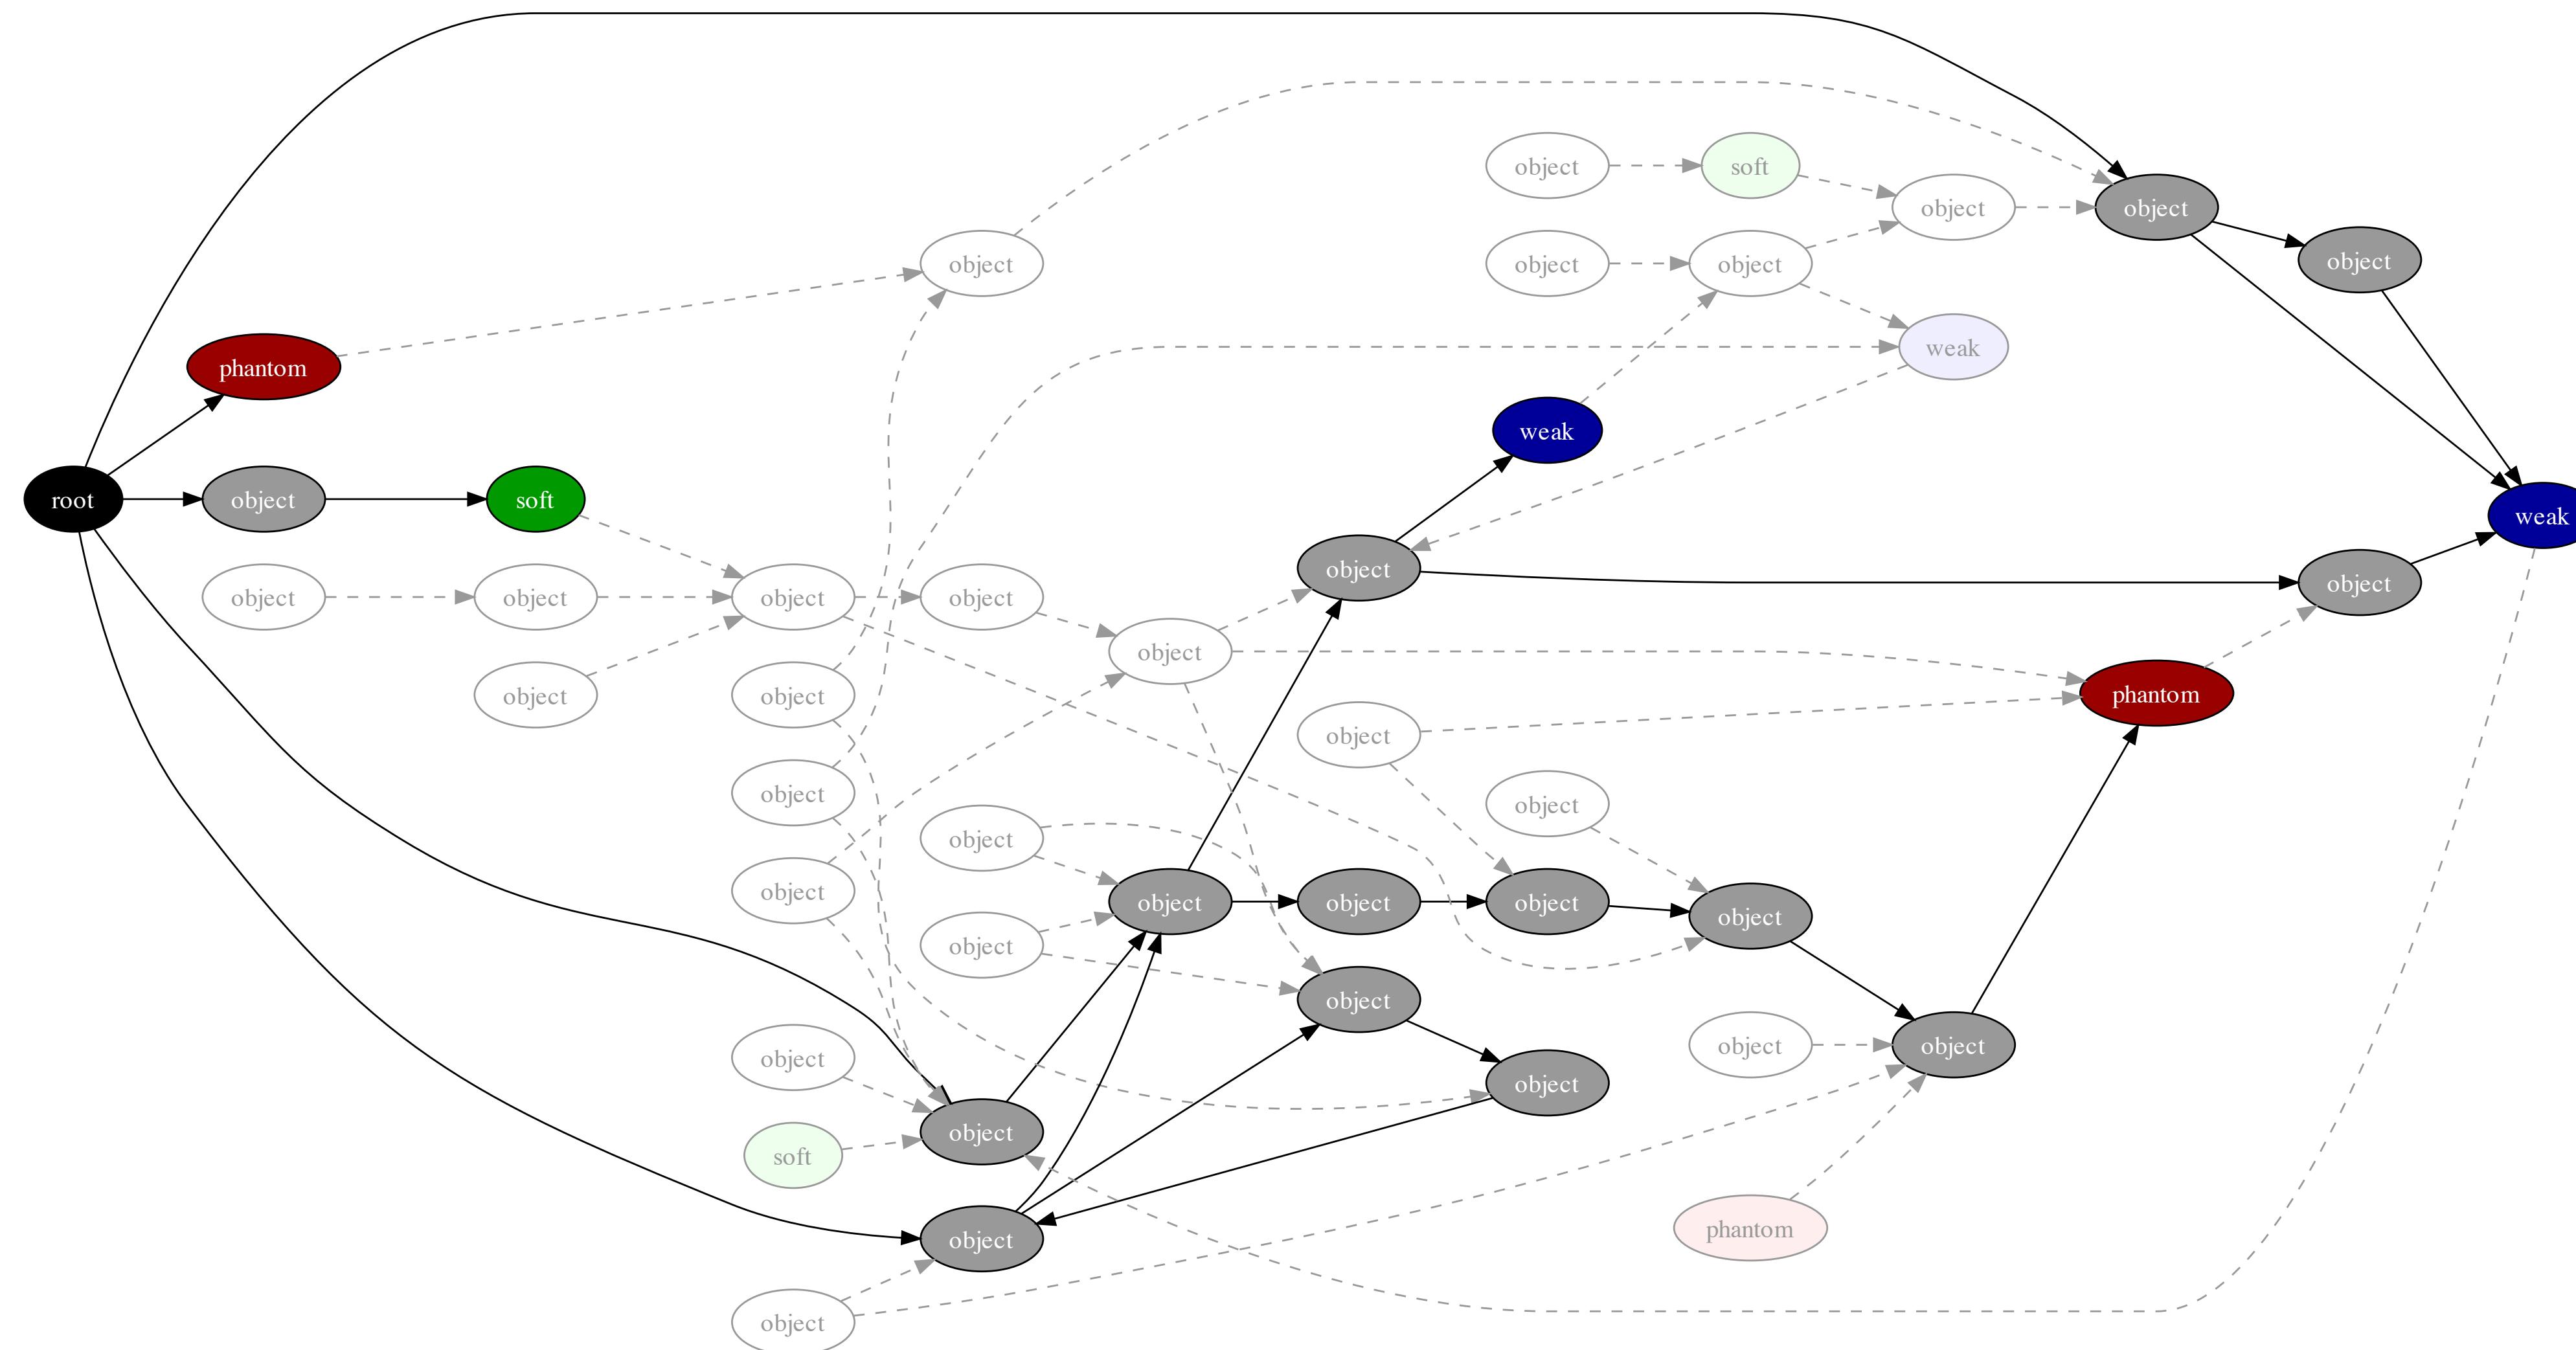
2. Trace and mark strongly-referenced objects.



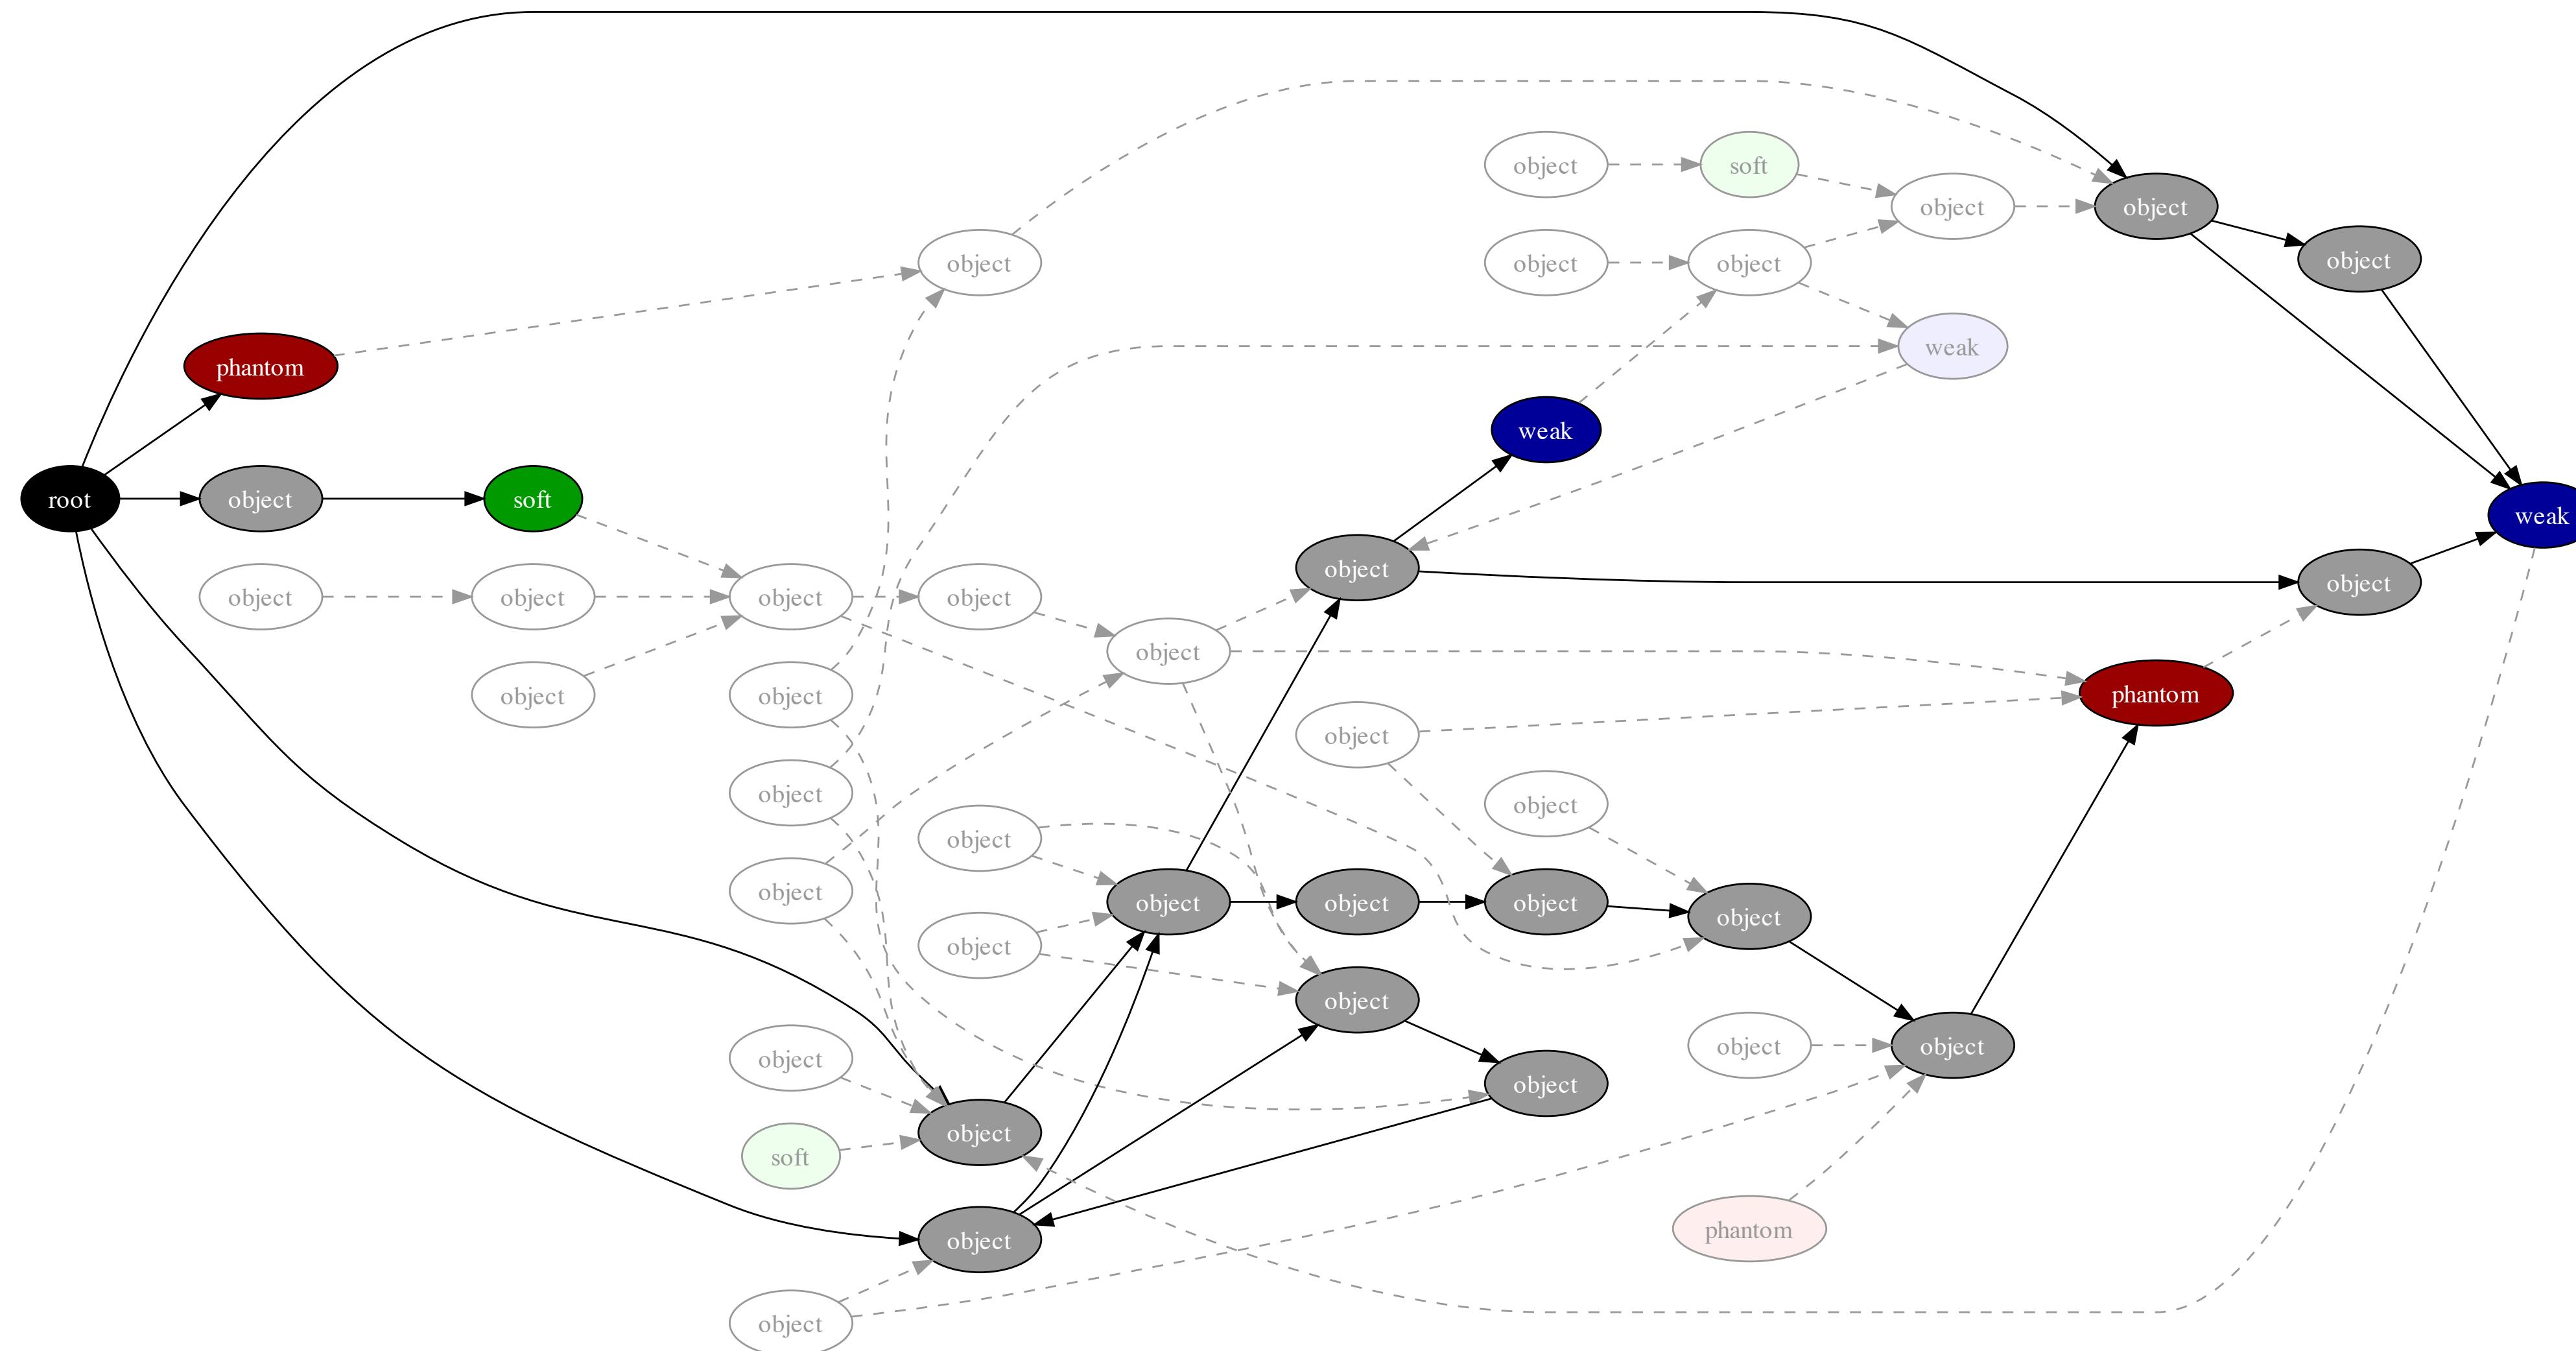
2. Trace and mark strongly-referenced objects.



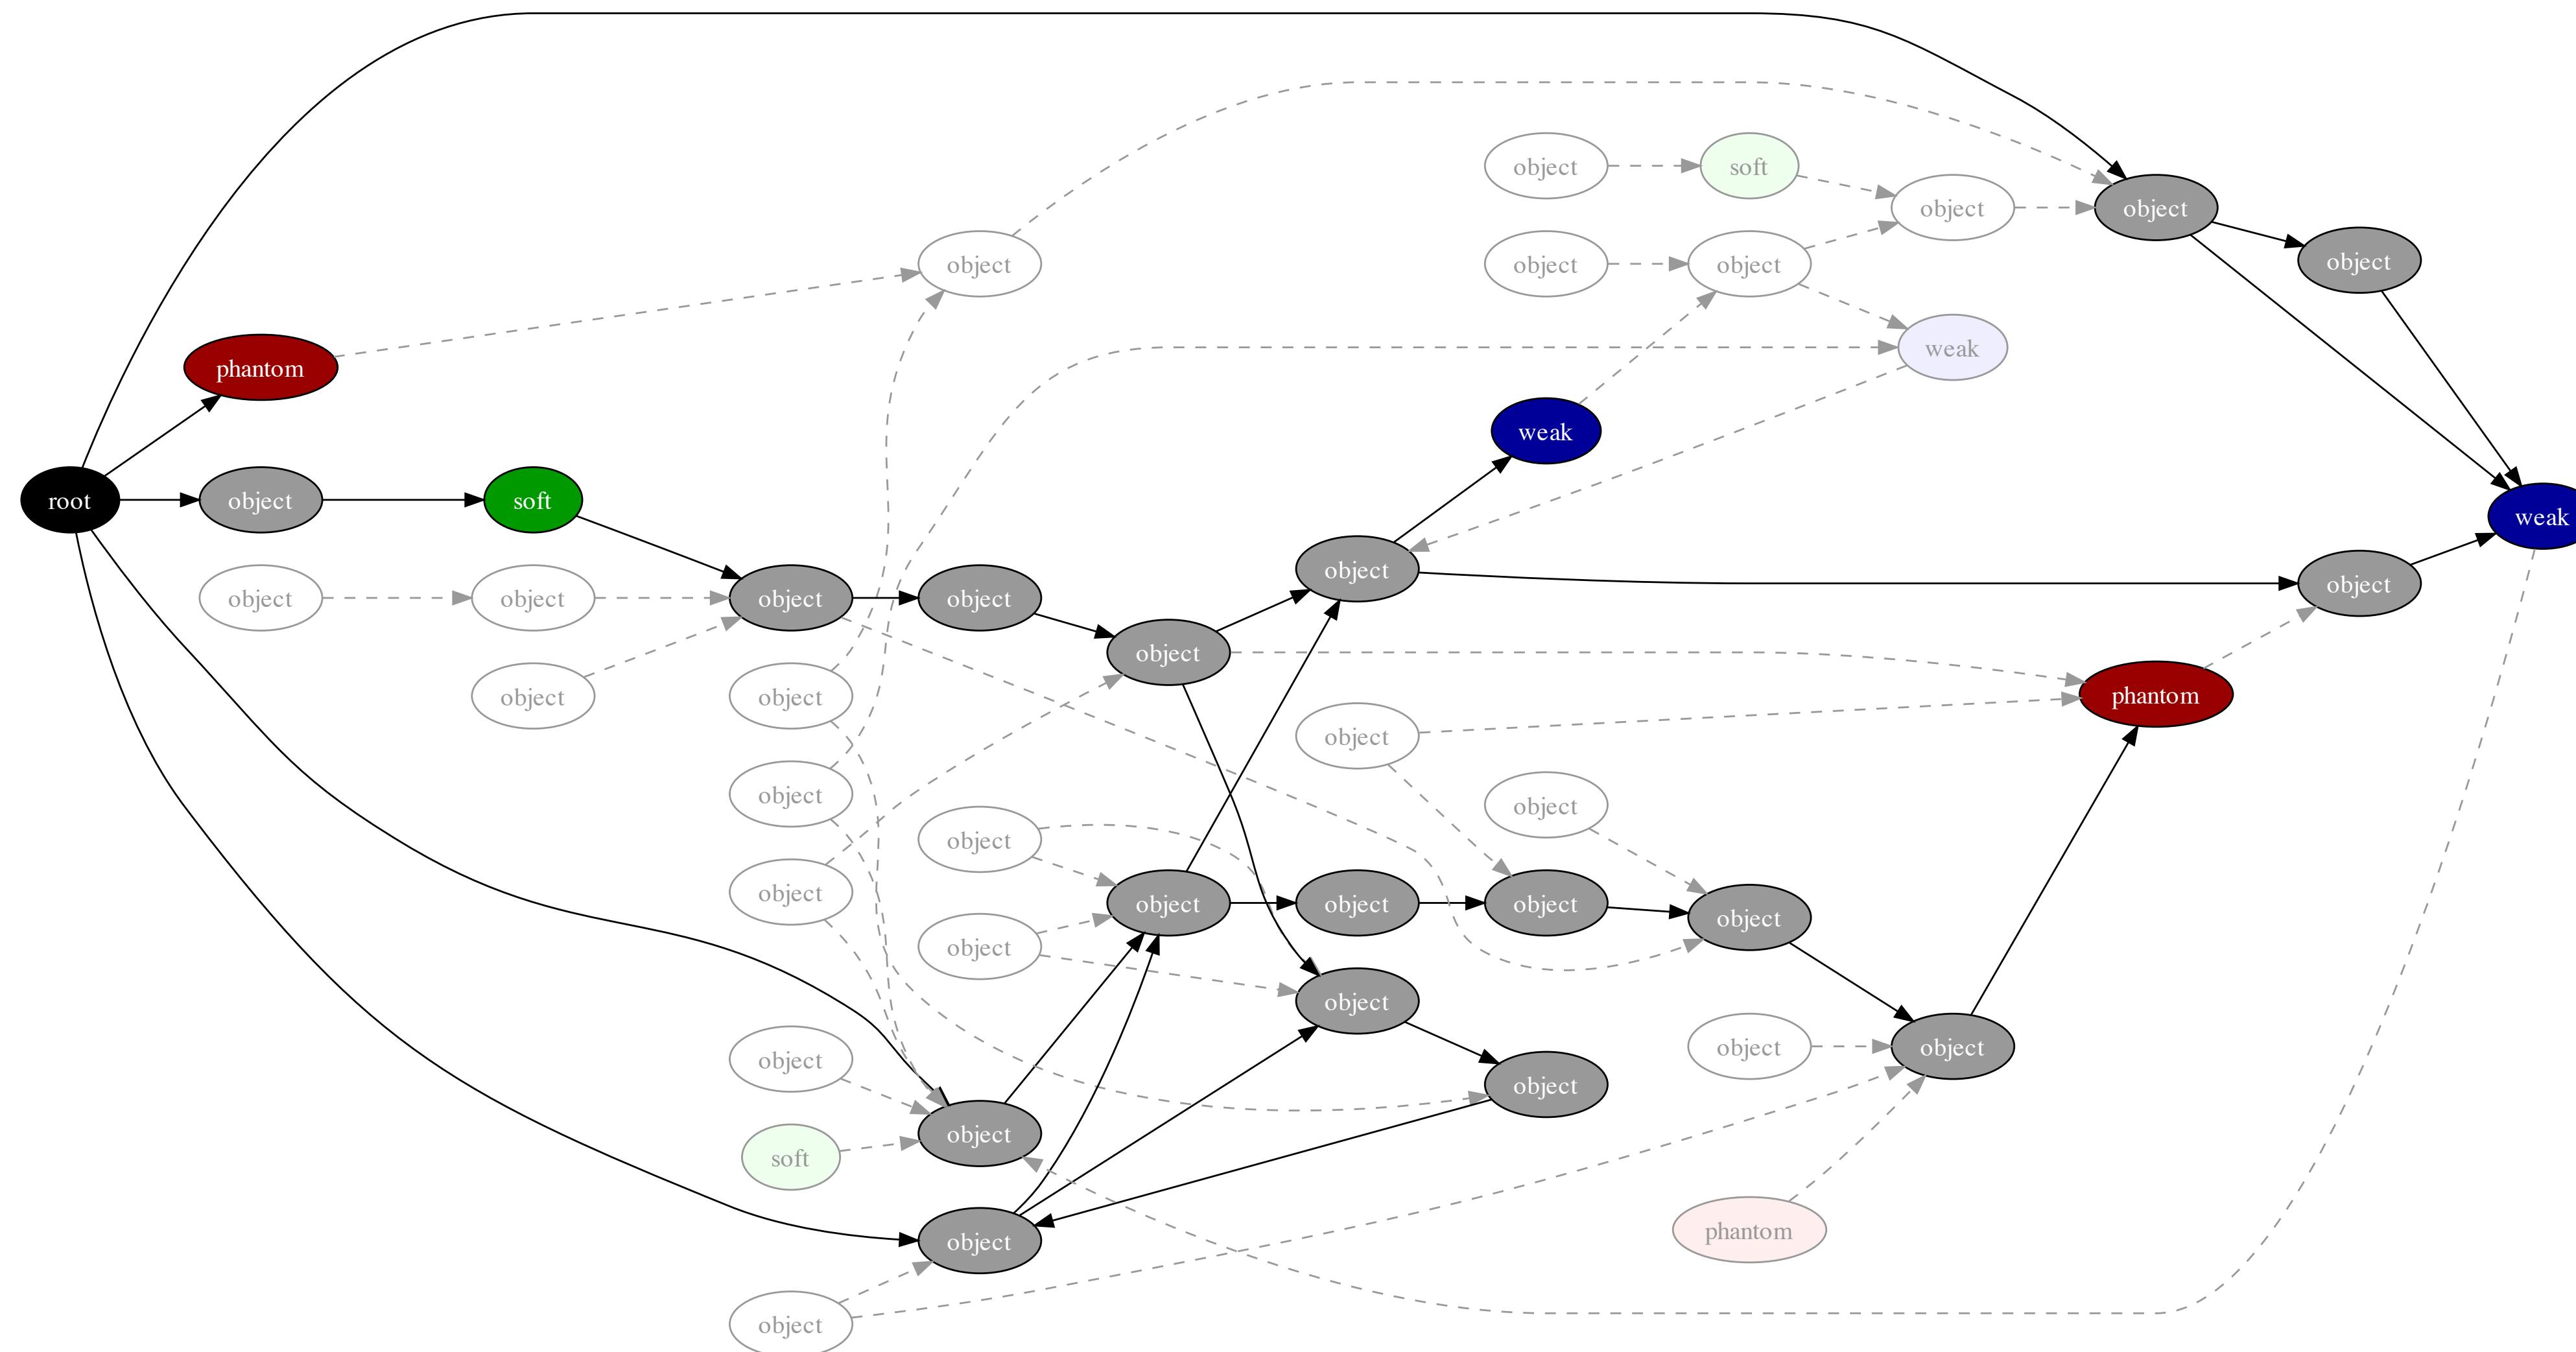
3. Optionally clear soft references.



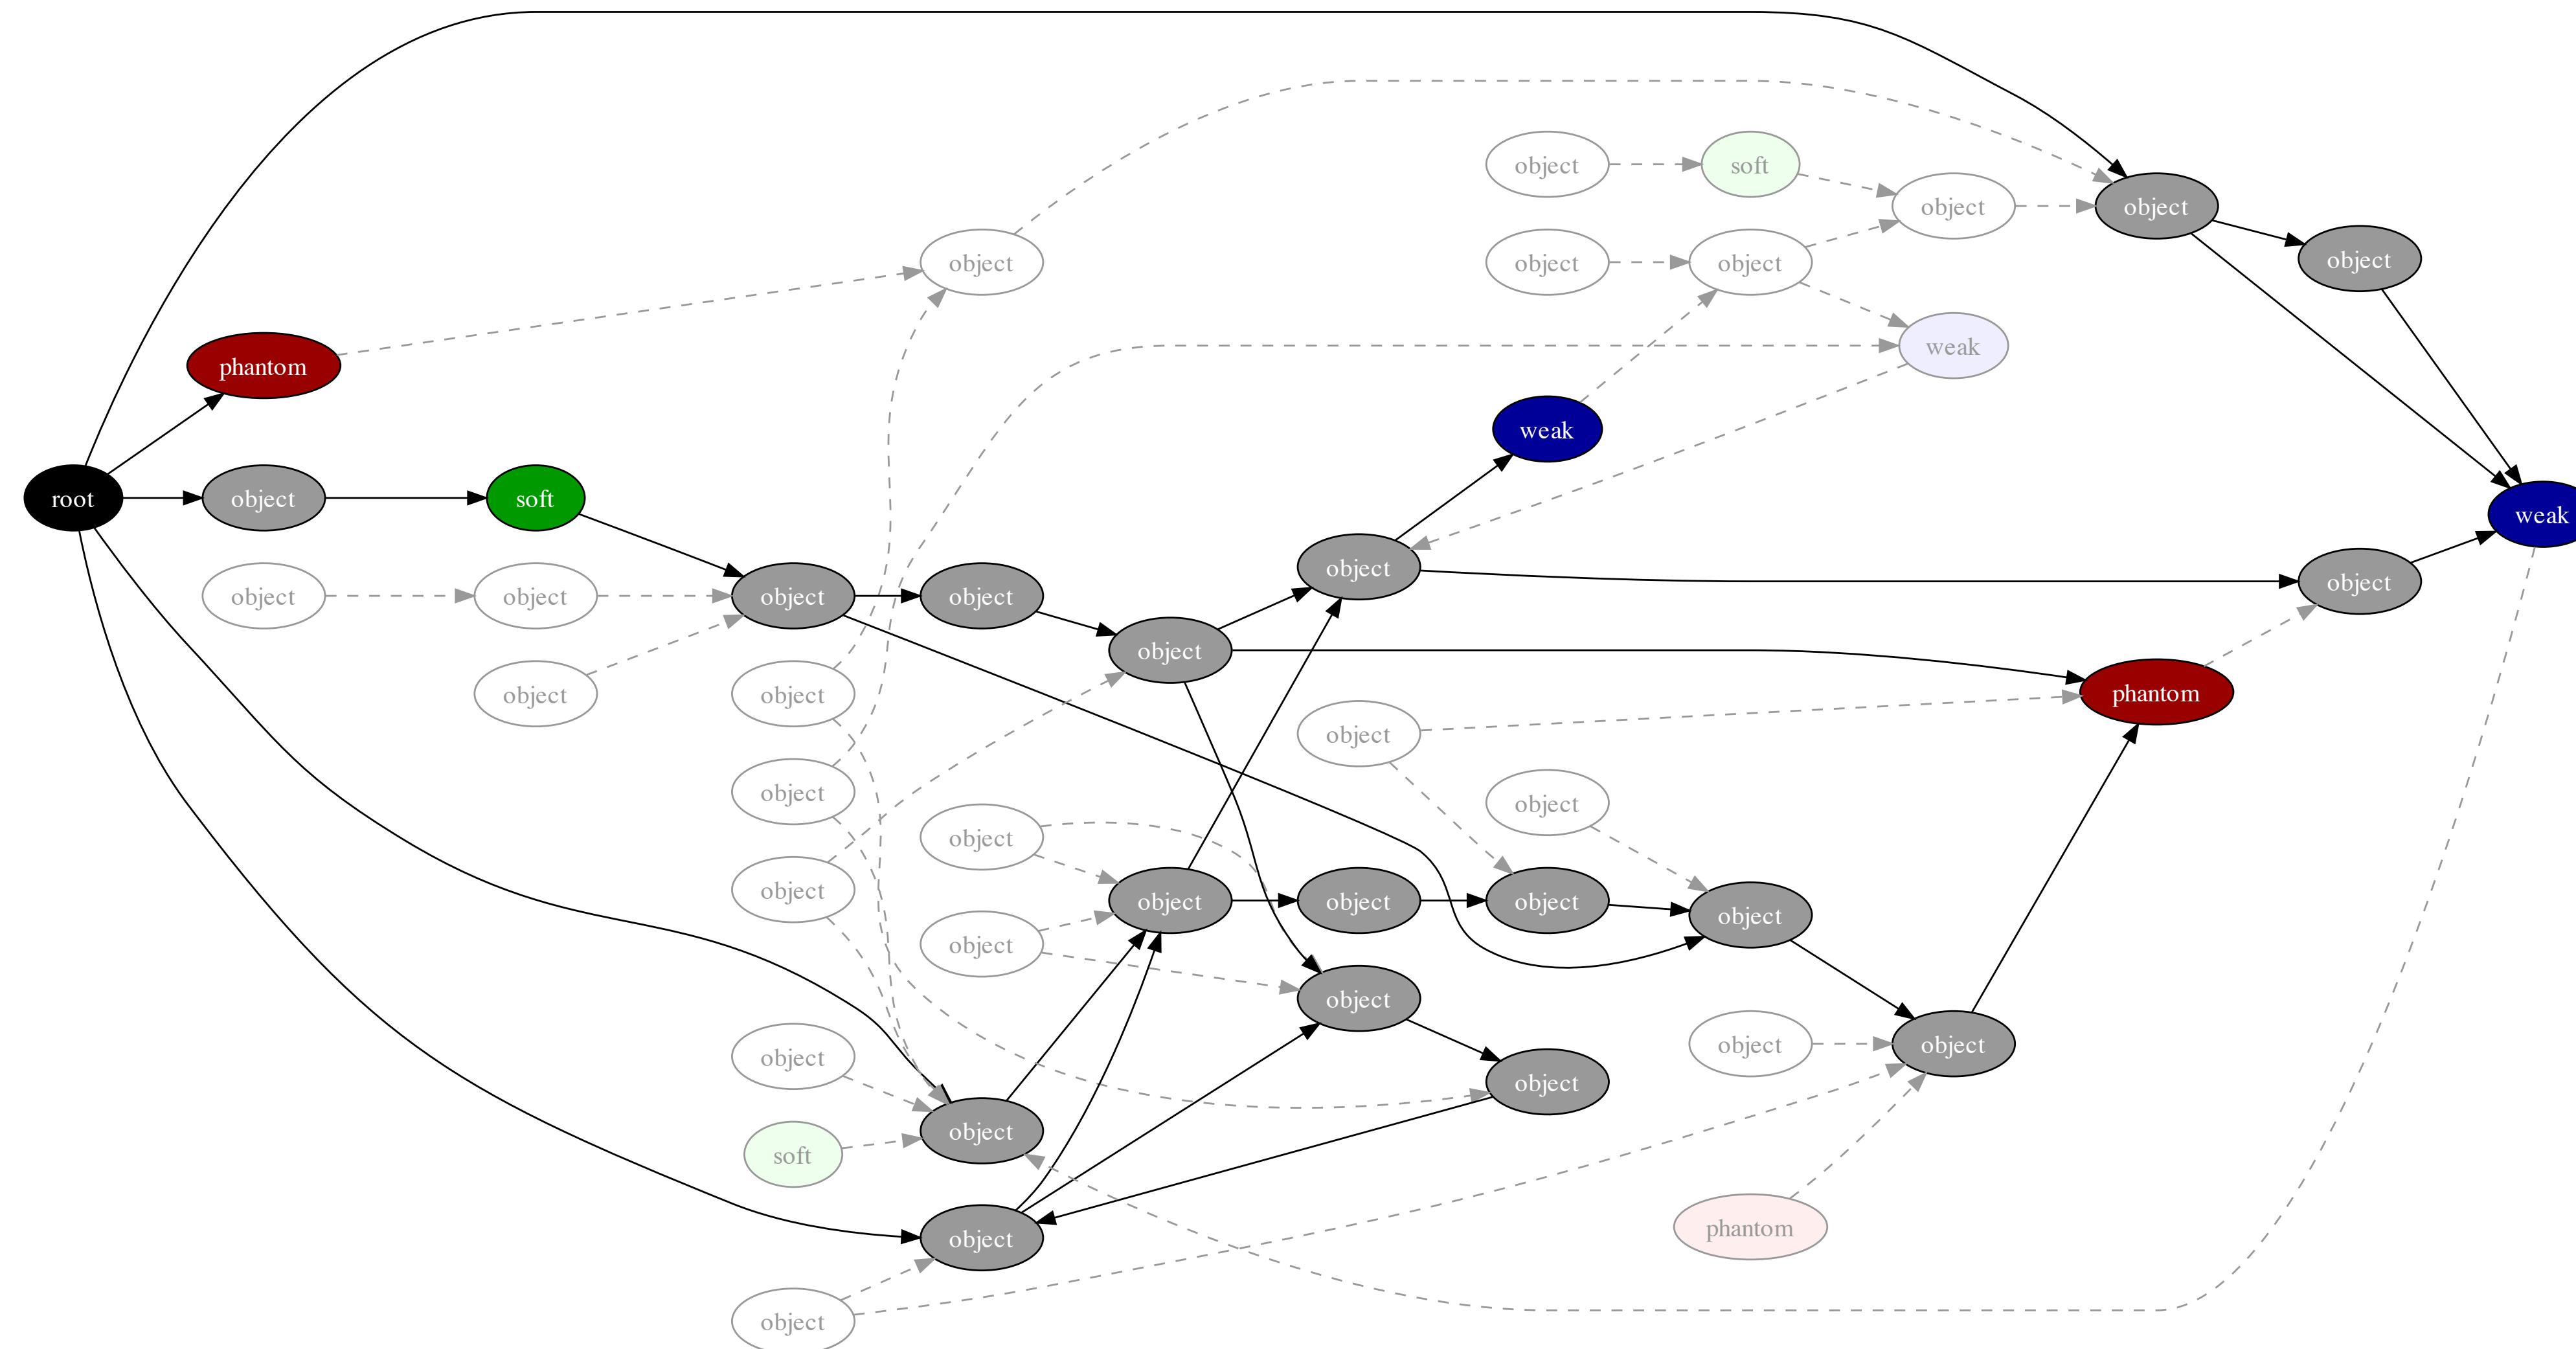
4. Trace and mark softly-referenced objects.



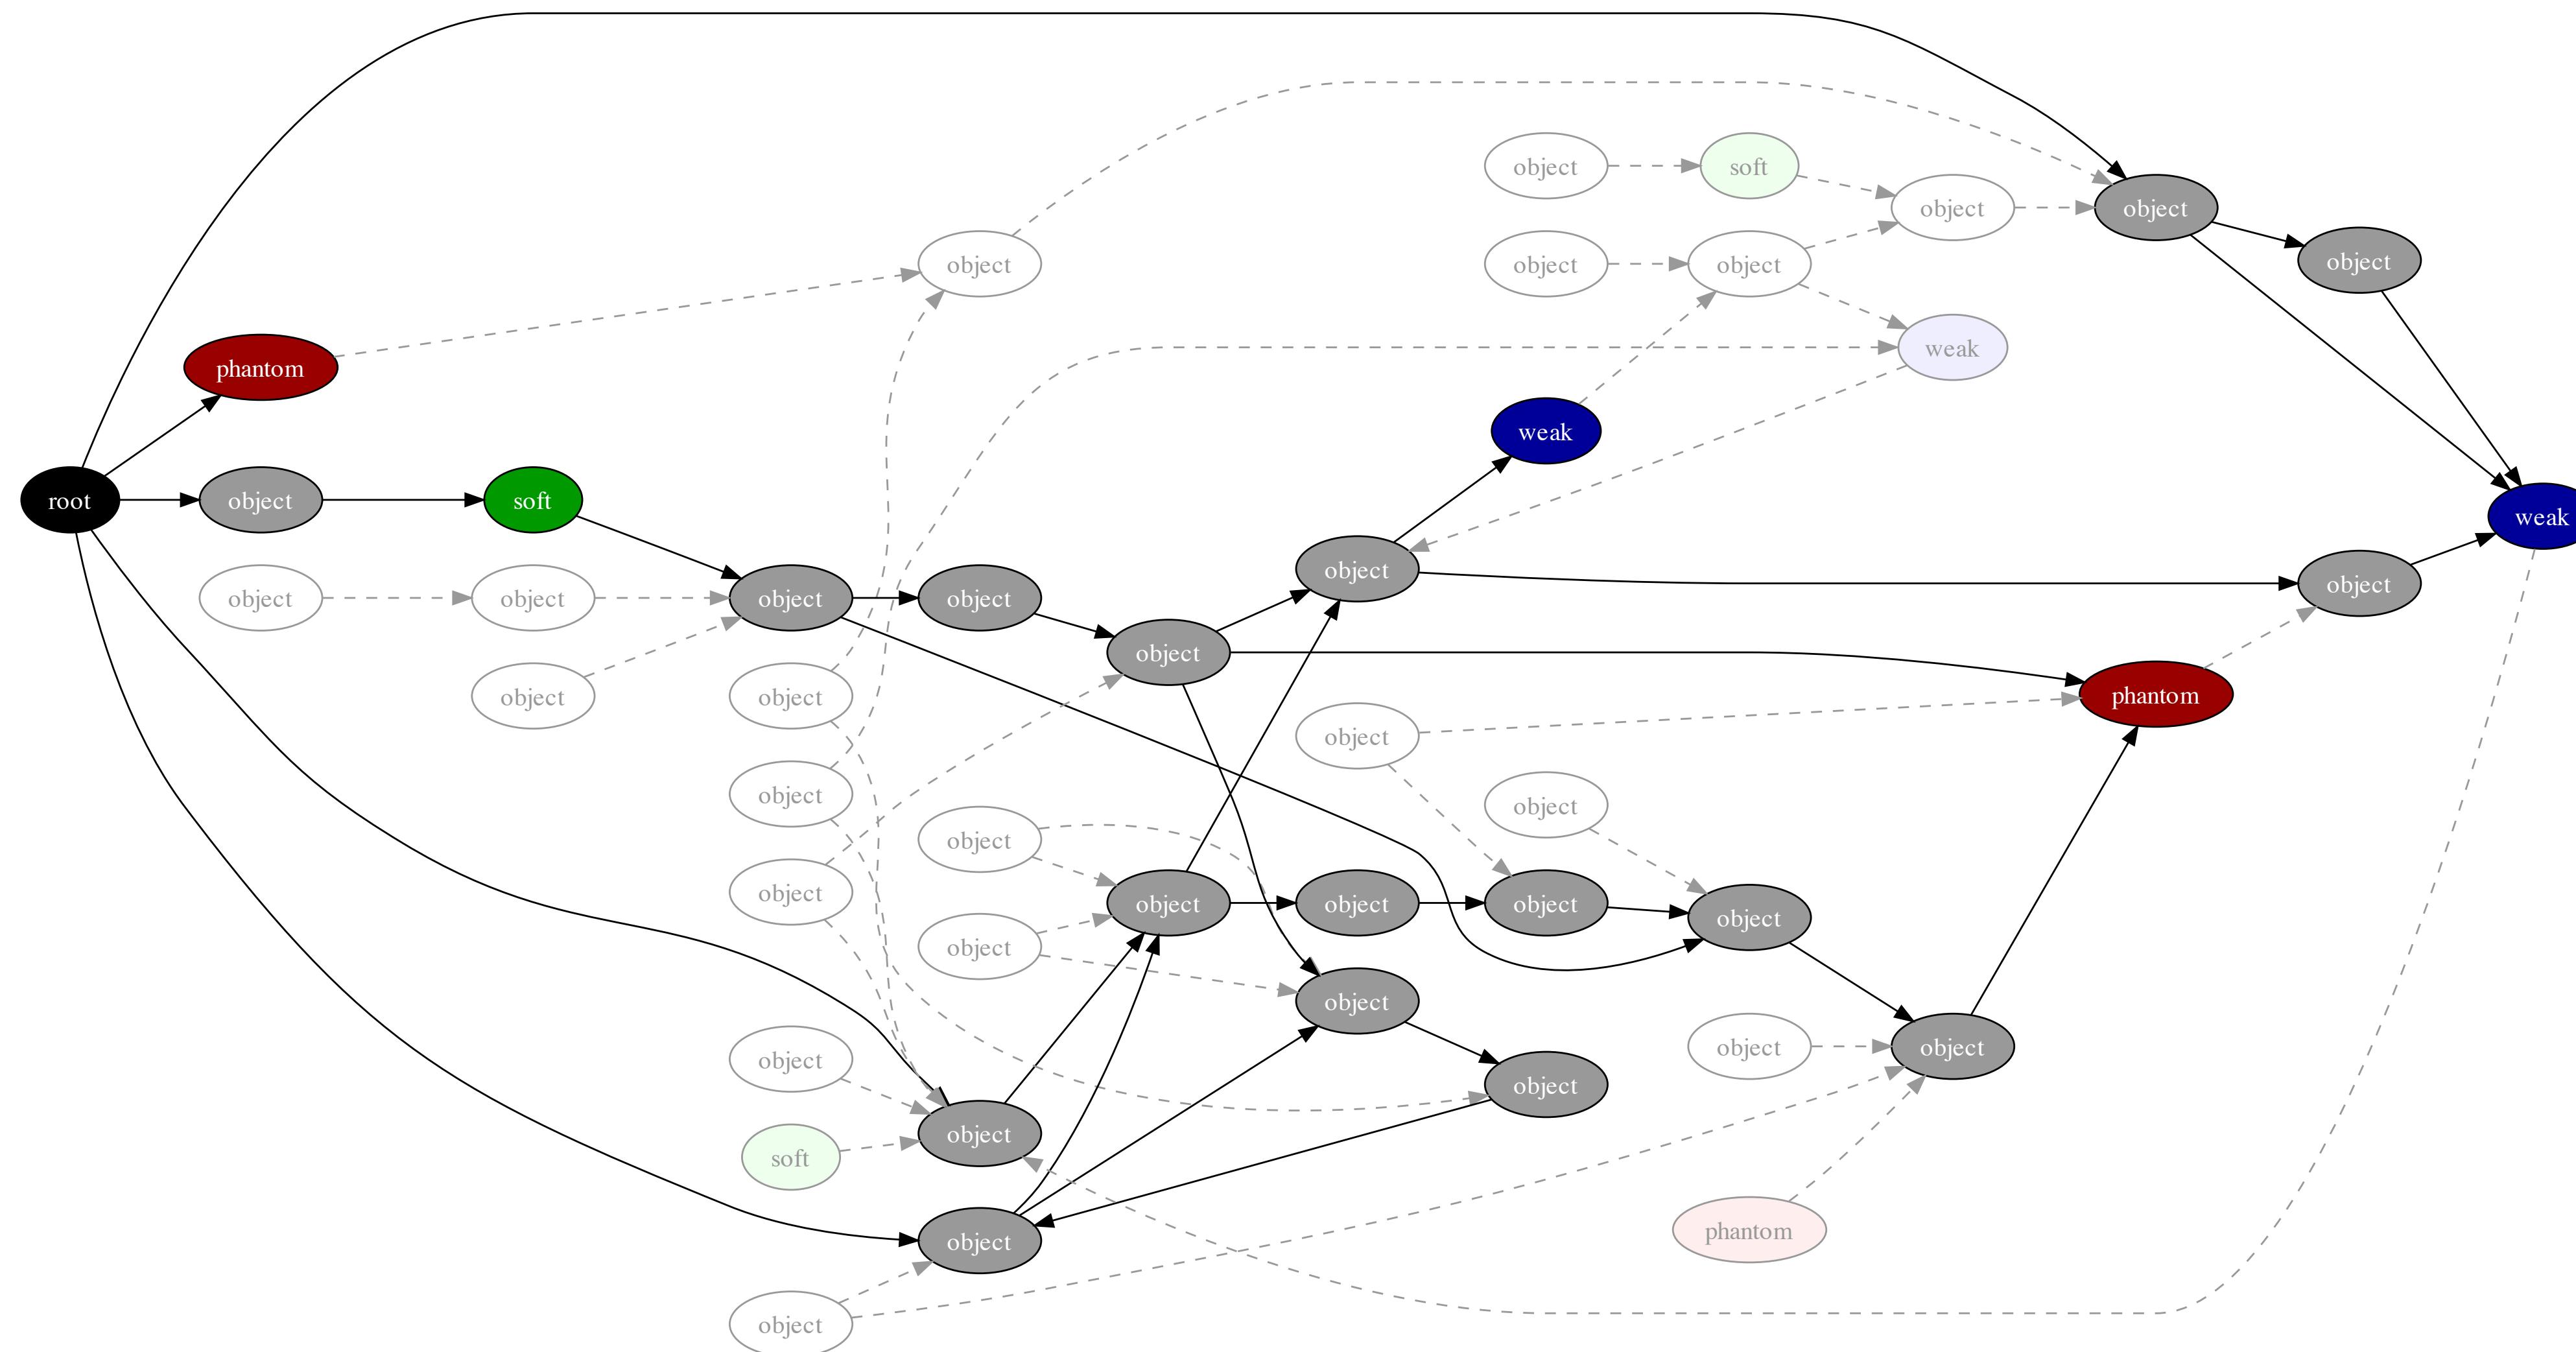
4. Trace and mark softly-referenced objects.



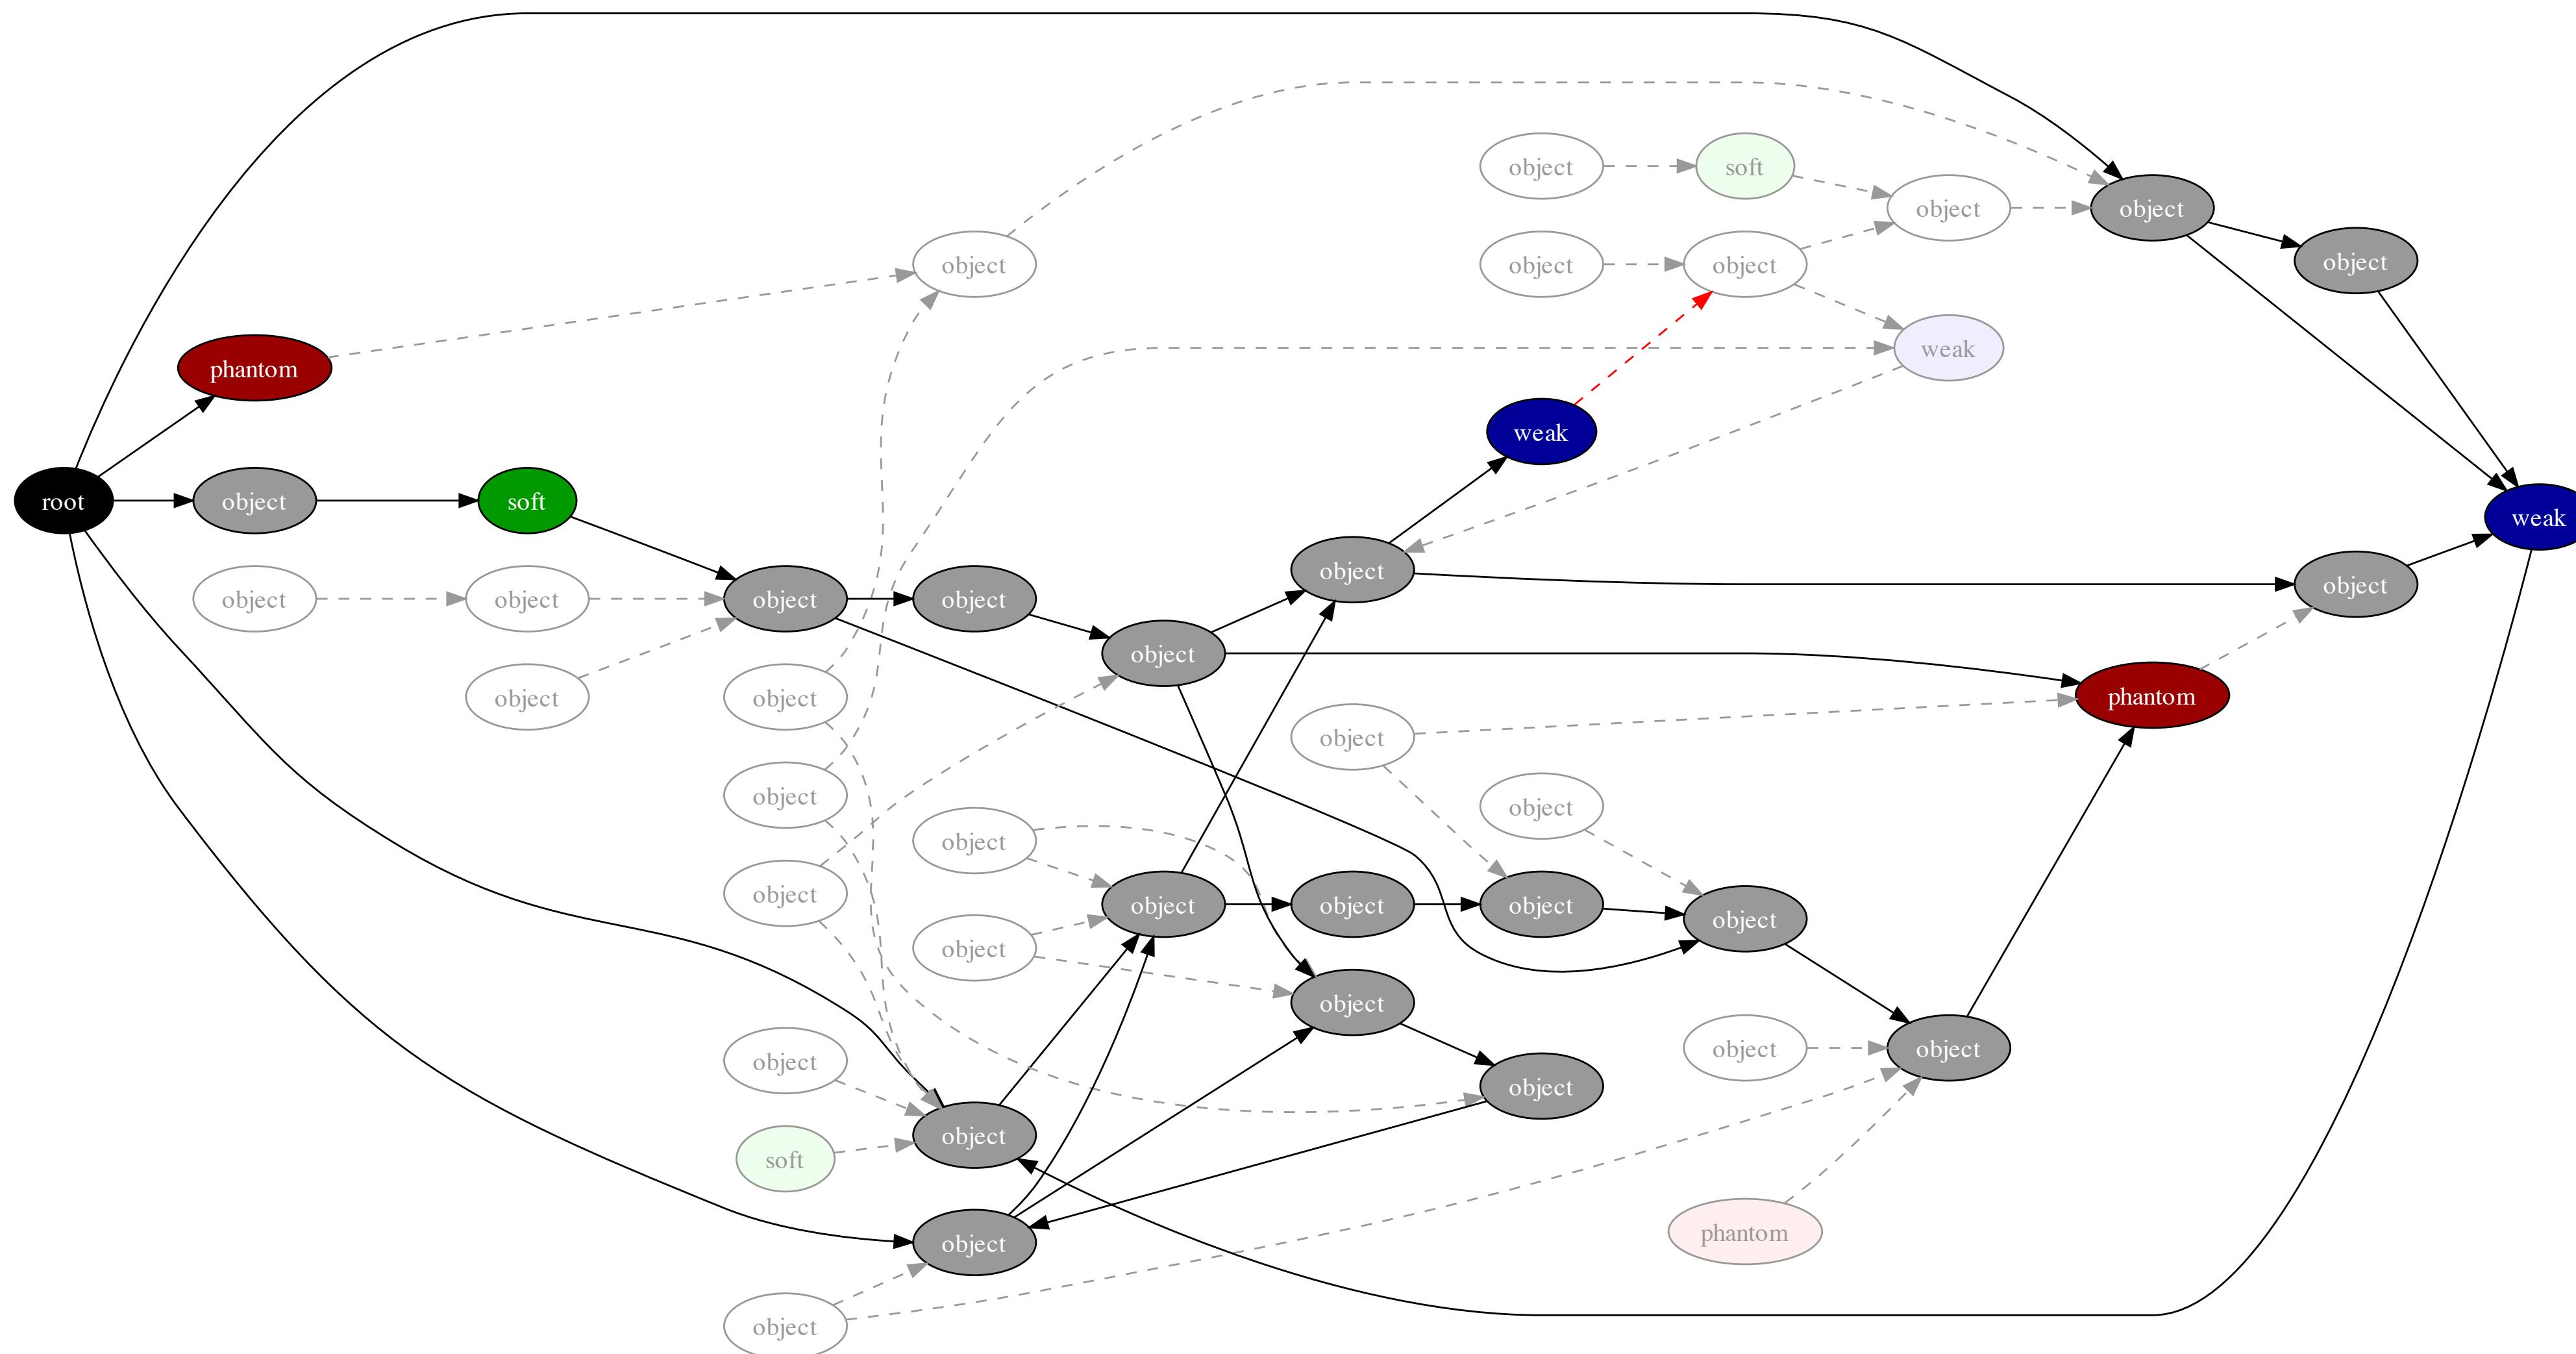
4. Trace and mark softly-referenced objects.



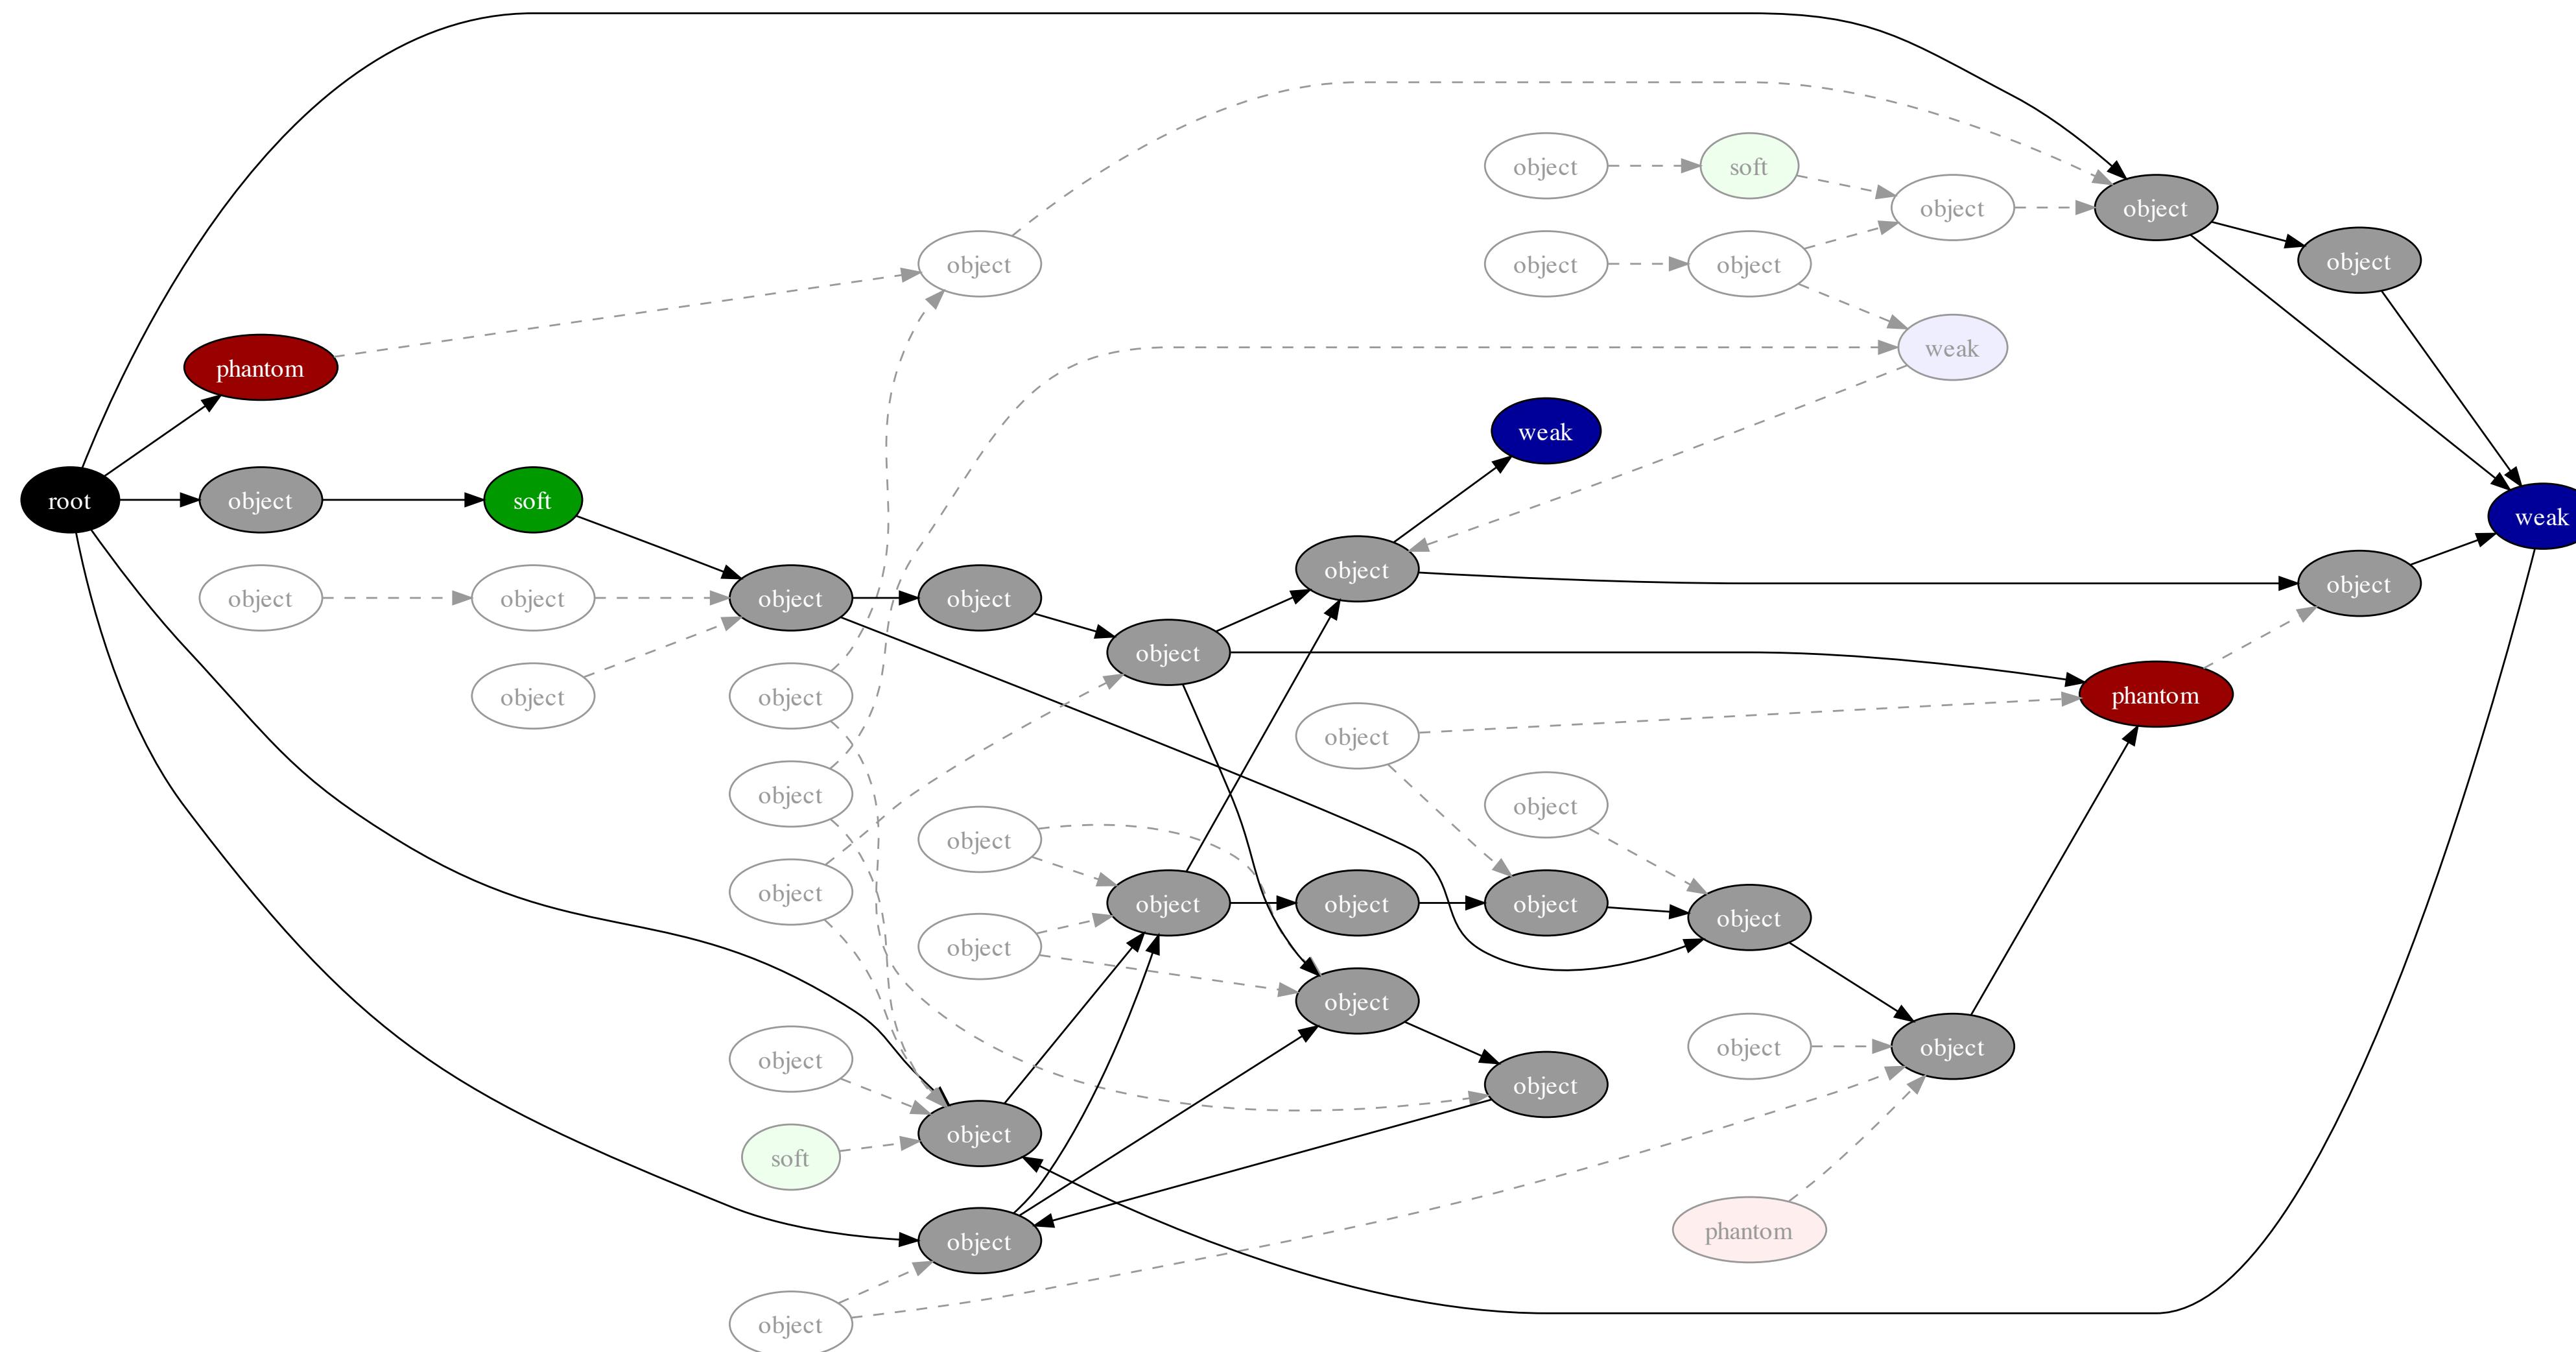
5. Clear weak references.



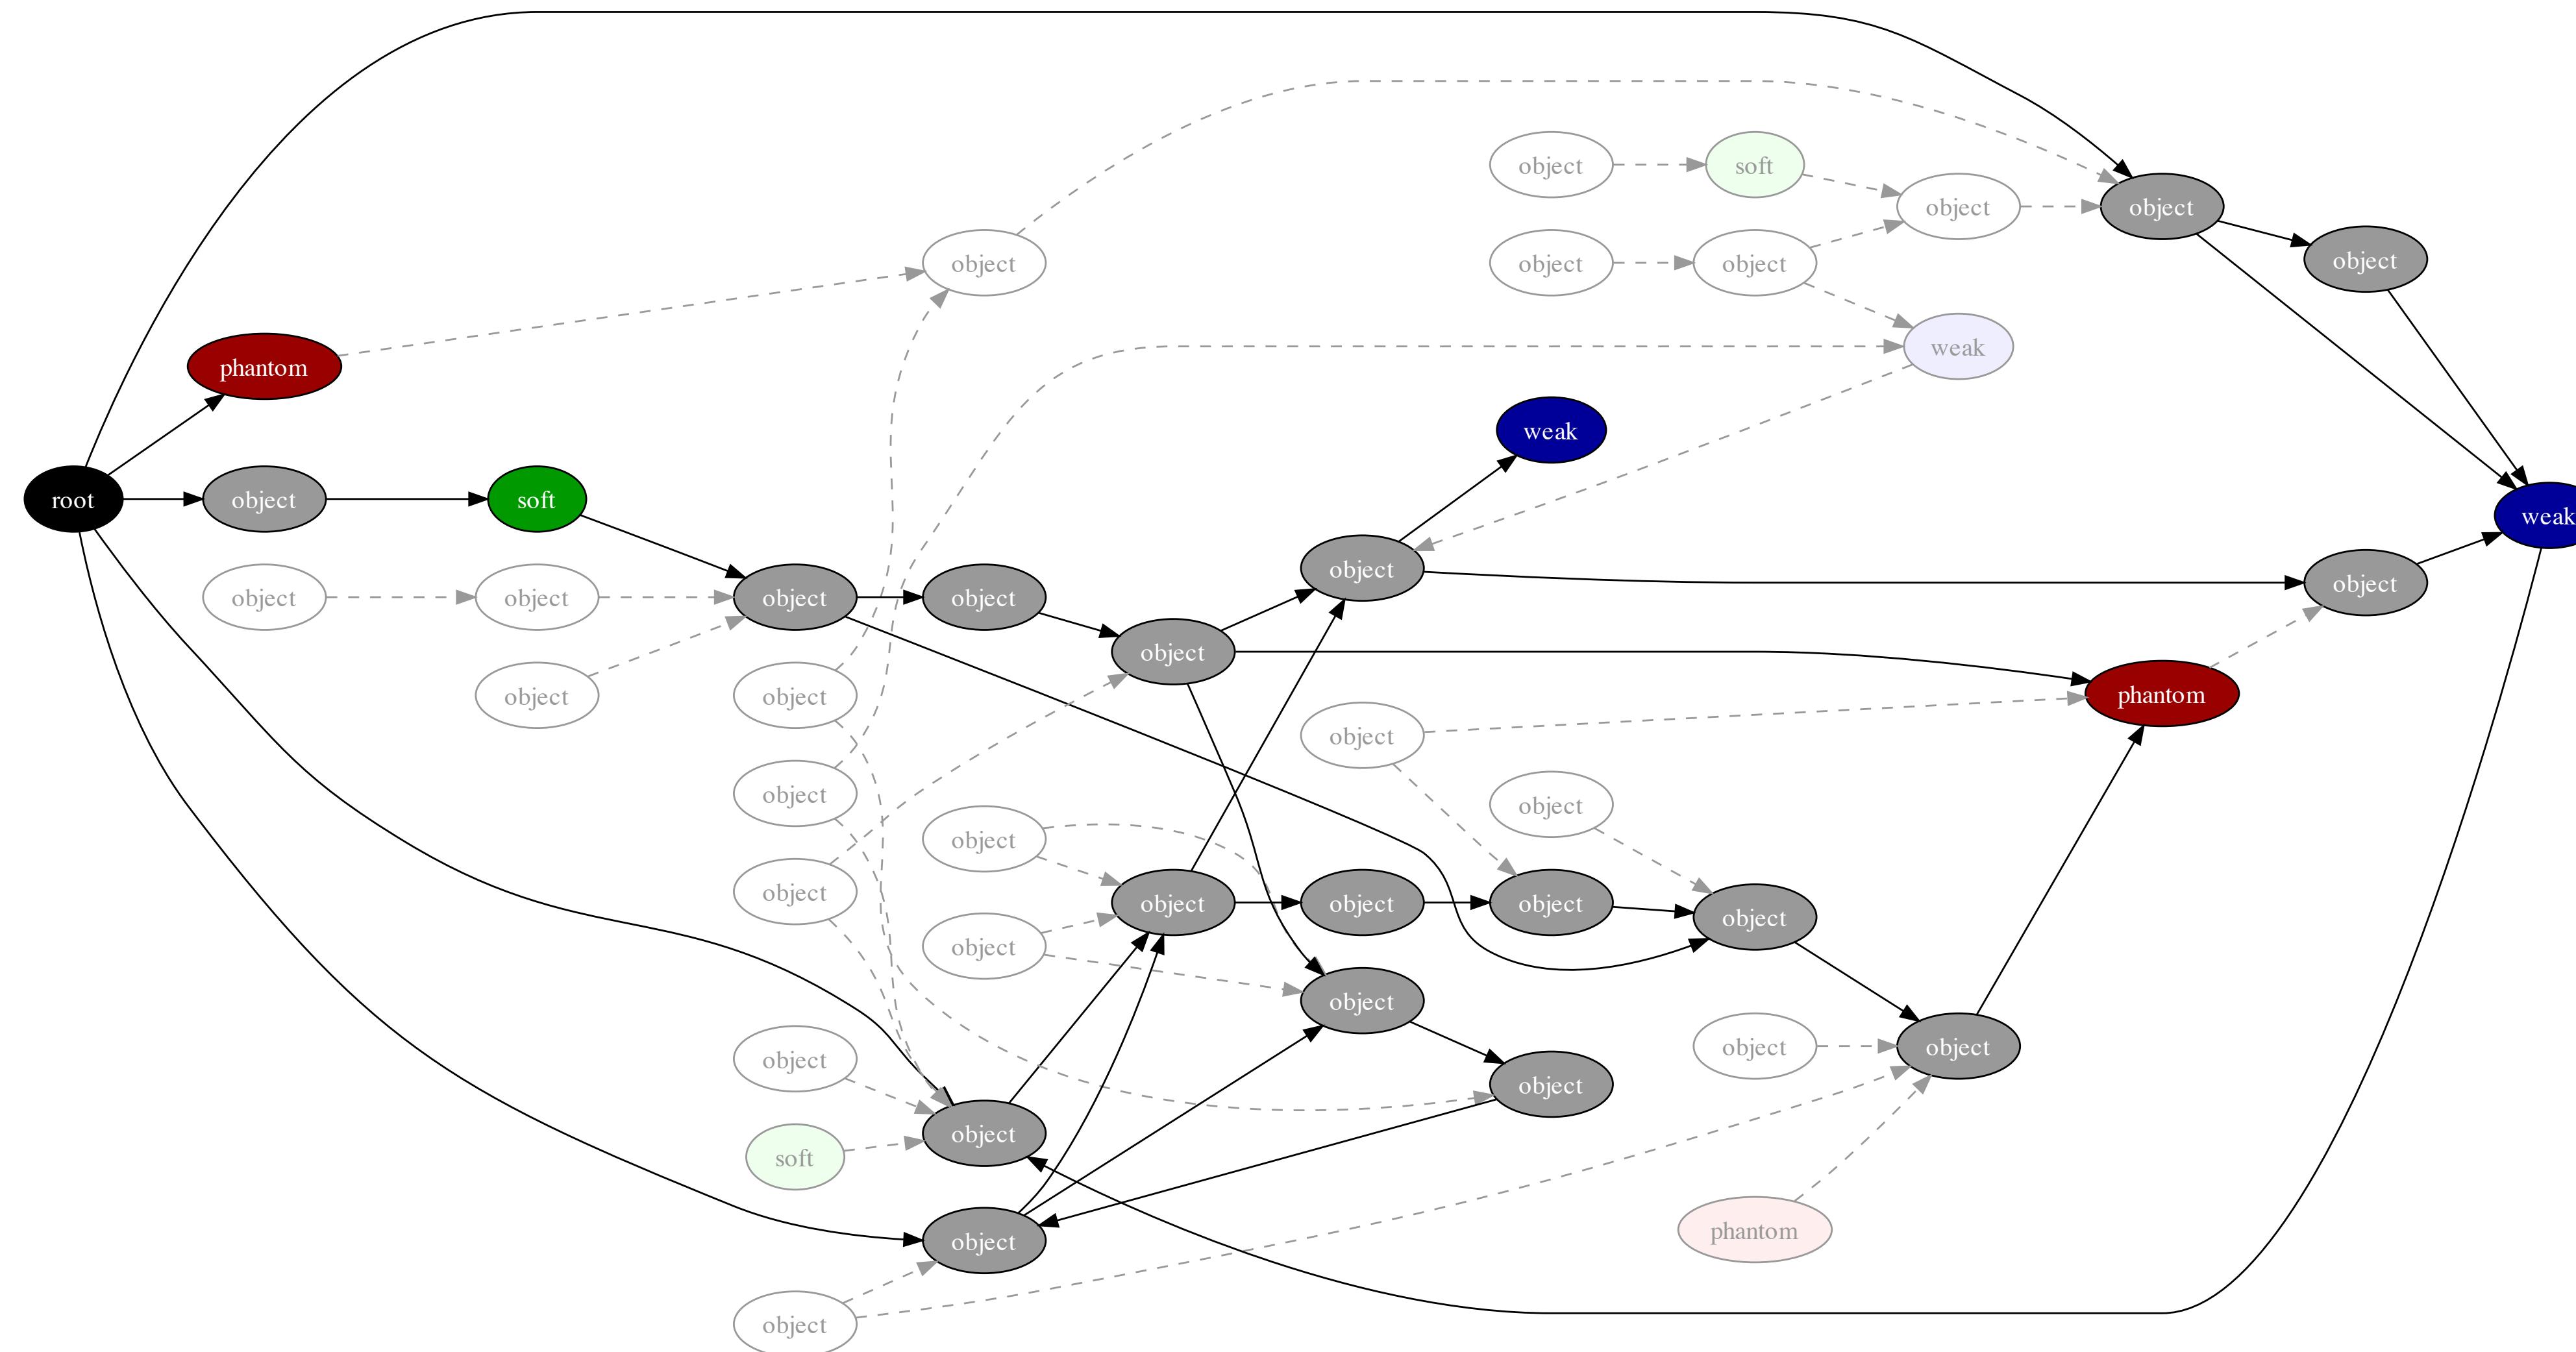
5. Clear weak references.



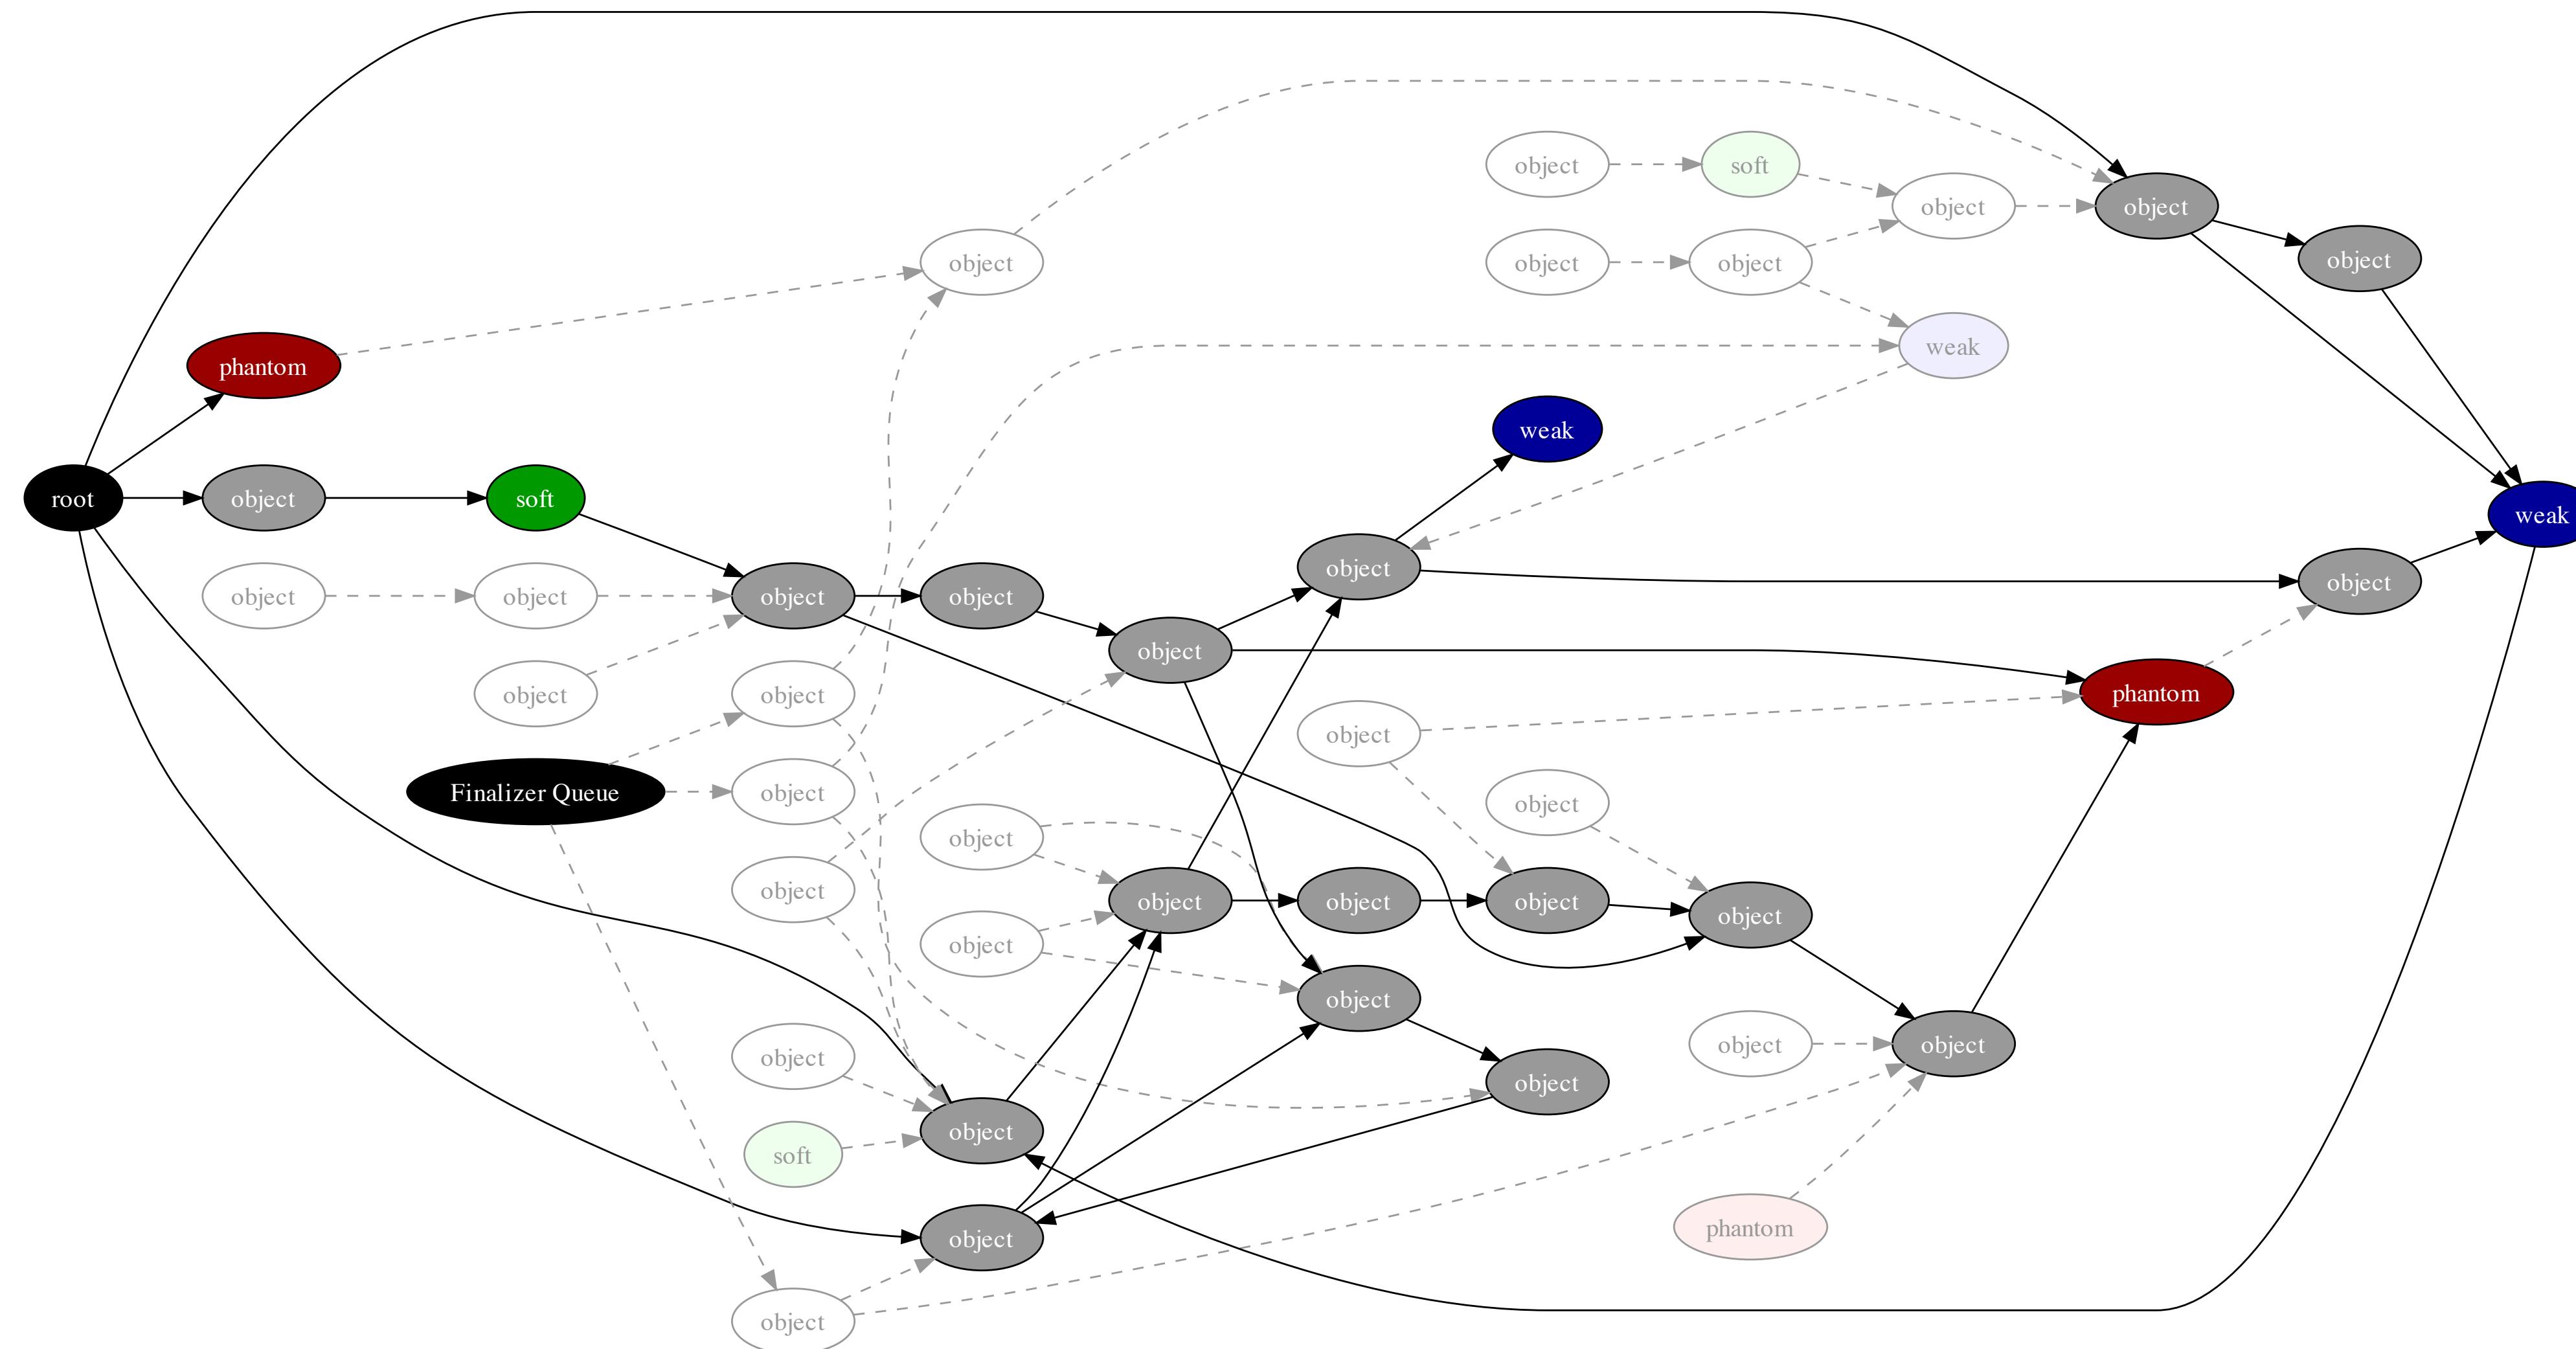
5. Clear weak references.



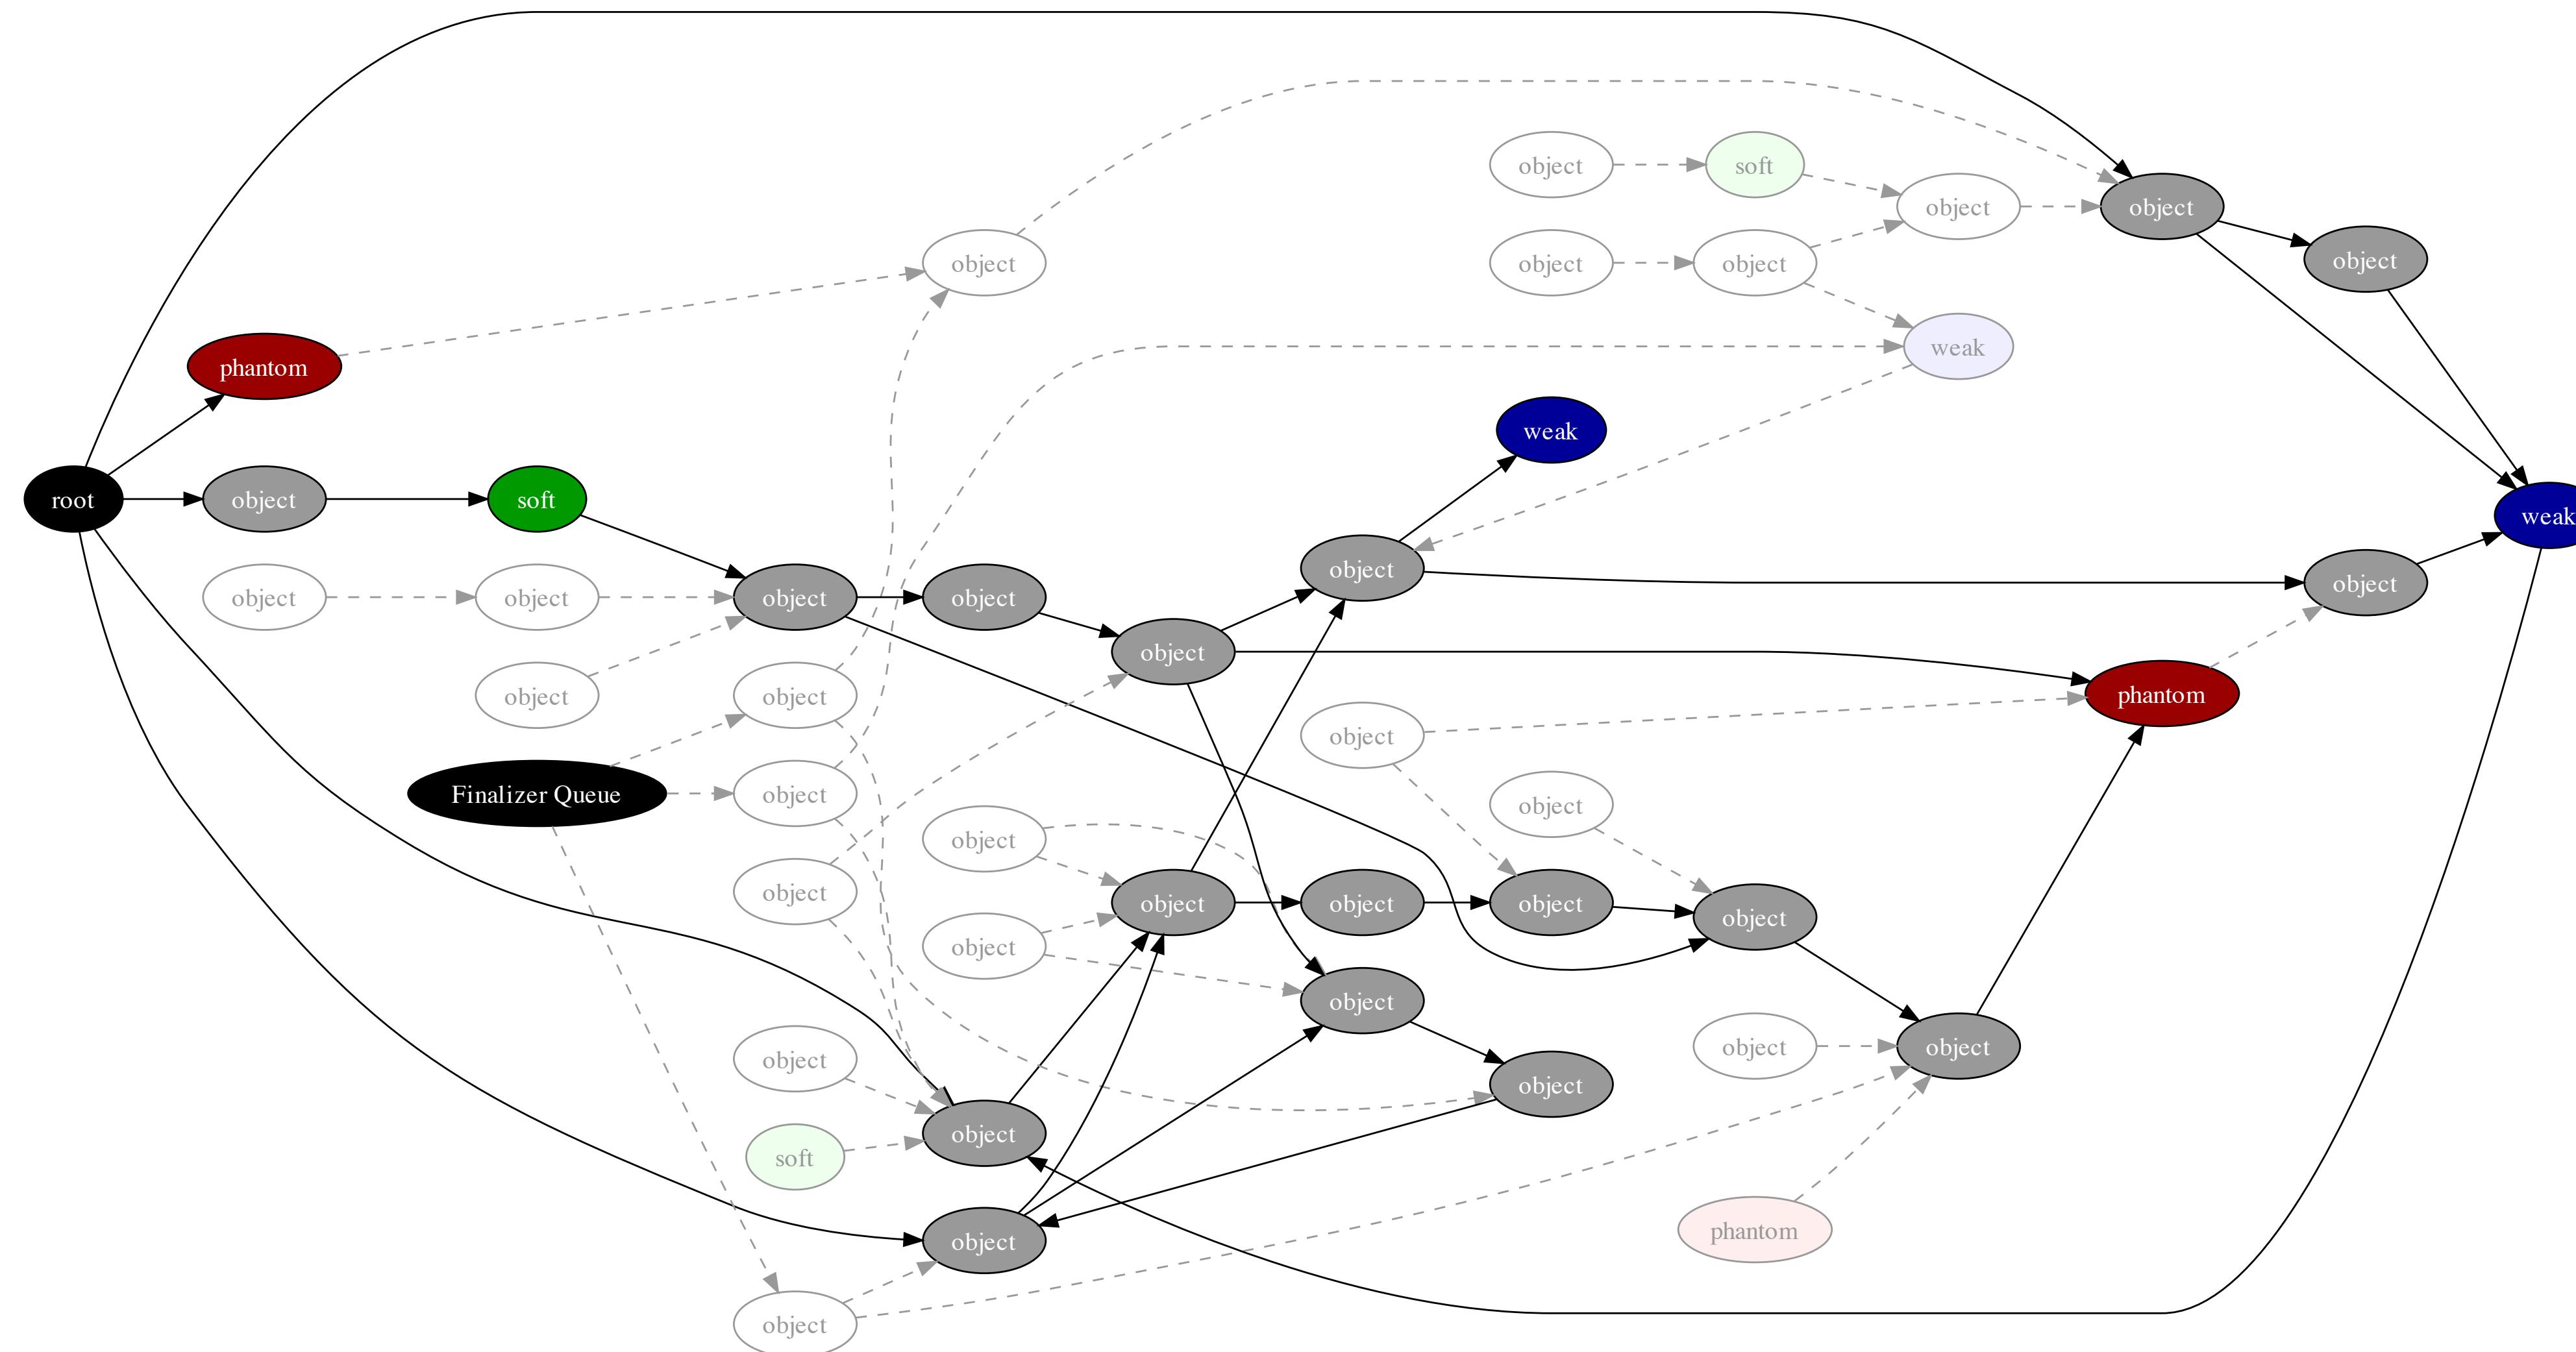
6. Enqueue finalizable objects.



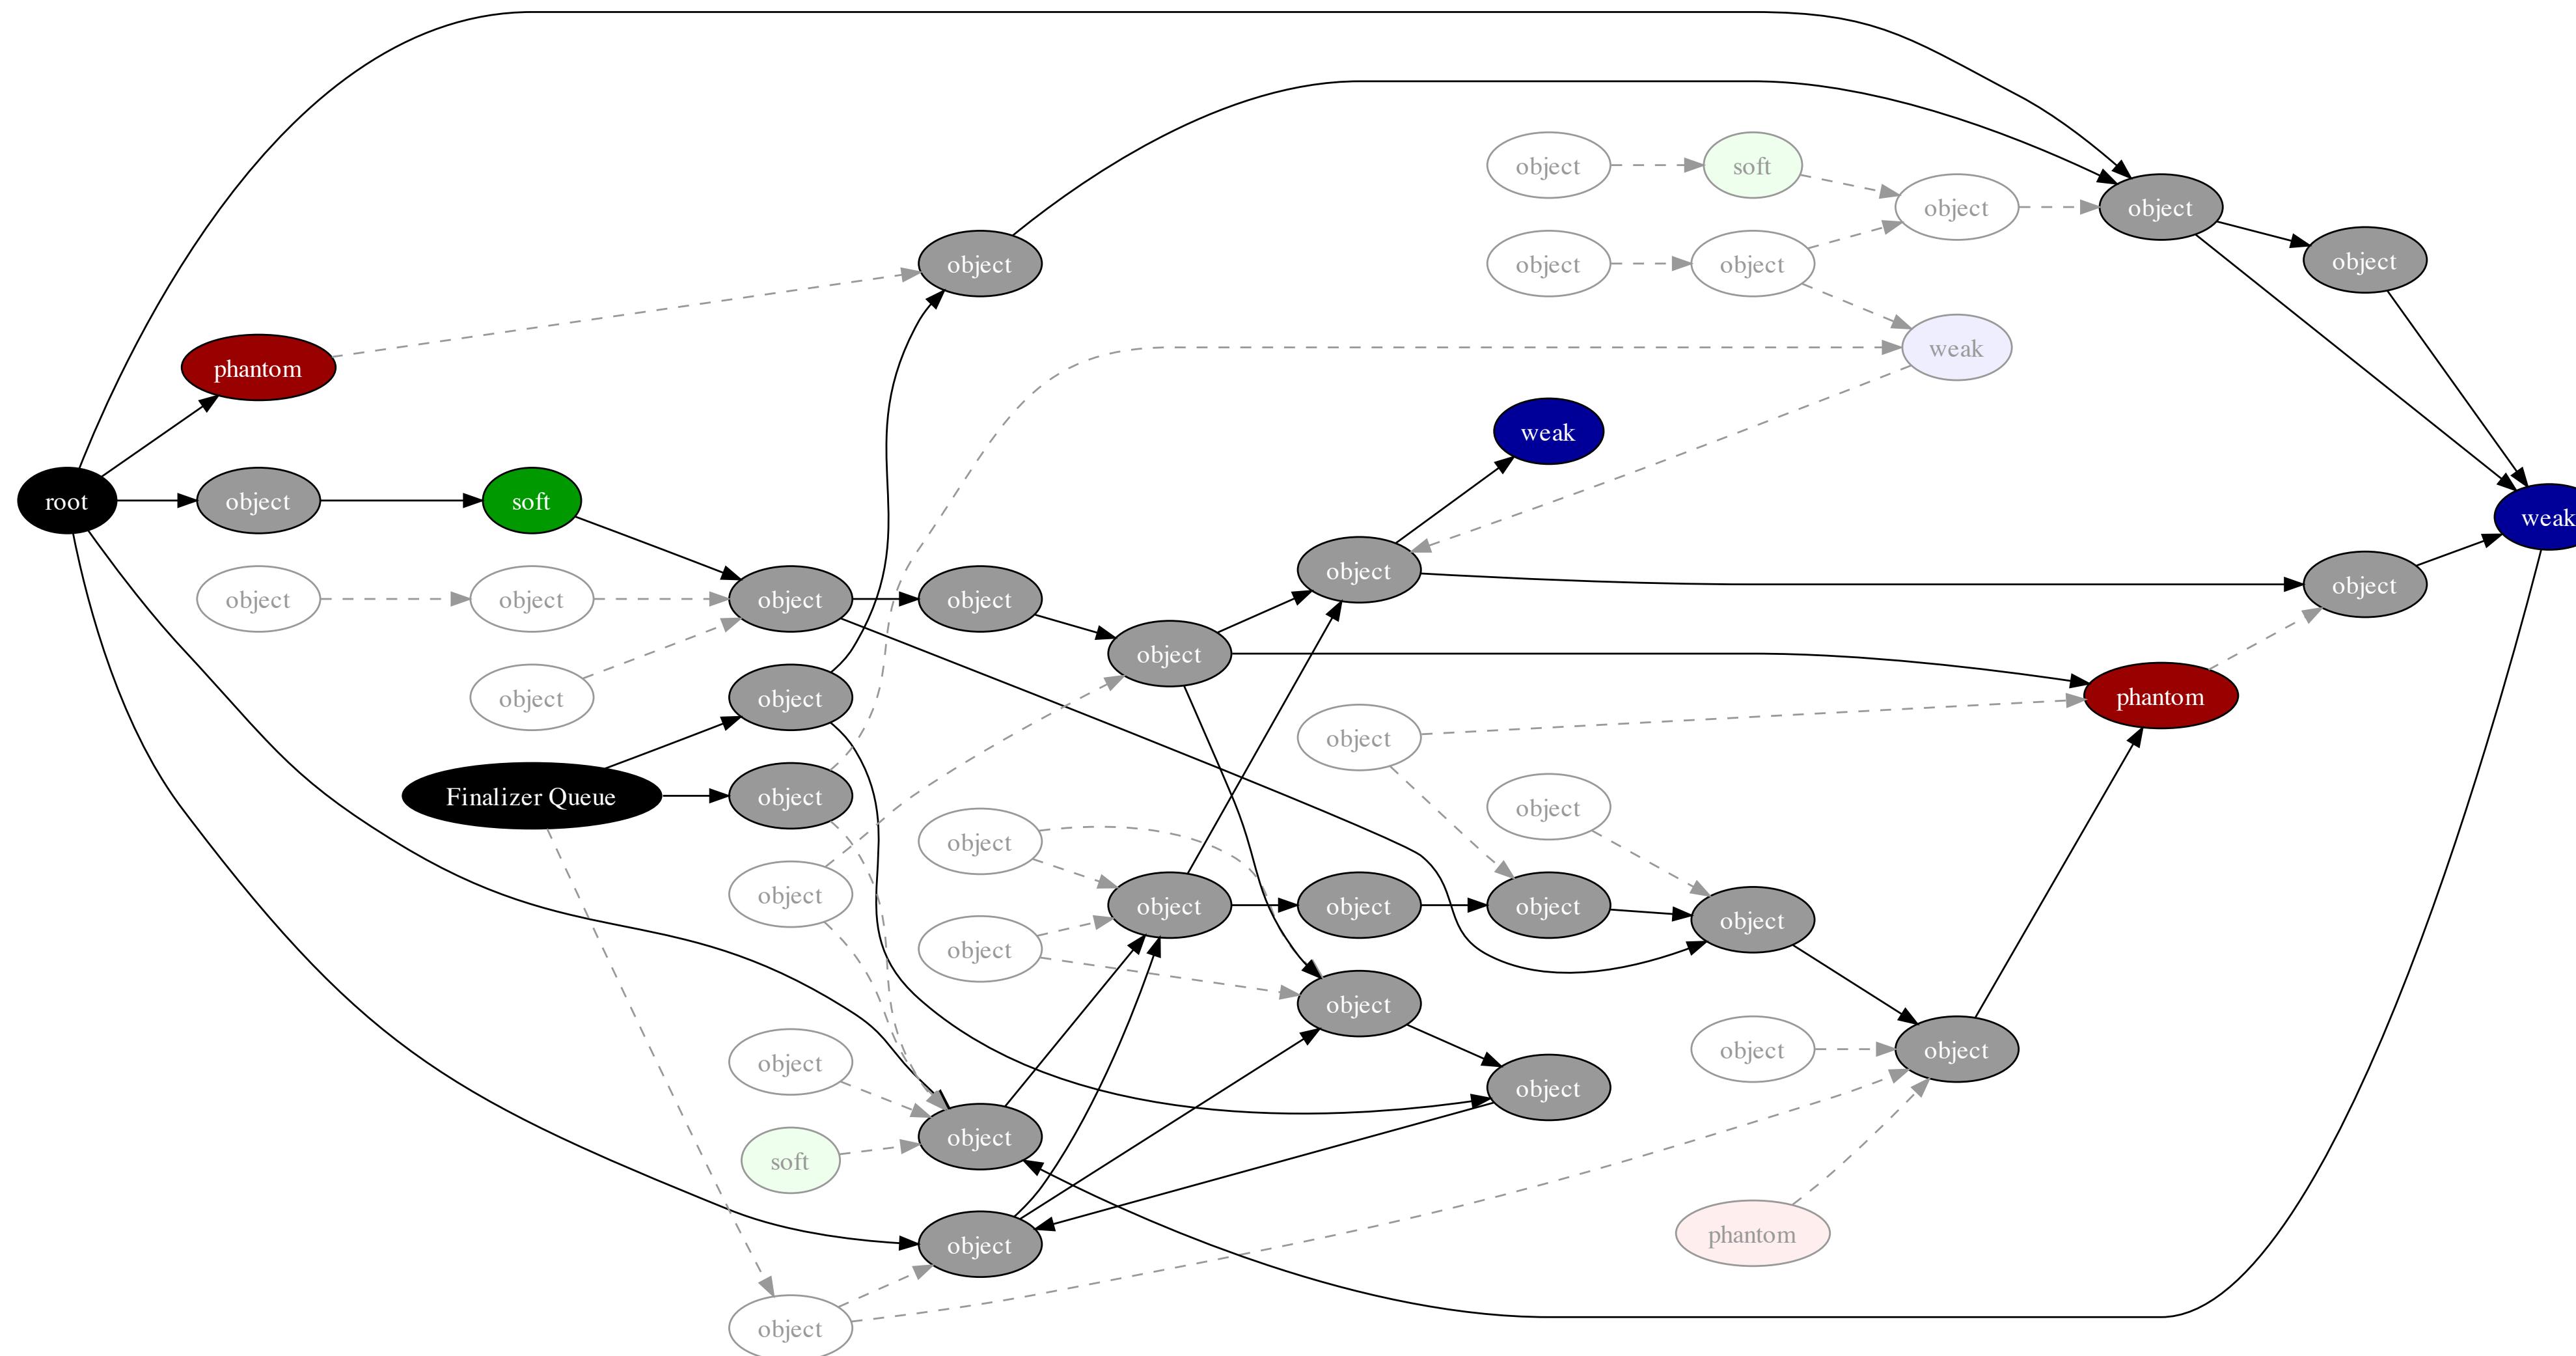
6. Enqueue finalizable objects.



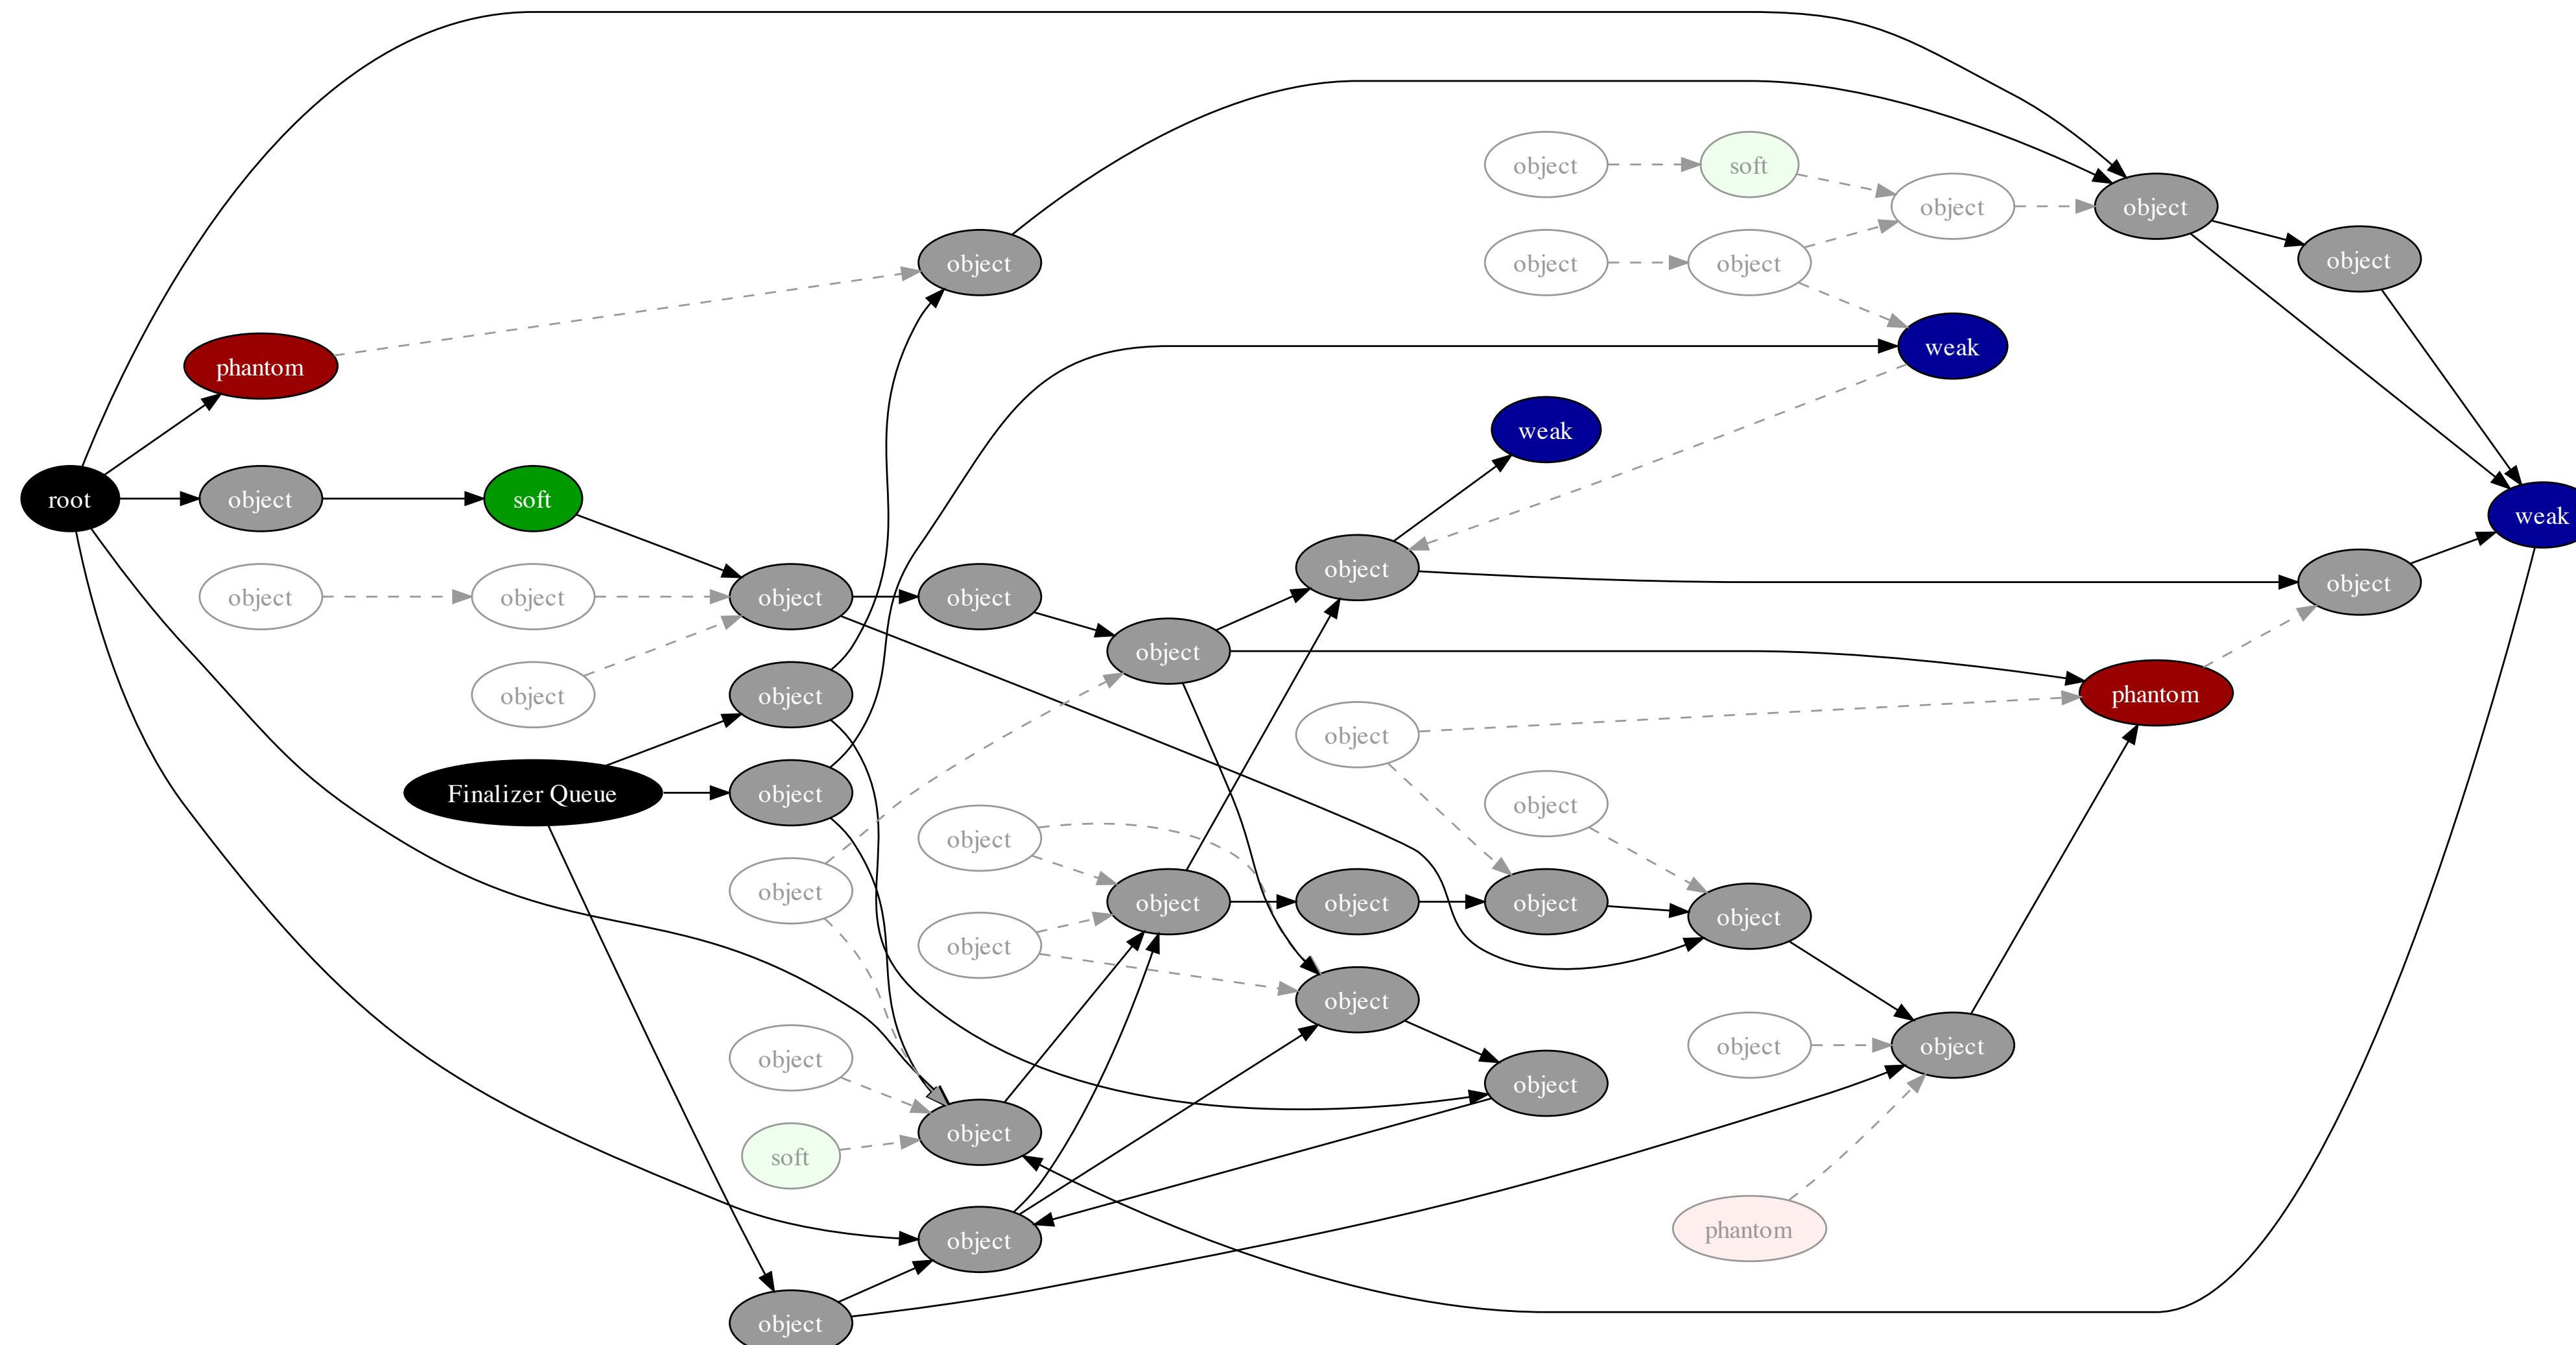
7. Repeat steps 1 through 5 for the queue.



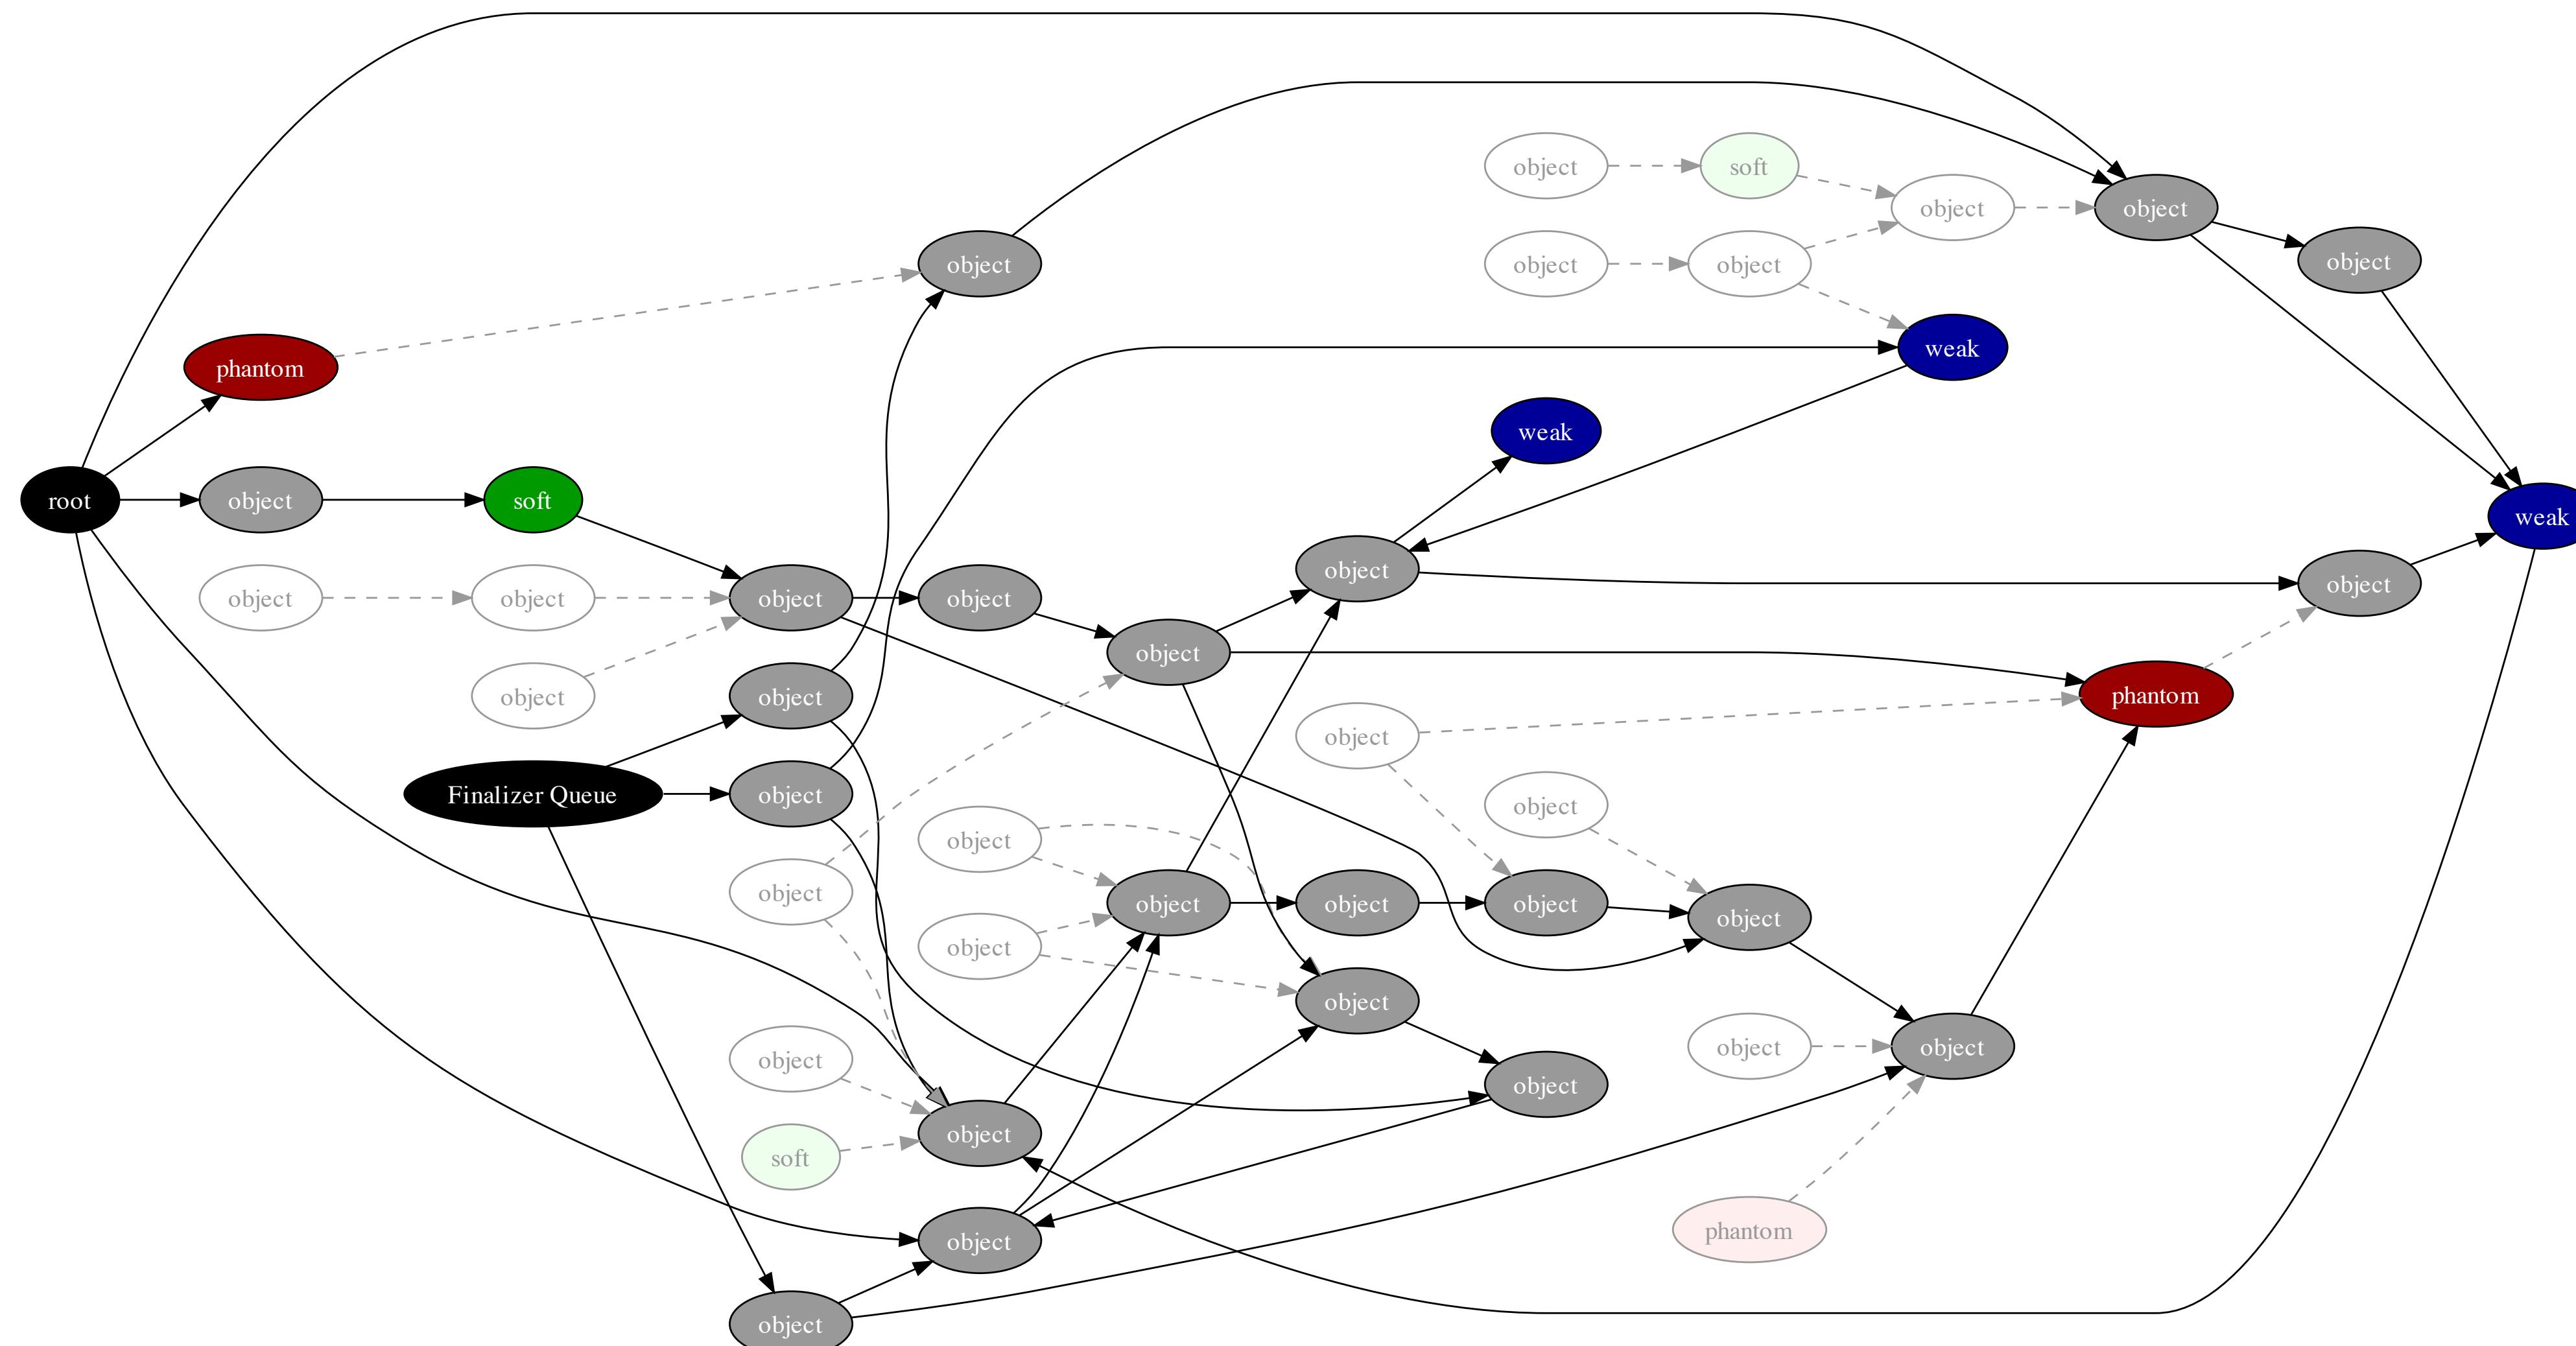
7. Repeat steps 1 through 5 for the queue.



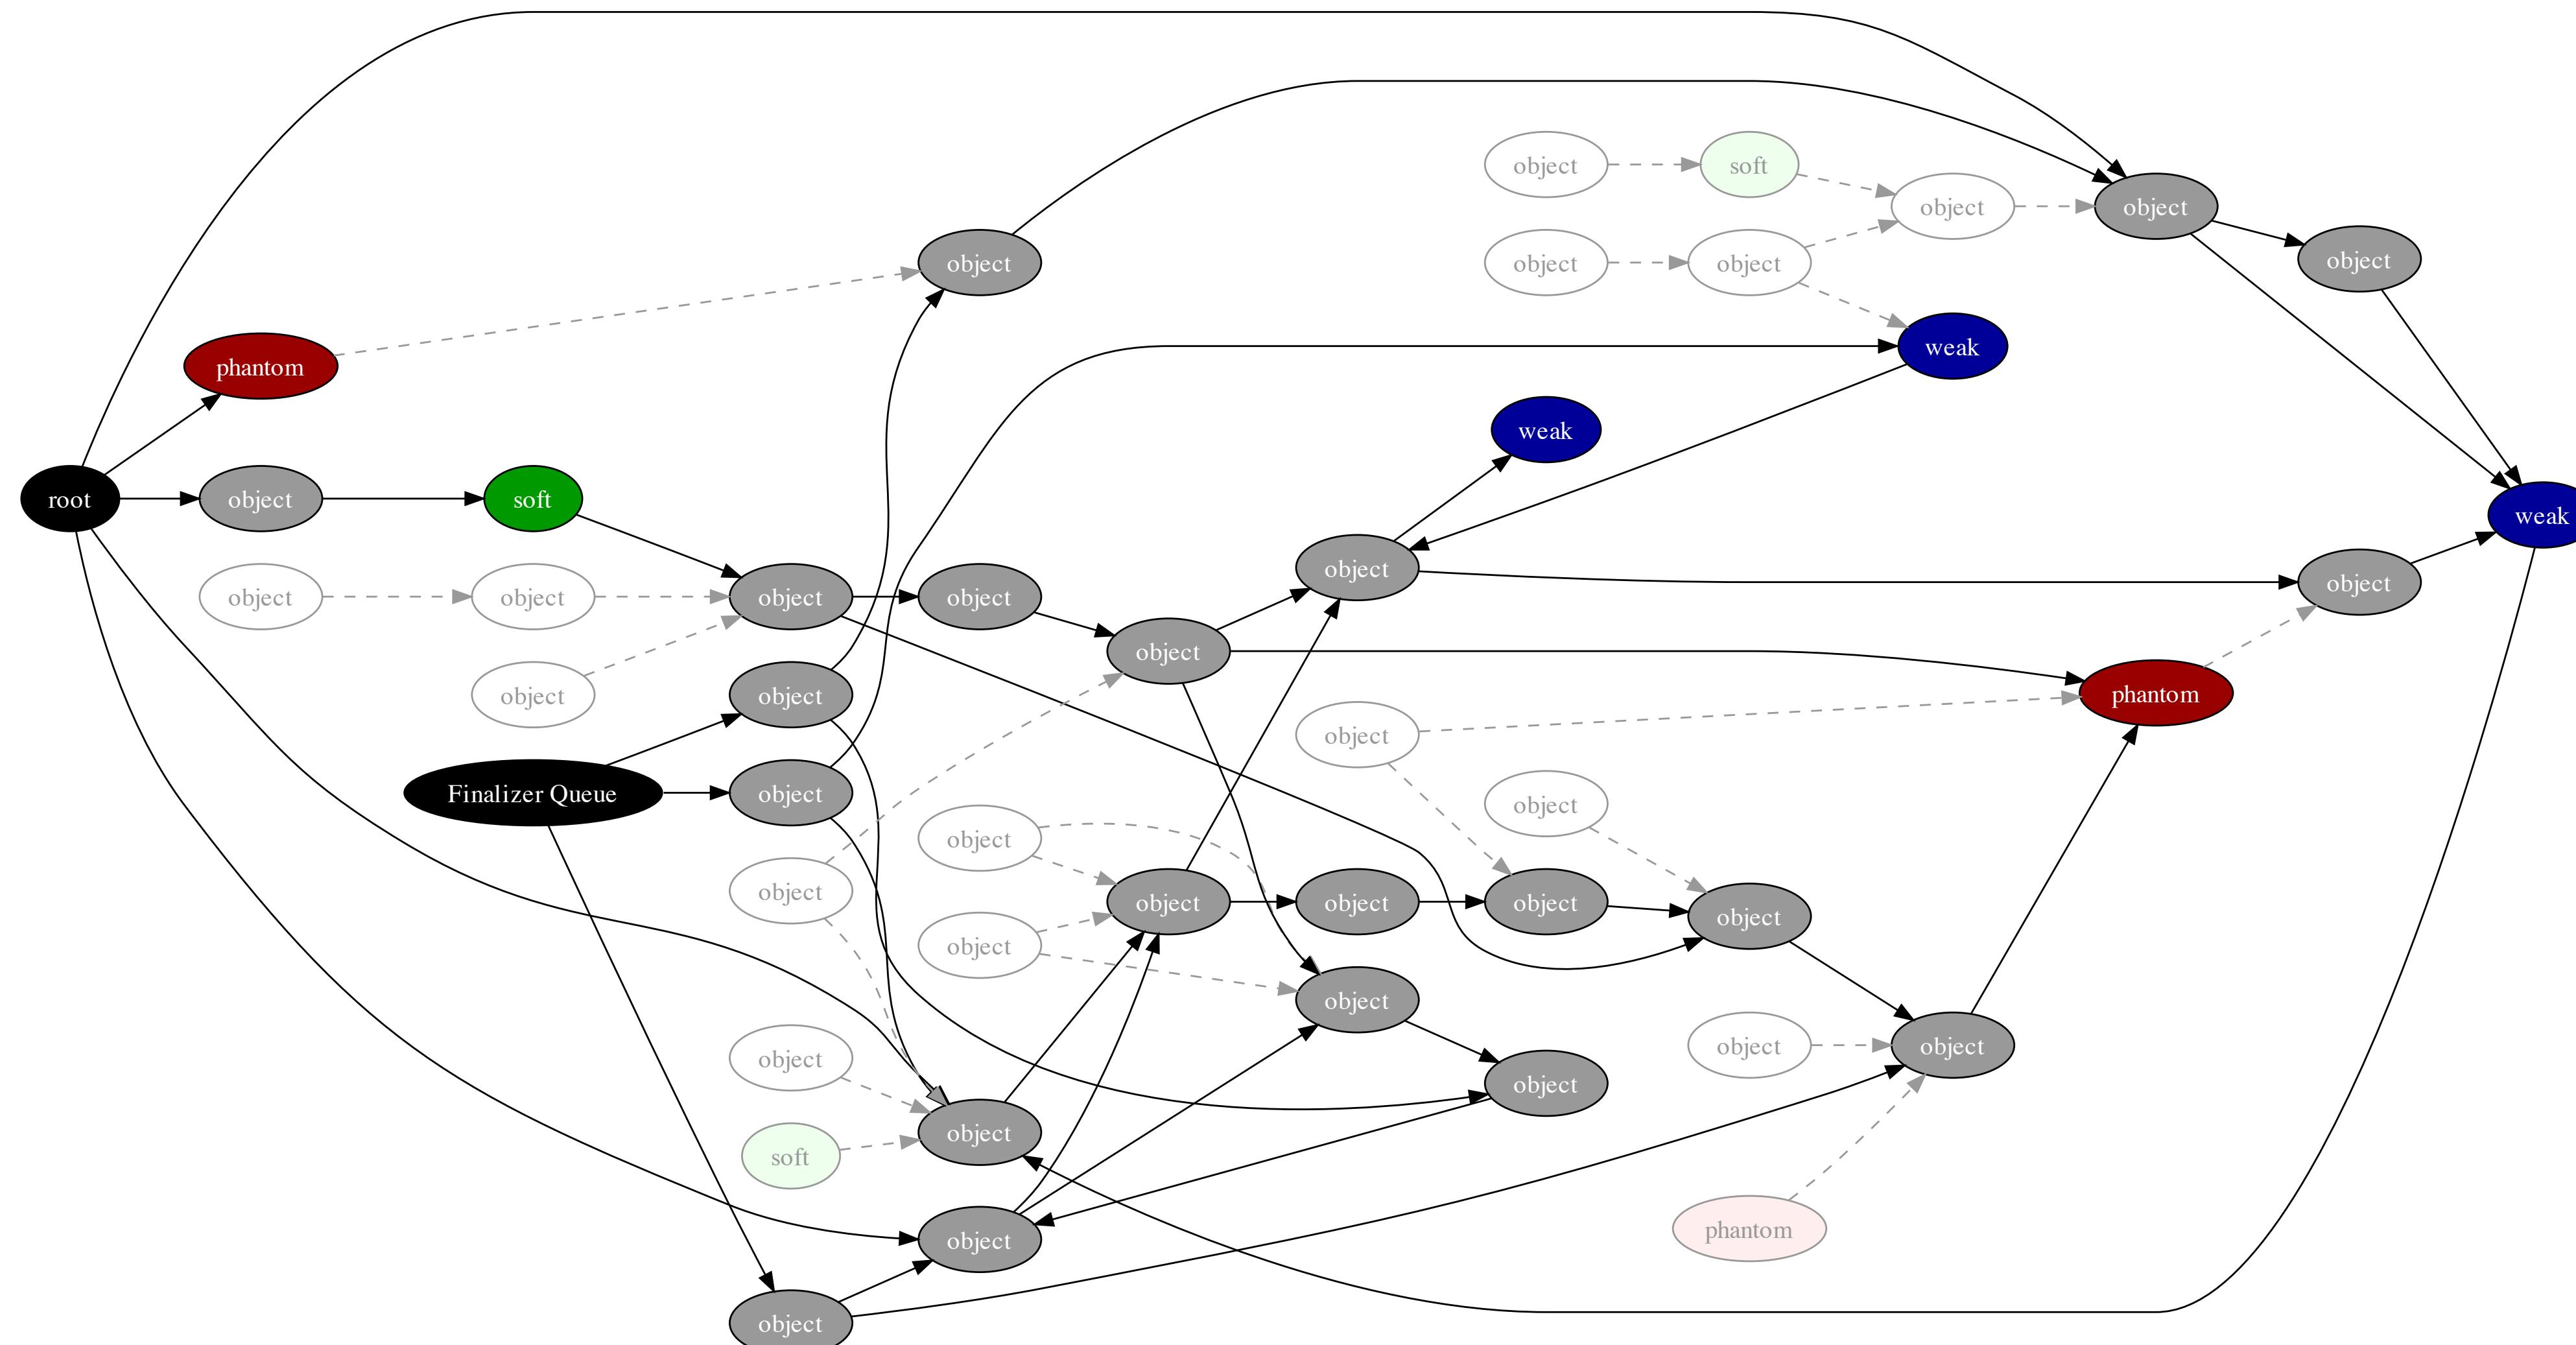
7. Repeat steps 1 through 5 for the queue.



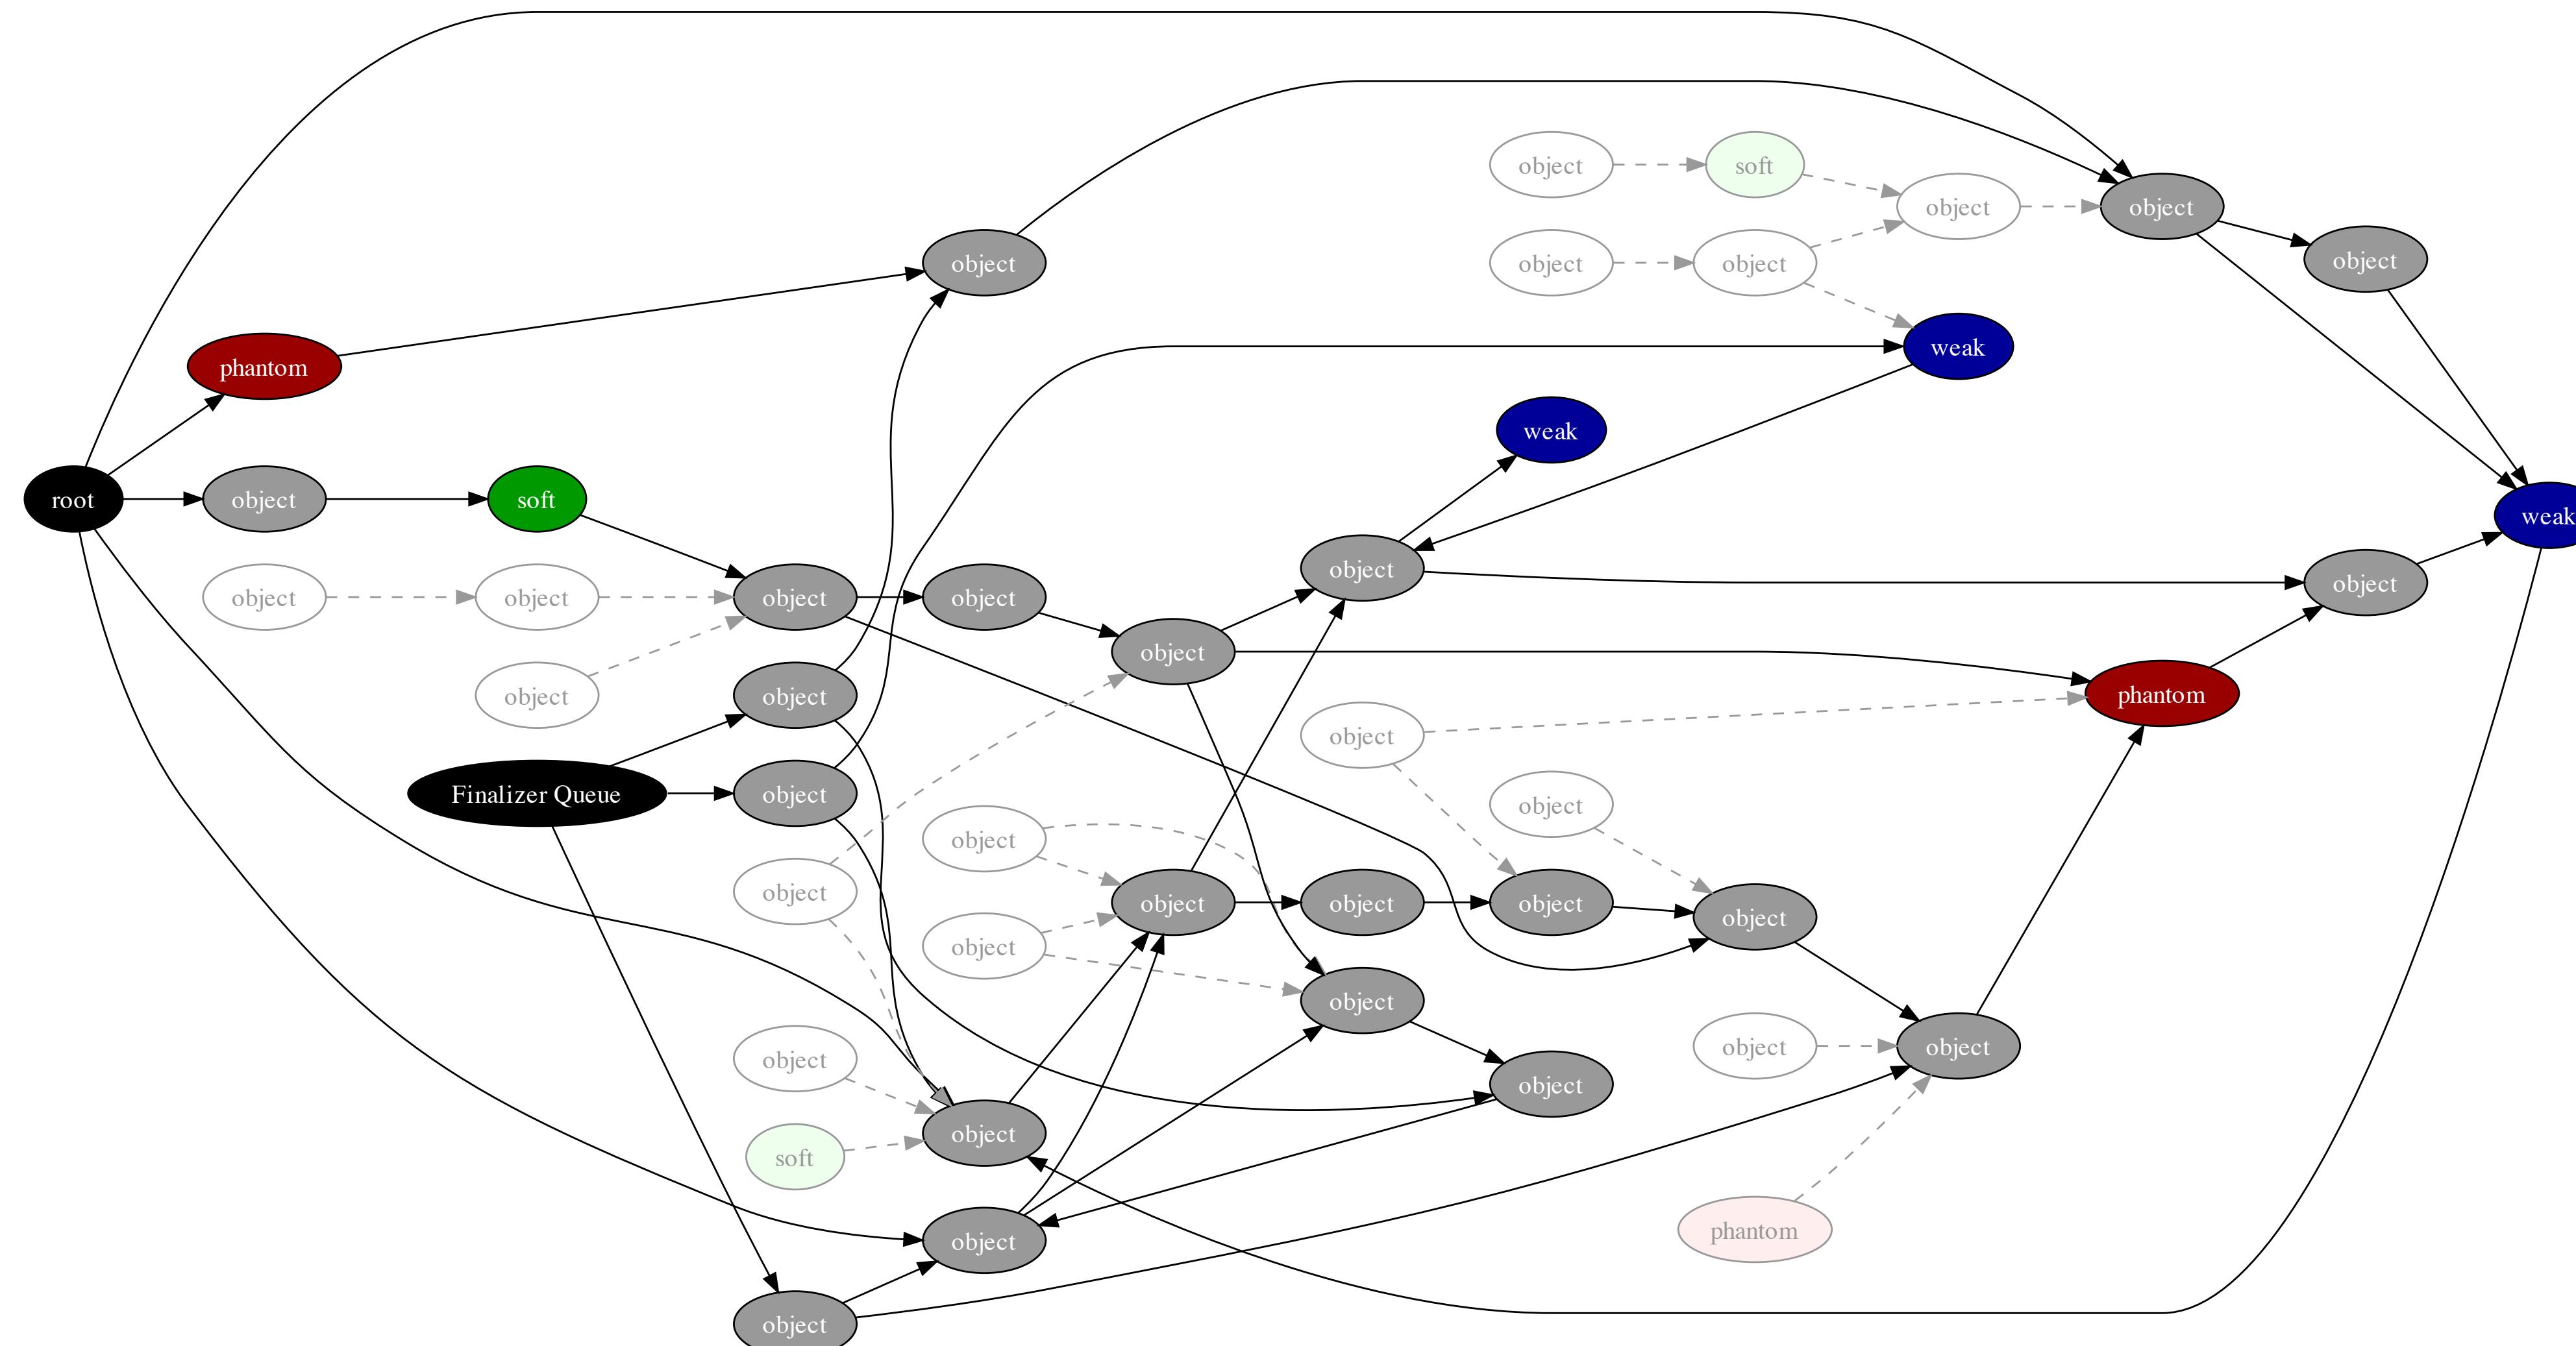
7. Repeat steps 1 through 5 for the queue.



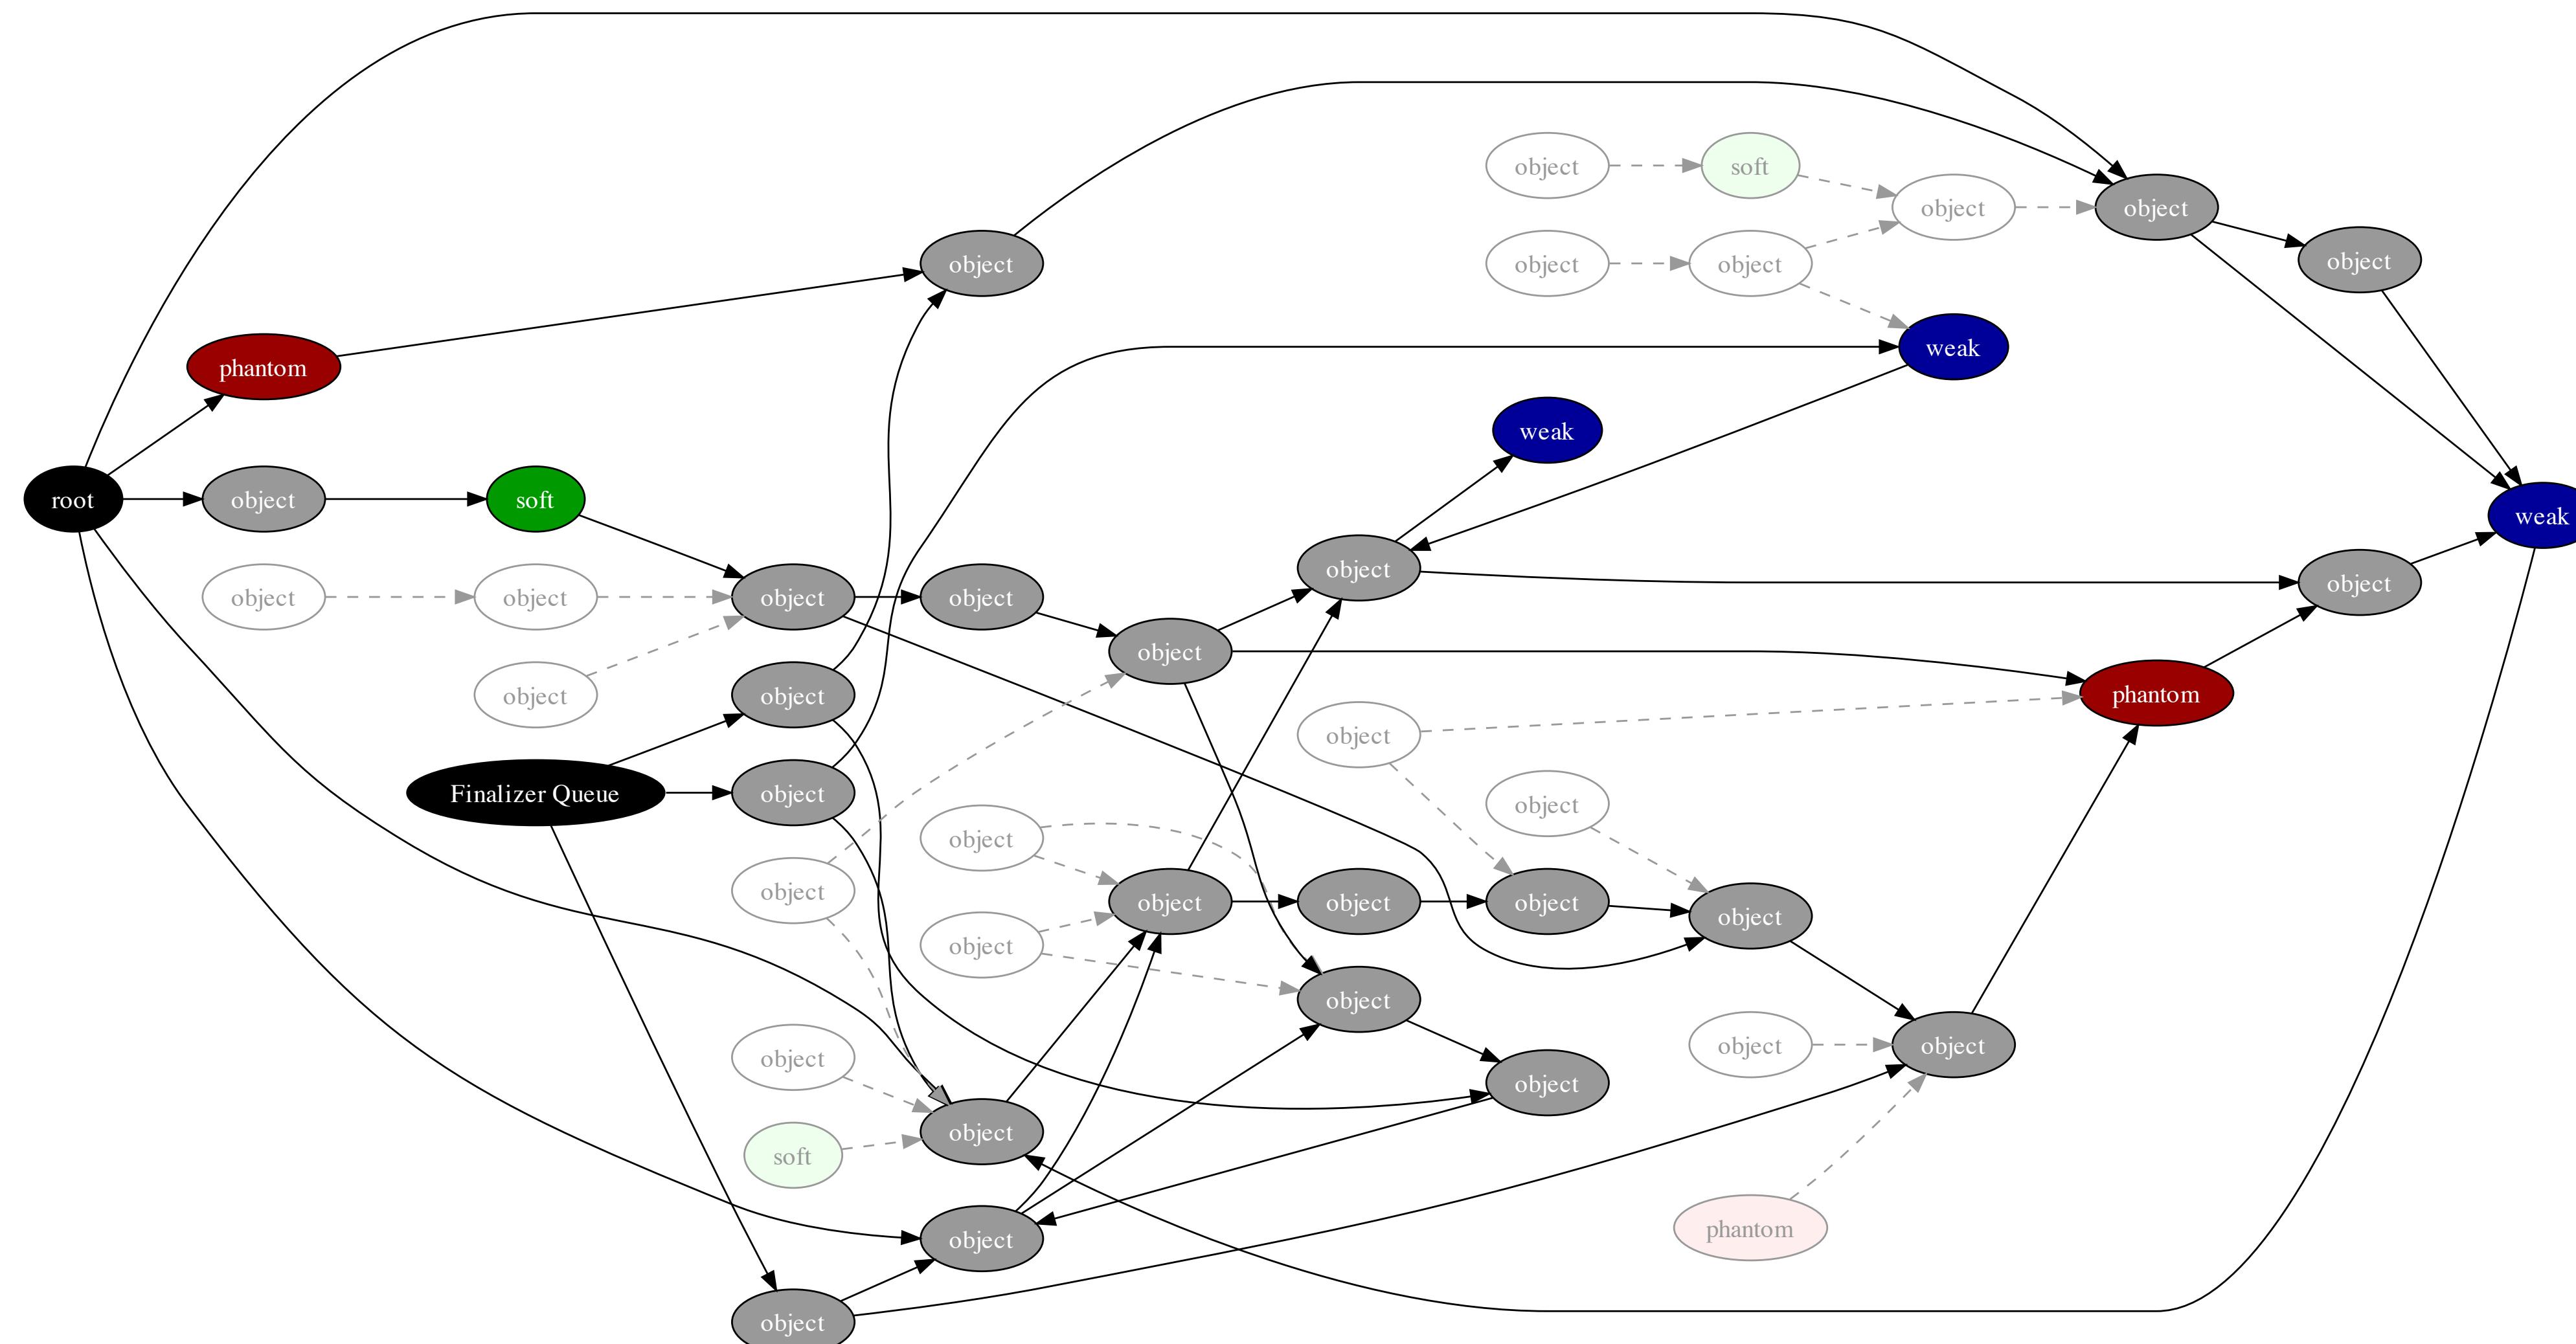
8. Possibly enqueue phantom references.



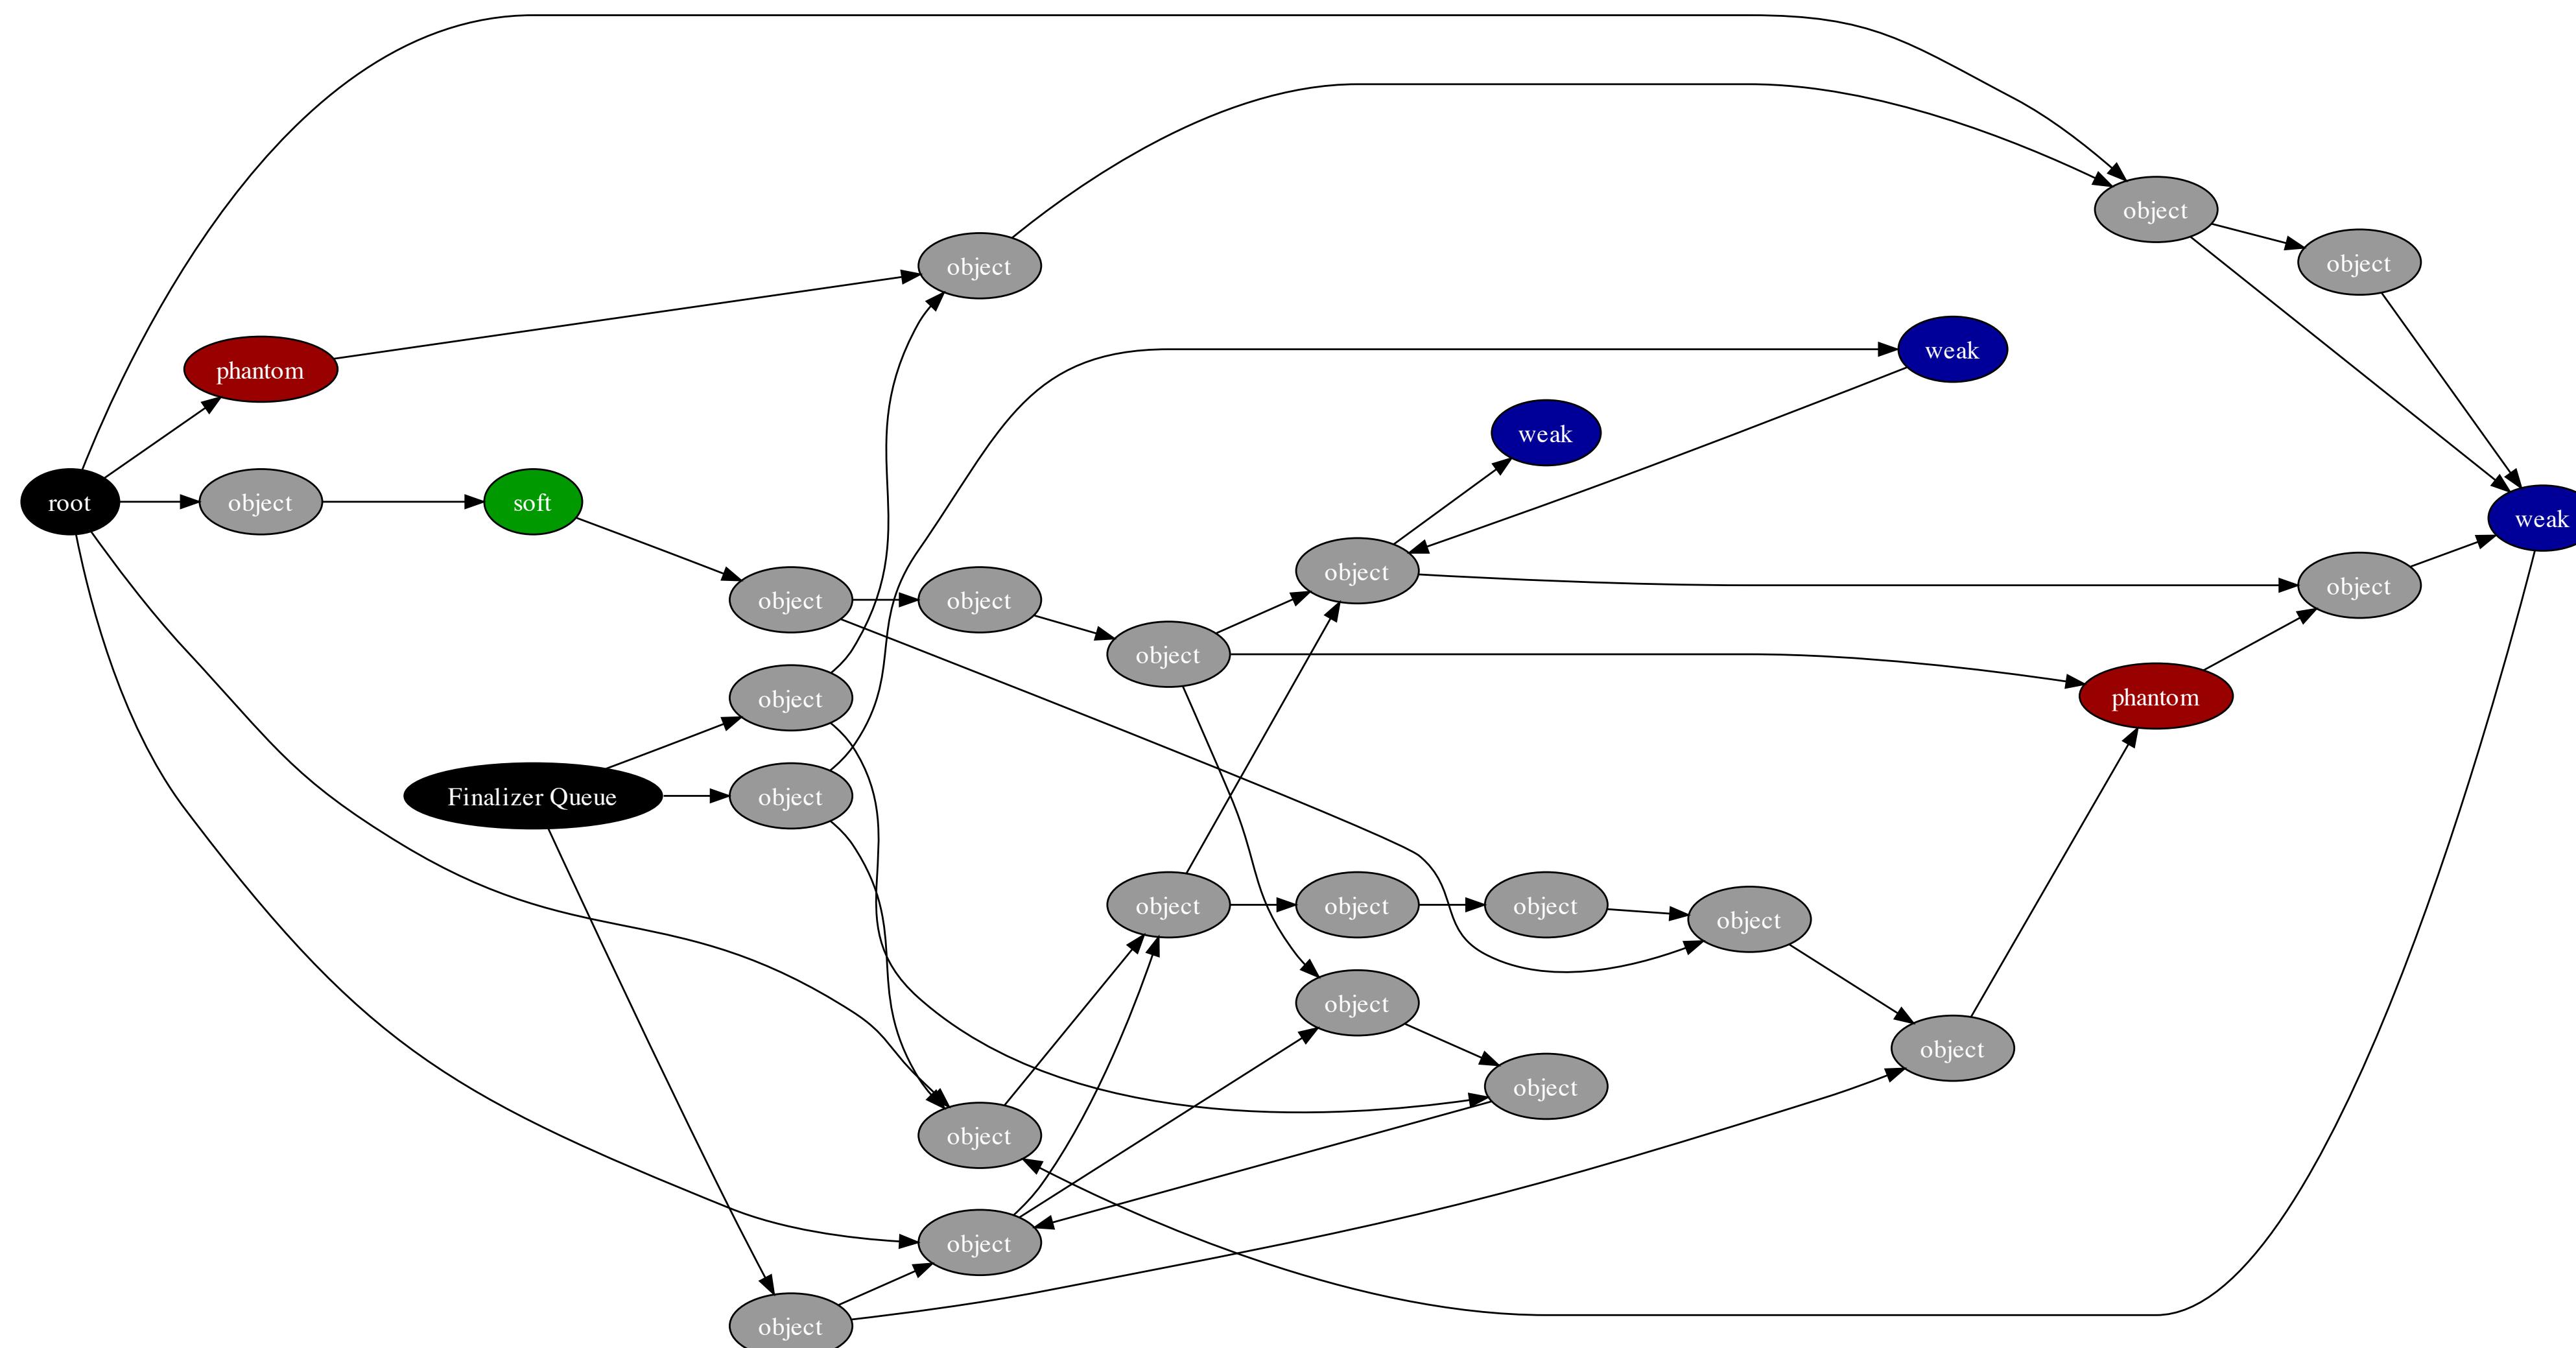
8. Possibly enqueue phantom references.



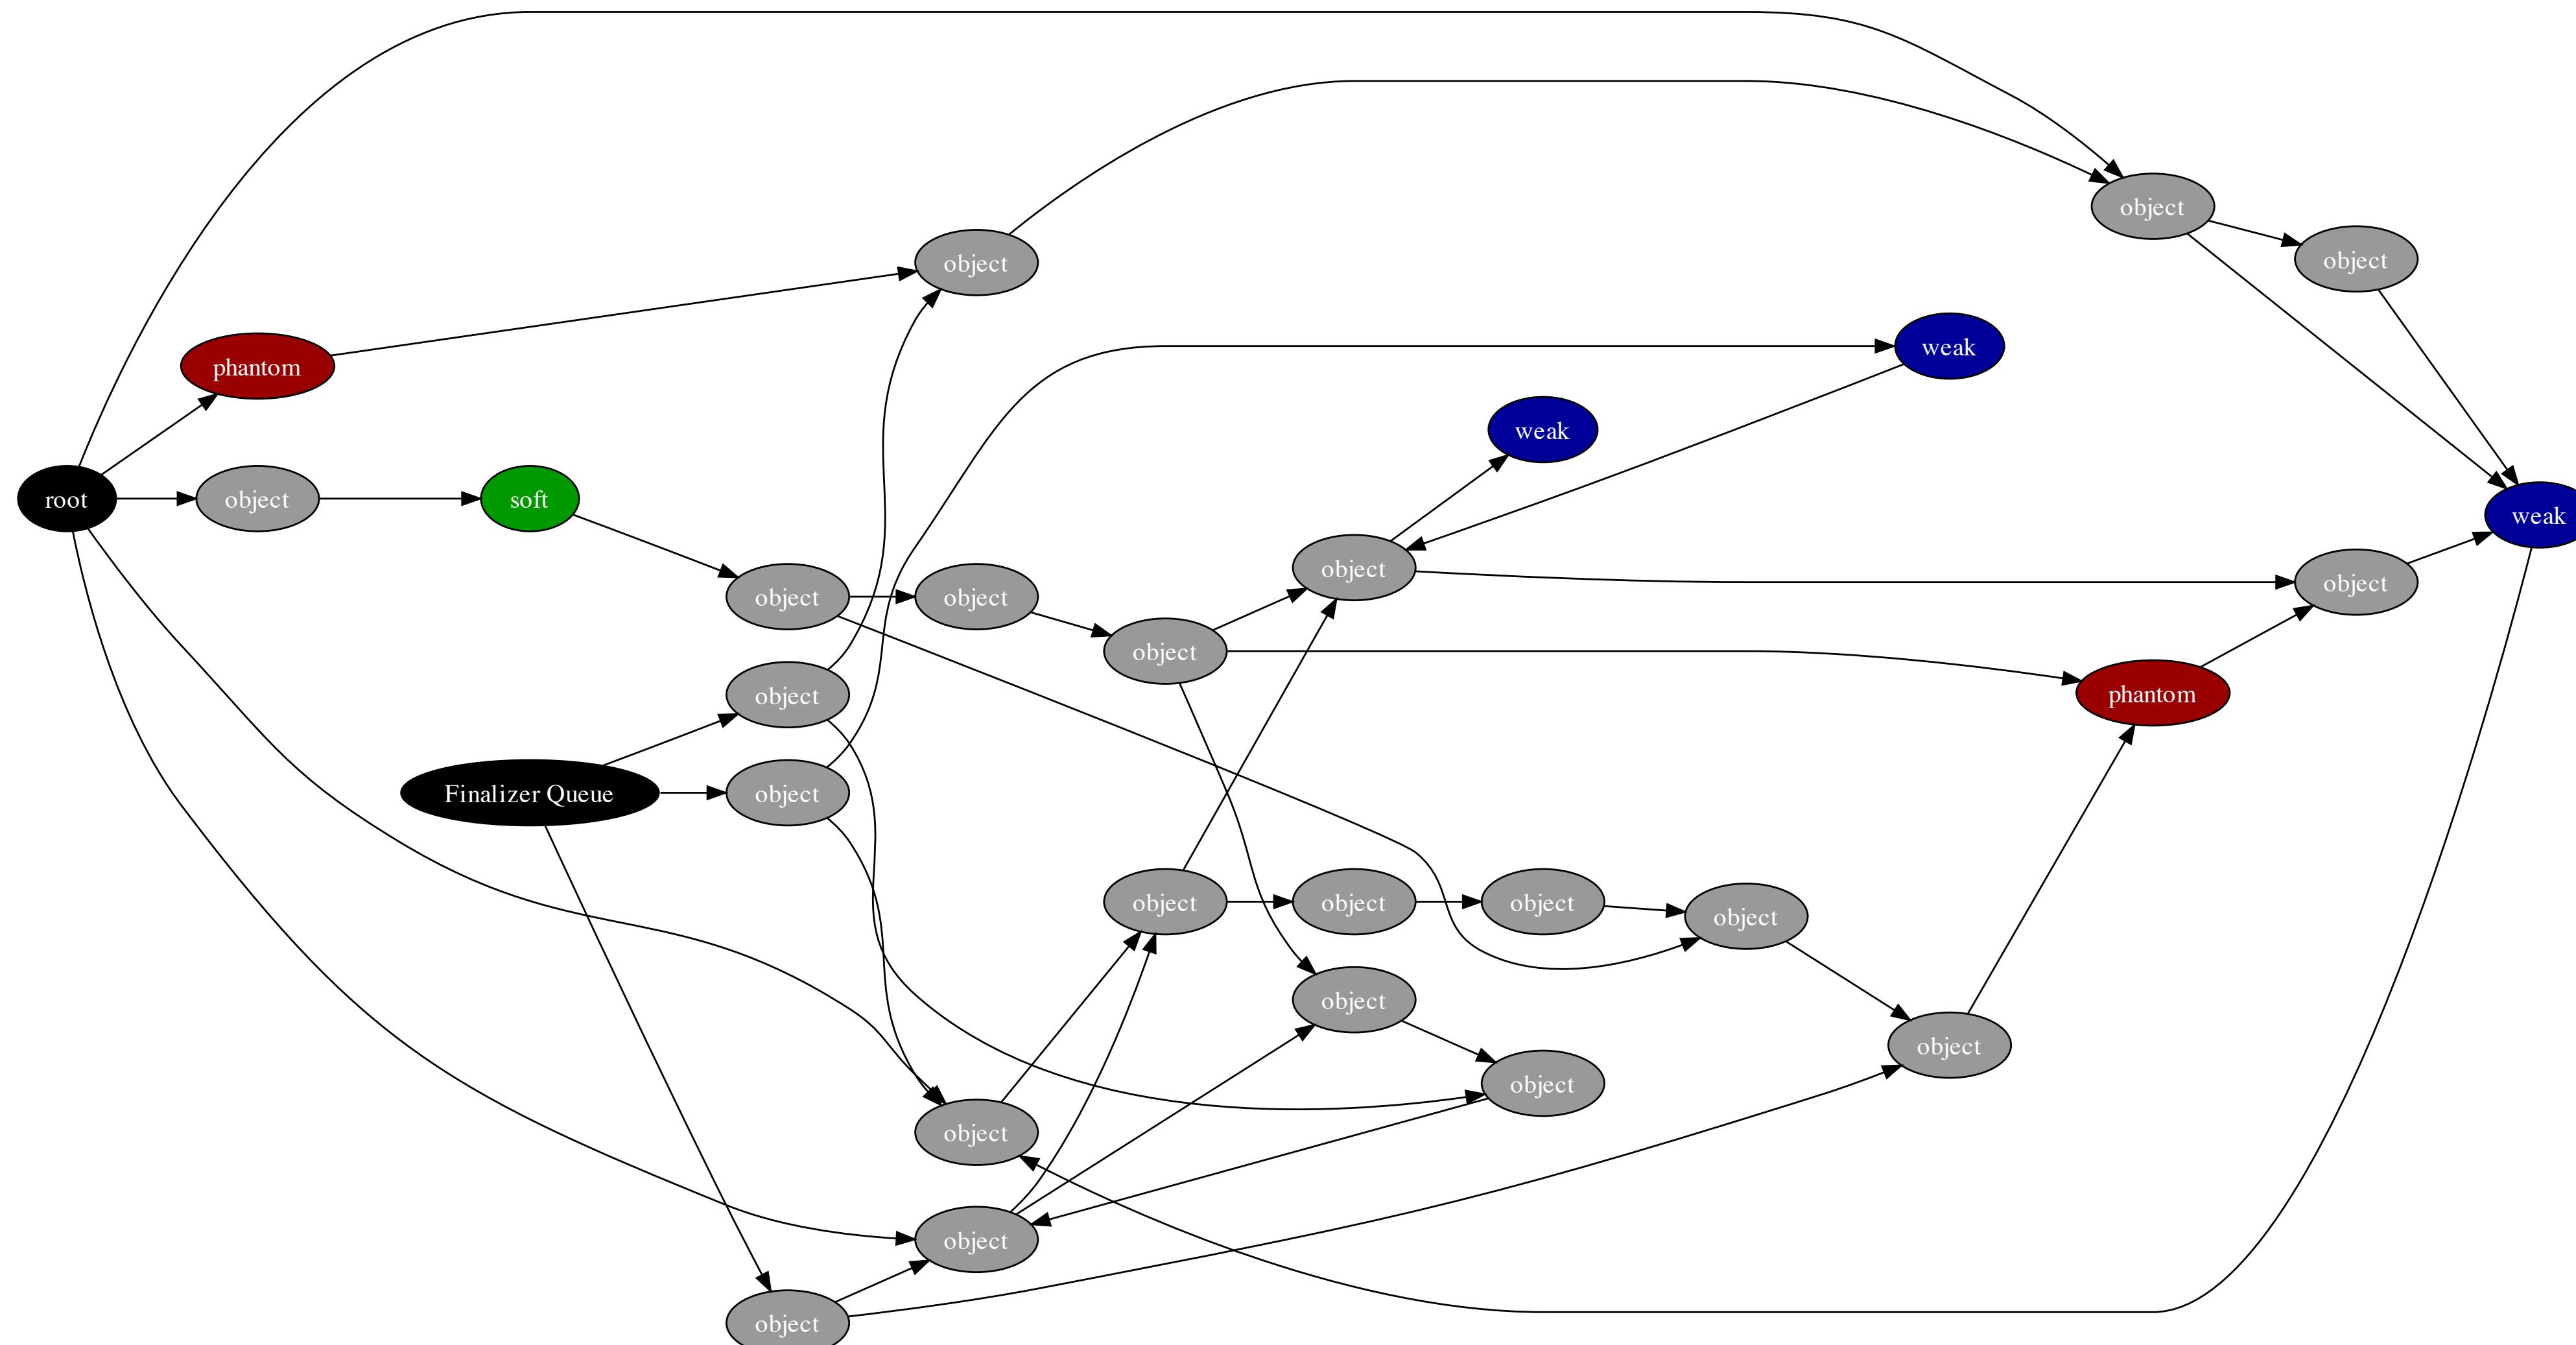
9. The remaining objects are dead.



9. The remaining objects are dead.



10. Repeat.



Recap

1. Start at a root.
2. Trace and mark strongly-referenced objects.
3. Optionally clear soft references.
4. Trace and mark softly-referenced objects.
5. Clear weak references.
6. Enqueue finalizable objects.
7. Repeat steps 1 through 5 for the queue.
8. Possibly enqueue phantom references.
9. The remaining objects are dead.
10. Repeat.

Thank You!

If you enjoyed this talk,
follow **@crazybob** on Twitter.

We're hiring! Email résumés
to **java-jobs@squareup.com**.

