



Java is a trademark of Sun Microsystems, Inc.



JavaOneSM

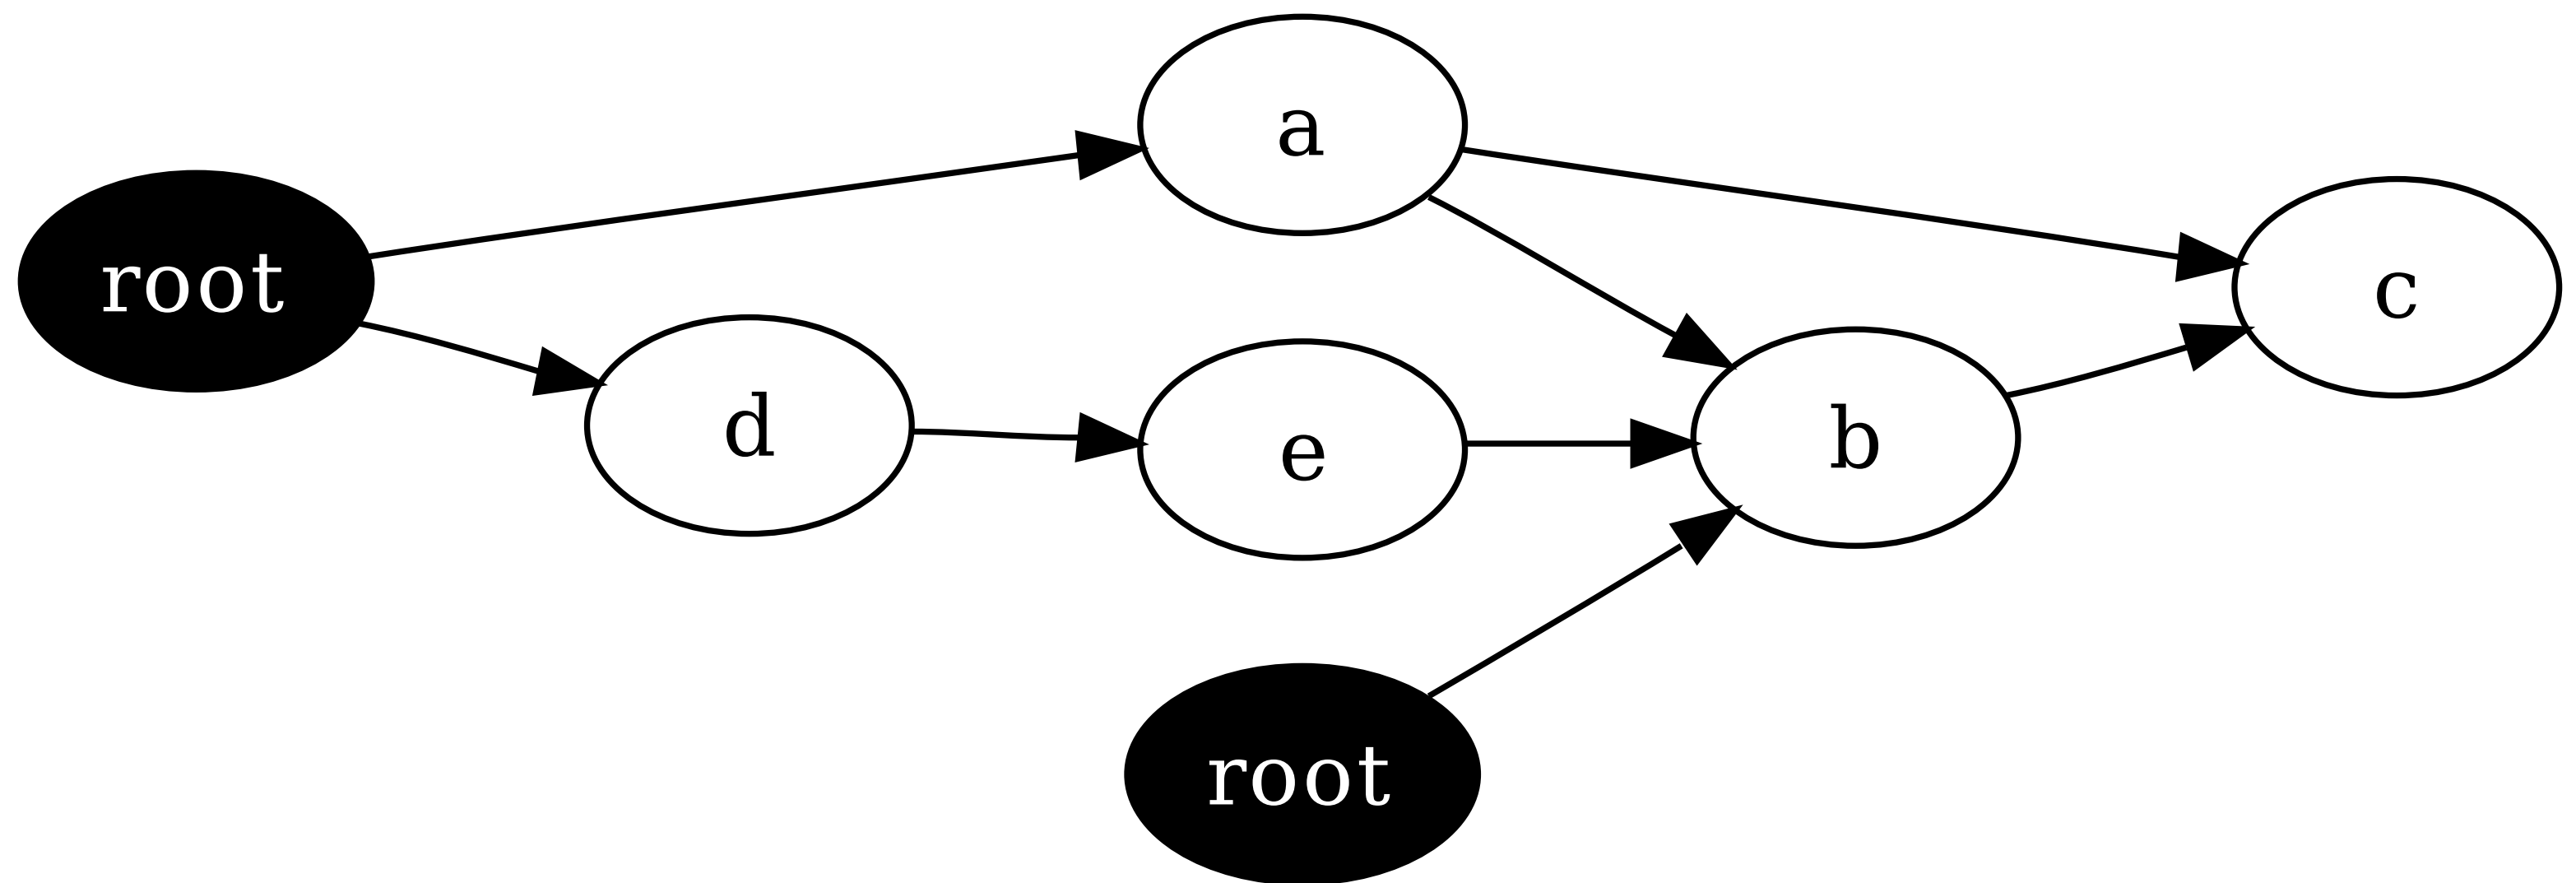
The Ghost in the Virtual Machine A Reference to References

Bob Lee
Google Inc.

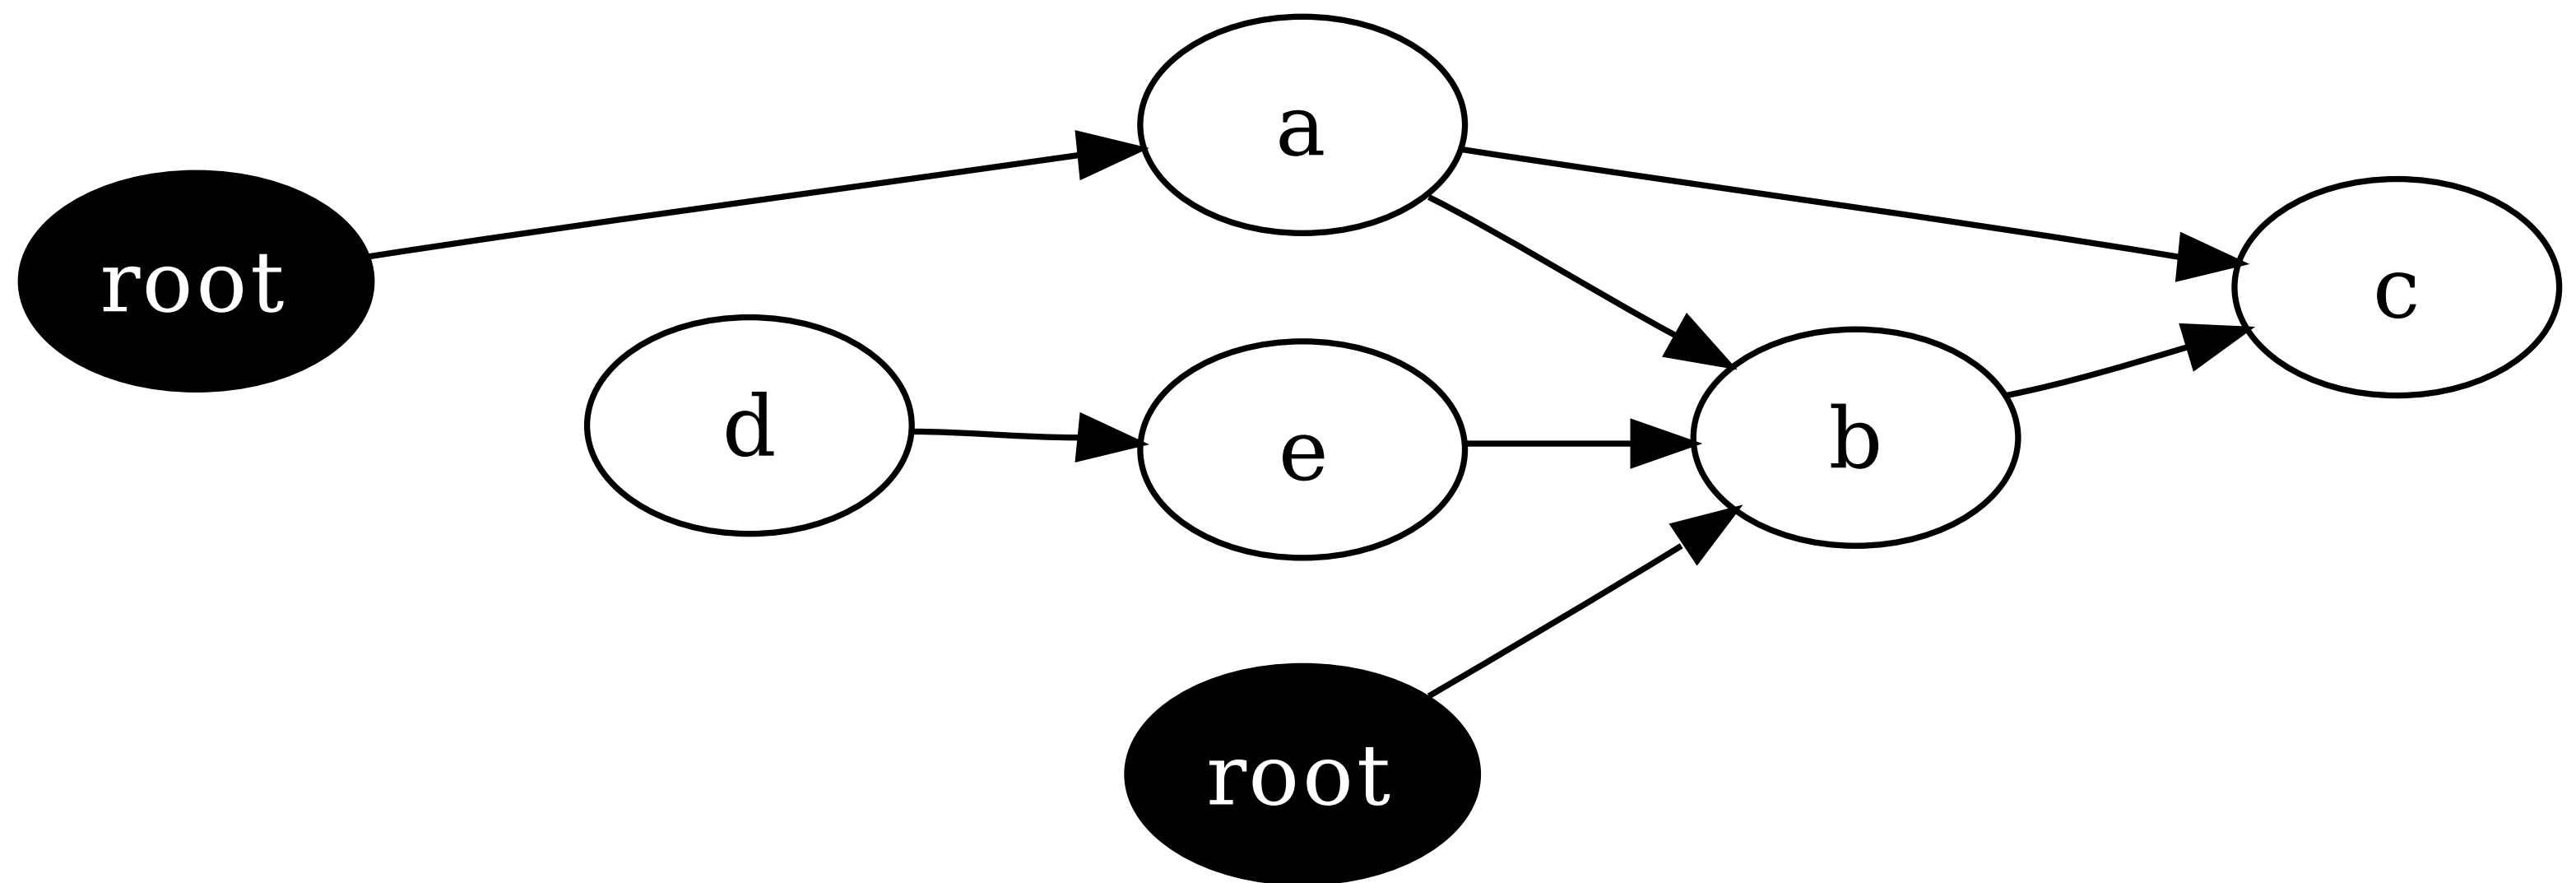
Goals

- > Take the mystery out of garbage collection.
- > Perform manual cleanup the right way.
- > Become honorary VM sanitation engineers.

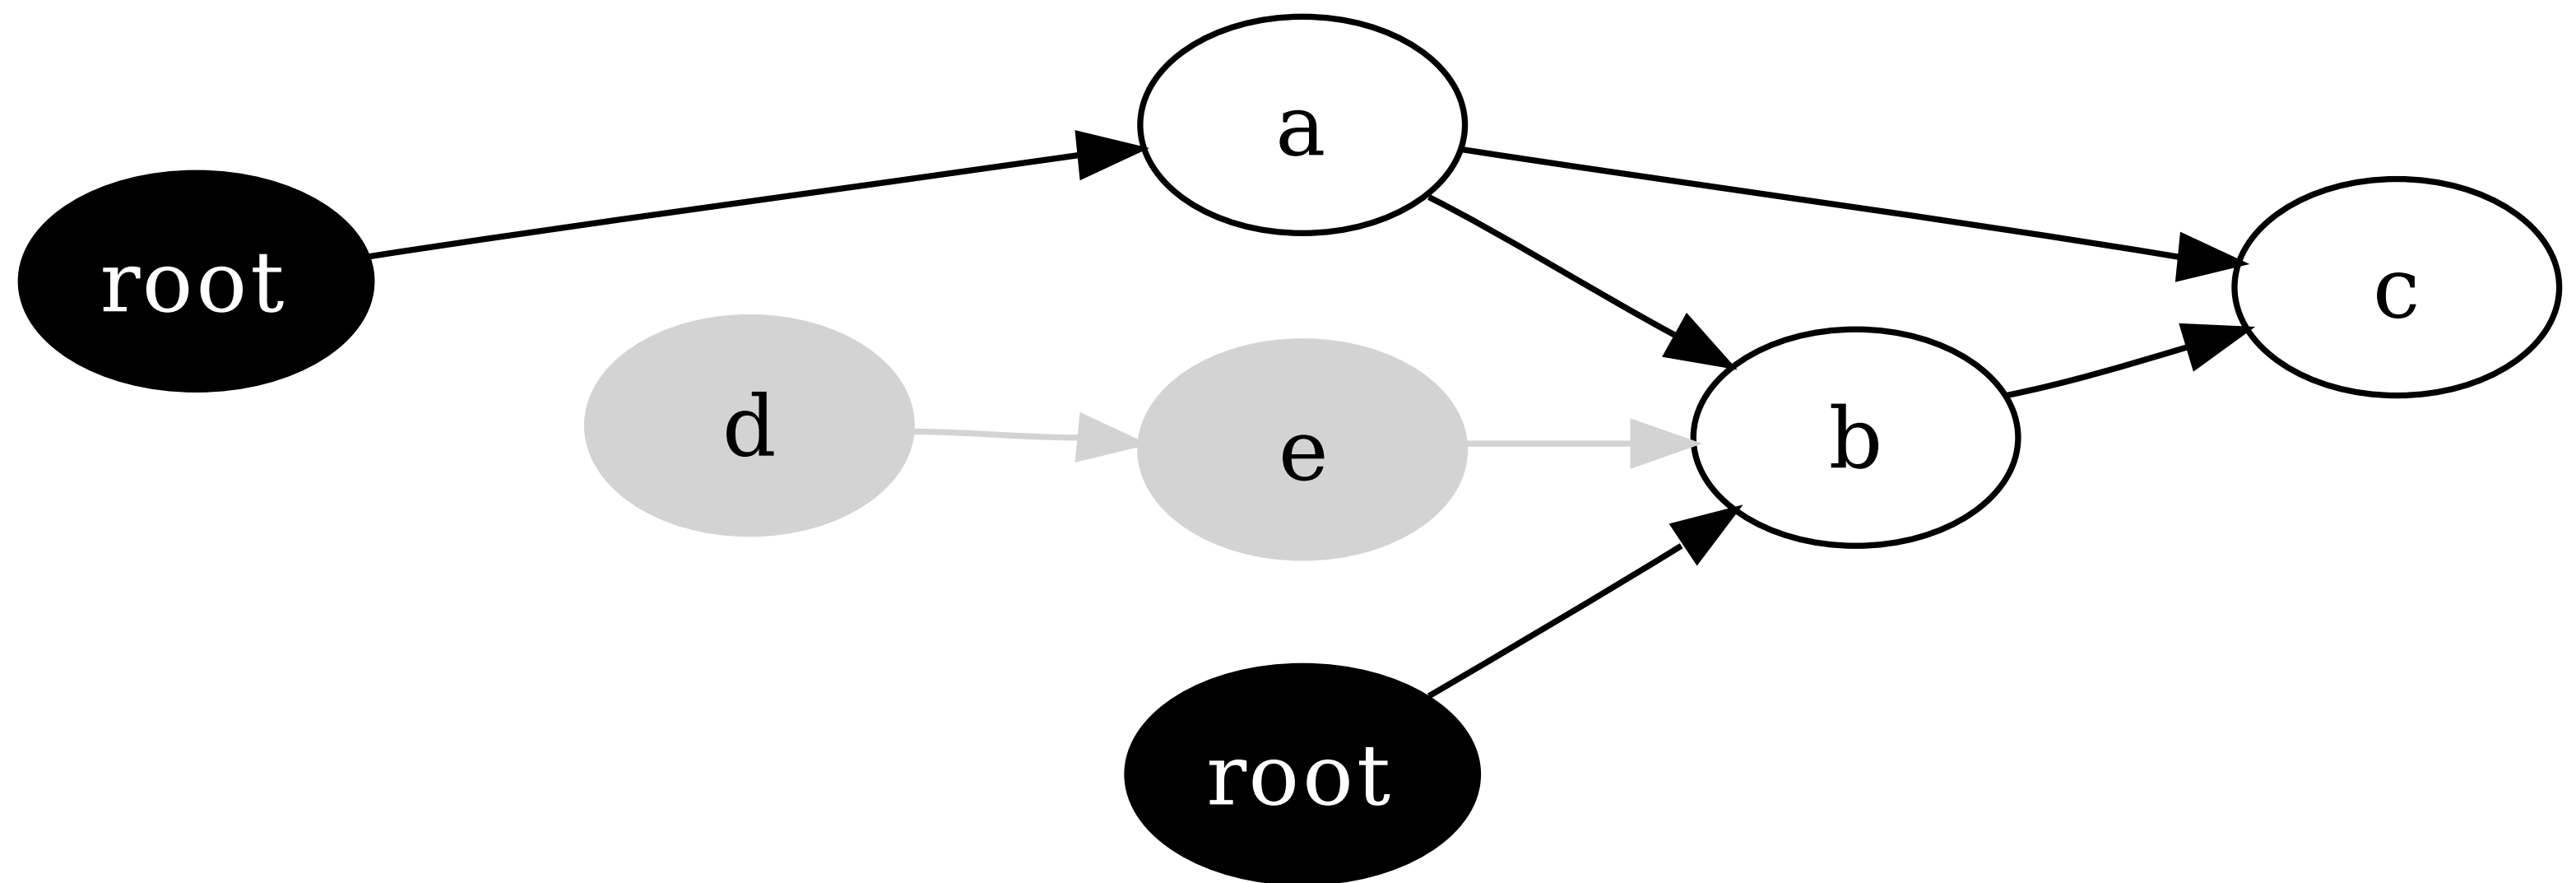
How does garbage collection work?



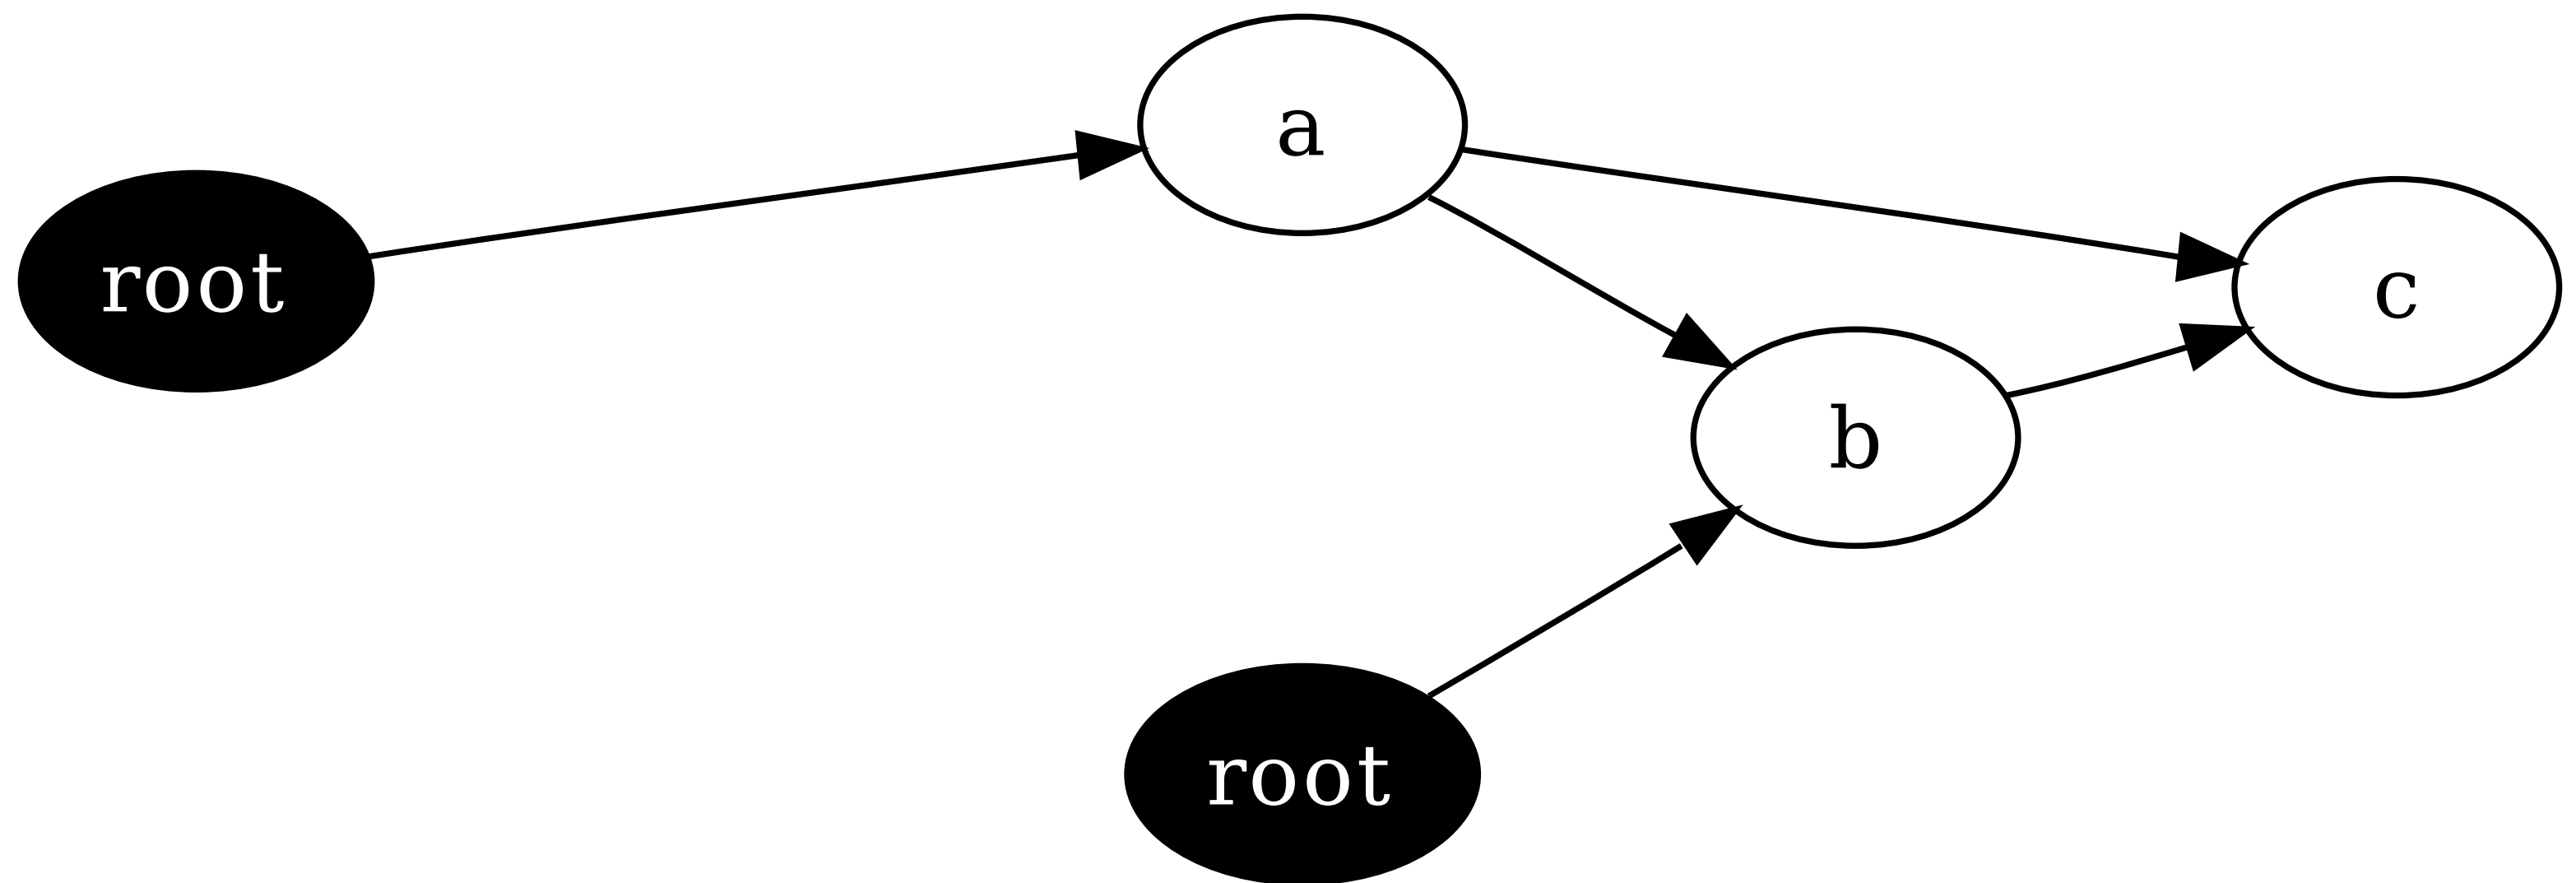
If the reference to D goes away...



We can no longer reach D or E.



So the collector reclaims them.



The GC can't do everything.

> Some things require manual cleanup.

- Listeners
- File descriptors
- Native memory
- External state (`IdentityHashMap`)

> Tools at your disposal:

- `finally`
- Overriding `Object.finalize()`
- Reference queues

What is a finalizer?

Extend `NativeResource` and make it safe.

```
public class SafeNativeResource extends NativeResource {  
    private boolean finalized;  
  
    @Override public synchronized void write(byte[] data) {  
        if (!finalized) super.write(data);  
        else /* do nothing? */;  
    }  
  
    @Override protected synchronized void finalize() {  
        finalized = true;  
        super.finalize();  
    }  
}
```


An external resource

```
public class NativeResource {  
    public NativeResource() { init(); }  
  
    /** Allocates native memory. */  
    private native void init();  
  
    /** Writes to native memory. */  
    public native void write(byte[] data);  
  
    /** Frees native memory. */  
    @Override protected native void finalize();  
}
```

Let's play War!

SegfaultFactory can cause a segfault if its finalizer executes after NativeResource's.

```
public class SegfaultFactory {  
    private final NativeResource nr;  
    public SegfaultFactory(NativeResource nr) {  
        this.nr = nr;  
    }  
  
    @Override protected void finalize() {  
        // 50/50 chance of failure  
        nr.write("I'm taking the VM with me!".getBytes());  
    }  
}
```



Use protection.

Extend `NativeResource` and make it safe.

```
public class SafeNativeResource extends NativeResource {  
    private boolean finalized;  
  
    @Override public synchronized void write(byte[] data) {  
        if (!finalized) super.write(data);  
        else /* do nothing? */;  
    }  
  
    @Override protected synchronized void finalize() {  
        finalized = true;  
        super.finalize();  
    }  
}
```

Reachability

- > An object is *reachable* if a live thread can access it.
- > Examples of heap roots:
 - System classes (which have static fields)
 - Thread stacks
 - In-flight exceptions
 - JNI global references
 - The finalizer queue
 - The interned String pool
 - etc. (VM-dependent)

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > **Strong**
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > **Soft**
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > **Weak**
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > **Finalizer**
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > **Phantom, JNI weak**
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > **Unreachable**

Two options for freeing native resources

- > Use a finalizer.
 - You must defend against subsequent use!
- > Or use a phantom reference.

Weak references aren't for caching!

- > Many collectors will reclaim weak refs immediately.
- > Use soft reference for caching, as intended:

“Virtual machine implementations are encouraged to bias against clearing recently-created or recently-used soft references.”

- The `SoftReference` documentation