



Java is a trademark of Sun Microsystems, Inc.



JavaOneSM

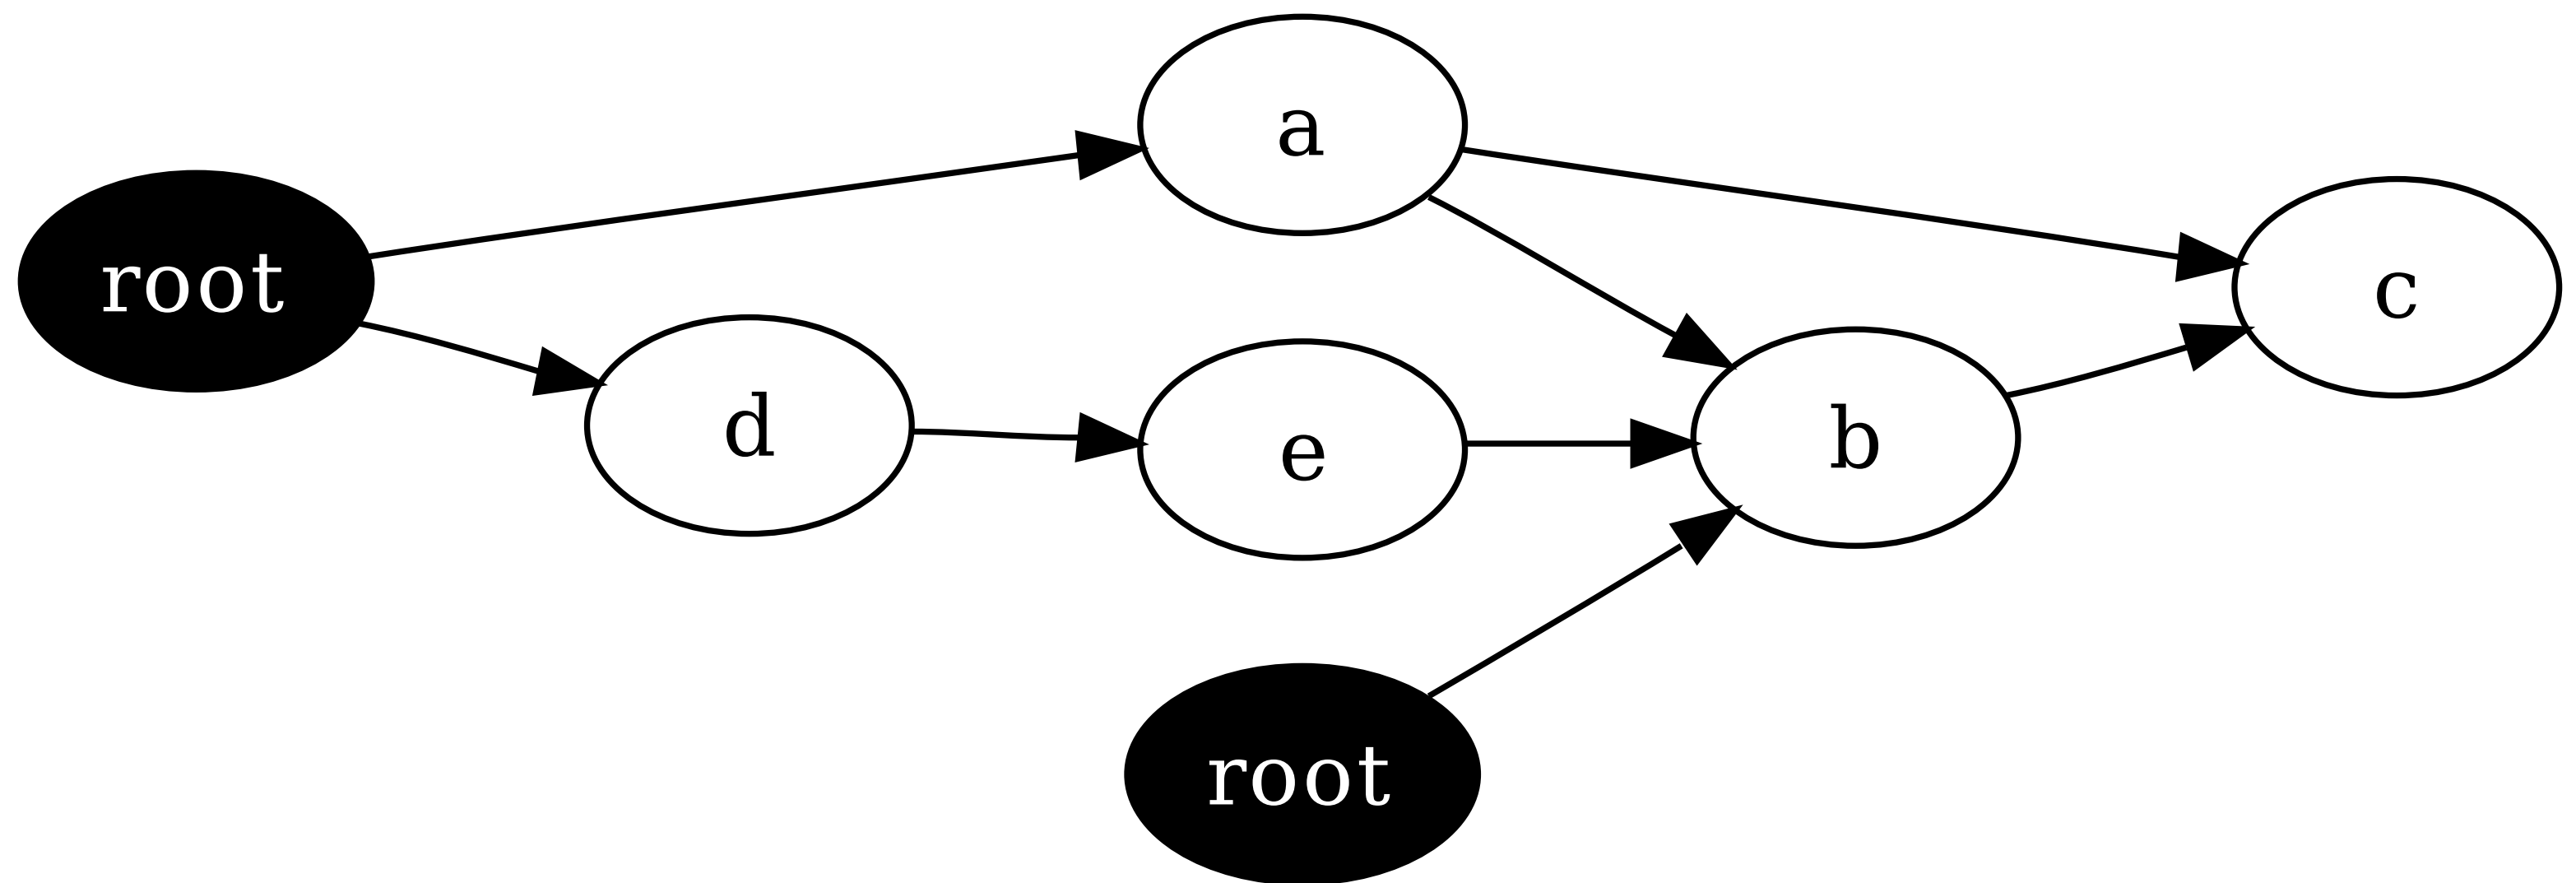
The Ghost in the Virtual Machine A Reference to References

Bob Lee
Google Inc.

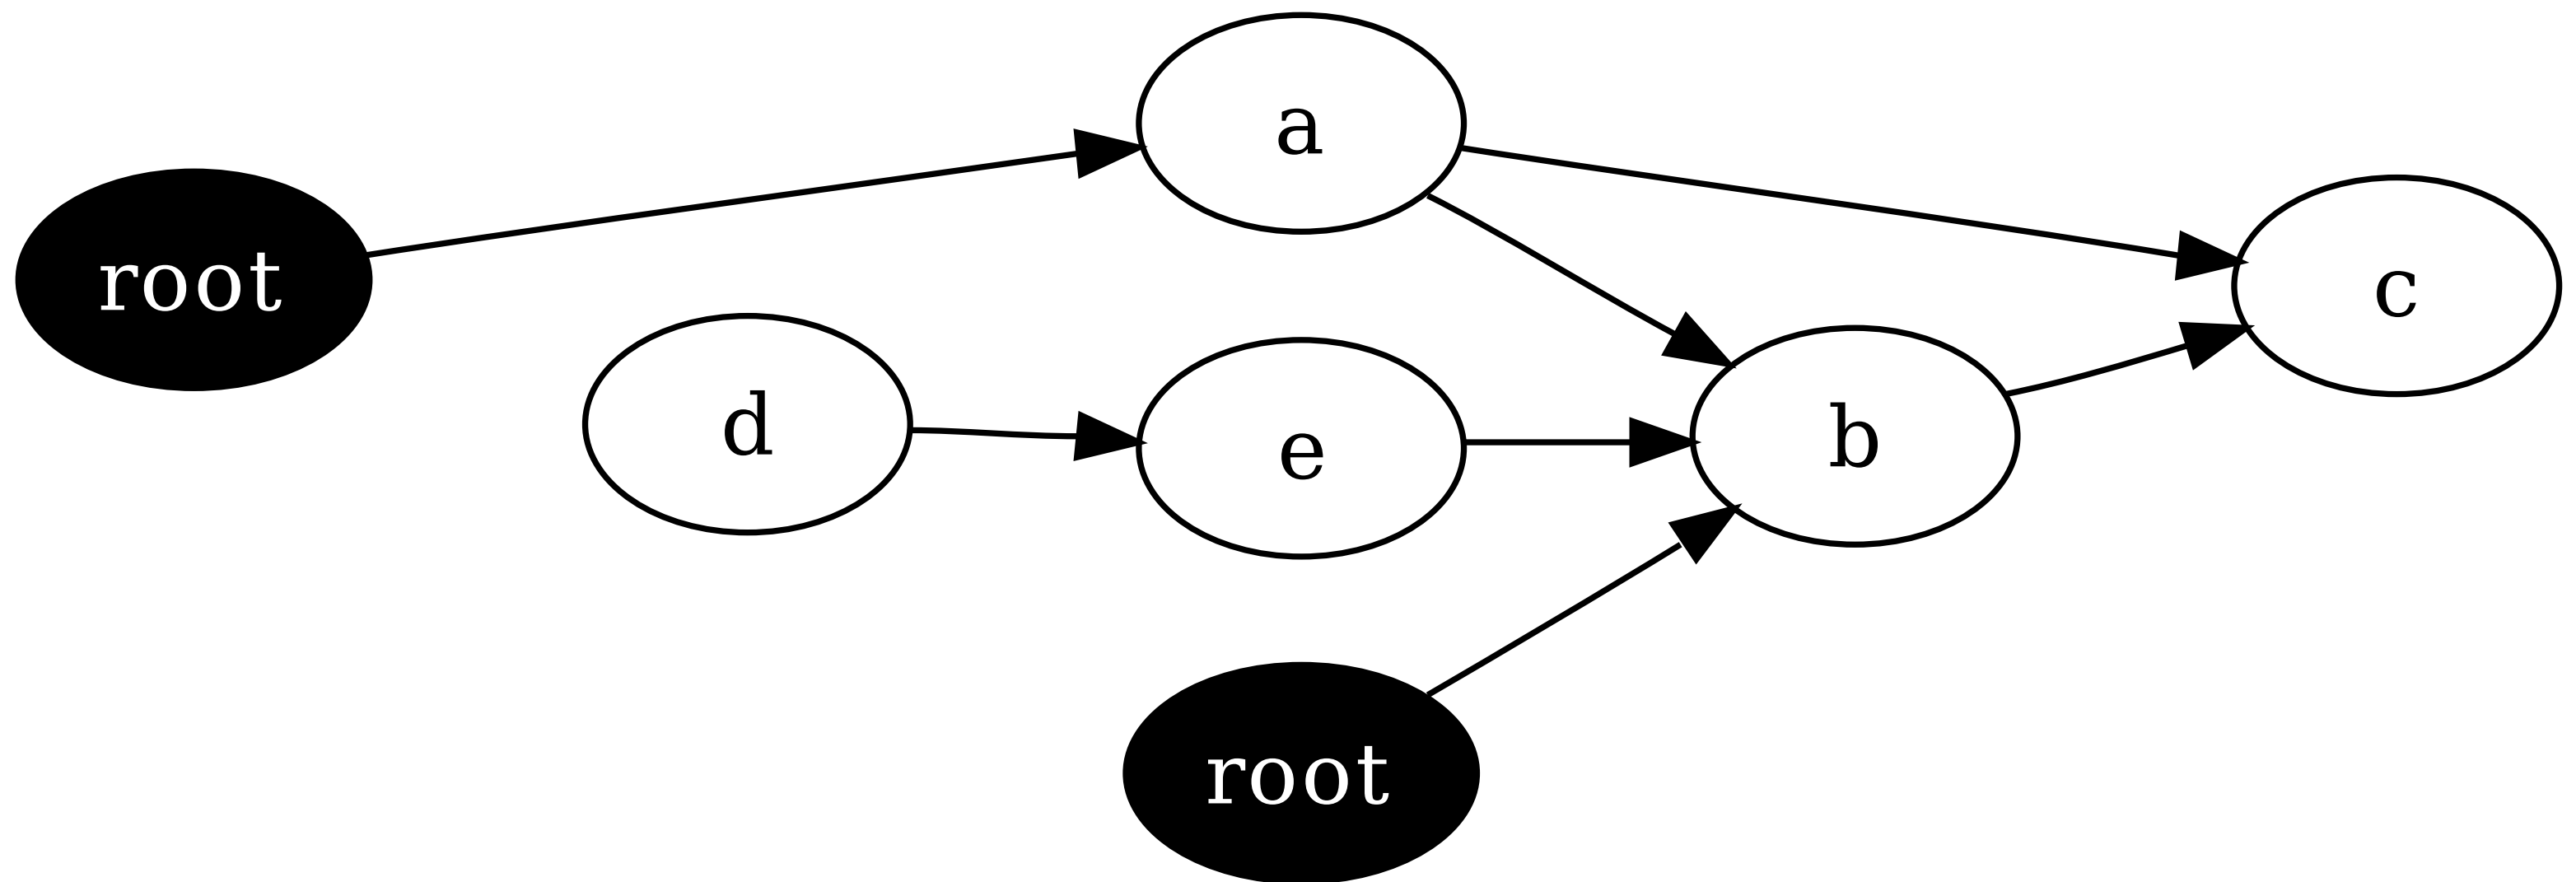
Goals

- > Take the mystery out of garbage collection.
- > Perform manual cleanup the Right way.
- > Become honorary VM sanitation engineers.

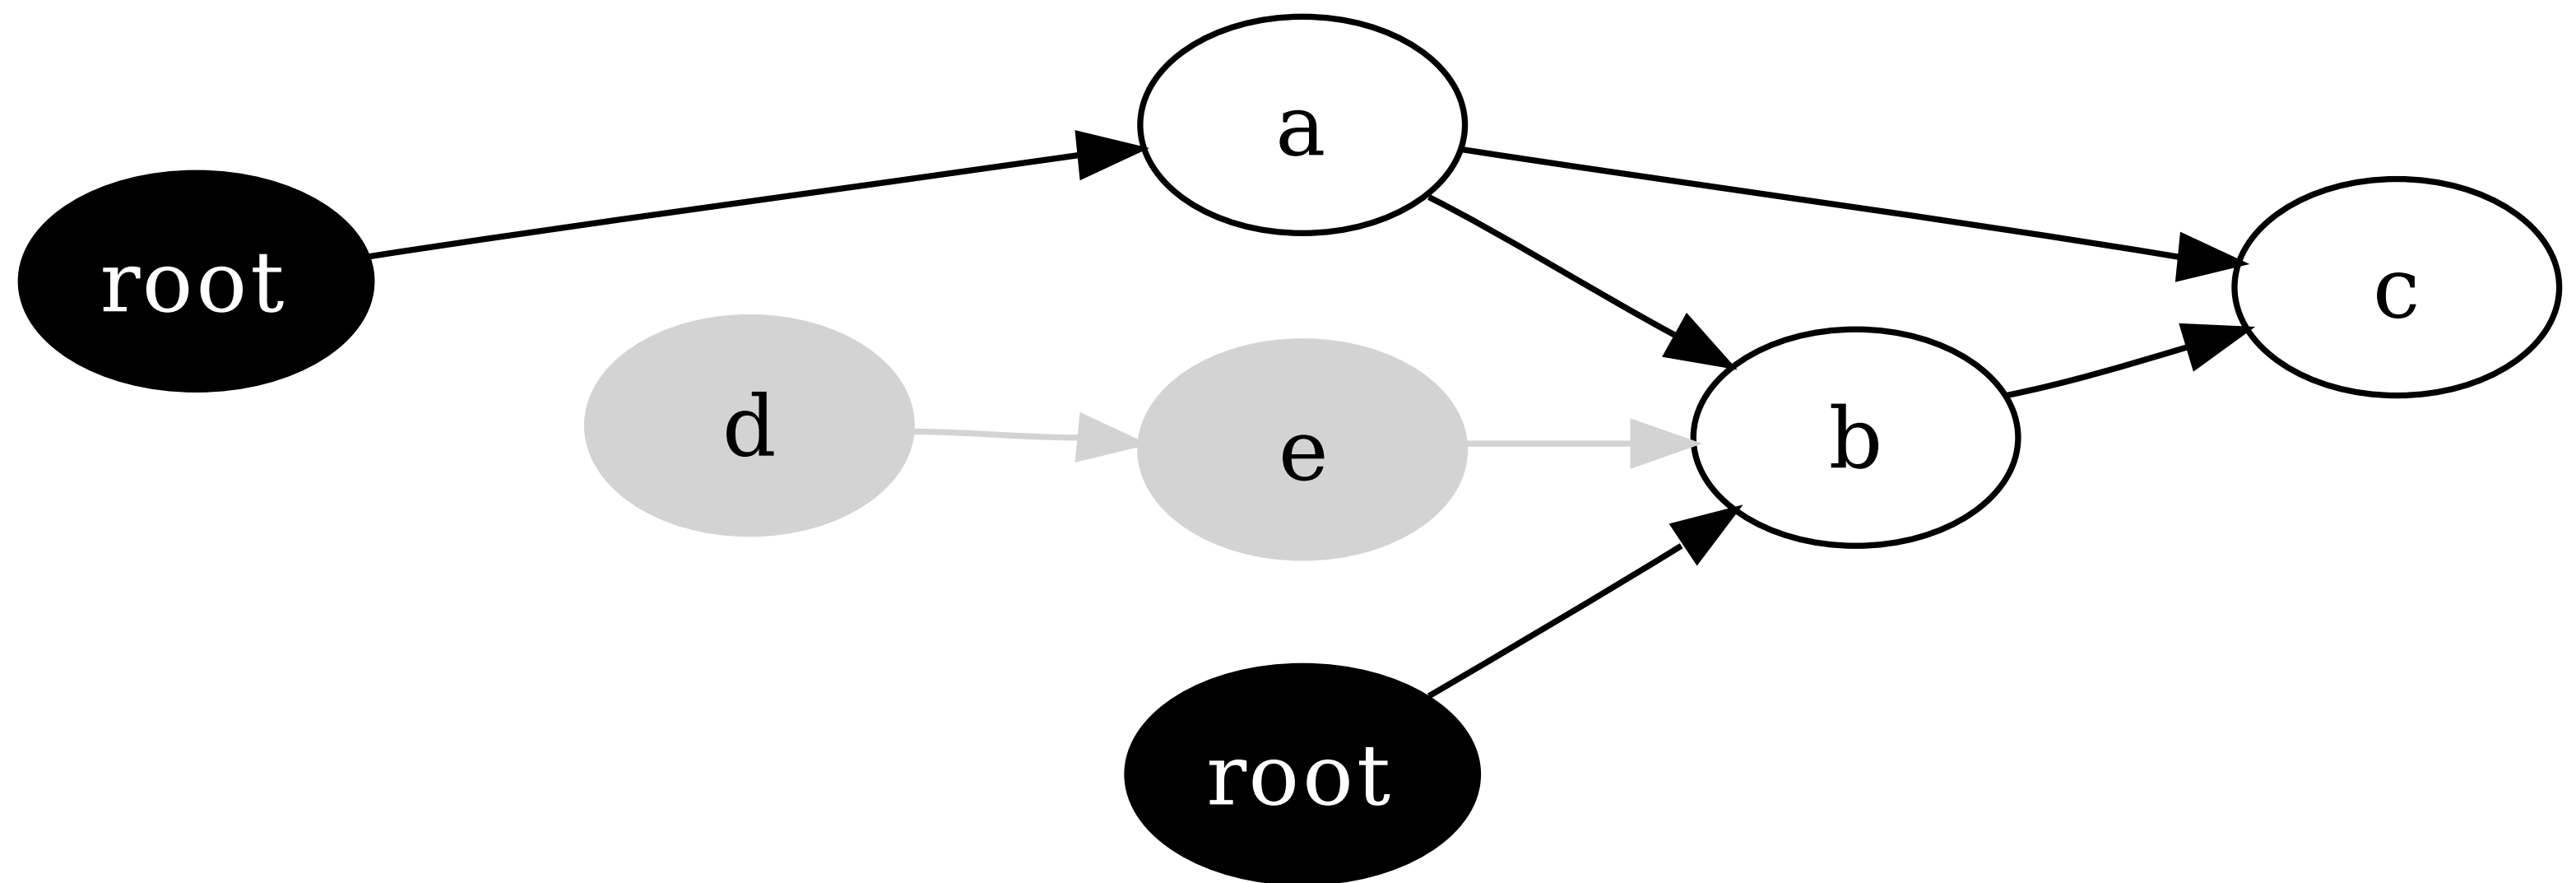
How does garbage collection work?



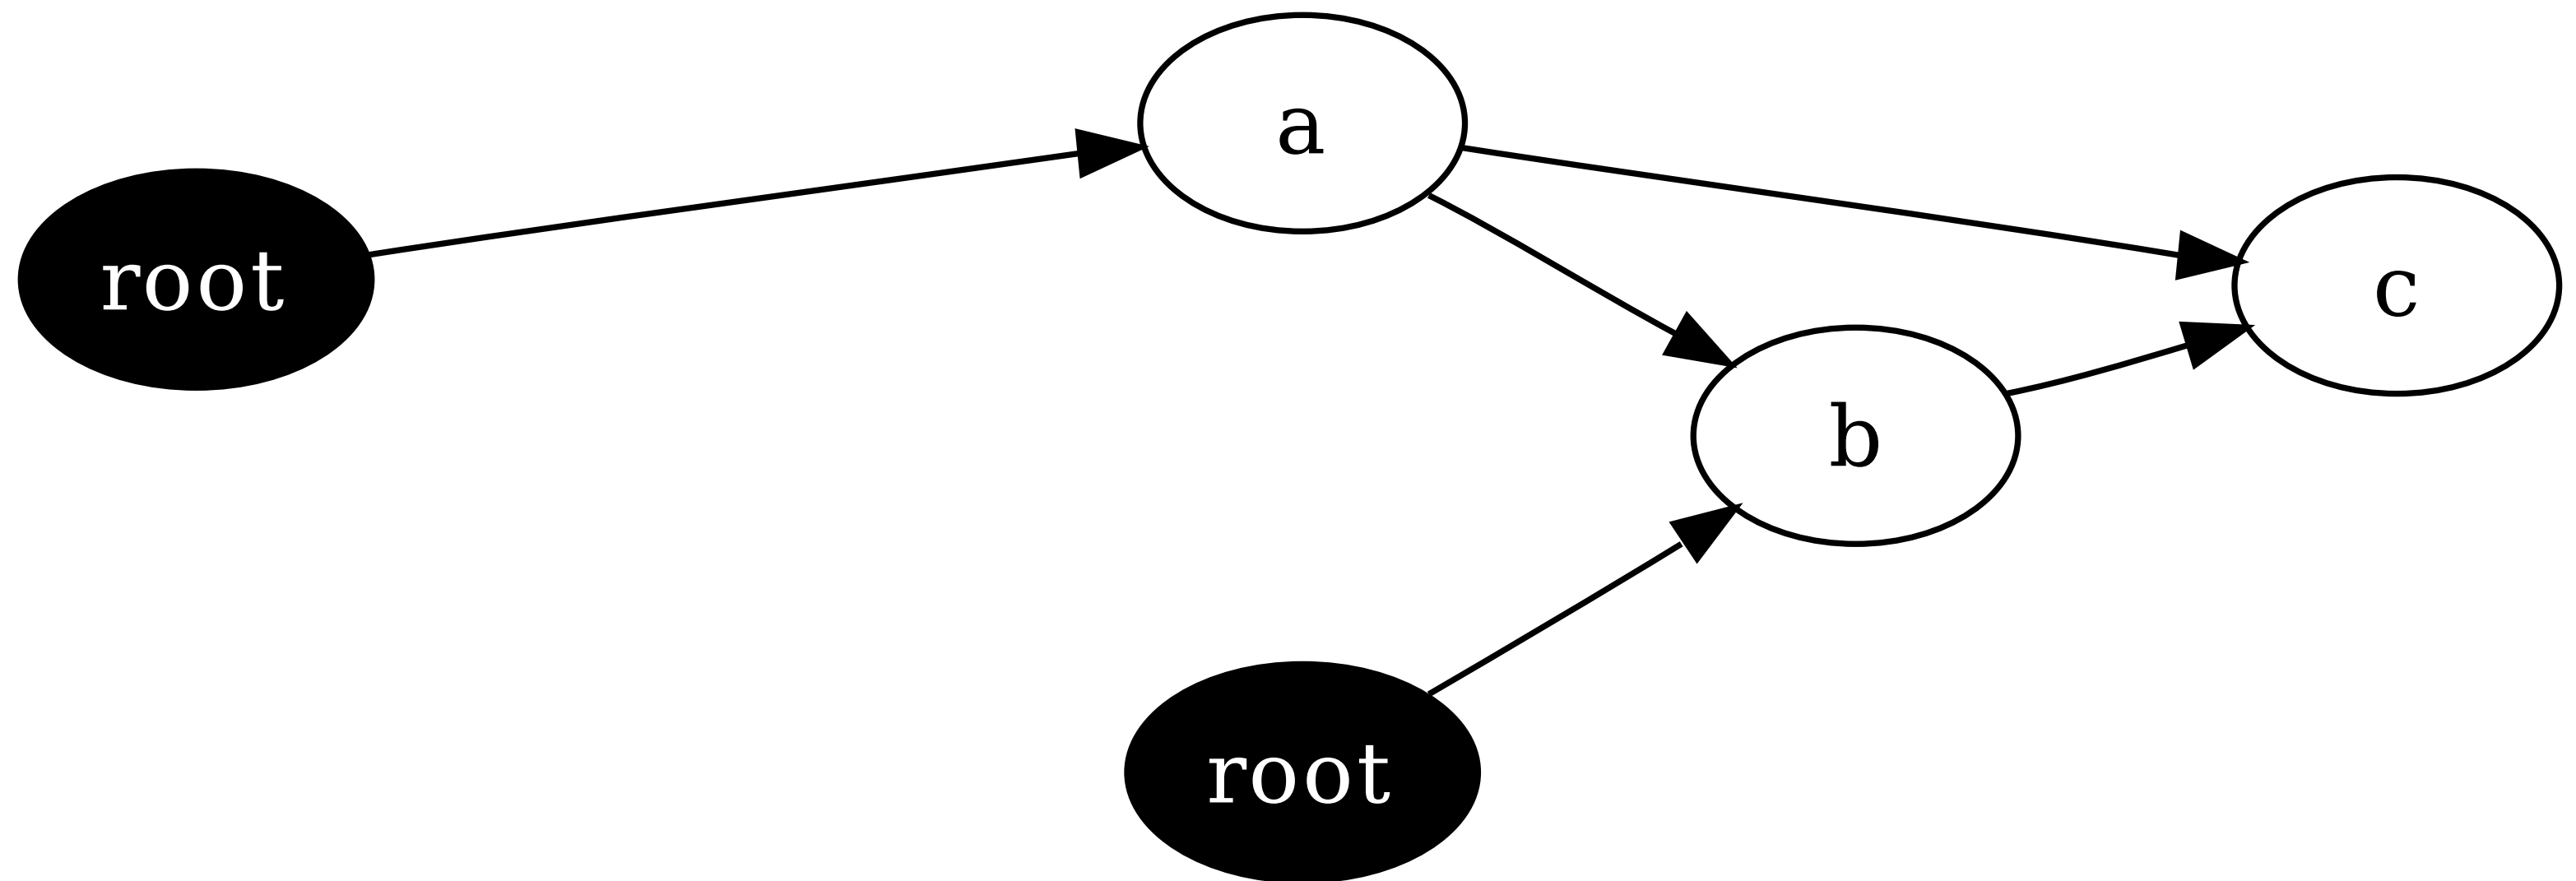
If the reference to D goes away...



We can no longer reach D or E.



So the collector reclaims them.



The GC can't do everything.

- > Some things require manual cleanup.
 - Listeners
 - File descriptors
 - Native memory
 - External state (`IdentityHashMap`)
- > Tools at your disposal:
 - `finally`
 - Overriding `Object.finalize()`
 - References (and reference queues)

Try `finally` first.

> Pros:

- More straightforward
- Handles exceptions in main thread
- Ensures cleanup keeps pace

> Cons:

- More work for programmers
- More error prone
- Cleanup happens in main thread

> ARM will help.

What is a finalizer?

```
public class Foo extends Bar {  
    @Override protected void finalize() throws Throwable {  
        try {  
            ... // Clean up Foo.  
        } finally {  
            super.finalize(); // Clean up Bar.  
        }  
    }  
}
```

Finalizers are seductively simple, but...

- > They're not guaranteed to run, especially not timely.
- > Avoid `System.runFinalizersOnExit()` and `runFinalization()`.
- > Undefined threading model, can run concurrently!
- > You must call `super.finalize()`.
- > Exceptions are ignored (per spec).
- > You can resurrect references.
- > Keeps objects alive longer.
- > Can make allocation/reclamation 430X slower
(Bloch, *Effective Java*)

Example

```
public class NativeMemory {
    final int address = allocate();
    /** Allocates native memory. */
    static native int allocate();

    /** Writes to native memory. */
    public void write(byte[] data) {
        write(address, data);
    }
    static native void write(int address, byte[] data);

    /** Frees native memory. */
    @Override protected void finalize() {
        free(address);
    }
    static native void free(int address);
}
```

Let's play War!

SegfaultFactory can cause a segfault if its finalizer executes after NativeMemory's:

```
public class SegfaultFactory {  
    private final NativeMemory nm;  
  
    public SegfaultFactory(NativeMemory nm) {  
        this.nm = nm;  
    }  
  
    @Override protected void finalize() {  
        // 50/50 chance of failure  
        nm.write("I'm taking the VM with me!".getBytes());  
    }  
}
```



Always use protection.

```
public class NativeMemory {
    final int address = allocate();
    /** Allocates native memory. */
    static native int allocate();

    /** Writes to native memory. */
    boolean finalized;
    public synchronized void write(byte[] data) {
        if (!finalized) write(address, data);
        else /* do nothing? */;
    }
    static native void write(int address, byte[] data);

    /** Frees native memory. */
    @Override protected synchronized void finalize() {
        finalized = true;
        free(address);
    }
    static native void free(int address);
}
```

Basically, finalizers are good for one thing.

Logging warnings

```
package java.lang.ref
```

```
public abstract class Reference<T> {  
    public T get() { ... }  
}
```

```
public class SoftReference<T> extends Reference<T> {  
    public SoftReference(T referent) { ... }  
    public SoftReference(T referent, ReferenceQueue<? super T> q) { ... }  
}
```

```
public class WeakReference<T> extends Reference<T> {  
    public WeakReference(T referent) { ... }  
    public WeakReference(T referent, ReferenceQueue<? super T> q) { ... }  
}
```

```
public class PhantomReference<T> extends Reference<T> {  
    public PhantomReference(T referent, ReferenceQueue<? super T> q) { ... }  
}
```

```
public class ReferenceQueue<T> {  
    public ReferenceQueue() { ... }  
    public Reference<? extends T> poll() { ... }  
    public Reference<? extends T> remove() { ... }  
}
```


Can you hear me now?

```
public class Button {
    public interface Listener {
        void onClick();
    }
    private final List<WeakReference<Listener>> listeners
        = new ArrayList<WeakReference<Listener>>();
    public void add(Listener l) {
        listeners.add(new WeakReference<Listener>(l));
    }
    public void click() {
        Iterator<WeakReference<Listener>> i
            = listeners.iterator();
        while (i.hasNext()) {
            Listener l = i.next().get();
            if (l == null) i.remove();
            else l.onClick();
        }
    }
}
```

Reachability

- > An object is *reachable* if a live thread can access it.
- > Examples of heap roots:
 - System classes (which have static fields)
 - Thread stacks
 - In-flight exceptions
 - JNI global references
 - The finalizer queue
 - The interned String pool
 - etc. (VM-dependent)

Making maps

```
public class BytecodeCache {  
    final static Map<Class<?>, byte[]> cache = new MapMaker()  
        .weakKeys()  
        .softValues()  
        .makeComputingMap(new Function<Class<?>, byte[]>() {  
            public byte[] apply(Class<?> clazz) {  
                ...  
            }  
        })  
};  
  
    public static byte[] bytesFor(Class<?> clazz) {  
        return cache.get(clazz);  
    }  
}
```

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > **Strong**
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > **Soft**
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > **Weak**
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > **Finalizer**
- > Phantom, JNI weak
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > **Phantom, JNI weak**
- > Unreachable

Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > **Unreachable**

Weak references aren't for caching!

- > Many collectors will reclaim weak refs immediately.
- > Use soft reference for caching, as intended:

“Virtual machine implementations are encouraged to bias against clearing recently-created or recently-used soft references.”

- The `SoftReference` documentation