



Java is a trademark of Sun Microsystems, Inc.



# JavaOne<sup>SM</sup>

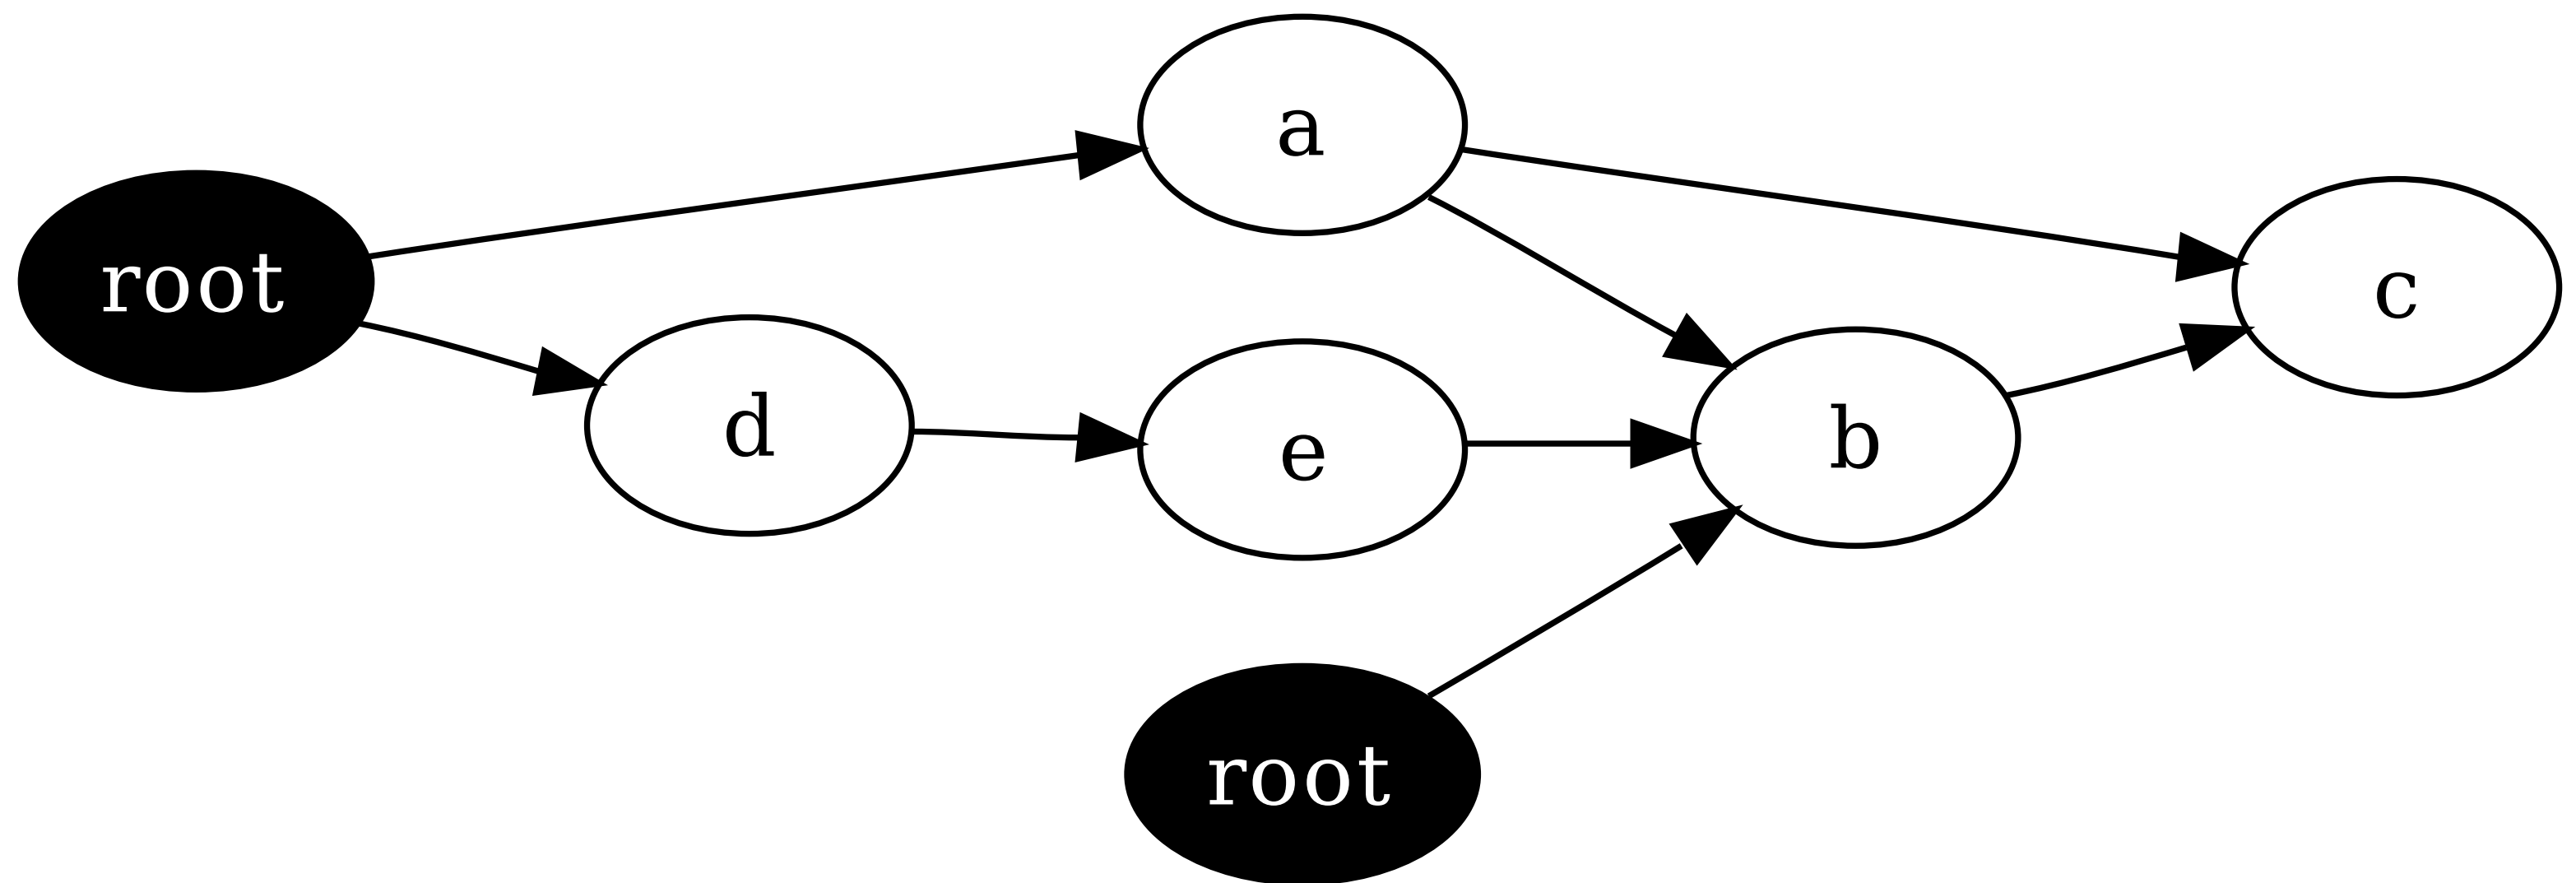
## The Ghost in the Virtual Machine A Reference to References

Bob Lee  
Google Inc.

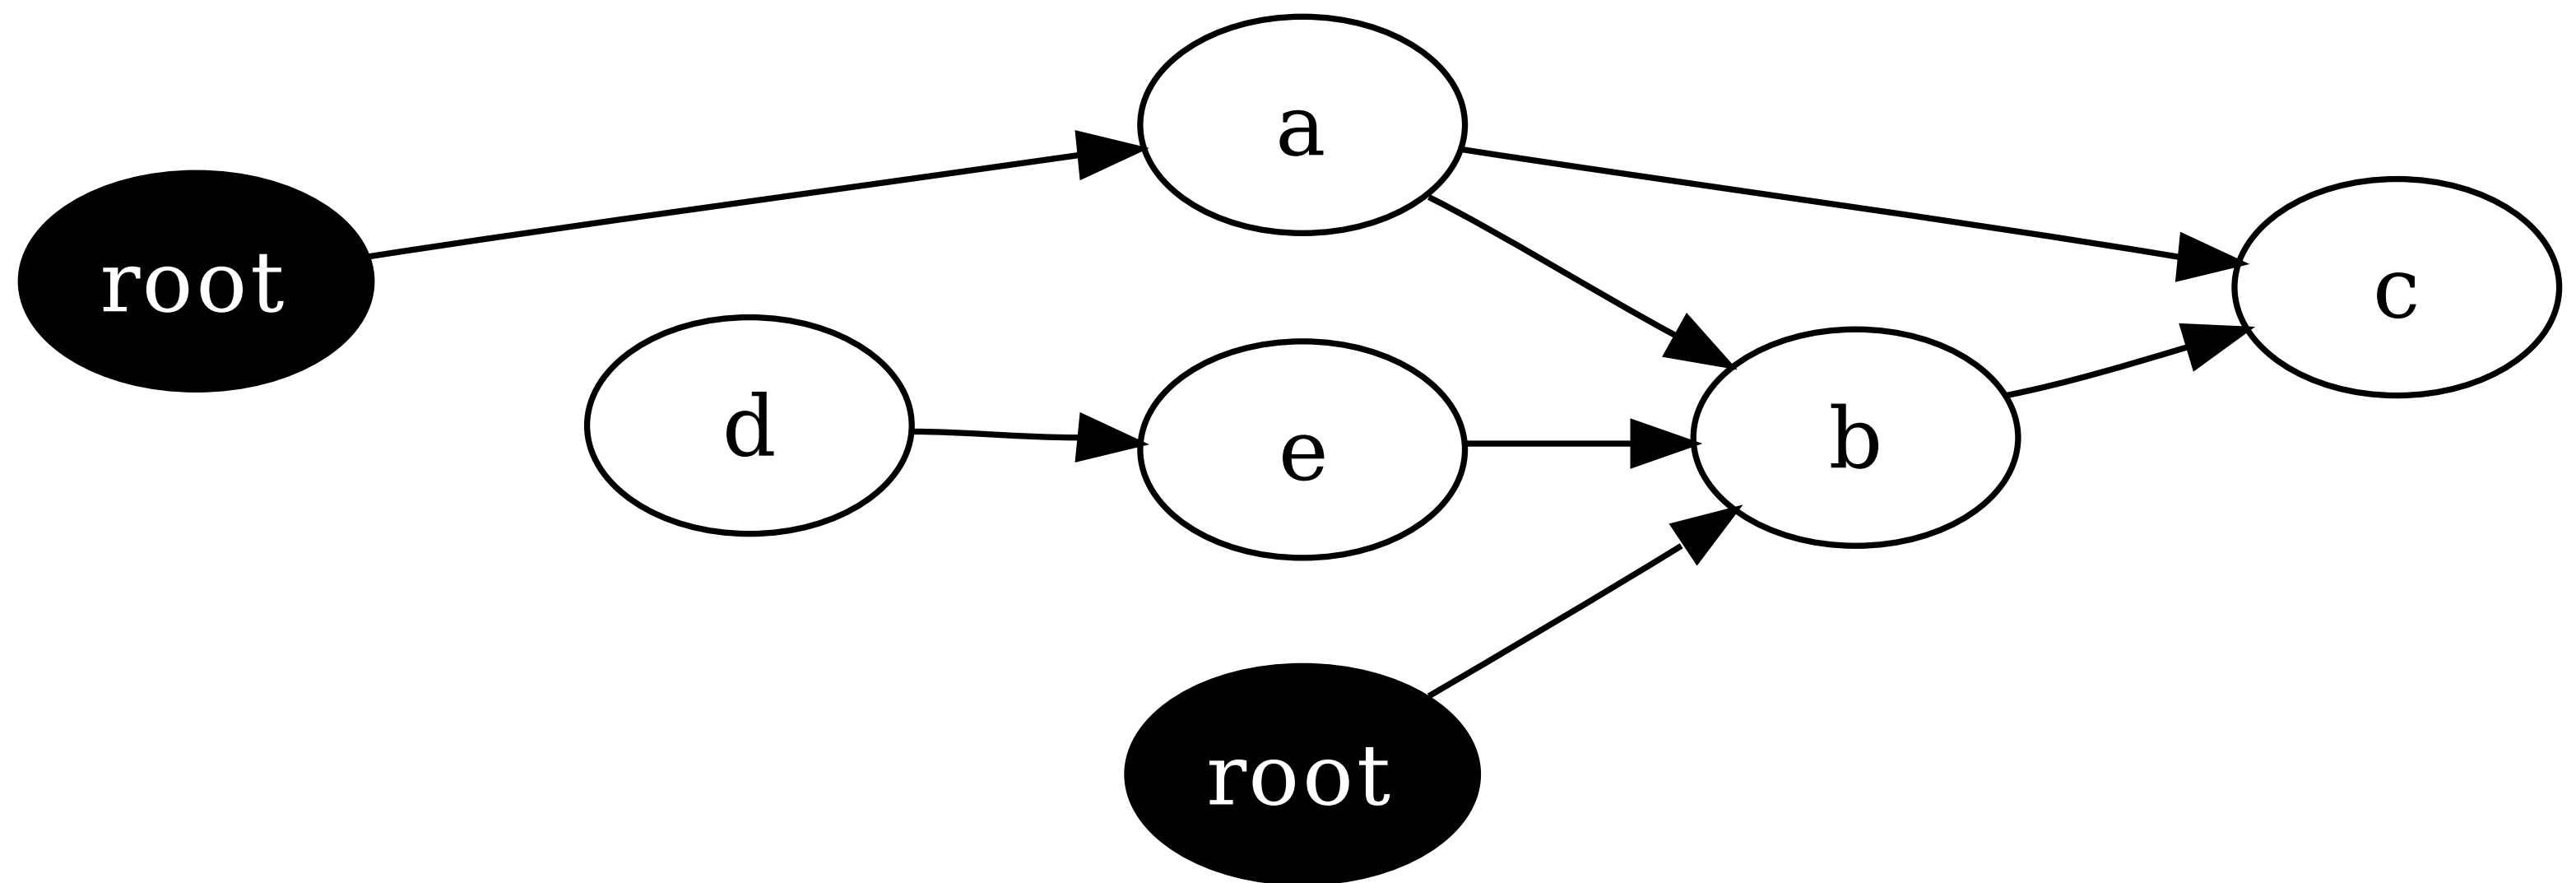
# Goals

- > Take the mystery out of garbage collection.
- > Perform manual cleanup the Right way.
- > Become honorary VM sanitation engineers.

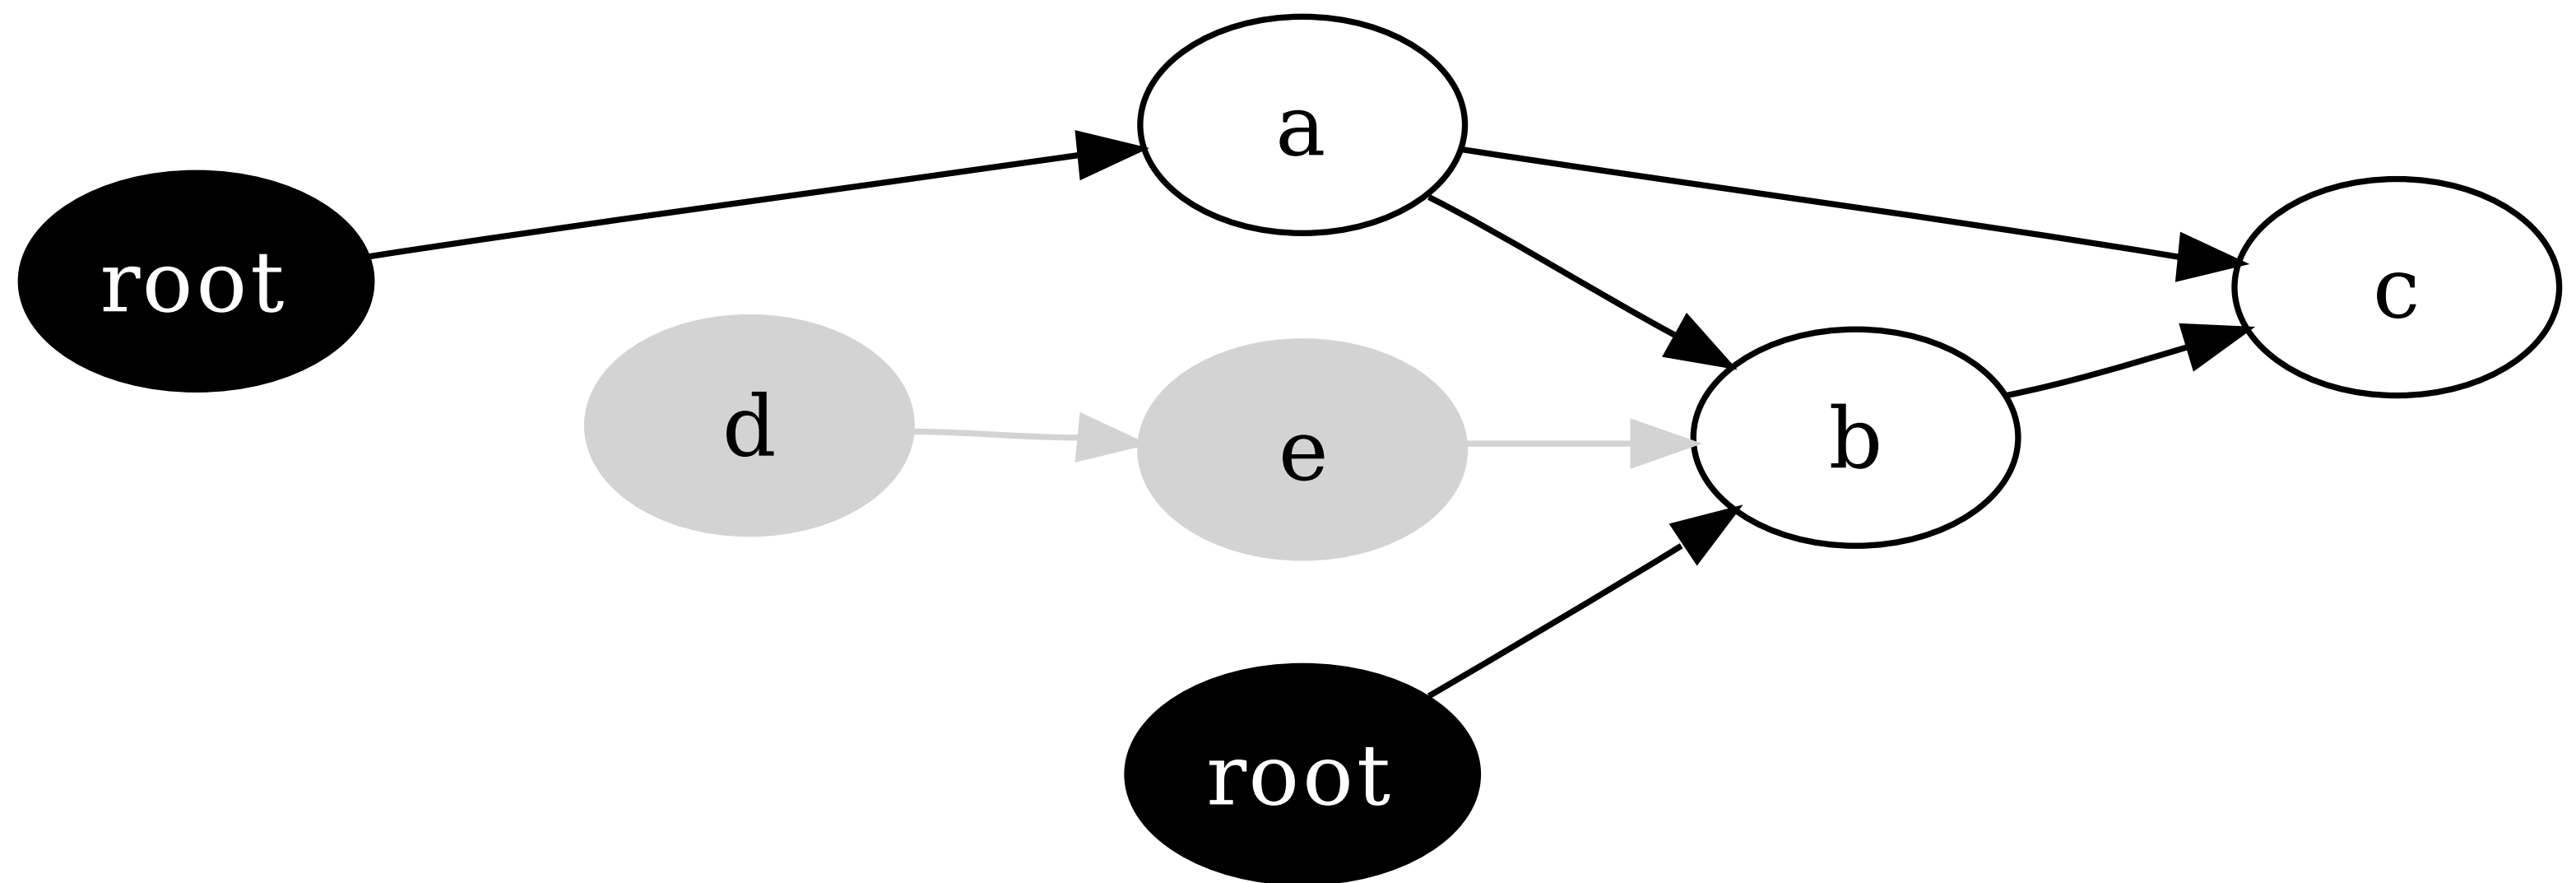
# How does garbage collection work?



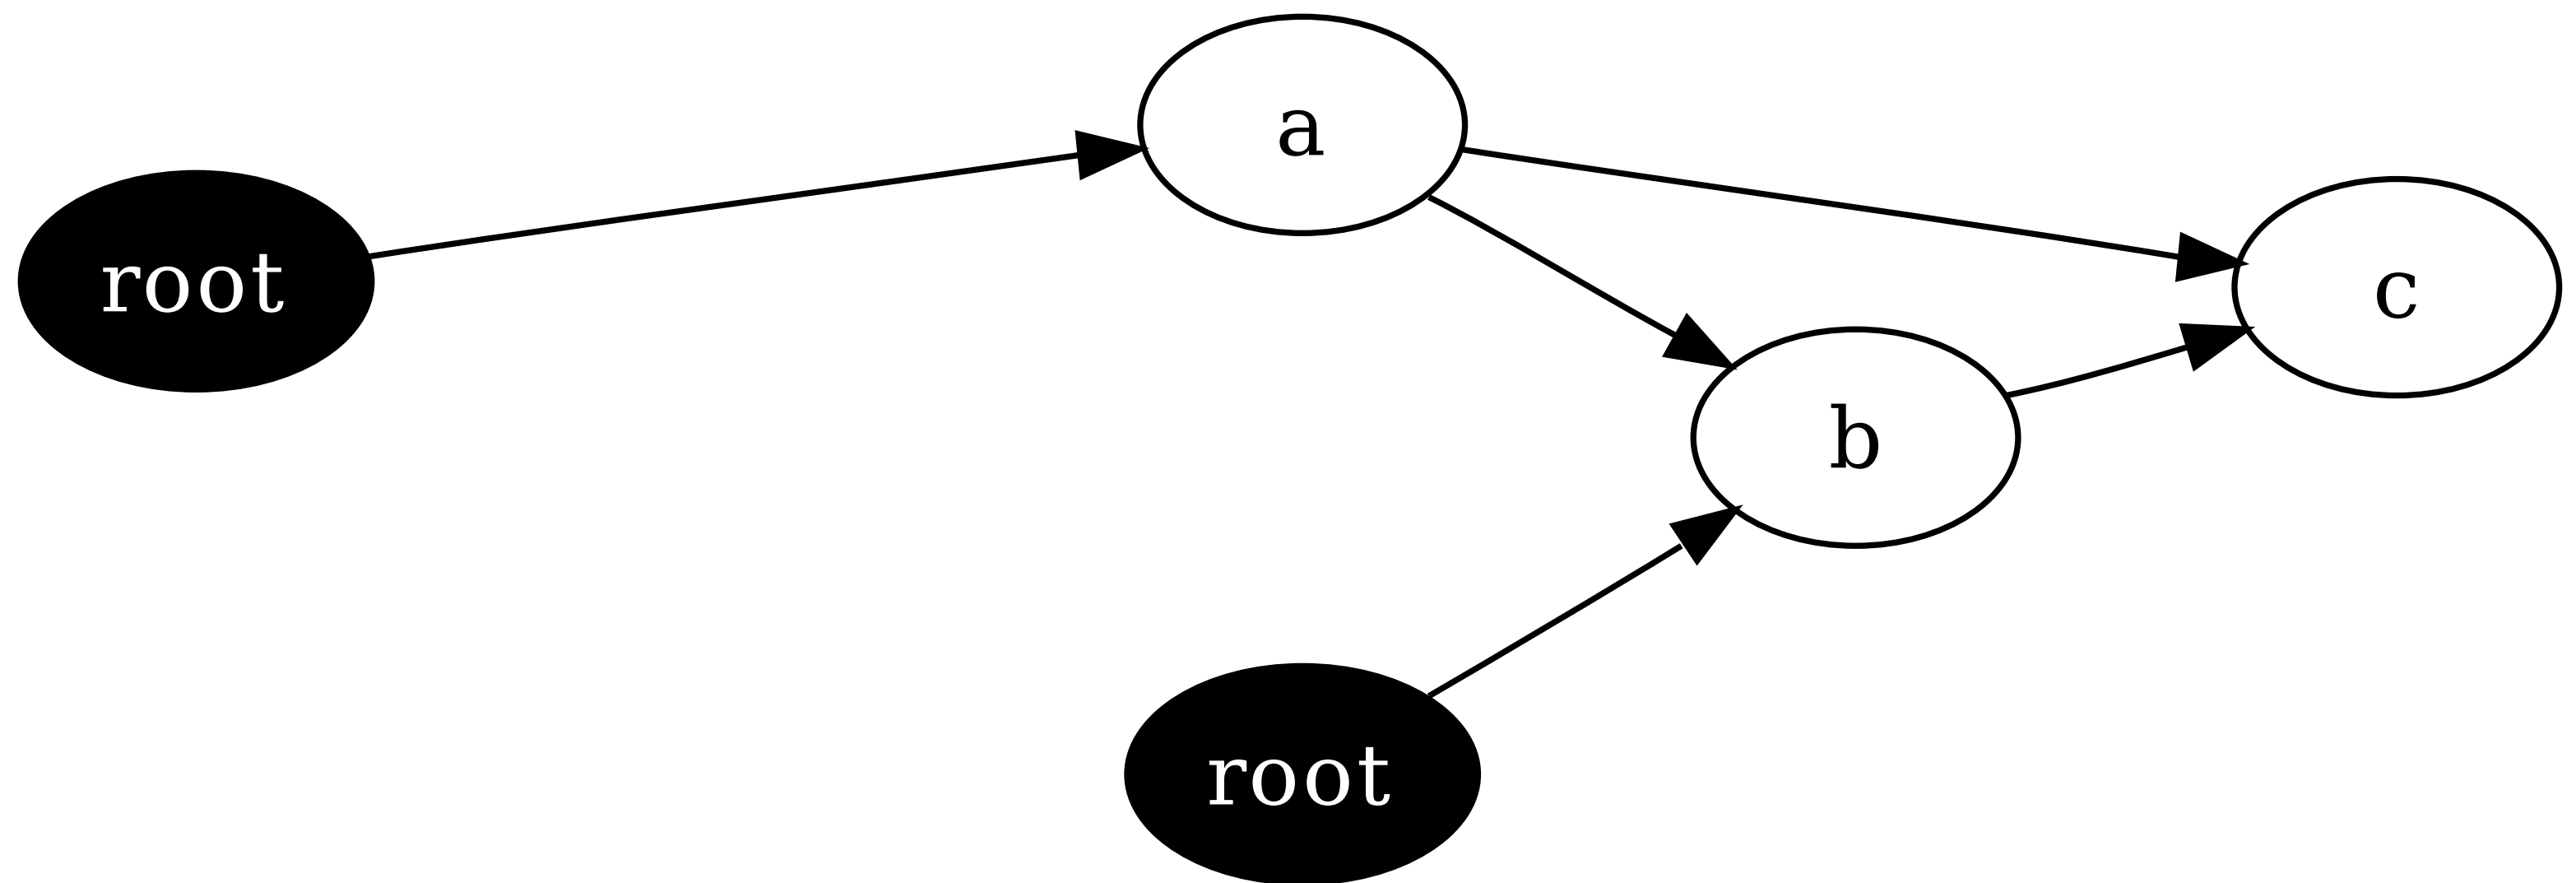
If the reference to D goes away...



We can no longer reach D or E.



So the collector reclaims them.





# The GC can't do everything.

## > Some things require manual cleanup.

- Listeners
- File descriptors
- Native memory
- External state (`IdentityHashMap`)

## > Tools at your disposal:

- `finally`
- Overriding `Object.finalize()`
- Reference queues

# Try `finally` first.

- > Reasons to not use `finally`:
  - More work for programmers
  - More error prone
  - Cleanup happens in main thread
- > ARM will help.



# What is a finalizer?

```
public class Foo extends Bar {  
    @Override protected void finalize() throws Throwable {  
        try {  
            ... // Clean up Foo.  
        } finally {  
            super.finalize(); // Clean up Bar.  
        }  
    }  
}
```

## Finalizers are seductively simple, but...

- > They're not guaranteed to run, especially not timely.
- > Avoid `System.runFinalizersOnExit()` and `runFinalization()`.
- > Undefined threading model, can run concurrently!
- > You must call `super.finalize()`.
- > Exceptions are ignored (per spec).
- > You can resurrect references.
- > Keeps objects alive longer.
- > Can make allocation/reclamation 430X slower (Bloch, Effective Java)

# An external resource

```
public class NativeResource {  
    public NativeResource() { init(); }  
  
    /** Allocates native memory. */  
    private native void init();  
  
    /** Writes to native memory. */  
    public native void write(byte[] data);  
  
    /** Frees native memory. */  
    @Override protected native void finalize();  
}
```

# Let's play War!

SegfaultFactory can cause a segfault if its finalizer executes after NativeResource's.

```
public class SegfaultFactory {  
    private final NativeResource nr;  
    public SegfaultFactory(NativeResource nr) {  
        this.nr = nr;  
    }  
  
    @Override protected void finalize() {  
        // 50/50 chance of failure  
        nr.write("I'm taking the VM with me!".getBytes());  
    }  
}
```



# Use protection.

Extend `NativeResource` and make it safe.

```
public class SafeNativeResource extends NativeResource {  
    private boolean finalized;  
  
    @Override public synchronized void write(byte[] data) {  
        if (!finalized) super.write(data);  
        else /* do nothing? */;  
    }  
  
    @Override protected synchronized void finalize() {  
        finalized = true;  
        super.finalize();  
    }  
}
```

# Finalizers are good for one thing.

## Logging warnings

# The alternative

An API-based approach:

```
public class WeakReference<T> {  
    public WeakReference(T referent) {  
        ...  
    }  
    public WeakReference(T referent,  
        ReferenceQueue<? super T> q) {  
        ...  
    }  
    public T get() {  
        ...  
    }  
    ...  
}
```



# The alternative

```
public class ReferenceQueue<T> {  
    public T poll() {  
        ...  
    }  
    public T remove() {  
        ...  
    }  
    public T remove(long timeout) {  
        ...  
    }  
}
```

# Reachability

- > An object is *reachable* if a live thread can access it.
- > Examples of heap roots:
  - System classes (which have static fields)
  - Thread stacks
  - In-flight exceptions
  - JNI global references
  - The finalizer queue
  - The interned String pool
  - etc. (VM-dependent)

# Let's make a map...

```
public class BytecodeCache {  
    final static Map<Class<?>, byte[]> cache = new MapMaker()  
        .weakKeys()  
        .softValues()  
        .makeComputingMap(new Function<Class<?>, byte[]>() {  
            public byte[] apply(Class<?> clazz) {  
                ...  
            }  
        })  
;  
  
    public static byte[] bytesFor(Class<?> clazz) {  
        return cache.get(clazz);  
    }  
}
```

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > **Strong**
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > **Soft**
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > **Weak**
- > Finalizer
- > Phantom, JNI weak
- > Unreachable



# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > **Finalizer**
- > Phantom, JNI weak
- > Unreachable

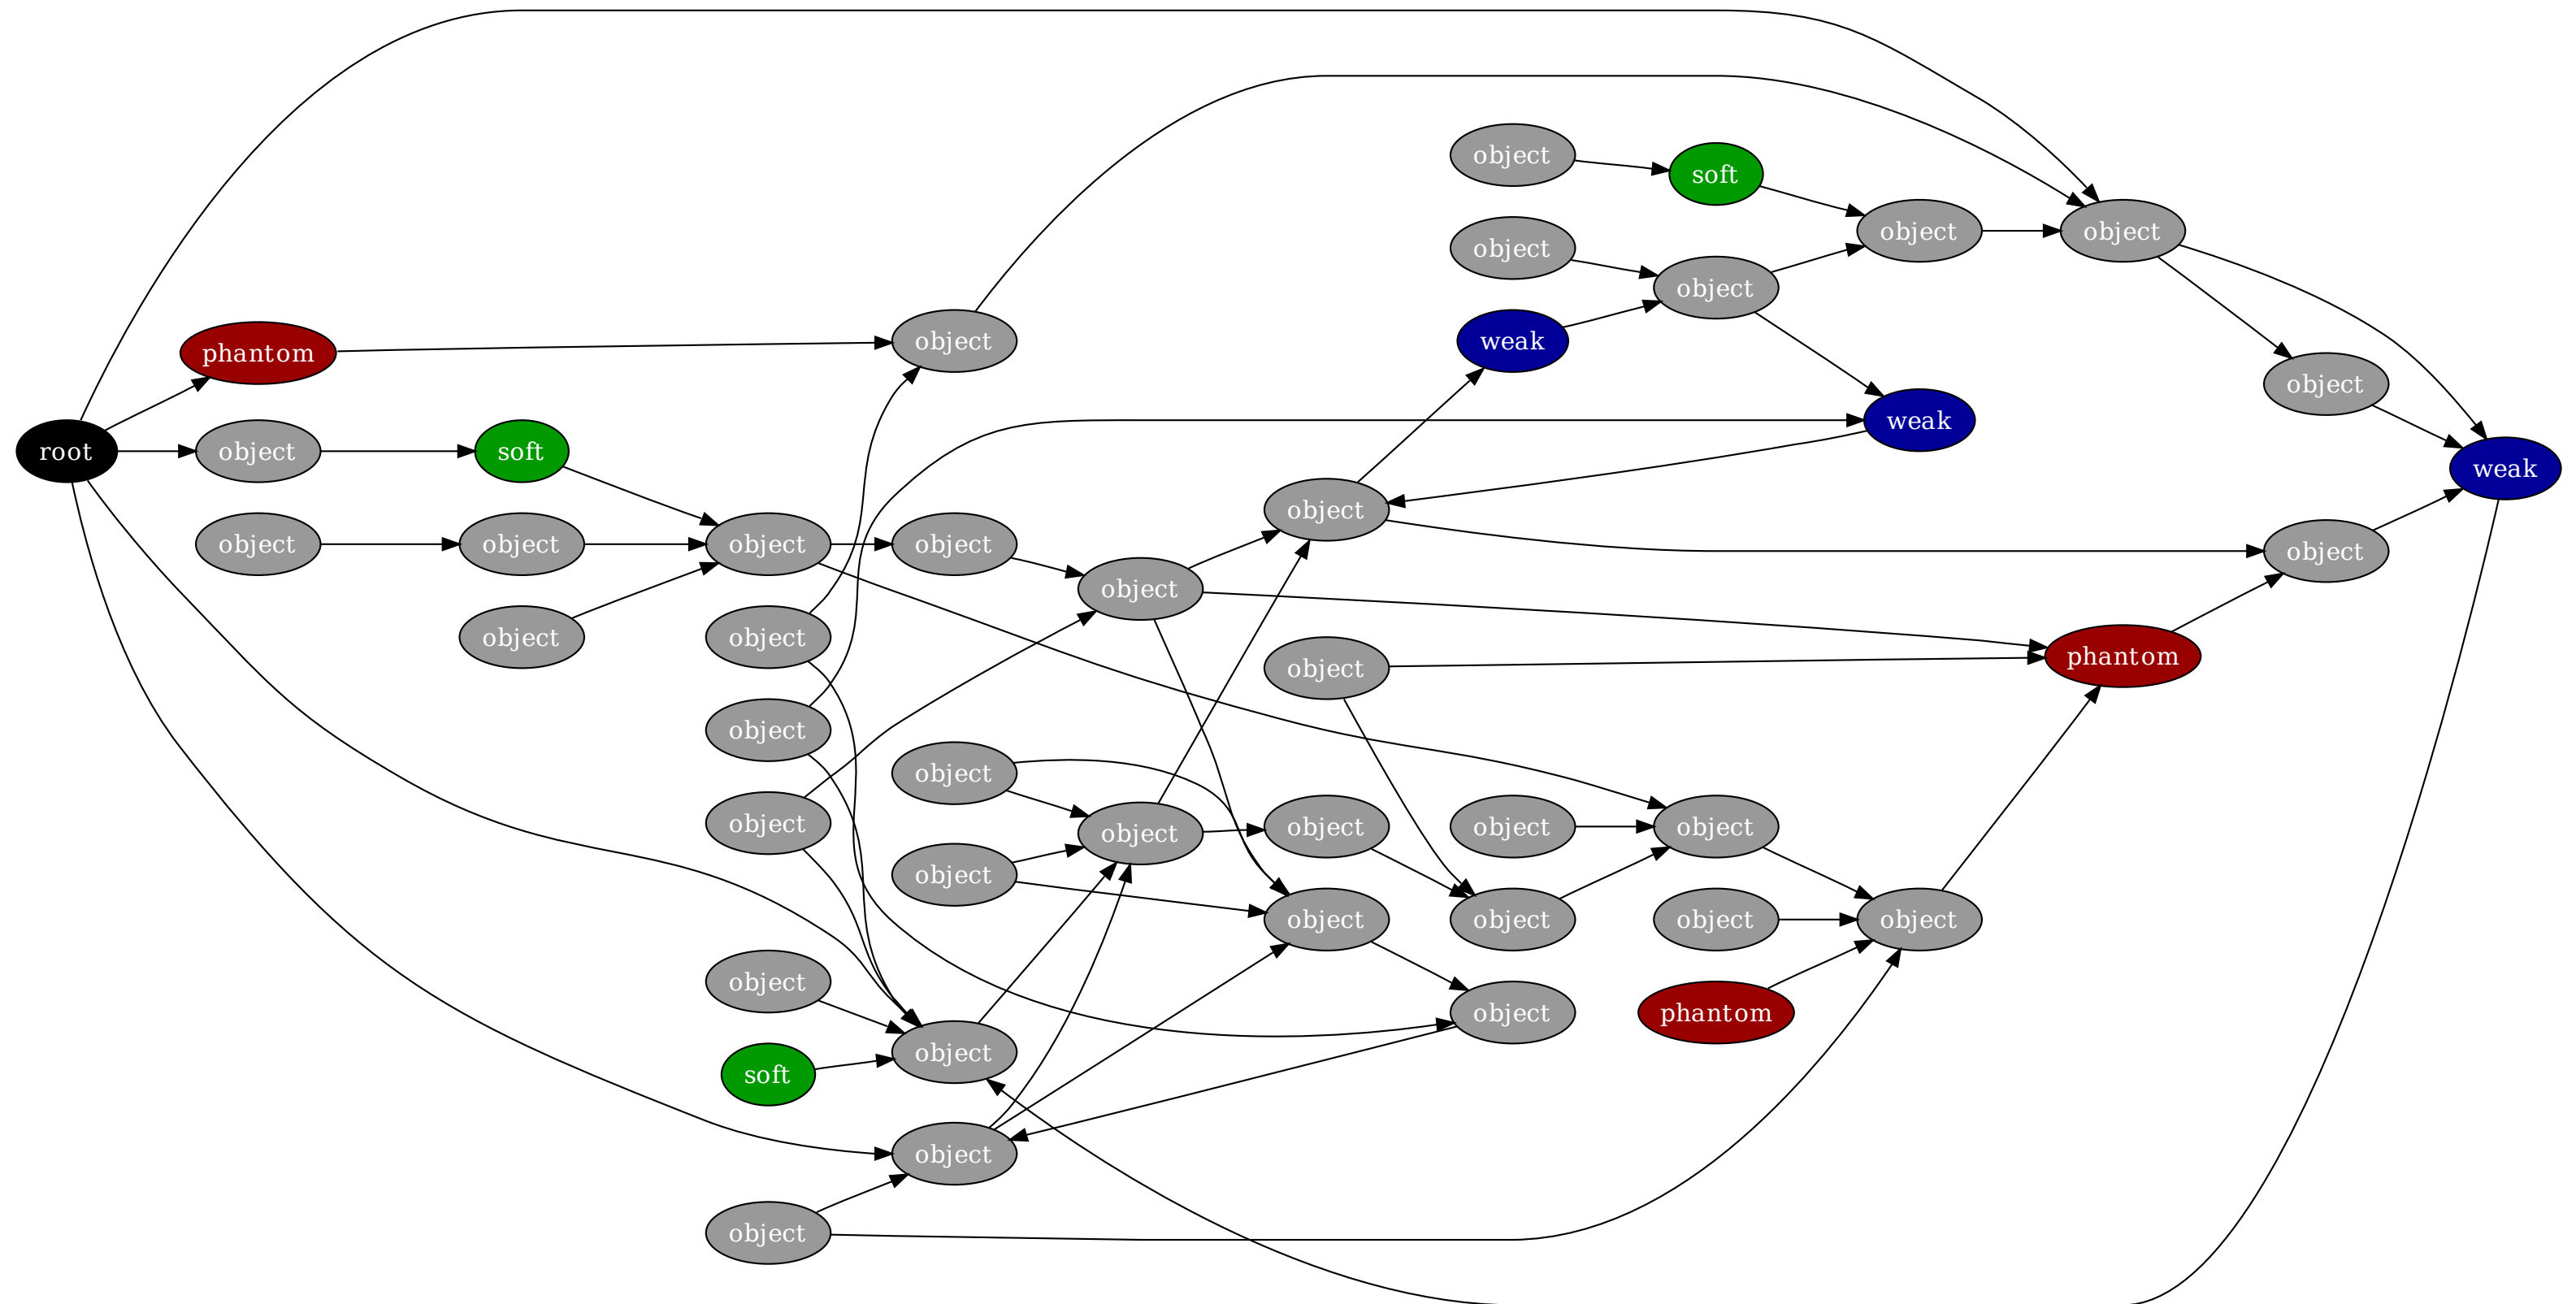
# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > **Phantom, JNI weak**
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > **Unreachable**


# Let's mark and sweep a heap!




# No objects are marked at first.



# 1. Start at a root.




## 2. Trace and mark strongly-referenced objects.







## 2. Trace and mark strongly-referenced objects.




## 2. Trace and mark strongly-referenced objects.




## 2. Trace and mark strongly-referenced objects.




## 2. Trace and mark strongly-referenced objects.



## 2. Trace and mark strongly-referenced objects.

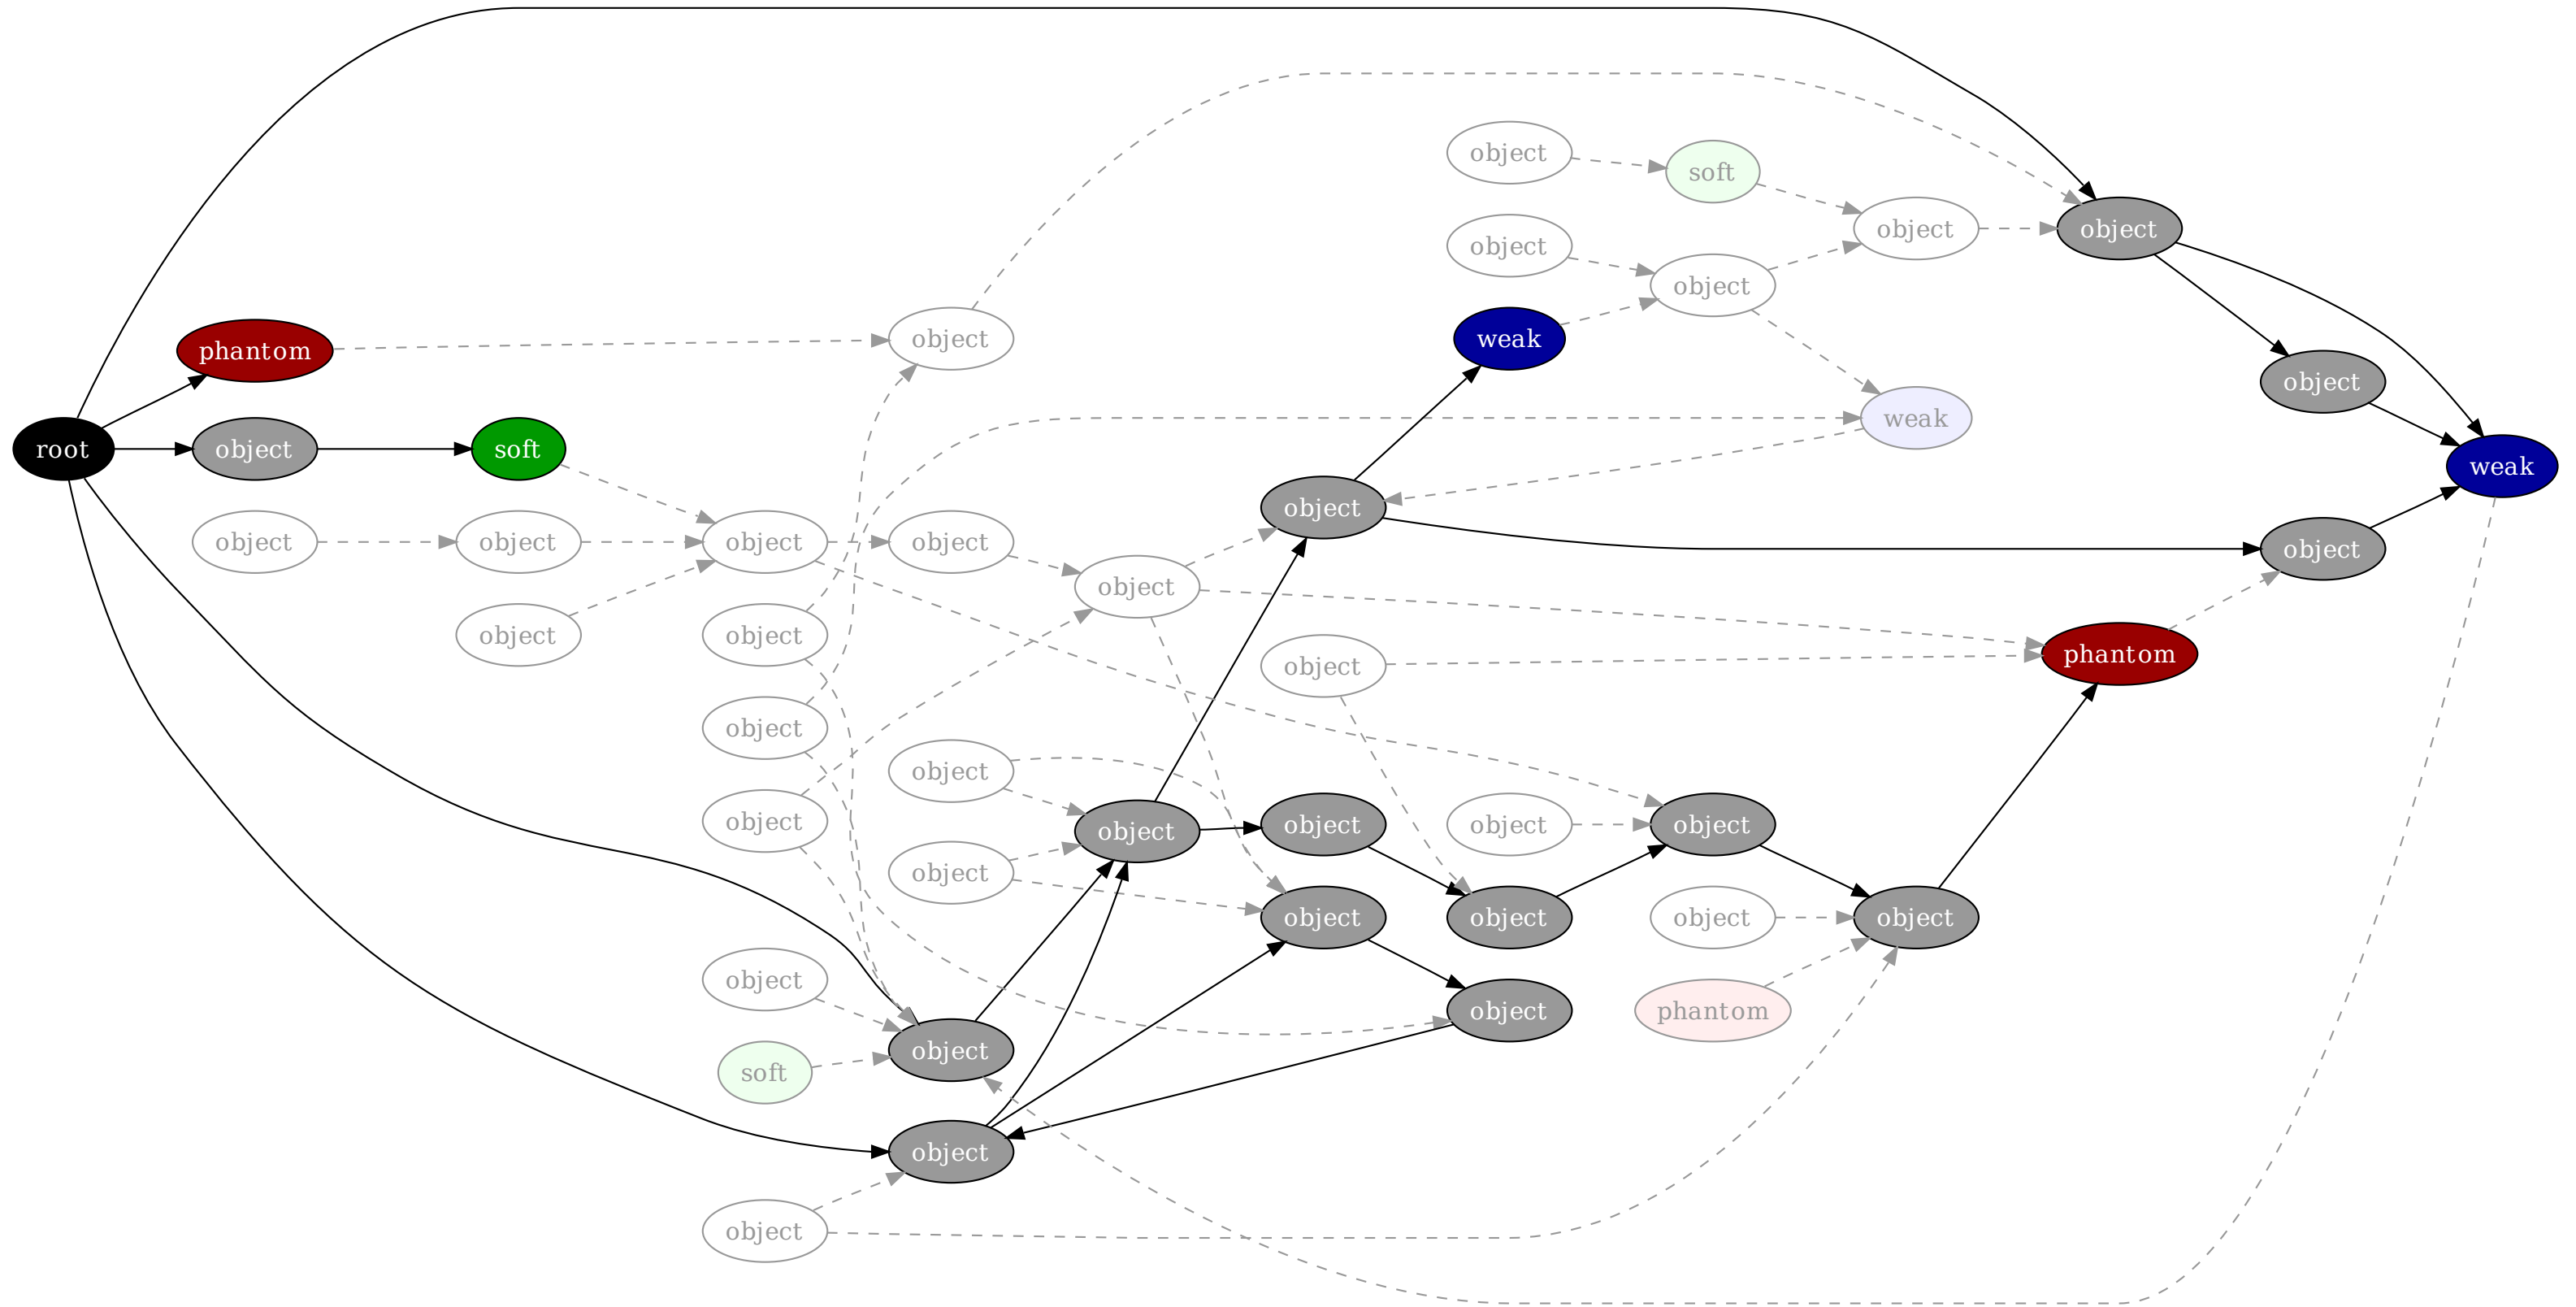


### 3. Optionally clear soft references.






## 4. Trace and mark softly-referenced objects.







## 4. Trace and mark softly-referenced objects.



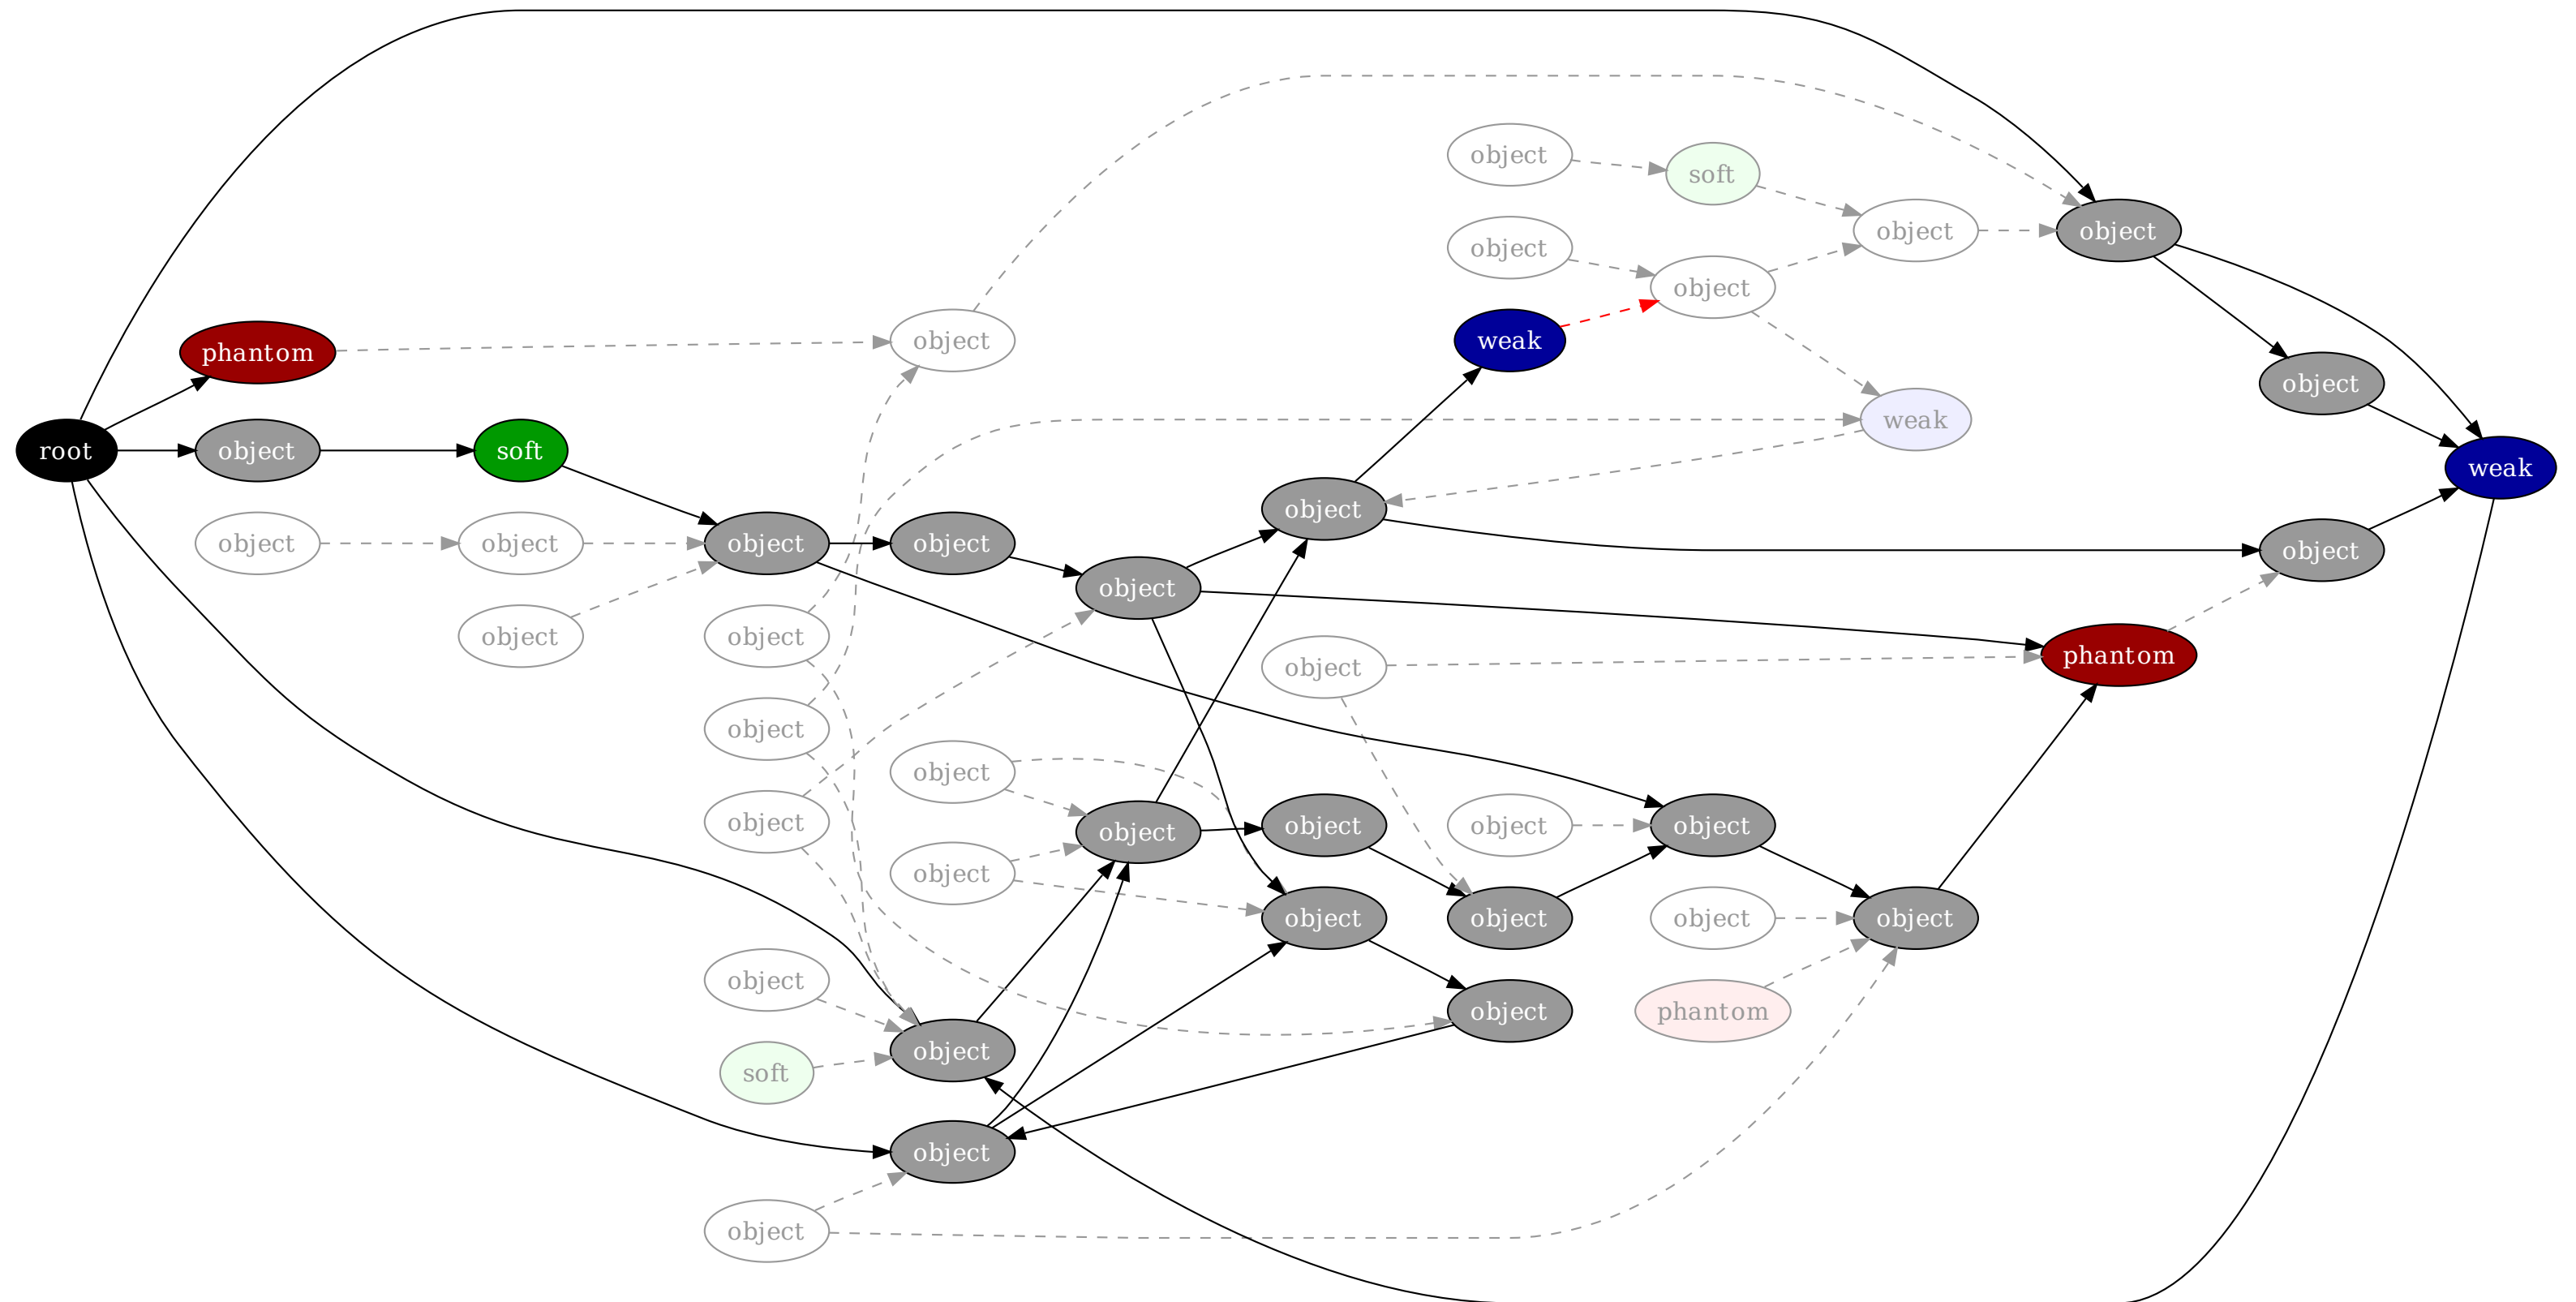
## 4. Trace and mark softly-referenced objects.




## 5. Clear weak references.




## 5. Clear weak references.



## 5. Clear weak references.




## 6. Enqueue finalizable objects.






## 6. Enqueue finalizable objects.




## 7. Repeat steps 1 through 5 for the queue.






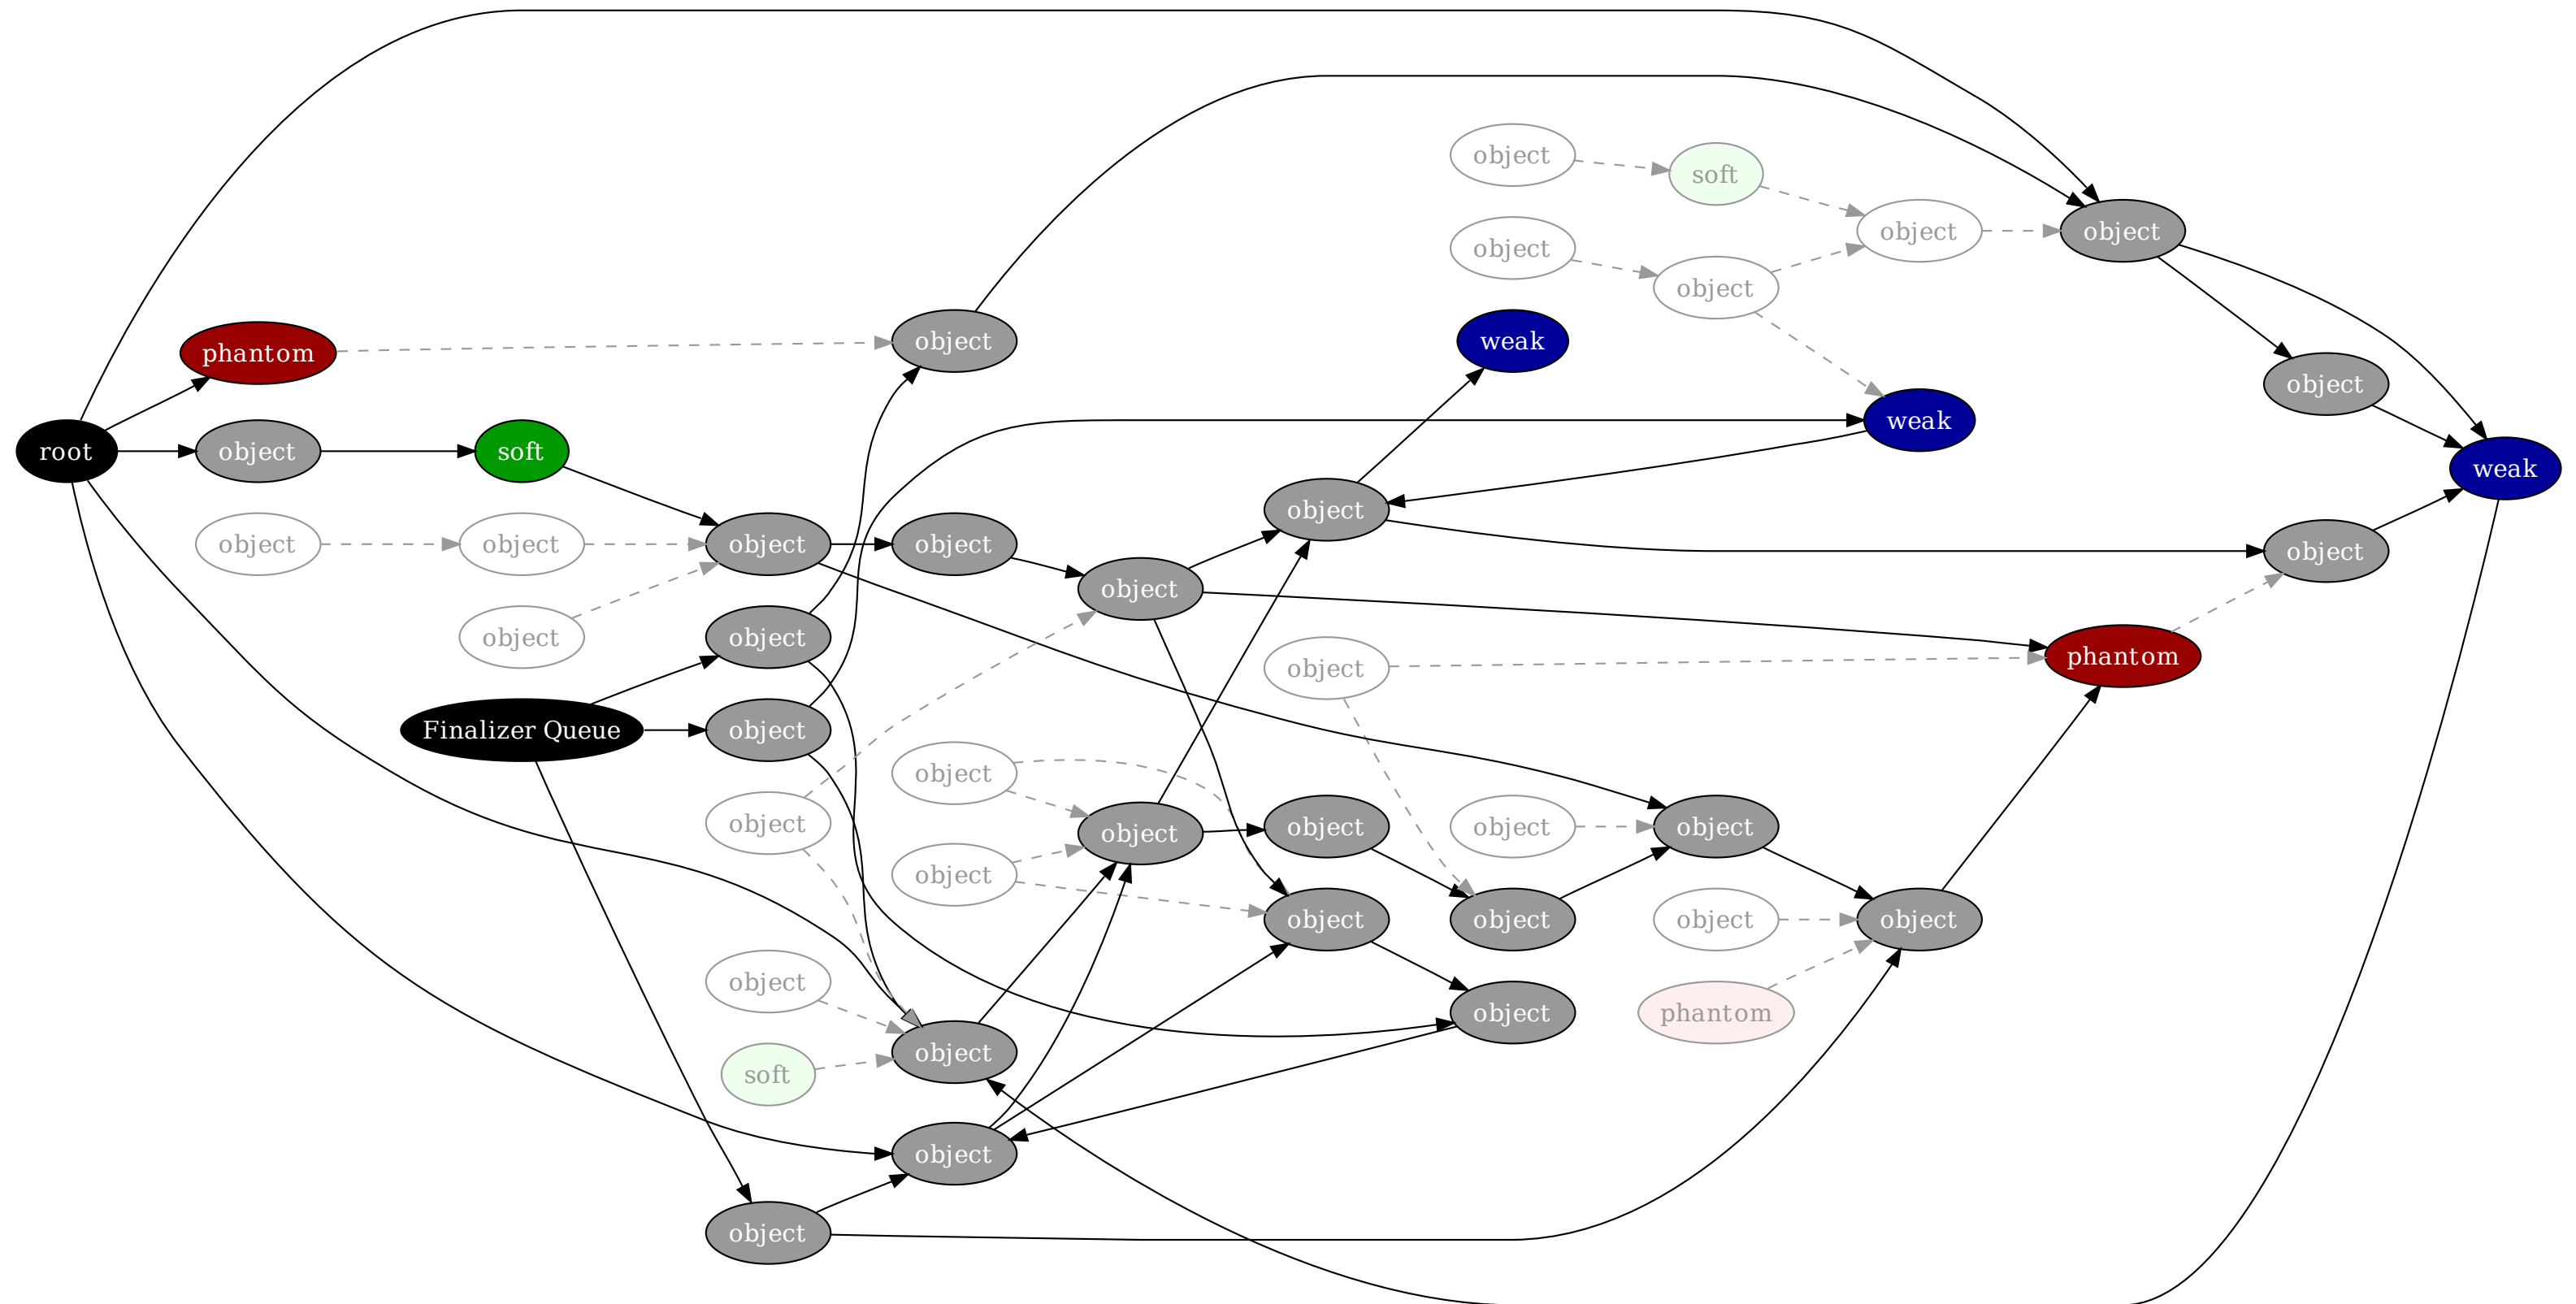
## 7. Repeat steps 1 through 5 for the queue.




## 7. Repeat steps 1 through 5 for the queue.




## 7. Repeat steps 1 through 5 for the queue.



## 8. Possibly enqueue phantom references.

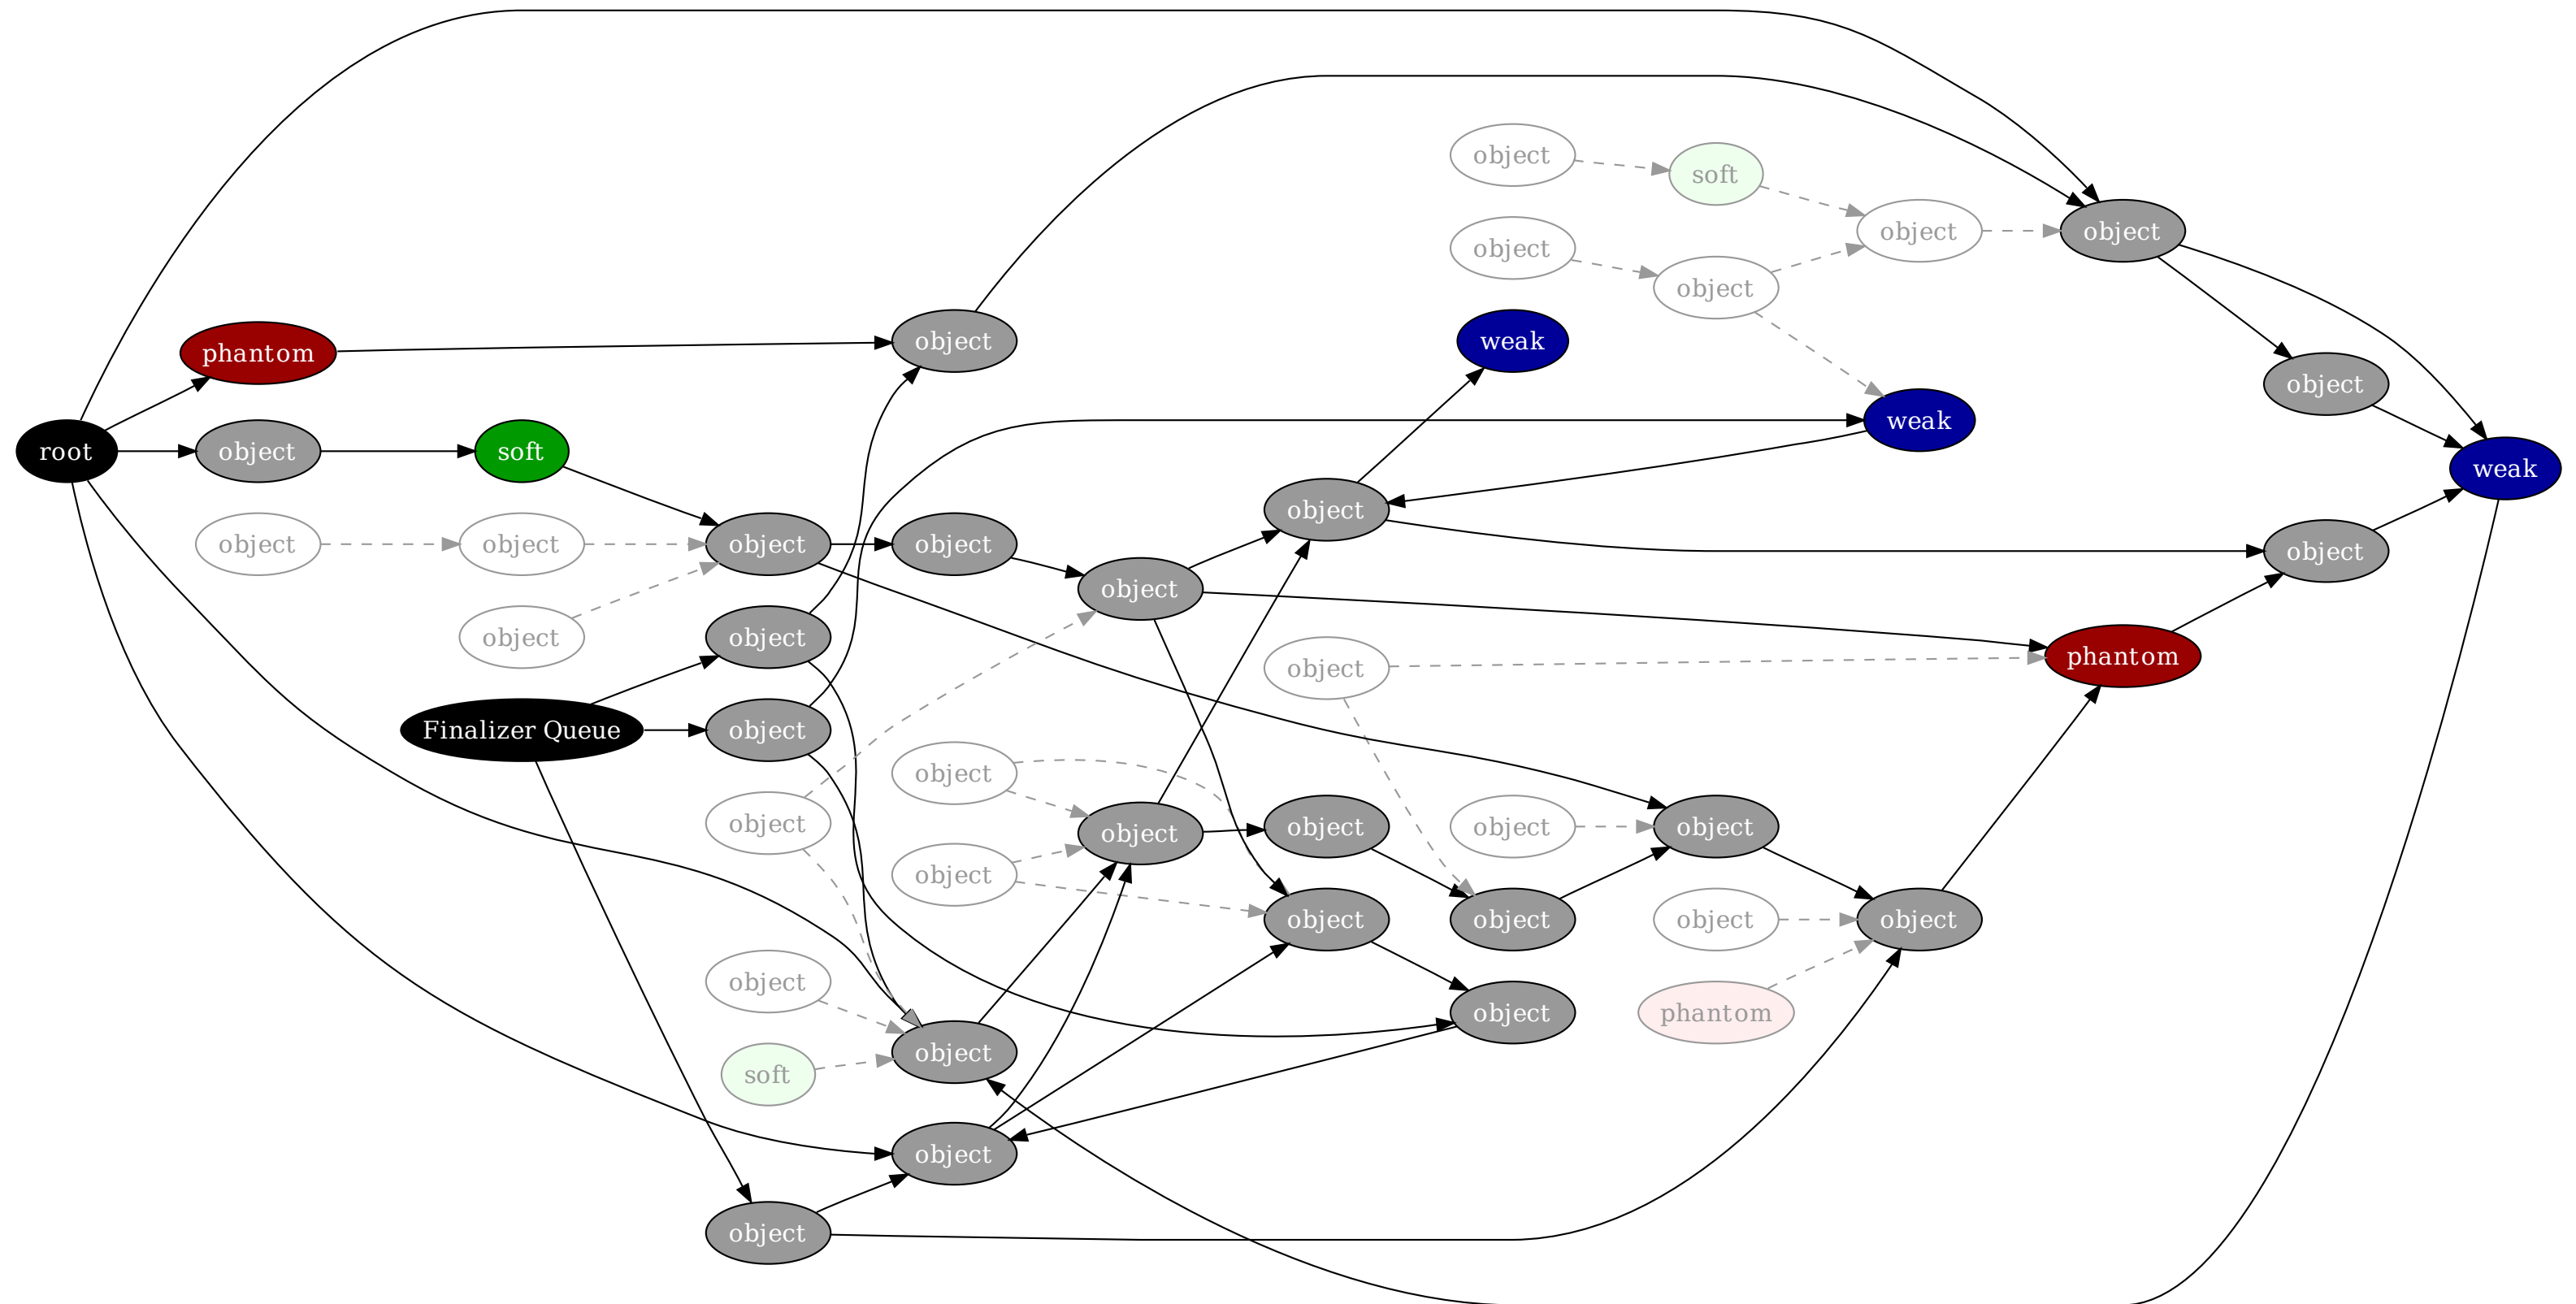


## 8. Possibly enqueue phantom references.






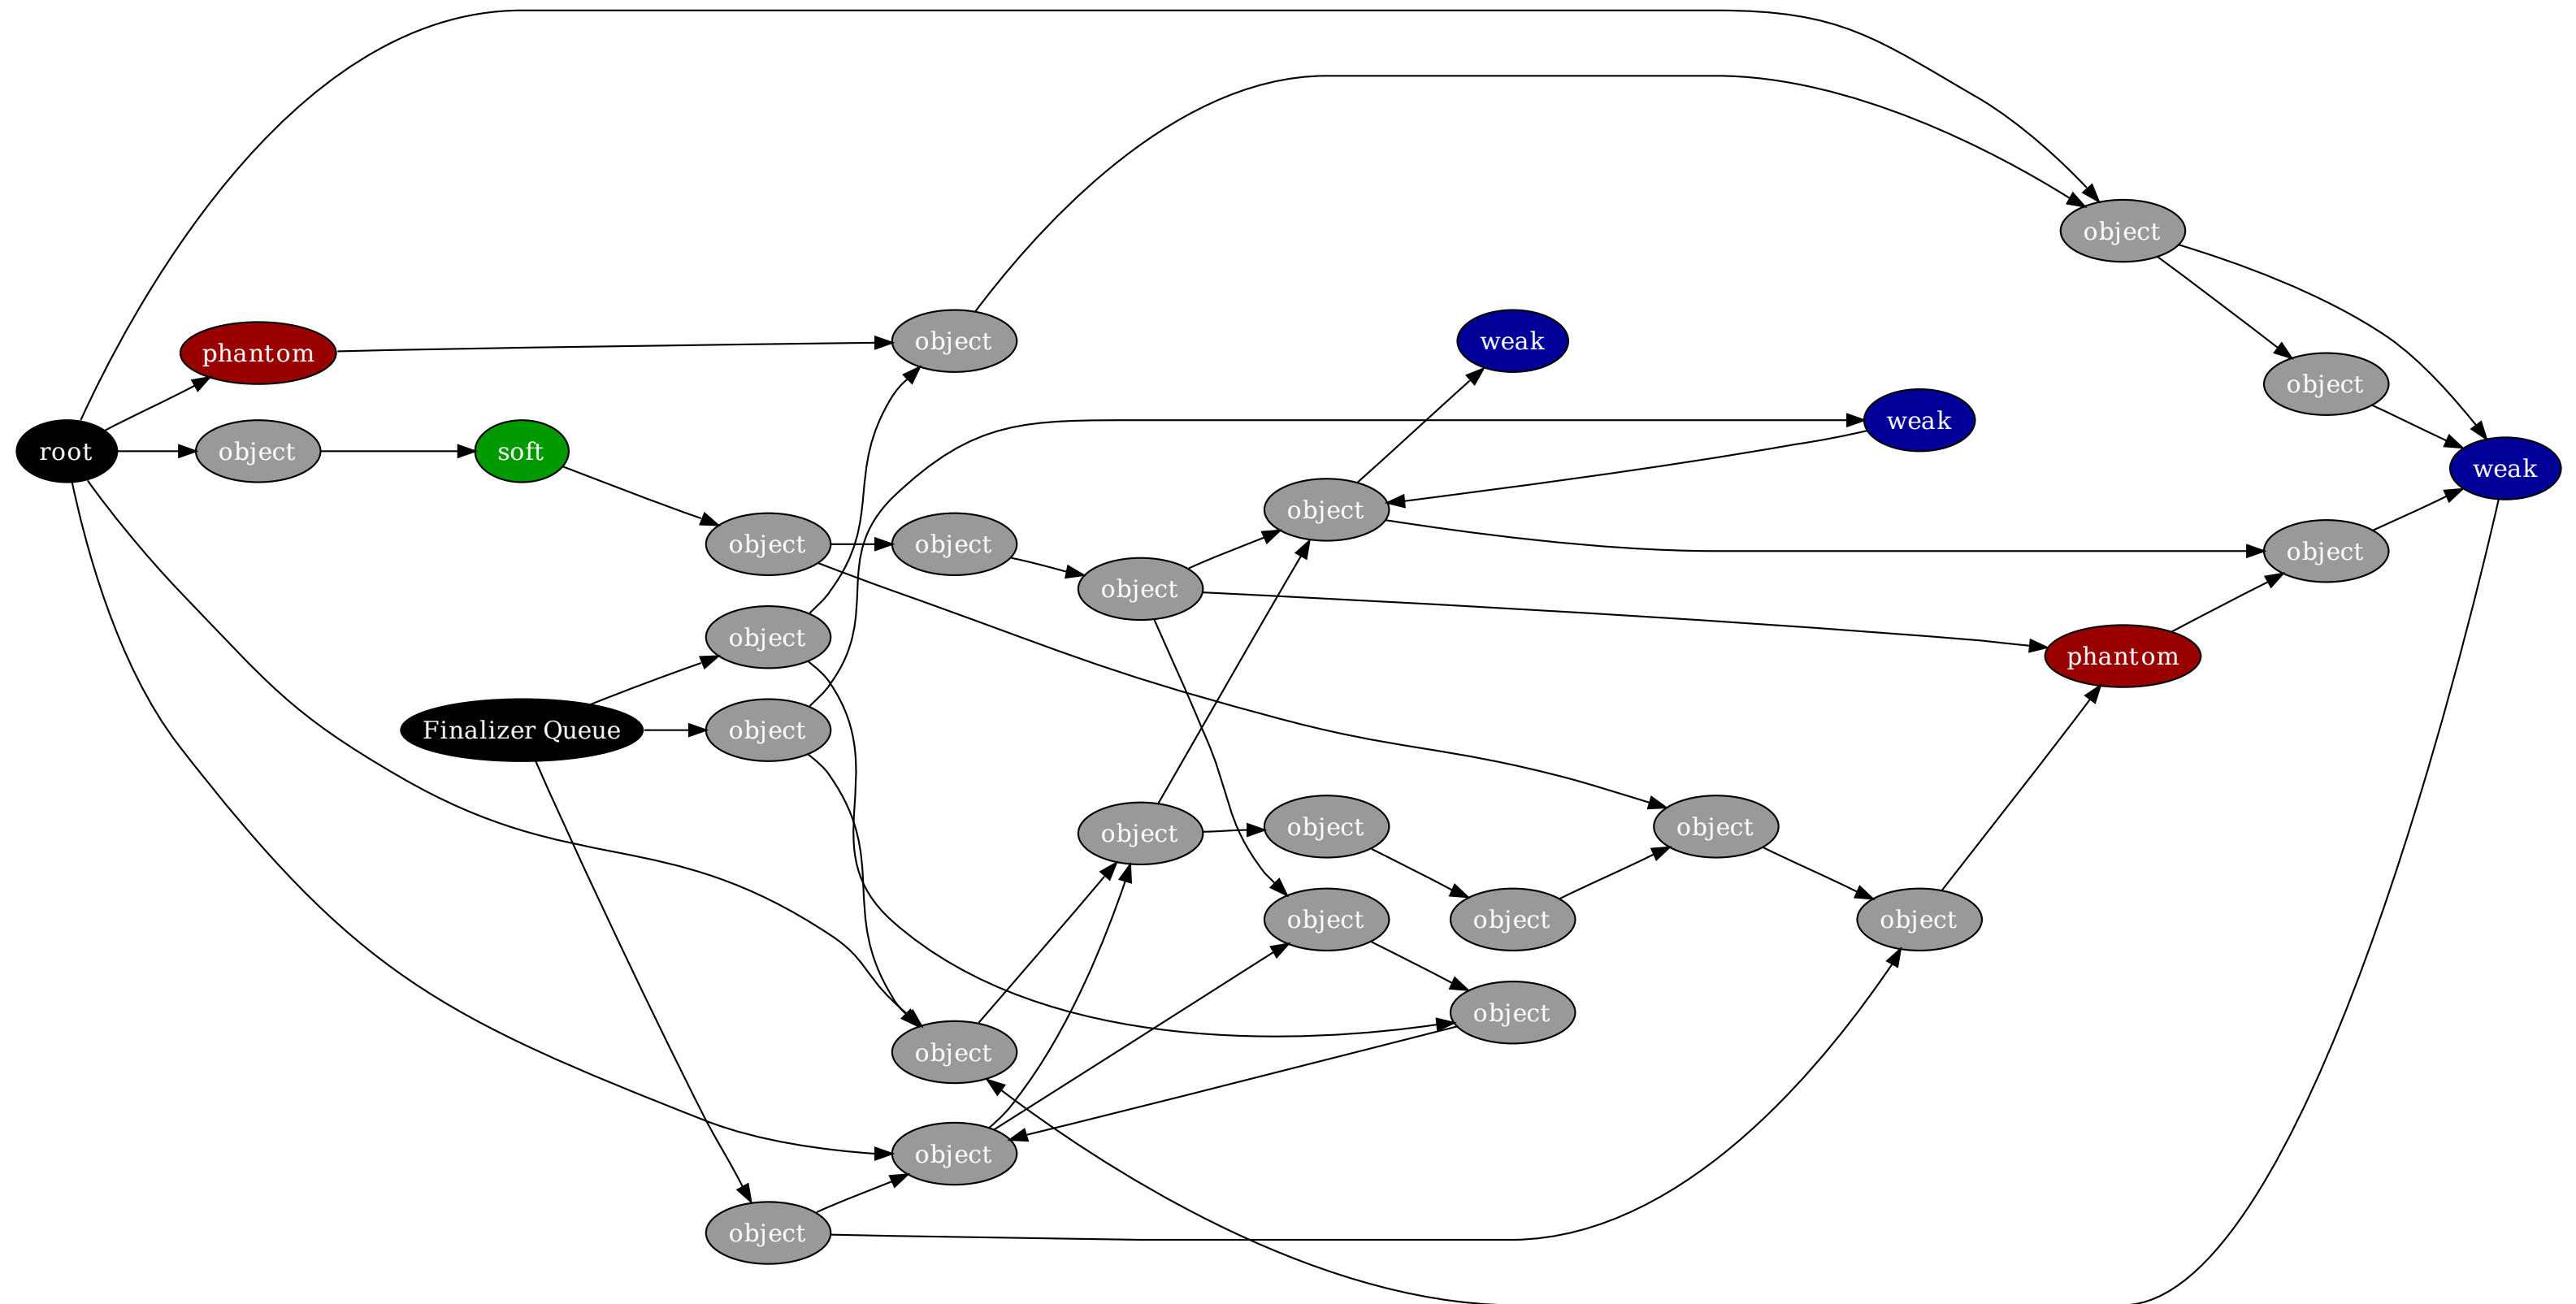
## 9. The remaining objects are dead.



## 9. The remaining objects are dead.



# 10. Repeat.





# Recap

1. Start at a root.
2. Trace and mark strongly-referenced objects.
3. Optionally clear soft references.
4. Trace and mark softly-referenced objects.
5. Clear weak references.
6. Enqueue finalizable objects.
7. Repeat steps 1 through 5 for the queue.
8. Possibly enqueue phantom references.
9. The remaining objects are dead.
10. Repeat.

# Weak references aren't for caching!

- > Many collectors will reclaim weak refs immediately.
- > Use soft reference for caching, as intended:

*“Virtual machine implementations are encouraged to bias against clearing recently-created or recently-used soft references.”*

- The `SoftReference` documentation