



Java is a trademark of Sun Microsystems, Inc.



# JavaOne<sup>SM</sup>

## The Ghost in the Virtual Machine A Reference to References

Bob Lee  
Google Inc.

# Goals

- > Perform manual cleanup the right way.
- > Take the mystery out of garbage collection.
- > Become honorary VM sanitation engineers.

# An external resource

```
public class NativeResource {  
    public NativeResource() { init(); }  
  
    /** Allocates native memory. */  
    private native void init();  
  
    /** Writes to native memory. */  
    public native void write(byte[] data);  
  
    /** Frees native memory. */  
    @Override protected native void finalize();  
}
```

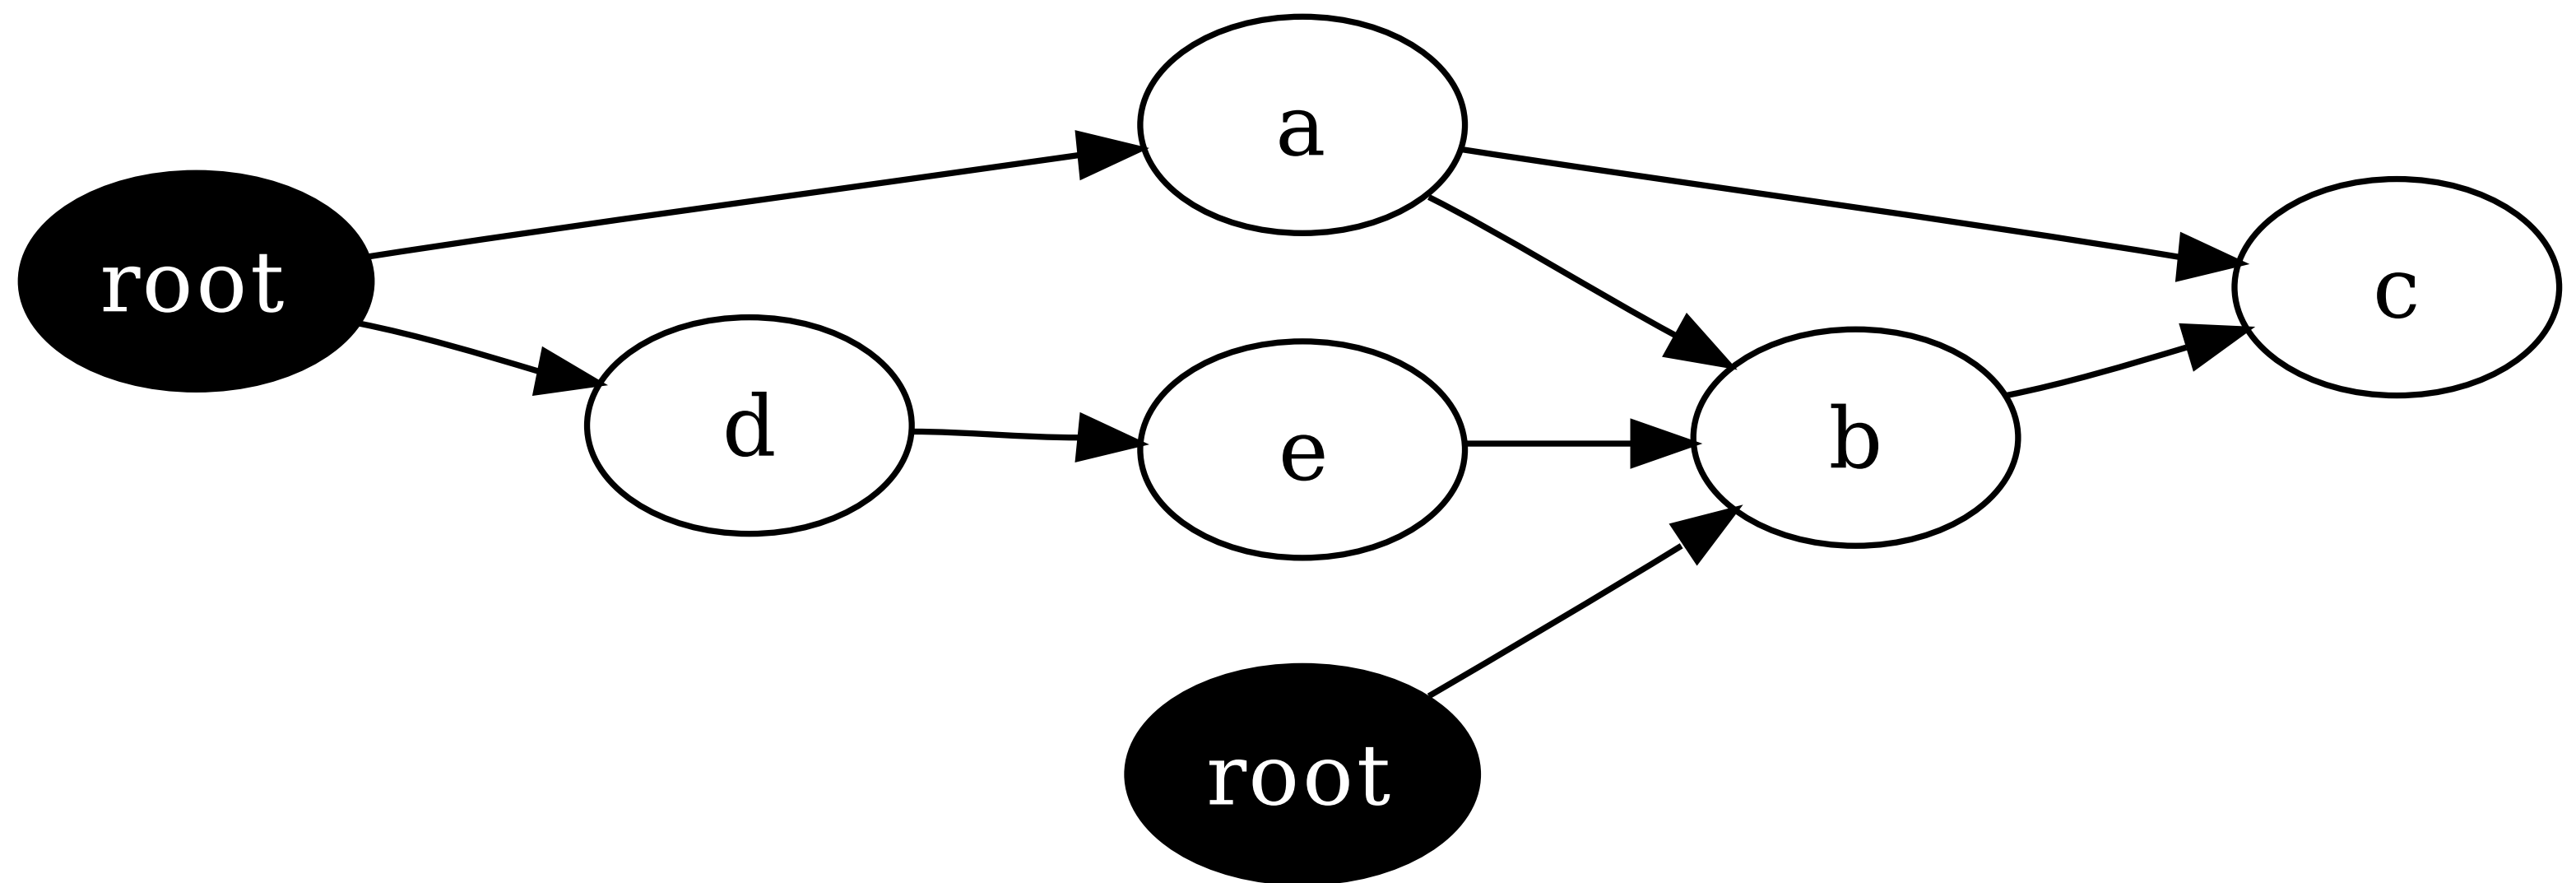
# Let's play War!

```
public class SegfaultFactory {  
    private final NativeResource nr;  
    public SegfaultFactory(NativeResource nr) {  
        this.nr = nr;  
    }  
  
    @Override protected void finalize() {  
        // 50/50 chance of failure  
        nr.write("I'm taking the VM with me!".getBytes());  
    }  
}
```

# Use protection.


```
public class SafeNativeResource extends NativeResource {  
    private boolean finalized;  
  
    @Override public synchronized void write(byte[] data) {  
        if (!finalized)  
            super.write(data);  
        else /* do nothing? */;  
    }  
  
    @Override protected synchronized void finalize() {  
        finalized = true;  
        super.finalize();  
    }  
}
```

# How does garbage collection work?






If the reference to D goes away...




We can no longer reach D or E.





So the collector reclaims them.



# Reachability

- > An object is *reachable* if a live thread can access it.
- > Examples of heap roots:
  - System classes (which have static fields)
  - Thread stacks
  - In-flight exceptions
  - JNI global references
  - The finalizer queue
  - The interned String pool
  - etc. (VM-dependent)

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > **Strong**
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > **Soft**
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > **Weak**
- > Finalizer
- > Phantom, JNI weak
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > **Finalizer**
- > Phantom, JNI weak
- > Unreachable



# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > **Phantom, JNI weak**
- > Unreachable

# Dante's Heap - The Levels of Reachability

- > Strong
- > Soft
- > Weak
- > Finalizer
- > Phantom, JNI weak
- > **Unreachable**

# Two options for freeing native resources

- > Use a finalizer.
  - You must defend against subsequent use!
- > Or use a phantom reference.

# Weak references aren't for caching!

- > Many collectors will reclaim weak refs immediately.
- > Use soft reference for caching, as intended:

*“Virtual machine implementations are encouraged to bias against clearing recently-created or recently-used soft references.”*

- The `SoftReference` documentation